

Introduction to MySQL Query Tuning

for Dev[Op]s

October 4, 2019

Sveta Smirnova



PERCONA

Table of Contents

- Basics
- When MySQL Uses Indexes
- Diagnostics
 - EXPLAIN: estimation on how Optimizer works
 - Real Numbers: Inside Storage Engine
 - Real Numbers: Inside the Server
- How to Affect Query Plans

Sveta Smirnova




- MySQL Support engineer
- Author of
 - **MySQL Troubleshooting**
 - JSON UDF functions
 - FILTER clause for MySQL
- Speaker
 - Percona Live, OOW, Fosdem, DevConf, HighLoad...



PERCONA


Basics

Troubleshooting Workflow



PERCONA

Troubleshooting Workflow: This Session



The Query

```
$system = System::factory()  
    ->setName($this->form->get(Field::NAME))  
    ->setDescription(  
        $this->form->get(Field::DESCRIPTION)  
    );  
DAO::system()->take($system);
```

The Query

```
$system = System::factory()  
  ->setName($this->form->get(Field::NAME))  
  ->setDescription(  
    $this->form->get(Field::DESCRIPTION)  
  );  
DAO::system()->take($system);
```


The Query

```
cursor = conn.cursor()
q = '''UPDATE 'foo' SET my_date=NOW(),
      subject = %s,
      msg = %s,
      address = %s,
      updated_at = NOW()
      WHERE id=%s
      , , ,
cursor.execute(q, [
    remote_resp.get('subject'),
    remote_resp.get('msg'),
    remote_resp.get('address'),
    my_id
])
```

The Query

```
cursor = conn.cursor()
q = '''UPDATE 'foo' SET my_date=NOW(),
      subject = %s,
      msg = %s,
      address = %s,
      updated_at = NOW()
      WHERE id=%s
      , , ,
cursor.execute(q, [
    remote_resp.get('subject'),
    remote_resp.get('msg'),
    remote_resp.get('address'),
    my_id
])
```

The Query

```
SELECT dept_name, title, gender,  
       min(salary) AS mins, max(salary) AS maxs  
FROM employees  
JOIN salaries USING(emp_no)  
JOIN titles USING(emp_no)  
JOIN dept_emp USING(emp_no)  
JOIN departments USING(dept_no)  
JOIN dept_manager USING(dept_no)  
WHERE dept_manager.to_date = '9999-01-01'  
GROUP BY dept_name, title, gender  
ORDER BY gender, maxs DESC;
```

The Query

```
SELECT dept_name, title, gender,  
       min(salary) AS mins, max(salary) AS maxs  
FROM employees  
JOIN salaries USING(emp_no)  
JOIN titles USING(emp_no)  
JOIN dept_emp USING(emp_no)  
JOIN departments USING(dept_no)  
JOIN dept_manager USING(dept_no)  
WHERE dept_manager.to_date = '9999-01-01'  
GROUP BY dept_name, title, gender  
ORDER BY gender, maxs DESC;
```



Allways Tune Raw Query

- PMM QAN



Allways Tune Raw Query

- PMM QAN
- Slow Query Log

Allways Tune Raw Query

- PMM QAN
- Slow Query Log
- Application log

Allways Tune Raw Query

- PMM QAN
- Slow Query Log
- Application log
- ...

Slow is relative

- Mind you data!
- 75,000,000 rows
 - (INT, INT)
 - $75,000,000 * (4 + 4) = 600,000,000$ bytes = 572 MB
 - (INT, INT, DATETIME, VARCHAR(255), VARCHAR(255))
 - $75,000,000 * (4 + 4 + 8 + 256 + 256) = 39,600,000,000$ bytes = 37 G
- $39,600,000,000 / 600,000,000 = 66$

Slow is relative

- Mind you data!
- Mind use case
 - Popular website
 - Admin interface
 - Weekly cron job

Slow is relative

- Mind you data!
- Mind use case
- Mind location
 - Server, used by multiple connections
 - Dedicated for OLAP queries

Why Query can be Slow

- MySQL performs a job to execute a query


Why Query can be Slow

- MySQL performs a job to execute a query
- In worst case scenario it will do a full table scan
 - CREATE INDEX
 - ANALYZE TABLE ... UPDATE HISTOGRAM ON


Why Query can be Slow

- MySQL performs a job to execute a query
- In worst case scenario it will do a full table scan
 - CREATE INDEX
 - ANALYZE TABLE ... UPDATE HISTOGRAM ON
- Incorrect index can be used

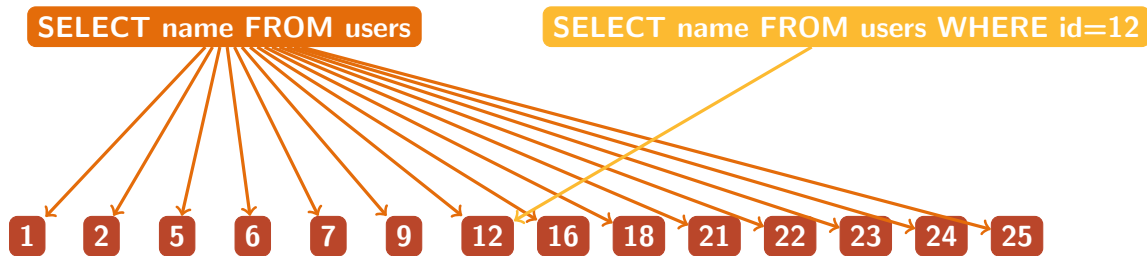
Query Execution Workflow




Full Table Scan



After Index Added



MySQL Indexes



- B-Tree (Mostly)
- Fractal Tree
- LSM Tree
- R-Tree (Spatial)
- Hash (Memory SE)
- Engine-dependent

How to Create an Index

- Single column

```
CREATE INDEX index_name ON  
the_table(the_column)
```

- Multiple columns

```
CREATE INDEX index_name ON  
the_table(column1, column2)
```

How to Create an Index

- Single column

```
ALTER TABLE table_name ADD INDEX  
[index_name] (the_column)
```

- Multiple columns

```
ALTER TABLE table_name ADD INDEX  
[index_name] (column1, column2)
```

When MySQL Uses Indexes



PERCONA

Conditions

- WHERE the_column = a_value
- WHERE the_column IN(value1, value2, value3)
- WHERE the_column LIKE 'value%'
- WHERE the_column LIKE '%value'

Conditions

- `WHERE left_part = value1 AND right_part = value2`
- `WHERE left_part = value1 OR right_part = value2`
- `WHERE right_part = value1 AND left_part = value2`
- ~~`WHERE right_part = value1 OR left_part = value2`~~

Joins

- `table1 JOIN table2 ON table1.column1 = table2.column2`

Joins

- `table1 JOIN table2 ON table1.column1 = table2.column2`
- Same as `FROM table1, table2 WHERE table1.column1 = table2.column2`

GROUP BY

- GROUP BY the_column
- GROUP BY left_part, right_part
- ~~GROUP BY right_part, left_part~~
- ~~GROUP BY the_index, another_index~~

ORDER BY

- ORDER BY the_column
- ORDER BY left_part, right_part
- ~~ORDER BY right_part, left_part~~
- ~~ORDER BY the_index, another_index~~

ORDER BY

~~5.7 ORDER BY left_part DESC, right_part ASC~~

8.0 ORDER BY left_part DESC, right_part
ASC

- left_part **must** be **descending**
- right_part **must** be **ascending**
- the_index(left_part DESC, right_part ASC)

Expressions

- Deterministic, **built-in**
 - Return same value for the same argument
 - WHERE the_column = FLOOR(123.45)

Expressions

- Deterministic, **built-in**
 - Return same value for the same argument
 - `WHERE the_column = FLOOR(123.45)`
- Non-deterministic
 - Return different values for different invocations
 - ~~`WHERE the_column = RAND() * 100`~~

Expressions

- Deterministic, **built-in**
 - Return same value for the same argument
 - `WHERE the_column = FLOOR(123.45)`
 - Non-deterministic
 - Return different values for different invocations
 - ~~`WHERE the_column = RAND() * 100`~~
 - Stored functions and UDFs
 - Indexes are not used
- Use generated column indexes



Diagnostics

Diagnostics

EXPLAIN: estimation on how Optimizer works

How to Find how MySQL Uses Indexes

- EXPLAIN
 - Estimates what happens during query execution

5.6- EXTENDED

5.6- PARTITIONS

5.6+ FORMAT=JSON

8.0+ FORMAT=TREE

How to Find how MySQL Uses Indexes

- EXPLAIN
 - Estimates what happens during query execution

5.6- EXTENDED

5.6- PARTITIONS

5.6+ FORMAT=JSON

8.0+ FORMAT=TREE

- INFORMATION_SCHEMA.OPTIMIZER_TRACE
 - Real data, collected after query was executed
 - **Advanced topic**

Effect of Indexes: Before

```
mysql> explain select * from t1\G
```

```
***** 1. row *****
```

```
...
```

```
rows: 12
```

```
Extra: NULL
```

```
mysql> explain select * from t1 where f2=12\G
```

```
***** 1. row *****
```

```
...
```

```
key: NULL
```

```
...
```

```
rows: 12
```

```
Extra: Using where
```

Same number of examined rows for both queries



PERCONA

Effect of Indexes: After

```
mysql> alter table t1 add index(f2);  
Query OK, 12 rows affected (0.07 sec)  
Records: 12  Duplicates: 0  Warnings: 0
```

```
mysql> explain select * from t1 where f2=12\G  
***** 1. row *****  
  
...  
key: f2  
key_len: 5  
ref: const  
rows: 1  
Extra: NULL  
1 row in set (0.00 sec)
```

**Much more effective!
Only 1 row examined**



PERCONA

EXPLAIN: overview

mysql> explain extended select * from t1 join t2 where t1.int_key=1;

Number of select		Tables, for which information is printed		Possible keys		Length of the key		Number of examined rows		Additional information	
id	select_type	table	type	possible_keys	key	key_len	ref	rows	f...	Extra	
1	SIMPLE	t1	ref	int_key,ik	int_key	5	const	4	100.	NULL	
1	SIMPLE	t2	index	NULL	pk	9	NULL	6	100.	Using in	

Using join buffer: Block nested loop join

2 rows in set (0.00 sec) 1 warning (0.00 sec)

Product of rows here: how many rows in all tables will be accessed
For this example estimated value is $4 \times 6 = 24$

% of filtered rows
rows x filtered / 100 — number of row
which will be joined with another table

Note (Code 1003): /* select#1 */ select 'test'.'t1'.'pk' AS 'pk','test'.'t1'.'int_key' AS 'int_key','test'.'t2'.'pk' AS 'pk','test'.'t2'.'int_key' AS 'int_key' from 'test'.'t1' join 'test'.'t2' where ('test'.'t1'.'int_key' = 'test'.'t2'.'int_key')

Actual (optimized) query as executed by MySQL Server

EXPLAIN in Details

```
mysql> explain extended select * from t1 join t2 where...
```

+----+-----+-----+-----+***				
id	select_type	table	type	***
+----+-----+-----+-----+***				
1	SIMPLE	t1	ref	***
1	SIMPLE	t2	index	***
+----+-----+-----+-----+***				

```
2 rows in set, 1 warning (0.00 sec)
```

SIMPLE;PRIMARY;UNION;DEPENDENT UNION;UNION RESULT;
SUBQUERY;DEPENDENT SUBQUERY;DERIVED;MATERIALIZED

system
const
eq_ref
ref
fulltext
ref_or_null
index_merge
unique_subquery
index_subquery
range
index
ALL



PERCONA

EXPLAIN in Details: keys

Keys, which can be used for resolving the query

Actual length of the key (Important for multiple-column keys)

```
mysql> explain extended select * from t1 join t2 where t1.int_key=1;
```

```
***+-----+-----+-----+-----+***
***| possible_keys | key      | key_len | ref      |***
***+-----+-----+-----+-----+***
***| int_key, ik   | int_key  | 5       | const    |***
***| NULL         | pk       | 9       | NULL     |***
***+-----+-----+-----+-----+***
```

```
2 rows in set, 1 warning (0.00 sec)
```

Constant
Numeric in our case

Index used
to resolve rows

Only one key was actually used

Which columns were compared with the index



PERCONA

EXPLAIN in Details: rows

Number of rows accessed

% of rows filtered

Additional information:
how query is resolved

Using filesort
Using temporary
etc.

```
mysql> explain extended select * from t1 join t2 where t1.int_key=1;
***+-----+-----+-----+
***| rows  | filtered | Extra |
***+-----+-----+-----+
***| 4X6   |         |       |
***| =     | 100.00  | NULL  |
***| 24    | 100.00  | Using index; Using join buffer (Block Nested Loop) |
***+-----+-----+-----+
2 rows in set, 1 warning (0.00 sec)
```

All rows used

EXPLAIN Type by Example: ALL

```
mysql> explain select count(*) from employees where hire_date > '1995-01-01'
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: employees
         type: ALL
possible_keys: NULL
         key: NULL
      key_len: NULL
         ref: NULL
         rows: 300157
  Extra: Using where
1 row in set (0.00 sec)
```

All rows in the table examined
Worst plan ever!

EXPLAIN Type by Example: range

- We need to add index to table employees first

```
mysql> alter table employees add index(hire_date);
```

```
Query OK, 0 rows affected (3.48 sec)
```

```
Records: 0 Duplicates: 0 Warnings: 0
```

EXPLAIN Type by Example: range

```
mysql> explain select count(*) from employees where hire_date>'1995-01-01'\G
***** 1. row *****
```

```
      id: 1
select_type: SIMPLE
      table: employees
      type: range
possible_keys: hire_date
       key: hire_date
    key_len: 3
       ref: NULL
      rows: 68654
```

Only rows from given range used

Compare with ALL:
 $300157 / 68654 = 4.3720$
4 times less rows examined!

Extra: Using where; Using index

1 row in set (0.00 sec)

Combined Indexes

- Consists of two or more columns

Combined Indexes

- Consists of two or more columns
- Only leftmost part used

```
mysql> alter table City add key  
      -> comb(CountryCode, District, Population),  
      -> drop key CountryCode;
```

Combined Indexes: example 1

```
mysql> explain select * from City where CountryCode = 'USA'\G
***** 1. row *****
      table: City
      type: ref
possible_keys: comb
      key: comb
      key_len: 3
           ref: const
           rows: 273
```

Uses first field from the comb key

Combined Indexes: example 2

```
mysql> explain select * from City where \  
-> District = 'California' and population > 10000\G  
***** 1. row *****  
      table: City  
      type: ALL  
possible_keys: NULL  
      key: NULL  
key_len: NULL  
      ref: NULL  
     rows: 3868
```

Can't use combined index:
not a leftmost part

Does not have the CountryCode
in the where clause
= can't use comb index

Combined Indexes: key_len

- Key_len = total size (in bytes)
- Index
 - comb(CountryCode, District, Population)

Explain:

Fields:

key: <u>comb</u>	CountryCode <u>char(3)</u>
key_len: <u>3</u>	District char(20)
	Population int(11)

3 -> Char(3) -> First field is used



PERCONA

EXPLAIN Type by Example: index

```
mysql> explain select count(*) from titles where title='Senior Engineer'\G
***** 1. row *****
```

```
      id: 1
  select_type: SIMPLE
        table: titles
         type: index
possible_keys: NULL
         key: emp_no
      key_len: 4
         ref: NULL
         rows: 444033
```

```
Extra: Using where; Using index
1 row in set (0.11 sec)
```

No row in the table was accessed to resolve the query!
Only index used
Still all records in the index were scanned

Covered Indexes

- Covered index = cover all fields in the query

```
select name from City  
where CountryCode = 'USA' and District = 'Alaska' and population > 10000
```

```
mysql> alter table City add key  
-> cov1(CountryCode, District, population, name);
```

1. Where part

2. Group By/Order (not used now)

3. Select part

Uses all fields in the query in particular order



PERCONA

EXPLAIN by Example: Covered Indexes

```
mysql> explain select name from City where CountryCode = 'USA' \
-> and District = 'Alaska' and population > 10000\G
***** 1. row *****
```

```
    table: City
    type: range
possible_keys: cov1
    key: cov1
key_len: 27
    ref: NULL
    rows: 1
Extra: Using where; Using index
```

Covered index is used
MySQL will only use index
Will not go to the data file

Diagnostics

Real Numbers: Inside Storage Engine

Handler_* Status Variables

- EXPLAIN is optimistic

```
mysql> explain select * from ol
      -> where thread_id=10432 and site_id != 9939 order by id limit 3\G
***** 1. row *****
           id: 1          |          ref: NULL
select_type: SIMPLE      |      rows: 33
           table: ol      |      filtered: 8.07
           partitions: NULL |      Extra: Using where
           type: index
possible_keys: thread_id
           key: PRIMARY
           key_len: 4
1 row in set, 1 warning (0,00 sec)
```

Handler_* Status Variables

- Status variables 'Handler_*' show truth

```
mysql> flush status; select * from ol
      -> where thread_id=10432 and site_id != 9939 order by id limit 3;
mysql> show status like 'Handler%';
```

```
+-----+-----+
| Variable_name          | Value |
+-----+-----+
...
| Handler_read_first     | 1     |
| Handler_read_key       | 1     |
| Handler_read_last      | 0     |
| Handler_read_next      | 100000 |
...
```

Diagnostics

Real Numbers: Inside the Server



PROCESSLIST

- `SHOW [FULL] PROCESSLIST`
- `INFORMATION_SCHEMA.PROCESSLIST`
- `performance_schema.THREADS`

PROCESSLIST

- `SHOW [FULL] PROCESSLIST`
- `INFORMATION_SCHEMA.PROCESSLIST`
- `performance_schema.THREADS`
- Your first alert in case of performance issue

PROCESSLIST

- `SHOW [FULL] PROCESSLIST`
- `INFORMATION_SCHEMA.PROCESSLIST`
- `performance_schema.THREADS`
- Your first alert in case of performance issue
- Shows all queries, running at the moment

Execution Stages

- Can be seen in PROCESSLIST

```
mysql> show processlist\G
```

```
***** 1. row *****
```

```
Id: 7
```

```
User: root
```

```
Host: localhost:48799
```

```
db: employees
```

```
Command: Query
```

```
Time: 2
```

```
State: Sending data
```

```
Info: select count(*) from employees join titles using(emp_no)  
      where title='Senior Engineer'
```

```
...
```



Execution Stages

- Can be seen in PROCESSLIST
 - Very useful when you need to answer on question: "What is my server doing now?"

Execution Stages

- PERFORMANCE_SCHEMA.EVENTS_STAGES_*

```
mysql> select eshl.event_name, substr(sql_text, 1, 15) as 'sql',  
-> eshl.timer_wait/1000000000000 w_s from events_stages_history_long  
-> eshl join events_statements_history_long esthl on  
-> (eshl.nesting_event_id = esthl.event_id) where  
-> esthl.current_schema='employees' and sql_text like  
-> 'select count(*) from employees%' order by eshl.timer_start asc;
```

event_name	sql	w_s
stage/sql/starting	select count(*)	0.0002
stage/sql/checking permissions	select count(*)	0.000

...



PERCONA

Execution Stages

- PERFORMANCE_SCHEMA.EVENTS_STAGES_*

...

stage/sql/checking permissions	select count(*)	0.0000	
stage/sql/Opening tables	select count(*)	0.0000	
stage/sql/init	select count(*)	0.0001	
stage/sql/System lock	select count(*)	0.0000	
stage/sql/optimizing	select count(*)	0.0000	
stage/sql/statistics	select count(*)	0.0001	
stage/sql/preparing	select count(*)	0.0000	
stage/sql/executing	select count(*)	0.0000	
stage/sql/Sending data	select count(*)	5.4915	
stage/sql/end	select count(*)	0.0000	

...

Temporary tables and other job

- Status variables

```
mysql> flush status;
```

```
Query OK, 0 rows affected (0,01 sec)
```

```
mysql> select count(*) from employees join titles using(emp_no)  
      -> where title='Senior Engineer';
```

```
+-----+
```

```
| count(*) |
```

```
+-----+
```

```
|    97750 |
```

```
+-----+
```

```
1 row in set (5,44 sec)
```


Temporary tables and other job

- Status variables

```
mysql> select * from performance_schema.session_status
      -> where variable_name in ('Created_tmp_tables',
      -> 'Created_tmp_disk_tables', 'Select_full_join',
      -> 'Select_full_range_join', 'Select_range',
      -> 'Select_range_check', 'Select_scan', 'Sort_merge_passes',
      -> 'Sort_range', 'Sort_rows', 'Sort_scan') and variable_value > 0;
```

VARIABLE_NAME	VARIABLE_VALUE
Select_scan	2

1 row in set (0,00 sec)

Temporary tables and other job

- PERFORMANCE_SCHEMA.EVENTS_STATEMENTS_*

```
mysql> select * from performance_schema.events_statements_history_long  
-> where sql_text like 'select count(*) from employees join %'\G  
***** 1. row *****
```

```
...  
          ROWS_SENT: 1          SELECT_RANGE_CHECK: 0  
          ROWS_EXAMINED: 541058      SELECT_SCAN: 1  
CREATED_TMP_DISK_TABLES: 0          SORT_MERGE_PASSES: 0  
      CREATED_TMP_TABLES: 0          SORT_RANGE: 0  
      SELECT_FULL_JOIN: 0            SORT_ROWS: 0  
SELECT_FULL_RANGE_JOIN: 0            SORT_SCAN: 0  
      SELECT_RANGE: 0                NO_INDEX_USED
```



PERCONA

Temporary tables and other job

- `sys.statement_analysis`

```
mysql> select * from statement_analysis where query like 'SELECT COUNT  
-> ( * ) FROM 'emplo%' and db='employees'\G
```

```
***** 1. row *****  
      query: SELECT COUNT ( * ) FROM 'emplo ... 'emp_no' ) WHE...  
      db: employees          max_latency: 5.59 s  
      full_scan:             avg_latency: 5.41 s  
      exec_count: 7           lock_latency: 2.24 ms  
      err_count: 0            rows_sent: 7  
      warn_count: 0           rows_sent_avg: 1  
      total_latency: 37.89 s   rows_examined: 3787406
```

Temporary tables and other job

- `sys.statement_analysis`

```
rows_examined_avg: 541058
  rows_affected: 0
rows_affected_avg: 0
  tmp_tables: 0
  tmp_disk_tables: 0
    rows_sorted: 0
sort_merge_passes: 0
      digest: 4086bc3dc6510a1d9c8f2fe1f59f0943
first_seen: 2016-04-14 15:19:19
last_seen: 2016-04-14 16:13:14
```

How to Affect Query Plans



PERCONA

What has Effect on Query Optimizer Plans?

- Index statistics
- Histogram statistics
- Optimizer switches
- Bugs in optimizer

Index Statistics

- Collected by storage engine

Index Statistics

- Collected by storage engine
- Used by Optimizer

Index Statistics

- Can be examined by SHOW INDEX command

```
mysql> show index from sbtest1;
```

Table	Key_name	Column_name	Cardinality
sbtest1	k_1	k	49142

```
mysql> select count(distinct id), count(distinct k) from sbtest1;
```

count(distinct id)	count(distinct k)
100000	17598

Index Statistics

- Can be updated
 - ANALYZE TABLE
 - If does not help: rebuild table
 - OPTIMIZE TABLE
 - ALTER TABLE ENGINE=INNODB; ANALYZE TABLE

Histogram Statistics

- Since version 8.0

Histogram Statistics

- Since version 8.0
- Collected and used by the Optimizer

Histogram Statistics

- Since version 8.0
- Collected and used by the Optimizer
- Can be examined in Information Schema

```
mysql> select HISTOGRAM from information_schema.column_statistics  
-> where table_name='example'\G
```

```
***** 1. row *****  
HISTOGRAM: {"buckets": [[1, 0.6], [2, 0.8], [3, 1.0]],  
"data-type": "int", "null-values": 0.0, "collation-id": 8,  
"last-updated": "2018-11-07 09:07:19.791470",  
"sampling-rate": 1.0, "histogram-type": "singleton",  
"number-of-buckets-specified": 3}  
1 row in set (0.00 sec)
```

Histogram Statistics

- Since version 8.0
- Collected and used by the Optimizer
- Can be examined in Information Schema



More details

Optimizer Switches

```
mysql> select @@optimizer_switch\G
***** 1. row *****
@@optimizer_switch: index_merge=on,index_merge_union=on,
index_merge_sort_union=on,index_merge_intersection=on,
engine_condition_pushdown=on,index_condition_pushdown=on,
mrr=on,mrr_cost_based=on,
block_nested_loop=on,batched_key_access=off,
materialization=on,semijoin=on,loosescan=on,firstmatch=on,
duplicateweedout=on,subquery_materialization_cost_based=on,
use_index_extensions=on,condition_fanout_filter=on,derived_merge=on
1 row in set (0,00 sec)
```

Optimizer Switches

- Turn ON and OFF particular optimization

Optimizer Switches

- Turn ON and OFF particular optimization
- Can be not helpful
 - Especially for queries, tuned for previous versions

Optimizer Switches

- Turn ON and OFF particular optimization
- Can be not helpful
- Work with them as with any other option
 - Turn OFF and try

```
SET optimizer_switch = 'use_index_extensions=off';  
SELECT ...  
EXPLAIN SELECT ...
```

Optimizer Switches

- Turn ON and OFF particular optimization
- Can be not helpful
- Work with them as with any other option
 - If helps implement in queries

```
SELECT /*+ SEMIJOIN(FIRSTMATCH, LOOSES SCAN) */ * FROM t1 ...;  
SELECT /*+ BKA(t1) NO_BKA(t2) */ * FROM t1 INNER JOIN t2 WHERE ...;
```

Bugs in Optimizer

- Optimizer choses wrong index for no reason

Bugs in Optimizer

- Optimizer choses wrong index for no reason
- Statistics is up to date
- Histograms are not usable

Bugs in Optimizer

- Optimizer choses wrong index for no reason
- Statistics is up to date
- Histograms are not usable
- Solution
 - Use index hints
 - `FORCE INDEX`
 - `IGNORE INDEX`

Bugs in Optimizer

- Optimizer choses wrong index for no reason
- Statistics is up to date
- Histograms are not usable
- Solution
- On every upgrade
 - Remove index hints
 - Test if query improved
 - You must do it even for minor version upgrades!

Summary

- EXPLAIN is essential for query tuning
- Real job is done by storage engine
- Index statistics affect query execution plan
- All index hints, optimizer hints and other workarounds must be validated on each upgrade

More information



EXPLAIN Syntax



EXPLAIN FORMAT=JSON is Cool! series



Troubleshooting Performance add-ons



Optimizer Statistics aka Histograms



Optimizer Hints



Tracing the Optimizer

Special thanks

Alexander Rubin for combined and covered index examples

Thank you!



www.slideshare.net/SvetaSmirnova



twitter.com/svetasmirnova



github.com/svetasmirnova

We're Hiring!

Percona's open source database experts are true superheroes, improving database performance for customers across the globe.

Percona's open source database experts are true superheroes, improving database performance for customers across the globe.

Discover what it means to have a Percona career with the smartest people in the database performance industries, solving the most challenging problems our customers come across.



glassdoor

4.8 ★★★★★ Rating Trends



Recommend to a Friend



Approve of CEO



Peter Zaitsev
45 Ratings

Jul 12, 2019



"Best company I have ever worked for"

★★★★★

Current Employee - Solutions Engi

■ Recommends

■ Positive Outlook



PERCONA

Approves of CEO



**DATABASE PERFORMANCE
MATTERS**