



MySQL 5.6 GTID in a nutshell

Miguel Ángel Nieto
Percona Live University - Toronto

Who am I?

- **Miguel Ángel**
- I live in the north of **Spain**
- Support Engineer at **Percona**
- Hobbies:
 - **Scuba Diving**
 - Videogames
 - American TV Series
 - Beers



Agenda

I'm going to answer the following questions and give a detailed overview that will let us to start working with it:

- What is GTID?
- What problems GTID solves?
- How can I implement it?
- How can I repair it?
- How can I use it for HA and Failover?
- Take in account...

What is GTID?

What is GTID?

8182213e-7c1e-11e2-a6e2-080027635ef5 : 1

What is GTID?

And that thing is a GTID?

8182213e-7c1e-11e2-a6e2-080027635e

Not impressed



What is GTID?

8182213e-7c1e-11e2-a6e2-080027635ef5 : 1



SID. This is the server's 128 bit identification number (**SERVER_UUID**). It identifies where the transaction was originated. Every server has its own **SERVER_UUID**.

What is GTID?

8182213e-7c1e-11e2-a6e2-080027635ef5 : 1



GNO. This is the transaction identification number. It is a sequence number that increments with every new transaction.

What is GTID?

- This is how we can see the GTID inside the binary logs:

```
# at 300
#130221 13:08:58 server id 101  end_log_pos 348 CRC32 0xc18cdbda      GTID [commit=yes]
SET @@SESSION.GTID_NEXT= '8182213e-7c1e-11e2-a6e2-080027635ef5:2'/*!*/;
# at 348
BEGIN
insert into t values(1)
COMMIT/*!*/;
# at 565
#130221 13:09:03 server id 101  end_log_pos 613 CRC32 0x5b25189e      GTID [commit=yes]
SET @@SESSION.GTID_NEXT= '8182213e-7c1e-11e2-a6e2-080027635ef5:3'/*!*/;
# at 697
BEGIN
insert into t values(100)
COMMIT/*!*/;
```

- The GTID is replicated to Slave servers.

What problems GTID solves?

```
INSERT INTO t VALUES(100);
```

bin-log.000407
10983



MASTER1



bin-log.000010
4



bin-log.000133
984



SLAVE 1

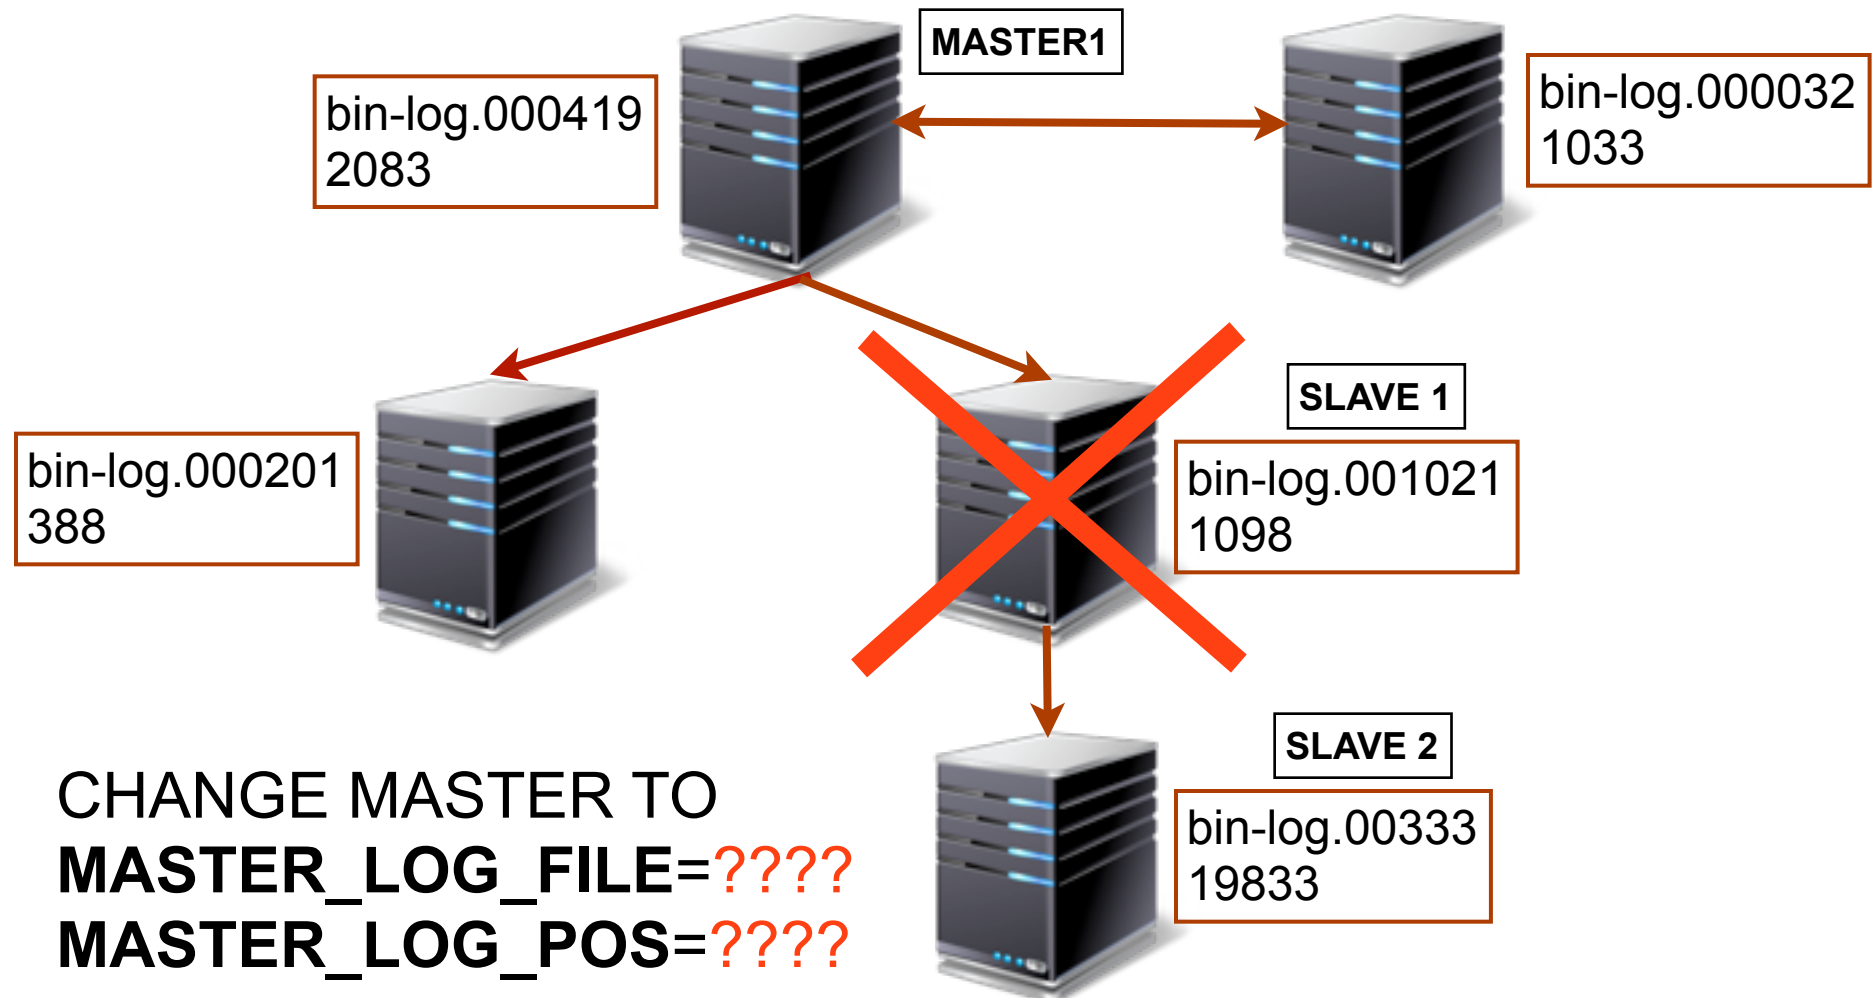
bin-log.001021
1098



SLAVE 2


bin-log.00333
19833

What problems GTID solves?

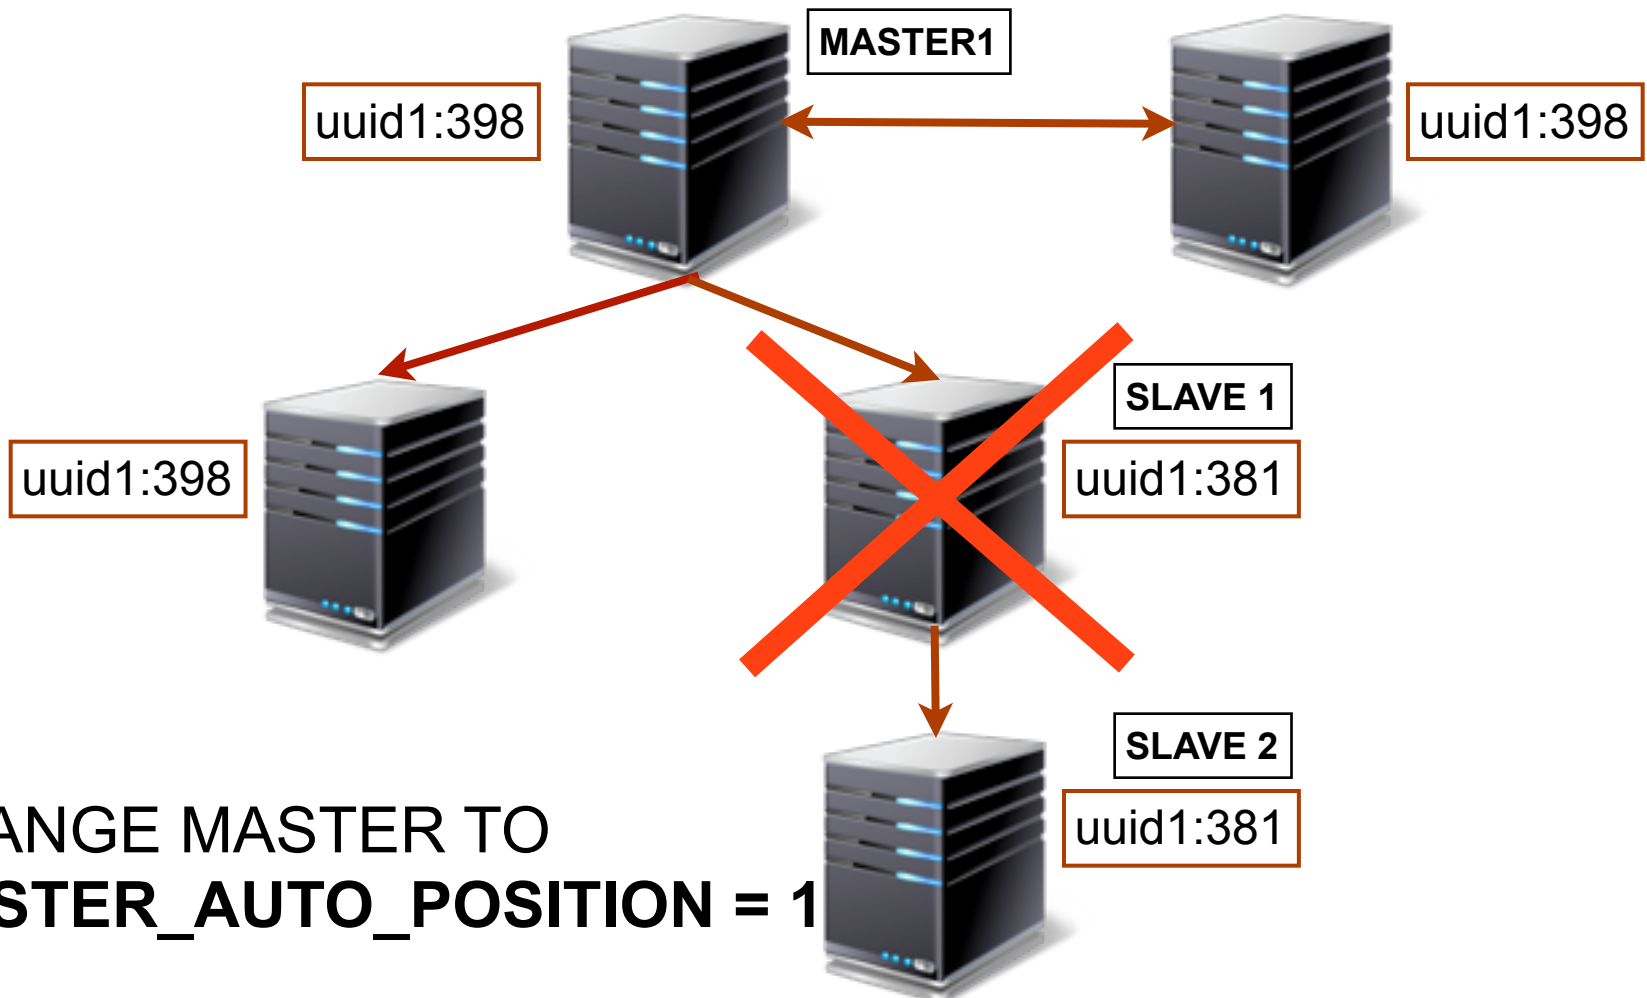


What problems GTID solves?

`INSERT INTO t VALUES(100);`



What problems GTID solves?



What problems GTID solves?

- It is possible to identify a transaction **uniquely** across the replication servers.
- Make the automation of failover process much easier. There is no need to do calculations, inspect the binary log and so on. Just **MASTER_AUTO_POSITION=1**.
- At application level it is easier to do WRITE/READ split. After a write on the MASTER you have a GTID so just check if that GTID has been executed on the SLAVE that you use for reads.
- Development of new automation tools isn't a pain now.

How can I implement it?

- Three variables are needed in ALL servers of the replication chain
- **gtid_mode:** It can be ON or OFF (not 1 or 0). It enables the GTID on the server.
- **log_bin:** Enable binary logs. Mandatory to create a replication environment.
- **log-slave-updates:** Slave servers must log the changes that comes from the master in its own binary log.
- **enforce-gtid-consistency:** Statements that can't be logged in a transactionally safe manner are denied by the server.

How can I implement it?

- **enforce-gtid-consistency**

- CREATE TABLE ... SELECT statements.

ERROR 1786 (HY000): CREATE TABLE ... SELECT is forbidden when ENFORCE_GTID_CONSISTENCY = 1.

- CREATE TEMPORARY TABLE inside transactions.


ERROR 1787 (HY000): When ENFORCE_GTID_CONSISTENCY = 1, the statements CREATE TEMPORARY TABLE and DROP TEMPORARY TABLE can be executed in a non-transactional context only, and require that AUTOCOMMIT = 1.

- Transactions that mixes updates on transactional and non-transactional tables.

ERROR 1785 (HY000): When ENFORCE_GTID_CONSISTENCY = 1, updates to non-transactional tables can only be done in either autocommitted statements or single-statement transactions, and never in the same statement as updates to transactional tables.

How can I implement it?

New replication from scratch

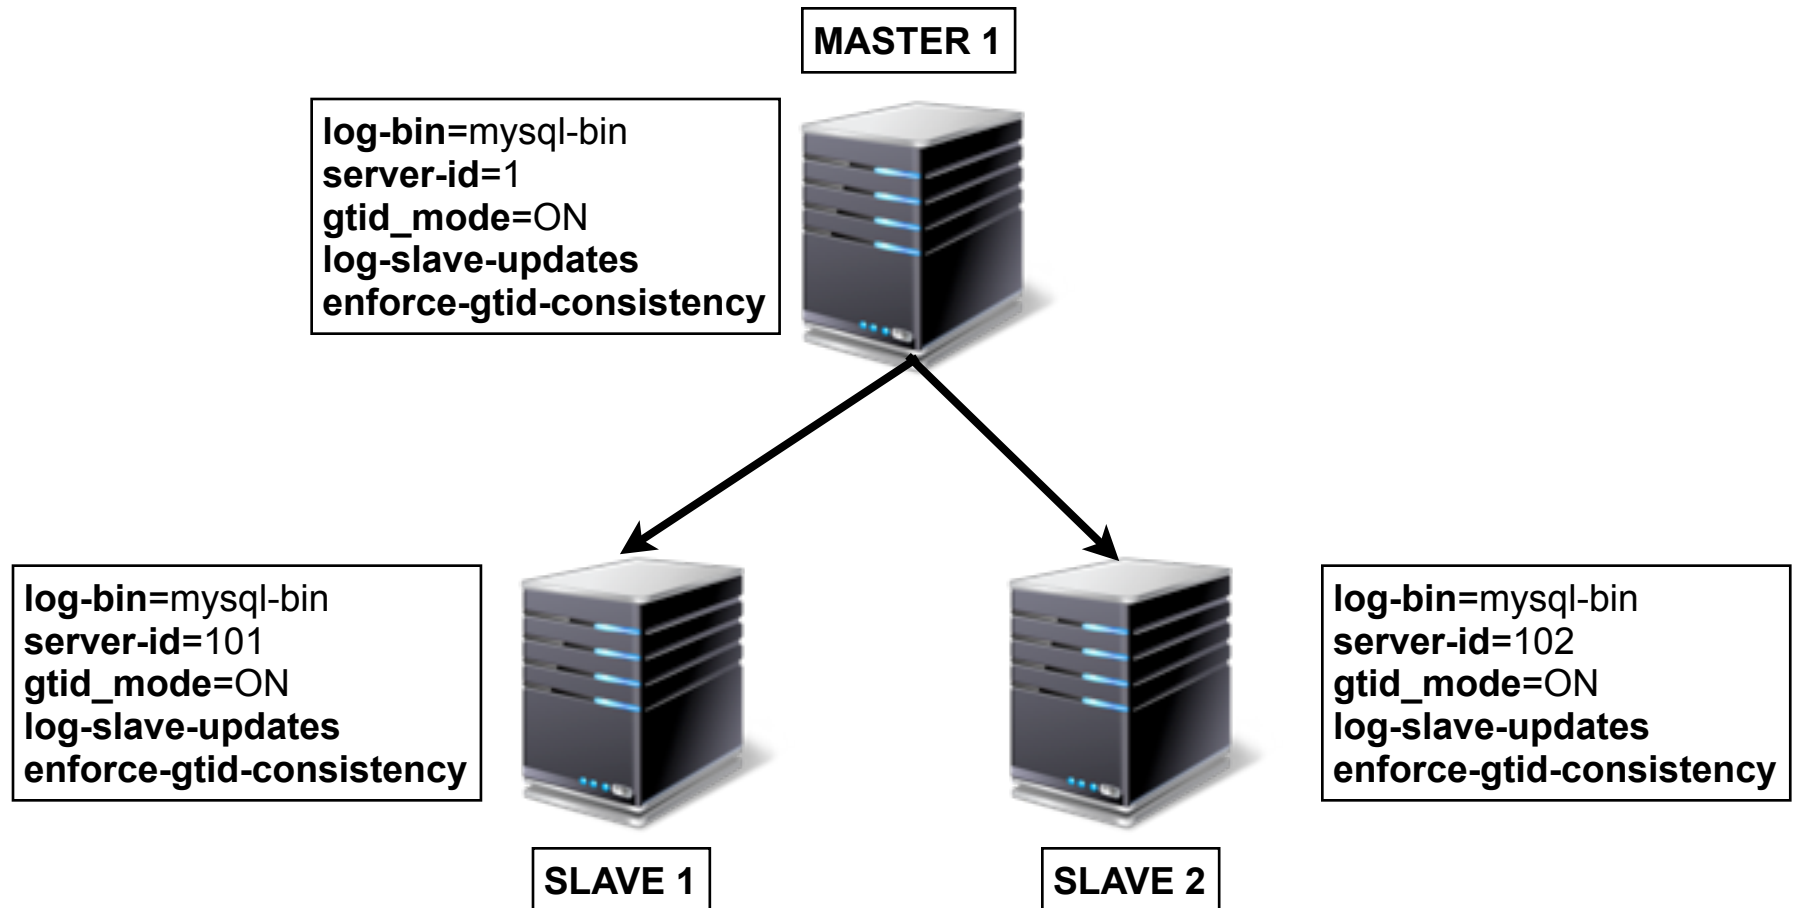


How can I implement it?


New replication from scratch

- 1) Create replication user on the master server.
- 2) Configure the parameters on all three servers:
 - `gtid_mode`
 - `log_bin`
 - `log-slave-updates`
 - `enforce-gtid-consistency`
 - `server_id`
- 3) Start all mysql services.
- 4) `CHANGE MASTER TO...` with **`MASTER_AUTO_POSITION=1`** on the two slave servers.

How can I implement it?



How can I implement it?



CHANGE MASTER TO **MASTER_HOST**="127.0.0.1", **MASTER_PORT**=18675,
MASTER_USER="msandbox", **MASTER_PASSWORD**="msandbox", **MASTER_AUTO_POSITION**=1;

How can I implement it?

Move already running replication to GTID

1. Set the master as **read_only** and wait until the slaves catch up.
2. Stop all servers.
3. Configure the GTID variables in my.cnf.
4. Start all the servers:
 - Master should start in **read_only** mode.
 - Slaves should start with **skip_slave_start**.
5. **CHANGE MASTER** with **MASTER_AUTO_POSITION=1** on the slaves.
6. **START SLAVE**; on slave servers.
7. **SET GLOBAL read_only=0**; on master server.

How can I implement it?

- Now we run two transactions on the master:

```
CREATE TABLE t (i INT);  
INSERT INTO t VALUES(1);
```

- This is the status of slaves:

```
slave1 > show slave status\G
```

```
[...]
```

```
      Master_Log_File: mysql-bin.000001  
Read_Master_Log_Pos: 550  
      Relay_Log_File: mysql_sandbox18676-relay-bin.000002  
      Relay_Log_Pos: 760  
Relay_Master_Log_File: mysql-bin.000001
```

```
[...]
```

```
      Master_Server_Id: 1  
      Master_UUID: 1c9cdcc8-7c33-11e2-a769-080027635ef5
```

```
[...]
```

```
Retrieved_Gtid_Set: 1c9cdcc8-7c33-11e2-a769-080027635ef5:1-2  
Executed_Gtid_Set: 1c9cdcc8-7c33-11e2-a769-080027635ef5:1-2  
      Auto_Position: 1
```

How can I implement it?

- Now we have new variables to check:
 - **gtid_executed (ro)**: shows the transactions that have been executed in this server.
`1c9cdcc8-7c33-11e2-a769-080027635ef5:1-3`
 - **gtid_purged (ro)**: shows the transactions that have been purged from the binary log (purge binary logs to...).
`1c9cdcc8-7c33-11e2-a769-080027635ef5:1-2`
 - **gtid_next**: the next GTID that will be used.
`SET @@SESSION.GTID_NEXT=`
`'8182213e-7c1e-11e2-a6e2-080027635ef5:2'/*!*/;`

How can I repair it?

- Even with GTID we have the same problem. MySQL replication can easily fail.
- The procedure to repair a replication is slightly different from the regular replication based on binary log position.
- There is a very good blog post written by a very good blogger that explains how to repair it:

<http://www.mysqlperformanceblog.com/2013/02/08/how-to-createrestore-a-slave-using-gtid-replication-in-mysql-5-6/>

How can I repair it?

- **ERROR!**

```
Slave_IO_Running: No
```

```
Slave_SQL_Running: Yes
```

```
Last_IO_Error: Got fatal error 1236 from master when reading data from binary log:  
'The slave is connecting using CHANGE MASTER TO MASTER_AUTO_POSITION = 1, but the  
master has purged binary logs containing GTIDs that the slave  
requires.'
```

- **mysqldump** supports GTID:

```
# mysqldump --all-databases --single-transaction --triggers --routines --  
host=127.0.0.1 --port=18675 --user=msandbox --password=msandbox > dump.sql
```

```
# grep PURGED dump.sql
```

```
SET @@GLOBAL.GTID_PURGED='9a511b7b-7059-11e2-9a24-08002762b8af:  
1-13';
```

- **Xtrabackup's** support for GTID is in **Work in Progress**.

How can I repair it?

- The server was already running as slave, so **GTID_EXECUTED** and **GTID_PURGED** has values:

```
slave1 > source test.sql;  
ERROR 1840 (HY000): GTID_PURGED can only be set when GTID_EXECUTED is  
empty.
```

- So, let's empty **GTID_EXECUTED**. But... How? It is a read only variable!

```
slave1 > reset master;  
slave1 > show global variables like 'GTID_EXECUTED';  
+-----+-----+  
| Variable_name | Value |  
+-----+-----+  
| gtid_executed |      |  
+-----+-----+  
slave1> source test.sql;  
slave1> start slave;
```

How can I repair it?

- Another way, injecting empty transactions

<http://dev.mysql.com/doc/refman/5.6/en/replication-gtids-failover.html#replication-gtids-failover-empty>

- **SQL_SLAVE_SKIP_COUNTER** doesn't work anymore with GTID
- We need to find what transaction is causing the replication to fail
 - From binary logs
 - From SHOW SLAVE STATUS (retrieved vs executed)

How can I repair it?

- Slave failed:

```
Last_SQL_Error: Error 'Duplicate entry '4' for key 'PRIMARY'' on
query. Default database: 'test'. Query: 'insert into t VALUES(NULL,'salazar')'
Retrieved_Gtid_Set: 7d72f9b4-8577-11e2-a3d7-080027635ef5:1-5
Executed_Gtid_Set: 7d72f9b4-8577-11e2-a3d7-080027635ef5:1-4
```

- So, this slave has retrieved transactions from 1 to 5 but only 1 to 4 has been applied. Seems that transaction 4 is the problem here.

```
STOP SLAVE;
SET GTID_NEXT="7d72f9b4-8577-11e2-a3d7-080027635ef5:5";
BEGIN; COMMIT;
SET GTID_NEXT="AUTOMATIC";
START SLAVE;
[...]
Retrieved_Gtid_Set: 7d72f9b4-8577-11e2-a3d7-080027635ef5:1-5
Executed_Gtid_Set: 7d72f9b4-8577-11e2-a3d7-080027635ef5:1-5
```

How can I repair it?

- mysqldump can be also used to create new slaves.
- It is a new slave so GTID_EXECUTED and GTID_PURGED are empty. No RESET MASTER is needed.
- Now we know how to create and repair a replication with GTID.



FEEL LIKE A SIR

How can I use it for HA and Failover?


- We have seen how to implement and repair a GTID based replication.
- Now we are going to see it can help us with HA and failover:

mysqlrpladmin: replication administration tool. For failover and switchover.

mysqlfailover: replication health check and automatic failover tool.

How can I use it for HA and Failover?

- Where can I download these tools?
- <https://launchpad.net/mysql-utilities>
- MySQL Workbench:



How can I use it for HA and Failover?

- **mysqlrpladmin** is a tool used to perform planned maintenance tasks in our replication environment:
 - **switchover**: a planned stop for a master. Slave is promoted to new master. No possibility of transaction loss.
 - **failover**: a non-planned stop of the master and a slave promoted to new master. Last transaction can be lost. The tool will chose the most up-to-date slave.

How can I use it for HA and Failover?

- Some pre-requisites:
 - To make the autodiscover work the **master-info-repository=TABLE** should be enabled.
 - Slaves should have the replication user created in order to be elected as new masters.

```
# mysqlrpladmin --master=root:msandbox@master:18675
--discover-slaves-login=root:msandbox health
# Discovering slaves for master at 127.0.0.1:18675
# Checking privileges.
# Replication Topology Health:
```

host	port	role	state	gtid_mode	health
master	18675	MASTER	UP	ON	OK
SBslave1	18676	SLAVE	UP	ON	OK
SBslave2	18677	SLAVE	UP	ON	OK
SBslave3	18678	SLAVE	UP	ON	OK

How can I use it for HA and Failover?

- **health**: shows the health status of the replication servers.
- **elect**: shows which slave server should be elected as new master in case of a failover.
- **failover**: performs a failover selecting the most up-to-date slave.
- **gtid**: shows GTID information from all nodes
- **reset, start, stop**: reset, start or stop command on all slaves.
- **switchover**: do a slave promotion using the --new-master parameter.

How can I use it for HA and Failover?

- We need to remove the master from the replication and promote a slave to a new master:

```
# mysqlrpladmin --demote-master --
master=msandbox:msandbox@master:18675 --new-
master=root:msandbox@sbslave1:18676 --
slaves=root:msandbox@sbslave1:18676,root:msandbox@sbslave2:
18677,root:msandbox@sbslave3:18678 switchover
# Performing switchover from master at master:18675 to slave at
sbslave1:18676.
# Switchover complete.
```

host	port	role	state	gtid_mode	health
sbslave1	18676	MASTER	UP	ON	OK
master	18675	SLAVE	UP	ON	OK
sbslave2	18677	SLAVE	UP	ON	OK
sbslave3	18678	SLAVE	UP	ON	OK

How can I use it for HA and Failover?

- **mysqlfailover** tool do a health check on the replication servers and run a failover automatically in case it is necessary.
- It monitors the master and in case of a failure on the health check it selects the best slave and performs the failover.
- You can give the tool a list of slaves that should be taken in account for master promotion. For example in case of different hardware characteristics.

How can I use it for HA and Failover?

- **auto**: performs an automatic failover.
- **elect**: the same as auto but if no candidates on the candidates list are viable it shows an error and exists.
- **fail**: doesn't perform any failover, just shows an error and exist.

```
mysqlfailover --master=msandbox:msandbox@sbslave1:18676 --  
slaves=root:msandbox@master:  
18675,root:msandbox@sbslave2:18677,root:msandbox@sbslave3:18678 auto
```

How can I use it for HA and Failover?

Failover Mode = auto Next Interval = Sun Feb 24 13:45:04 2013

Master Information

Binary Log File	Position	Binlog_Do_DB	Binlog_Ignore_DB
mysql-bin.000002	552		

GTID Executed Set

ce40779f-7e7b-11e2-b64d-080027635ef5:1-2

Replication Health Status

host	port	role	state	gtid_mode	health
sbslave1	18676	MASTER	UP	ON	OK
master	18675	SLAVE	UP	ON	OK
sbslave2	18677	SLAVE	UP	ON	OK
sbslave3	18678	SLAVE	UP	ON	OK

Q-quit R-refresh H-health G-GTID Lists U-UUIDs

How can I use it for HA and Failover?

- We kill the slave server and the automatic failover starts:

```
Failover starting in 'auto' mode...  
# Candidate slave master:18675 will become the  
new master.  
# Preparing candidate for failover.  
# Creating replication user if it does not exist.  
# Stopping slaves.  
# Performing STOP on all slaves.  
# Switching slaves to new master.  
# Starting slaves.  
# Performing START on all slaves.  
# Checking slaves for errors.  
# Failover complete.
```

How can I use it for HA and Failover?

- Failover done:

```
Failover Mode = auto      Next Interval = Sun Feb 24 13:51:31 2013
```

Master Information

```
Binary Log File    Position  Binlog_Do_DB  Binlog_Ignore_DB
mysql-bin.000002   552
```

GTID Executed Set

```
ce40779f-7e7b-11e2-b64d-080027635ef5:1-2
```

Replication Health Status

host	port	role	state	gtid_mode	health
master	18675	MASTER	UP	ON	OK
sbslave2	18677	SLAVE	UP	ON	OK
sbslave3	18678	SLAVE	UP	ON	OK

How can I use it for HA and Failover?

- **--exec-after**
- **--exec-before**
- You can use your own -pre and -post scripts to do a variety of different tasks:
 - Send a mail.
 - Move a virtual IP.
 - Electrocute the DBA.
 - ...

Take in account...

- 5.6 is a new GA release so there can be bugs...
<http://bugs.mysql.com/bug.php?id=68460>
- There have been some problems to make GTID compatible with MyISAM. From 5.6.9 it is possible to run single statements updating MyISAM tables. Be cautious.
- `mysql_upgrade --write-binlog=ON` can't connect to a server with GTID enabled. So from 5.6 by default `mysql_upgrade` has `write-binlog` disabled.

That's all Folks!

miguel.nieto@percona.com



That's all Folks!