

# MySQL 8 *Tips and Tricks*



Dave Stokes

@stoker

[david.stokes@oracle.com](mailto:david.stokes@oracle.com)

[Elephantdolphin.blogger.com](http://Elephantdolphin.blogger.com)

[OpensourceDBA.wordpress.com](http://OpensourceDBA.wordpress.com)

ORACLE®

# What This Talk Is About??

# MySQL 8 Features

This is not a simple talk on performance tuning a database or a cookbook where you set X to Y and get Z percent better performance.

Instead this a talk about developments that have the potential to make big changes in the way you use MySQL Instances.

Simple Answer:

Set

**INNODB\_BUFFER\_POOL\_SIZE**

To ~ 80% of RAM

# Quick c

```
mysql> SELECT  
@@innodb_buf
```



Daniel Storj {turnoff.us}  
Thanks to [jtomaszon](https://www.jtomaszon.com)

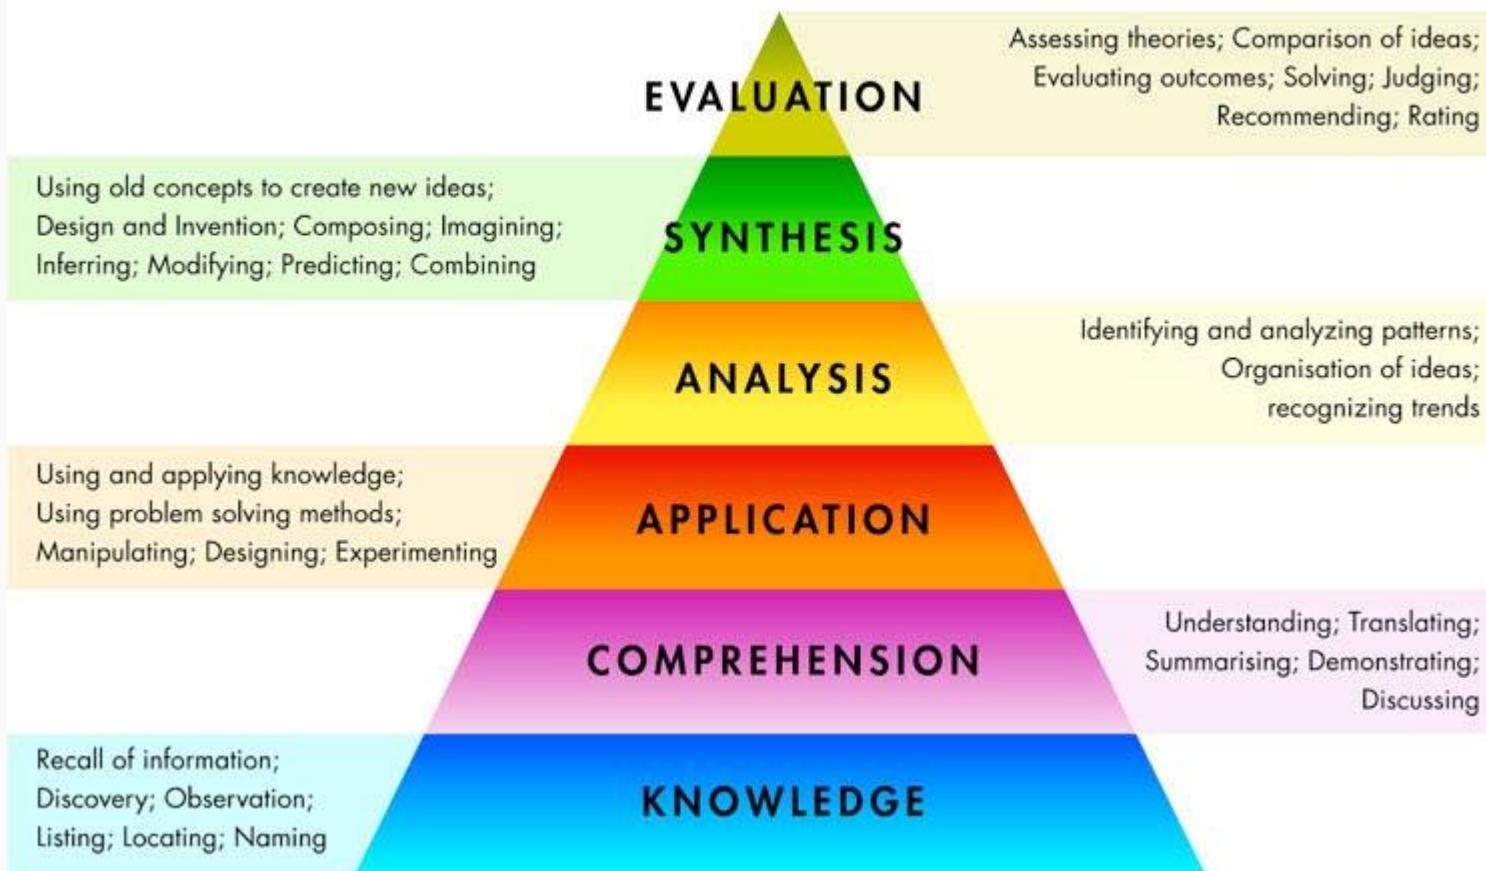
# setting:

```
/1024;
```

**Simple answers are great if**

**... you only live in a  
simple world!**

# BLOOMS TAXONOMY



# Quiet Database Revolution

Cloud

NoSQL

Security

Self-tuning



# 1. Upgrade

## **Minor Interruption**

**Please excuse  
this small rant  
about help  
forums!**

Hi!

I know nothing about brain surgery but ....

I popped the top of the skull off my coworker in an attempt to adjust their attitudes.

How do I do make those adjustments?

And what is the red stuff leaking on the carpet?

I have an Ikea allen wrench, a screwdriver, and some duct tape!

Please advise ASAP as the coworker is vital to production.

And how do you clean a carpet??

# End of Rant

# Big Changes Behind the Scenes

<https://stackoverflow.com/questions/50505236/mysql-8-0-group-by-performance>

To compare MySQL 5.7 and 8.0 I created a table using sysbench. And I tried the test. The performance of the server is exactly the same As a result, oltp\_point\_select showed almost similar performance.

However, when doing the group by tests below, MySQL 8.0 showed *10 times better performance*.

**But I do not know why it is fast.**

I do not know if I can find the MySQL 8.0 Release Notes. In 8.0, who will tell me why group by are faster?

# Oystein Answers

MySQL 8.0 uses a new storage engine, TempTable, for internal temporary tables. (See MySQL Manual for details.) This engine does not have a max memory limit per table, but a common memory pool for all internal tables. It also has its own overflow to disk mechanism, and does not overflow to InnoDB or MyISAM as earlier versions.

The profile for 5.7 contains "converting HEAP to ondisk". This means that the table reached the max table size for the MEMORY engine (default 16 MB) and the data is transferred to InnoDB. Most of the time after that is spent accessing the temporary table in InnoDB. In MySQL 8.0, the default size of the memory pool for temporary tables is 1 GB, so there will probably not be any overflow to disk in that case.

# **Please Upgrade**

**Besides the  
obvious security  
and bug updates  
there are some  
major  
improvements  
waiting for you in  
MySQL 8**

## 2. Data Dictionary

# Metadata before 8

MySQL Server incorporates a transactional data dictionary that stores information about database objects. In previous MySQL releases, dictionary data was stored in metadata files, non transactional tables, and storage engine-specific data dictionaries.

Metadata was kept in a series of files --- eating up inodes, getting damaged or deleted at the wrong time, and hard to fix

# Data Dictionary

Benefits of the MySQL data dictionary include:

- Simplicity of a centralized data dictionary schema that uniformly stores dictionary data.
- Removal of file-based metadata storage.
- Transactional, crash-safe storage of dictionary data. Uniform and centralized caching for dictionary objects. A simpler and improved implementation for some INFORMATION\_SCHEMA tables.
- Atomic DDL.

# Big Change

Good news: You can now have millions of tables within a schema

Bad news: You can now have millions of tables within a schema

# Instant Add Column

This INSTANT ADD COLUMN patch was contributed by the Tencent Games DBA Team. We would like to thank and acknowledge this important and timely contribution by Tencent Games.

# Bye Bye Bug #199

**No more Innodb auto\_increment stats loss**

# 3. CATS



# Contention Aware Transaction Schedule

<https://arxiv.org/pdf/1602.01871.pdf>

Identifying the Major Sources of Variance in Transaction Latencies: Towards More Predictable Databases -- University of Michigan

The CATS algorithm is based on a simple intuition: **not all transactions are equal**, and not all objects are equal. **When a transaction already has a lock on many popular objects, it should get priority when it requests a new lock.** In other words, unblocking such a transaction will indirectly contribute to unblocking many more transactions in the system, which means higher throughput and lower latency overall.

# Indexes Versus Histograms

Indexes are great but have a *cost at insert update, delete, and at statistic gathering time.*

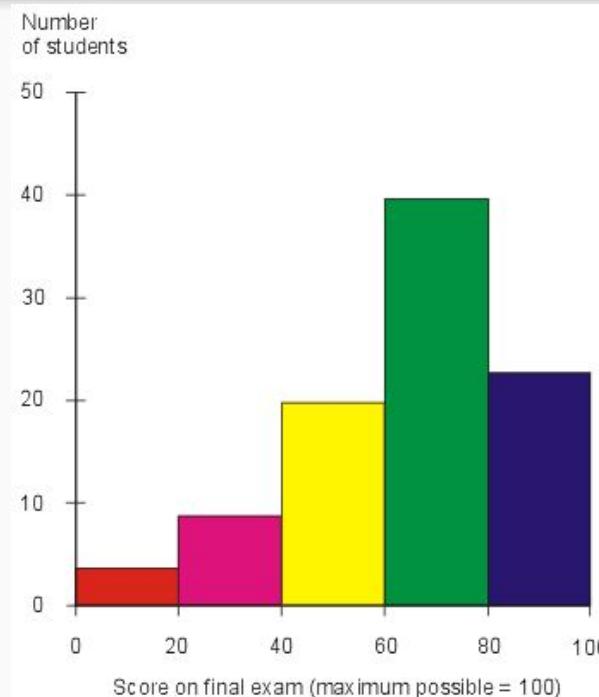
Histograms can be run after major changes to data or at slack times.

# Histograms

**The query optimizer needs statistics to create a query plan.**

- How many rows are there in each table?
- How many distinct values are there in each column?
- How is the data distributed in each column?

# What is a Histogram?



# What is a Histogram?

A *histogram* is an approximation of the data distribution for a column. It can tell you with a reasonable accuracy whether your data is skewed or not, which in turn will help the database server understand the nature of data it contains.

MySQL has chosen to support two different types: The “singleton” histogram and the “equi-height” histogram. Common for all histogram types is that they split the data set into a set of “buckets”, and MySQL automatically divides the values into buckets, and will also automatically decide what type of histogram to create.

# Syntax

```
ANALYZE TABLE tbl_name
```

```
UPDATE HISTOGRAM ON col_name [, col_name]  
WITH N BUCKETS;
```

```
ANALYZE TABLE tbl_name DROP HISTOGRAM ON col_name [, col_name];
```

# 4. Invisible Indexes

# What is an Invisible Index?

Indexes can be marked as ‘invisible’ to the optimizer

Use EXPLAIN to see query plan and tell if index aids or hinders query

```
ALTER TABLE t1 ALTER INDEX i_idx INVISIBLE;  
ALTER TABLE t1 ALTER INDEX i_idx VISIBLE;
```

# 5. Replacing Many-to-many joins with JSON

# Relational Database + JSON Fields (hybrid)

## Leverage power of RDMS but augmented with JSON fields

- Use JSON to eliminate one of the issues of traditional relational databases
  - the many-to-many join
- Allows more freedom to store unstructured data (data with pieces missing)
- You can still use SQL to work with the data via a database connector but the JSON documents can be manipulated directly in code

# JSON Document Tips

- Minimize joins - *reducing how many joins you need can speed up queries.*  
Faster access over data denormalization
- Plan for mutability - Schema-less design are based mutability. Build your applications with the ability to change the document as needed (and within reason)
- Use embedded arrays and lists to store relationship among documents
  - This can be as simple as embedding the data in document or embedding an array of document ids in the document. In the first case the data is available when you read the document. In the second, it takes only one more step to get the data.
  - In cases of seldom read (used) relationships the array of ids is more efficient as there is less data to read on the first pass

# Quick Example

**Customer table -- ID**

**Address -- Address1 .. n**

**Phone -- Phone1..n**

**Payment -- Bank1...n**

**4 or more reads to process an order**

**Customer table -- ID**

**JSON docs -- Address, Phone, Payment**

**1 read**

# 6. Resource Groups

# Resource Groups

MySQL supports creation and management of resource groups, and permits assigning threads running within the server to particular groups so that threads execute according to the resources available to the group.

Group attributes enable control over its resources, to enable or restrict resource consumption by threads in the group. DBAs can modify these attributes as appropriate for different workloads.

# Resource Groups

Currently, CPU time is a manageable resource, represented by the concept of “virtual CPU” as a term that includes CPU cores, hyperthreads, hardware threads, and so forth.

*The server determines at startup how many virtual CPUs are available, and database administrators with appropriate privileges can associate these CPUs with resource groups and assign threads to groups.*

# Create a Resource Group

CREATE RESOURCE GROUP Batch

TYPE = USER

VCPU = 2-3 -- assumes a system with at least 4 CPUs

THREAD\_PRIORITY = 10;

# Using a Resource Group

```
INSERT /*+ RESOURCE_GROUP(Batch) */ INTO t2 VALUES(2);
```

# 9. Autonomy

# Self Tuning Databases

Databases are getting better at realizing their environments (cores, disks, busses, virtual, container, buffers), loads, query patterns, and networks.

You will see much more of this much sooner than you would expect.



*"I expect you all to be independent, innovative, critical thinkers who will do exactly as I say!"*

# The Payoff is less ...

## Human Labor

## Human Error

## No Manual Labor

# 10. JSON Updates

# JSON Data Type Extremely Popular

Introduced in MySQL 5.7, the JSON data type provides a 1GB document store in a column of a row in a table.

Over thirty functions to support JSON data types

The foundation on the MySQL Document Store, a NoSQL JSON document store

# Inplace Update of JSON columns

In MySQL 8.0, the optimizer can perform a partial, in-place update of a JSON column instead of removing the old document and writing the new document in its entirety to the column.

# New JSON Functions

**JSON\_PRETTY**

**JSON array and object aggregations**

**JSON\_SIZE and JSON\_FREE**

**Change in JSON\_MERGE : JSON\_MERGE\_PRESERVE and JSON\_MERGE\_PATCH**

# The JSON Functions

Name	Description
JSON_ARRAY()	Create JSON array
JSON_ARRAY_APPEND()	Append data to JSON document
JSON_ARRAY_INSERT()	Insert into JSON array
->	Return value from JSON column after evaluating path; equivalent to JSON_EXTRACT().
JSON_CONTAINS()	Whether JSON document contains specific object at path
JSON_CONTAINS_PATH()	Whether JSON document contains any data at path
JSON_DEPTH()	Maximum depth of JSON document
JSON_EXTRACT()	Return data from JSON document
>>	Return value from JSON column after evaluating path and unquoting the result; equivalent to JSON_UNQUOTE(JSON_EXTRACT()).
JSON_INSERT()	Insert data into JSON document
JSON_KEYS()	Array of keys from JSON document
JSON_LENGTH()	Number of elements in JSON document
JSON.Merge() (deprecated 8.0.3)	Merge JSON documents, preserving duplicate keys. Deprecated synonym for JSON.Merge_Preserve()
JSON.Merge_Patch()	Merge JSON documents, replacing values of duplicate keys
JSON.Merge_Preserve()	Merge JSON documents, preserving duplicate keys
JSON_OBJECT()	Create JSON object
JSON_PRETTY()	Prints a JSON document in human-readable format, with each array element or object member printed on a new line, indented two spaces with respect to its parent.
JSON_QUOTE()	Quote JSON document
JSON_REMOVE()	Remove data from JSON document
JSON_REPLACE()	Replace values in JSON document
JSON_SEARCH()	Path to value within JSON document
JSON_SET()	Insert data into JSON document
JSON_STORAGE_FREE()	Freed space within binary representation of a JSON column value following a partial update
JSON_STORAGE_SIZE()	Space used for storage of binary representation of a JSON document; for a JSON column, the space used when the document was inserted, prior to any partial updates
JSON_TABLE()	Returns data from a JSON expression as a relational table
JSON_TYPE()	Type of JSON value
JSON_UNQUOTE()	Unquote JSON value
JSON_VALID()	Whether JSON value is valid

# **JSON\_TABLE**

**JSON\_TABLE takes schema-less JSON documents and turn it into a temporary relational table that can be processed like any other relational table.**

```
mysql> select country_name,
    IndyYear
        from countryinfo,
        json_table(doc,"$" columns
                    (country_name char(20) path "$.Name",
                     IndyYear int path "$.IndepYear"))
        as stuff
        where IndyYear > 1992;
```

country_name	IndyYear
Czech Republic	1993
Eritrea	1993
Palau	1994
Slovakia	1993

```
mysql> select country_name,
    IndyYear
    from countryinfo,
    json_table(doc,"$" columns
                (country_name char(20) path "$.Name",
                 IndyYear int path "$.IndepYear"))
    as stuff
    where IndyYear > 1992;
```

country_name	IndyYear
Czech Republic	1993
Eritrea	1993
Palau	1994
Slovakia	1993

# 12. Sys Schema

# What is in the SYS Schema

MySQL 8.0 includes the `sys` schema, a set of objects that helps DBAs and developers interpret data collected by the Performance Schema. `sys` schema objects can be used for typical tuning and diagnosis use cases. Objects in this schema include:

- Views that summarize Performance Schema data into more easily understandable form.
- Stored procedures that perform operations such as Performance Schema configuration and generating diagnostic reports.
- Stored functions that query Performance Schema configuration and provide formatting services.

## Statements in Highest 5 Percent by Runtime

List all statements whose average runtime, in microseconds is in highest 5 percent

Query	Full T...	Executed (#)	Errors (#)	Warnings (#)	Total Time (	Maximum Ti...	Avg Time (us)
CREATE TABLE SYSTEM_USER ( `username` VARCHA...		1	0	0	354508.33	354508.33	354508.33
CREATE TABLE `testx` ( `id` INTEGER , `c1` JSON )		1	0	0	351993.09	351993.09	351993.09
CREATE TABLE `UserRecords` ( `id` INTEGER (?) NO...		1	0	0	286686.84	286686.84	286686.84
CREATE TABLE `test` ( `id` INTEGER , `c1` JSON )		2	1	0	425183.98	419627.39	212591.84
INSERT INTO `testx` VALUES (...) /* , ... */		1	0	0	111718.35	111718.35	111718.35

## Tables with Full Table Scans

Find tables that are being accessed by full table scans ordering by the number of rows scanned descending

Schema	Object	Full Scanne...
world_x	city	4079
test	userrecords	55
test	testx	24
test	user	8

## Top I/O Time by User/Thread

Show the top IO time consumers by User/thread

User	Thread Id	Process List Id	Total IOs (#)	Total Time (	Min Time (us)	Avg Time (us)	Max Time (us)
main	1	0	52116	27199168.48	0.21	4286.21	690767.80
log_flusher_thread	18	0	100	2300070.77	18.37	23000.71	131842.48
page_flush_coordinator_...	14	0	394	1031708.76	9.48	2618.55	129726.92
io_write_thread	10	0	20	620911.77	7232.35	31045.59	116960.38
io_write_thread	11	0	8	258066.16	8405.60	32258.27	93117.47
dict_stats_thread	27	0	20	221030.76	231.19	11051.54	62549.93
io_write_thread	13	0	5	166138.96	8025.01	33227.79	110270.01
worker	29	0	84	160630.76	8.66	1912.27	28306.08
log_writer_thread	17	0	362	140328.05	8.14	387.65	101744.19
root@localhost	52	13	35	115720.94	1.32	3110.60	21231.09
root@localhost	50	11	1	15043.91	15043.91	15043.91	15043.91
log_checkpointer_thread	15	0	9	11583.21	39.75	1287.02	10889.76

## Statement Statistics

Shows statement execution statistics for each user

User	Statement	Total Event...	Total Time (	Max Time (us)	Lock Time (us)	Rows Sent (#)	Rows Examined	Rows Affected	Full Scans (#)
background	select	1	82664.71	82664.71	0.00	1	0	0	0
root	create_table	5	1418372.24	419627.39	1211446.00	0	0	0	0
root	select	117	1155600.71	269190.47	547608.00	8009	61155	0	74
root	insert	8	713038.15	131034.97	3523.00	0	0	13	0
root	show_status	241	621996.64	12009.79	60422.00	101440	206320	0	241
root	show_fields	109	173671.66	34202.20	125162.00	1013	5412	0	0
root	show_...ables	29	135444.66	34890.88	44192.00	1753	31726	0	29
root	show_tables	6	112752.87	73131.69	28860.00	411	2479	0	0
root	show_...bases	6	61244.50	36198.52	53072.00	38	212	0	6
root	set_option	44	51612.63	21146.78	0.00	0	0	0	0
root	show_...status	2	38312.63	36273.94	23829.00	0	4	0	0
root	update	2	34036.64	27277.57	11679.00	0	1225	939	0
root	show_keys	2	26894.47	25873.77	26075.00	0	20	0	0
root	Ping	184	26136.68	314.59	0.00	0	0	0	0
root	show_plugins	1	23417.73	23417.73	263.00	43	43	0	1
root	show_...gines	2	18428.13	18140.58	14618.00	18	18	0	2
root	stmt	1	14415.05	14415.05	0.00	0	1	0	0
root	set	2	14368.78	14367.87	0.00	0	1	0	0
root	Init DB	1	14247.69	14247.69	5972.00	0	0	0	0
root	freturn	152	14079.12	13412.08	0.00	0	0	0	0
root	show_grants	4	13623.91	13326.45	0.00	0	0	0	0
root	jump_if_not	777	9730.42	8495.19	0.00	0	0	0	0
root	drop_table	1	8491.28	8491.28	8171.00	0	0	0	0
root	show_...ations	2	4224.91	2215.06	740.00	540	1784	0	2
root	show_...rsets	2	3990.54	3219.54	2989.00	82	246	0	2
root	show_...status	2	1701.56	999.66	872.00	0	4	0	0
root	error	4	688.98	253.60	0.00	0	0	0	0
root	change_db	4	442.29	145.13	50.00	0	0	0	0
root	Quit	2	191.10	103.66	0.00	0	0	0	0

# 13. Set Persist

# Saving Configuration Changes

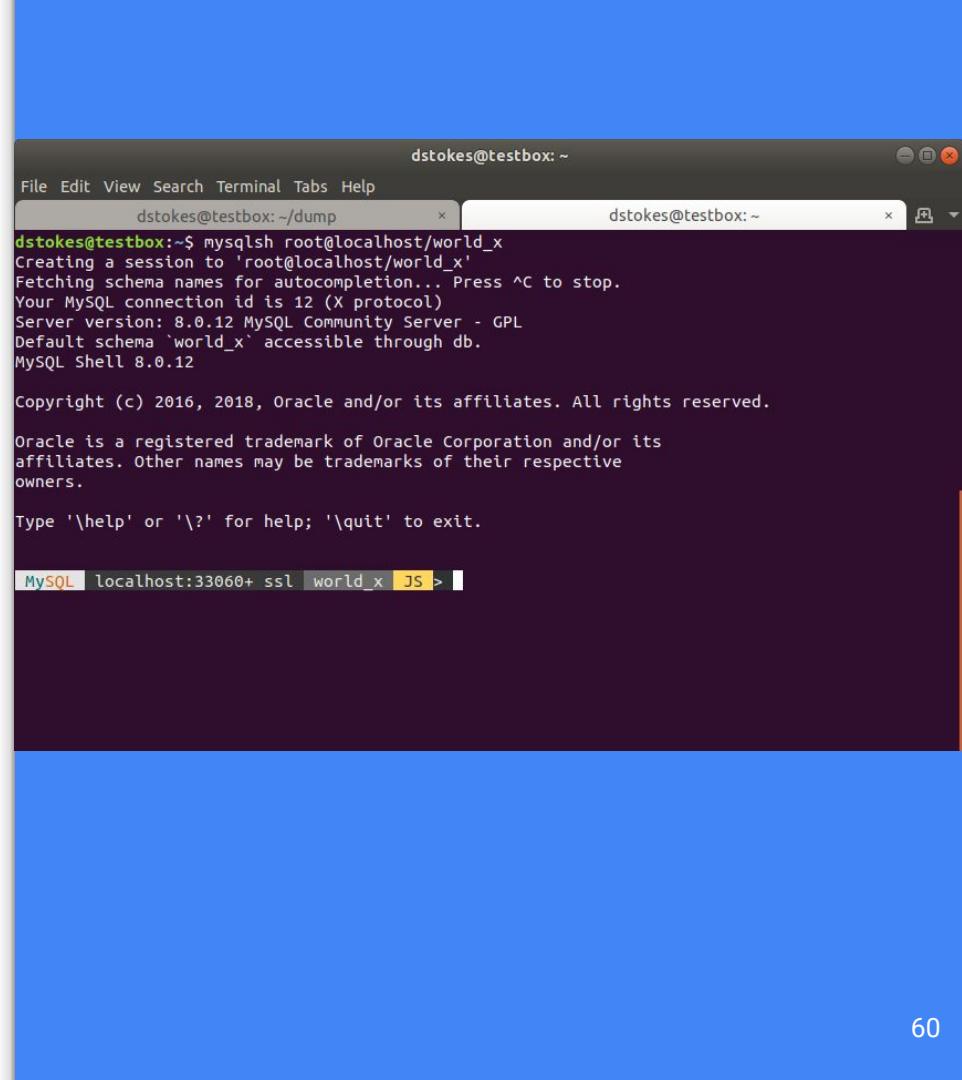
```
SET PERSIST innodb_buffer_pool_size = 512 * 1024 * 1024;
```

The file `mysqld-auto.cnf` is created the first time a `SET PERSIST` statement is executed. Further `SET PERSIST` statement executions will append the contents to this file.

# 14. New Shell

# MySQL Shell

Query tool, administration tool,  
cluster manager, and supports  
Python, JavaScript & SQL



The screenshot shows a terminal window titled "dstokes@testbox: ~" with two tabs open. The left tab is titled "dstokes@testbox:~/dump" and contains the MySQL Shell startup logs. The right tab is also titled "dstokes@testbox: ~". The logs show the connection to the 'world\_x' schema, the MySQL version (8.0.12), and the default schema. It also includes copyright information and a note about Oracle trademarks. The bottom of the window shows the MySQL prompt: "MySQL | localhost:33060+ ssl | world\_x | JS >".

```
dstokes@testbox:~$ mysqlsh root@localhost/world_x
Creating a session to 'root@localhost/world_x'
Fetching schema names for autocompletion... Press ^C to stop.
Your MySQL connection id is 12 (X protocol)
Server version: 8.0.12 MySQL Community Server - GPL
Default schema 'world_x' accessible through db.
MySQL Shell 8.0.12

Copyright (c) 2016, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type '\help' or '\?' for help; '\quit' to exit.

MySQL | localhost:33060+ ssl | world_x | JS >
```

# MySQL Shell

```
MySQL localhost:33060+ ssl world_x JS > \s
MySQL Shell version 8.0.12

Session type: X
Connection Id: 12
Default schema: world_x
Current schema: world_x
Current user: root@localhost
SSL: Cipher in use: ECDHE-RSA-AES128-GCM-SHA256 TLSv1.2
Using delimiter: ;
Server version: 8.0.12 MySQL Community Server - GPL
Protocol version: X protocol
Client library: 8.0.12
Connection: localhost via TCP/IP
TCP port: 33060
Server characterset: utf8mb4
Schema characterset: utf8mb4
Client characterset: utf8mb4
Conn. characterset: utf8mb4
Uptime: 4 hours 34 min 54.0000 sec
```

```
MySQL localhost:33060+ ssl world_x JS > |
```

# MySQL Shell

```
MySQL localhost:33060+ ssl world_x JS > \s
MySQL Shell version 8.0.12

Session type: X
Connection Id: 12
Default schema: world_x
Current schema: world_x
Current user: root@localhost
SSL: Cipher in use: ECDHE-RSA-AES128-GCM-SHA256 TLSv1.2
Using delimiter: ,
Server version: 8.0.12 MySQL Community Server - GPL
Protocol version: X protocol
Client library: 8.0.12
Connection: localhost via TCP/IP
TCP port: 33060
Server characterset: utf8mb4
Schema characterset: utf8mb4
Client characterset: utf8mb4
Conn. characterset: utf8mb4
Uptime: 4 hours 34 min 54.0000 sec
```

```
MySQL localhost:33060+ ssl world_x JS > |
```

# MySQL Shell

```
MySQL localhost:33060+ ssl world_x JS > \s
MySQL Shell version 8.0.12

Session type: X
Connection Id: 12
Default schema: world_x
Current schema: world_x
Current user: root@localhost
SSL: Cipher in use: ECDHE-RSA-AES128-GCM-SHA256 TLSv1.2
Using delimiter: ;
Server version: 8.0.12 MySQL Community Server - GPL
Protocol version: X protocol
Client library: 8.0.12
Connection: localhost via TCP/IP
TCP port: 33060
Server characterset: utf8mb4
Schema characterset: utf8mb4
Client characterset: utf8mb4
Conn. characterset: utf8mb4
Uptime: 4 hours 34 min 54.0000 sec

MySQL localhost:33060+ ssl world_x JS > 
```

# MySQL Shell

Python, JavaScript & SQL modes

Management

`util.checkForServerUpgrade('user@host.com:3306')`

`dba.configureLocalInstance`

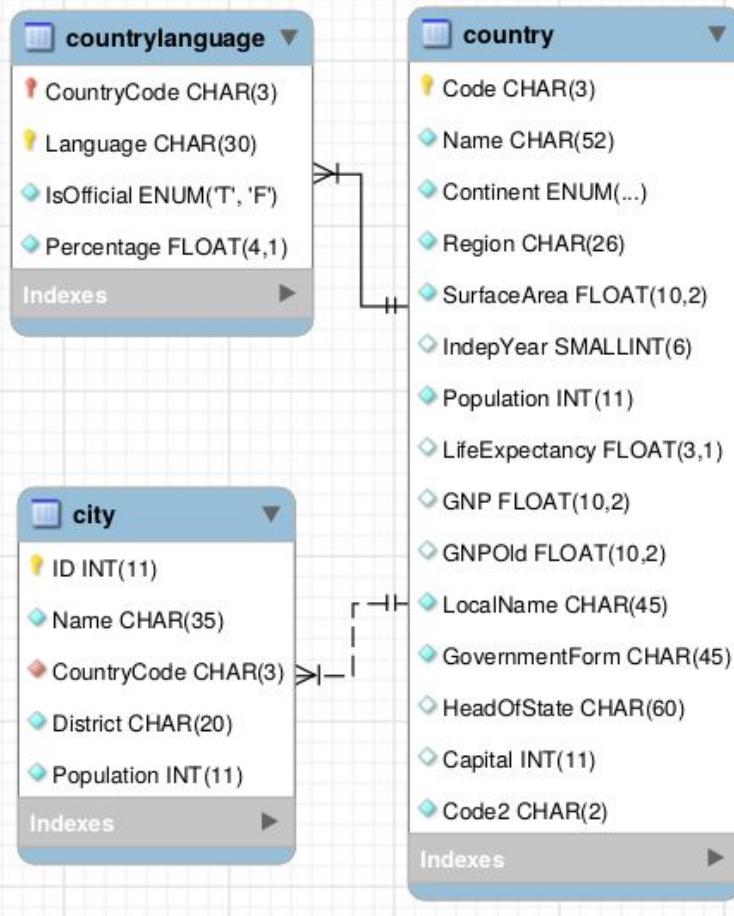
`dba.createCluster`

```
Mysqlx.Crud.Find {
    collection { name: "collection_name", schema: "test" }
    data_model: DOCUMENT
    criteria {
        type: OPERATOR
        operator {
            name: "=="
            param {
                type: IDENT,
                identifier { name: "_id" }
            }
            param {
                type: LITERAL,
                literal {
                    type: V_STRING,
                    v_string: { value: "some_string" }
                }
            }
        }
    }
}
```

# 15. MySQL Document Store

# NoSQL or Document Store

- Schemaless
  - No schema design, no normalization, no foreign keys, no data types, ...
  - Very quick initial development
- Flexible data structure
  - Embedded arrays or objects
  - Valid solution when natural data can not be modeled optimally into a relational model
  - Objects persistence without the use of any ORM - \*mapping object-oriented\*
- JSON
- close to frontend
- native in JS
- easy to learn



```
{
    "GNP" : 249704,
    "Name" : "Belgium",
    "government" : {
        "GovernmentForm" :
            "Constitutional Monarchy, Federation",
        "HeadOfState" : "Philippe I"
    },
    "_id" : "BEL",
    "IndepYear" : 1830,
    "demographics" : {
        "Population" : 10239000,
        "LifeExpectancy" : 77.8000030517578
    },
    "geography" : {
        "Region" : "Western Europe",
        "SurfaceArea" : 30518,
        "Continent" : "Europe"
    }
}
```

**What if there was a way to provide both SQL and NoSQL on one stable platform that has proven stability on well known technology with a large Community and a diverse ecosystem ?**

**With the MySQL Document Store, SQL is now optional!**

- ★ Provides a schema flexible JSON Document Store
- ★ No SQL required
- ★ No need to define all possible attributes, tables, etc.
- ★ Uses new X DevAPI
- ★ Can leverage generated column to extract JSON values into materialized columns that can be indexed for fast SQL searches.
- ★ Document can be ~1GB
  - It's a column in a row of a table
- ★ Allows use of modern programming styles
  - No more embedded strings of SQL in your code
  - Easy to read
- ★ Also works with relational Tables
- ★ Proven MySQL Technology

## ★ Connectors for

- C++, Java, .Net, Node.js, Python, PHP
- working with Communities to help them supporting it too

## ★ New MySQL Shell

- Command Completion
- Python, JavaScripts & SQL modes
- Admin functions
- New Util object
- A new high-level session concept that can scale from single MySQL Server to a multiple server environment

## ★ Non-blocking, asynchronous calls follow common language patterns

## ★ Supports CRUD operations

```
MySQL ➤ JS ➤ \c root@localhost
Creating a session to 'root@localhost'
Enter password:
Fetching schema names for autocomplete... Press ^C to stop.
Your MySQL connection id is 13 (X protocol)
Server version: 8.0.11 MySQL Community Server - GPL
No default schema selected; type \use <schema> to set one.
```

```
MySQL ➤ 🖥 localhost:33060+ 🔒 ➤ JS ➤ session.createSchema('docstore')
<Schema:docstore>
```

```
MySQL ➤ 🖥 localhost:33060+ 🔒 ➤ JS ➤ \use docstore
Default schema `docstore` accessible through db.
```

```
MySQL ➤ 🖥 localhost:33060+ 🔒 ➤ 🗁 docstore ➤ JS ➤ □
```

# Starting using MySQL in few seconds

# Migration from MongoDB to MySQL Document Store

For this example, I will use the well known restaurants collection:

We need to dump the data to a file and

we will use the MySQL Shell

with the Python interpreter to load the data.



# Dump and load using MySQL Shell & Python

This example is inspired by @datacharmer's work: <https://www.slideshare.net/datacharmer/mysql-documentstore>

```
$ mongo quiet eval 'DBQuery.shellBatchSize=30000;  
    db.restaurants.find().shellPrint()' \  
| perl -pe 's/(?:ObjectId|ISODate)\(\"[^\"]+\")/\$1/g' > all_recs.json
```

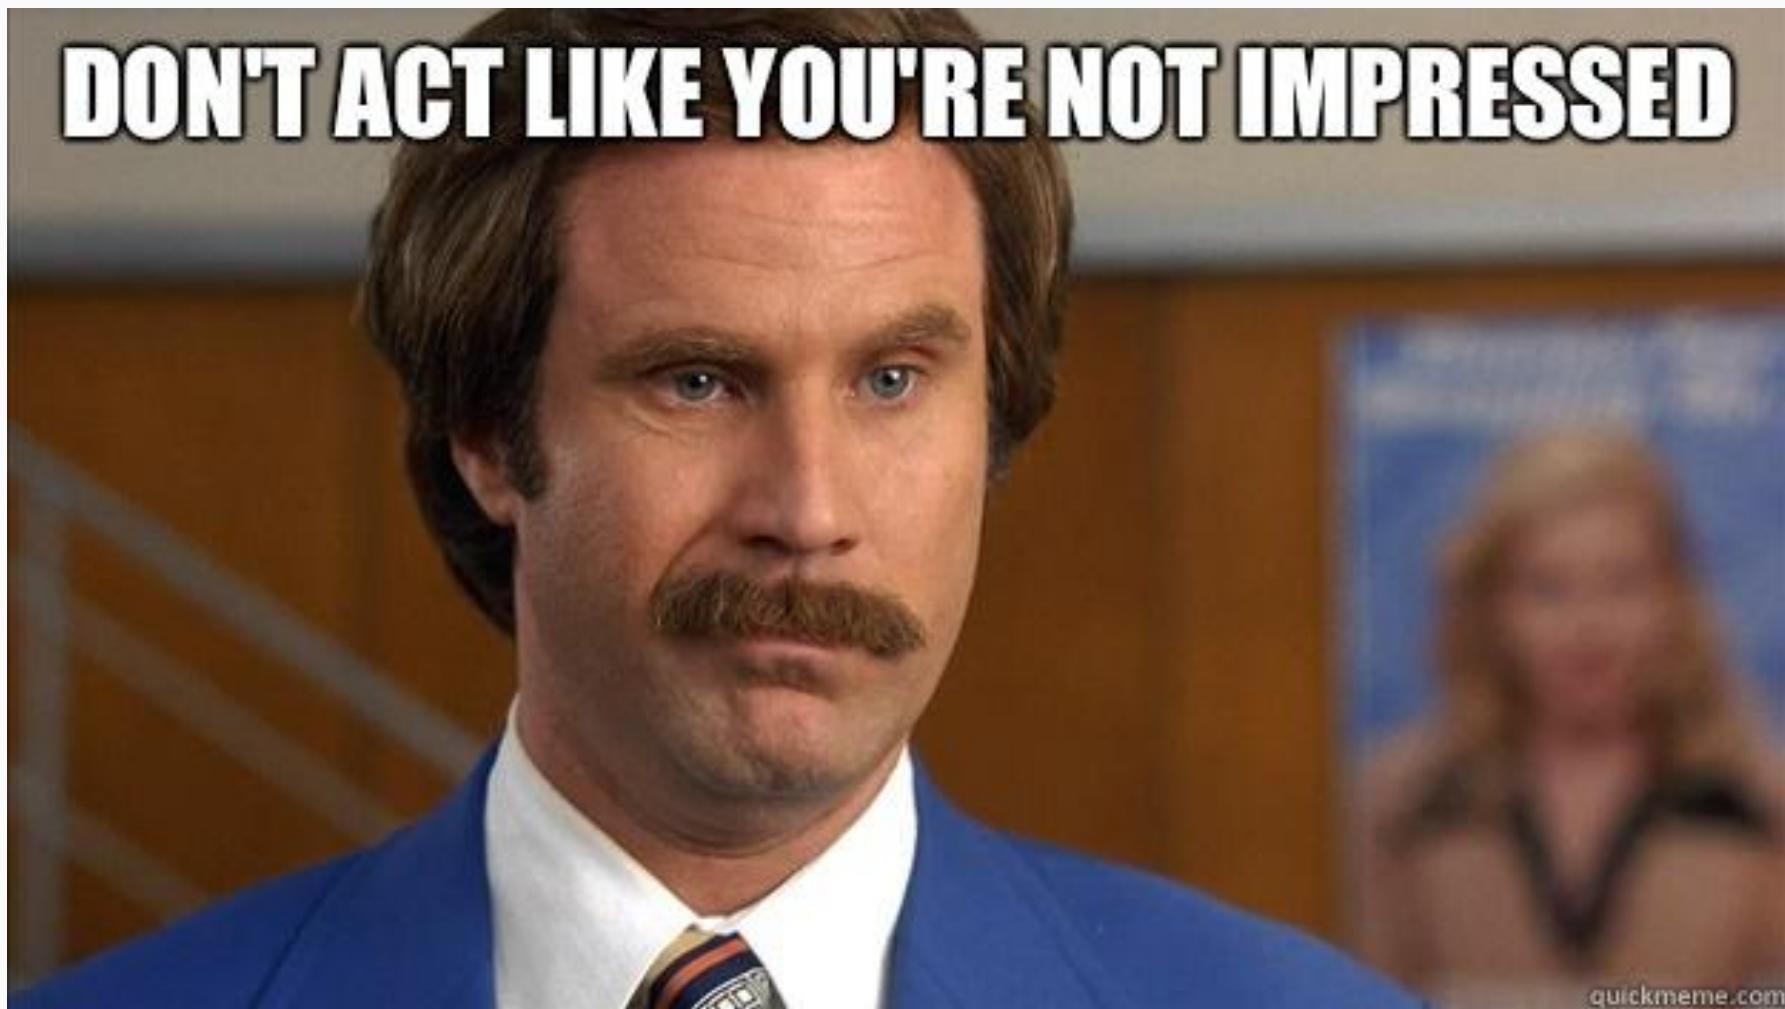
The screenshot shows the MySQL Shell interface with three command-line sessions:

- Session 1:** Shows the connection to `localhost:33060+ docstore` in JS mode. The command `\py` is entered, followed by the message "to Python mode...".
- Session 2:** Shows the connection to `localhost:33060+ docstore` in Py mode. The command `import json` is entered.
- Session 3:** Shows the connection to `localhost:33060+ docstore` in Py mode. The command `import re` is entered.

The code block below shows the full Python script used for the dump and load process:

```
import json  
import re  
with open('all_recs.json', 'r') as json_data:  
    for line in json_data:  
        skip = re.match('Type', line)  
        if not skip:  
            rec = json.loads(line)  
            db.restaurants.add(rec).execute()
```

**DON'T ACT LIKE YOU'RE NOT IMPRESSED**



# Let's query

MySQL ➔ 127.0.0.1:3306 ➔ JS ➔ restaurants.find()

Too many records to show here ... let's limit it!

MySQL ➔ 127.0.0.1:3306 ➔ JS ➔ restaurants.find().limit(1)

```
[  
  {  
    "_id": "5943c83d1adc26055941640c",  
    "address": {  
      "building": "351",  
      "coord": [  
        -73.9851,  
        40.7677  
      ],  
      "street": "West 57 Street",  
      "zipcode": "10019"  
    },  
    "borough": "Manhattan",  
    "cuisine": "Irish",  
    "grades": [  
      {  
        "date": "2014-09-06T00:00:00Z",  
        "grade": "A",  
        "score": 2  
      },  
      {  
        "date": "2013-07-22T00:00:00Z",  
        "grade": "A",  
        "score": 11  
      },  
      {  
        "date": "2012-07-31T00:00:00Z",  
        "grade": "A",  
        "score": 12  
      },  
      {  
        "date": "2011-12-29T00:00:00Z",  
        "grade": "A",  
        "score": 12  
      }  
    ],  
    "name": "Dj Reynolds Pub And Restaurant",  
    "restaurant_id": "30191841"  
  }  
]  
1 document in set (0.08 sec)
```

```
MySQL ➔ 127.0.0.1:33060 ➔ JS ➔ restaurants.find().fields(["name","cuisine"]).limit(2)
[
  {
    "cuisine": "Irish",
    "name": "Dj Reynolds Pub And Restaurant"
  },
  {
    "cuisine": "American",
    "name": "Riviera Caterer"
  }
]
2 documents in set (0.00 sec)
```

```
MySQL ➔ 127.0.0.1:33060 ➔ JS ➔ restaurants.find("cuisine='Italian").fields(["name","cuisine"]).limit(2)
[
  {
    "cuisine": "Italian",
    "name": "Philadelphia Grille Express"
  },
  {
    "cuisine": "Italian",
    "name": "Isle Of Capri Restaurant"
  }
]
2 documents in set (0.00 sec)
```

Let's add a selection criteria

```
MySQL ➤ 127.0.0.1:3306 ➤ JS ➤ restaurants.find("cuisine='French' AND borough!='Manhattan'").  
fields(["name", "cuisine", "borough"]).limit(2)  
[  
  {  
    "borough": "Queens",  
    "cuisine": "French",  
    "name": "La Baraka Restaurant"  
  },  
  {  
    "borough": "Queens",  
    "cuisine": "French",  
    "name": "Air France Lounge"  
  }  
]  
2 documents in set (0.00 sec)
```

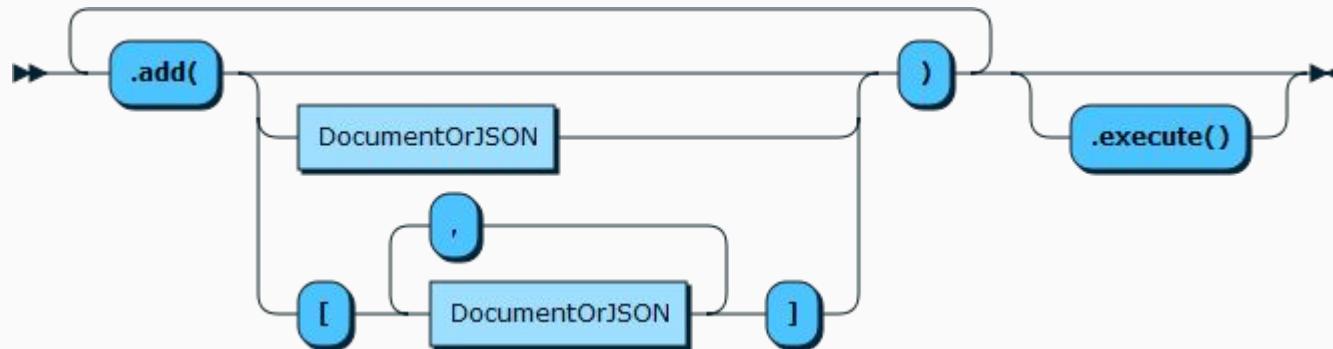
```
> db.restaurants.find({ "cuisine": "French",  
  "borough": { $not: /^Manhattan/ } },  
  { "_id": 0, "name": 1, "cuisine": 1, "borough": 1 }).limit(2)  
{ "borough" : "Queens", "cuisine" : "French",  
  "name" : "La Baraka Restaurant" }  
{ "borough" : "Queens", "cuisine" : "French",  
  "name" : "Air France Lounge" }
```

***Syntax is slightly  
different than  
MongoDB***

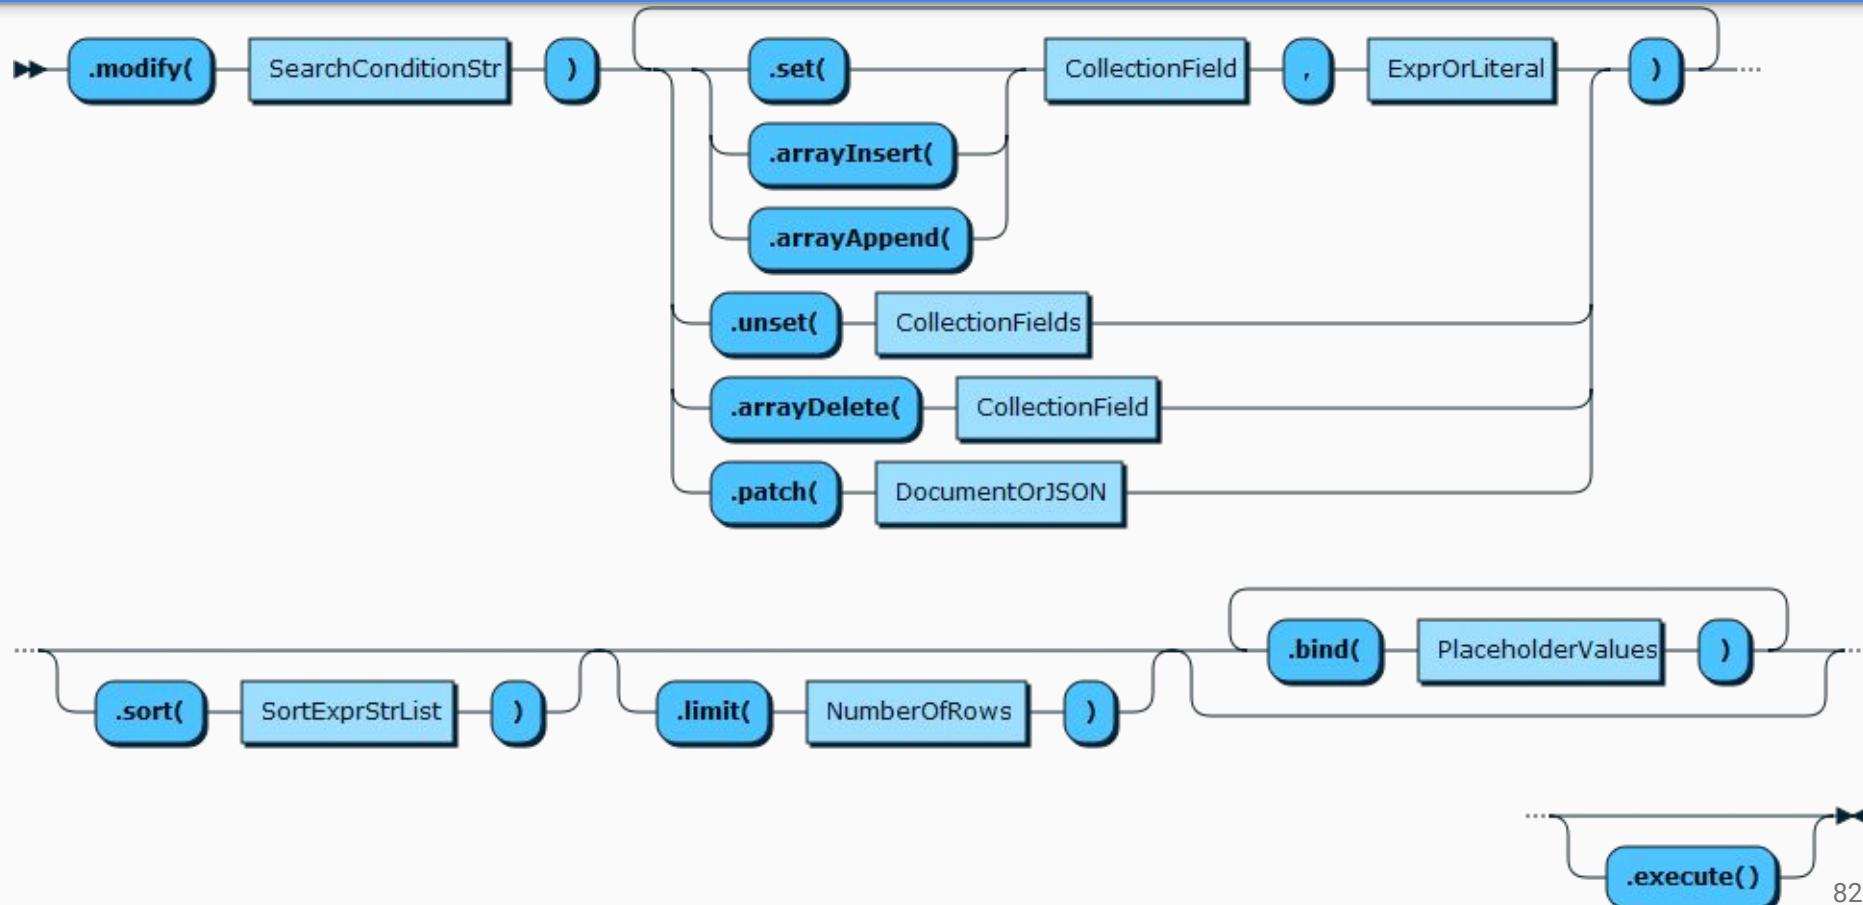
```
MySQL ➔ 127.0.0.1:33060 ➔ JS ➔ restaurants.remove("cuisine='French' AND borough!='Manhattan'").limit(2)
Query OK, 2 items affected (0.16 sec)
```

```
MySQL ➔ 127.0.0.1:33060 ➔ JS ➔ restaurants.find("cuisine='French' AND borough!='Manhattan'").fields(["name","cuisine","_id"]).limit(2)
[
  {
    "_id": "5943c83e1adc2605594170aa",
    "cuisine": "French",
    "name": "Bar Tabac"
  },
  {
    "_id": "5943c83e1adc260559417255",
    "cuisine": "French",
    "name": "Tournesol"
  }
]
2 documents in set (0.01 sec)
```

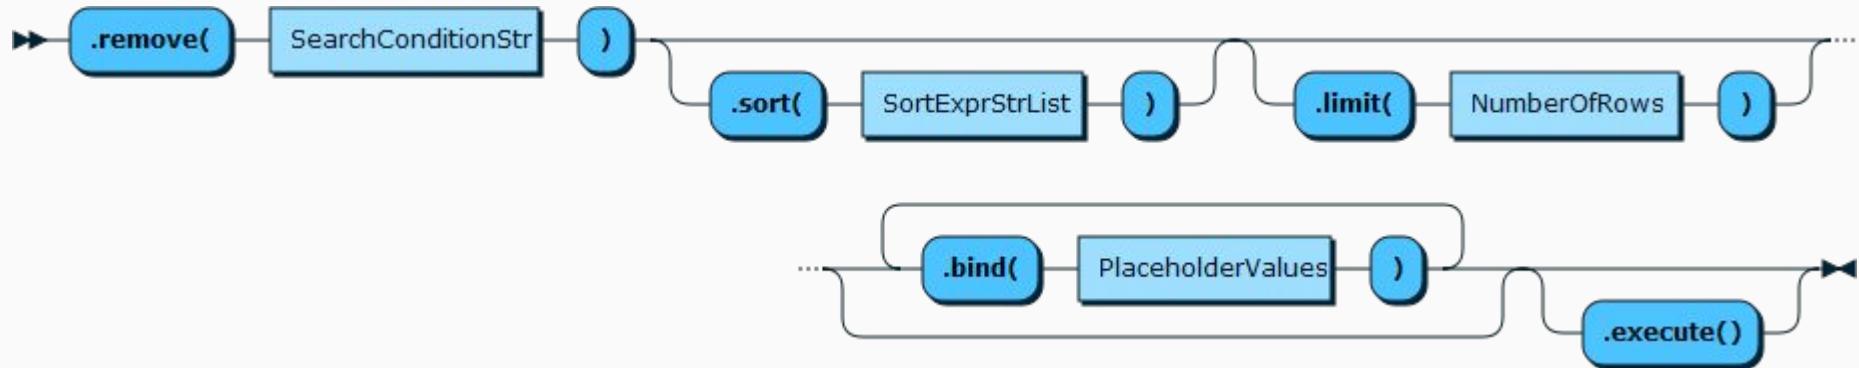
# Add a Document



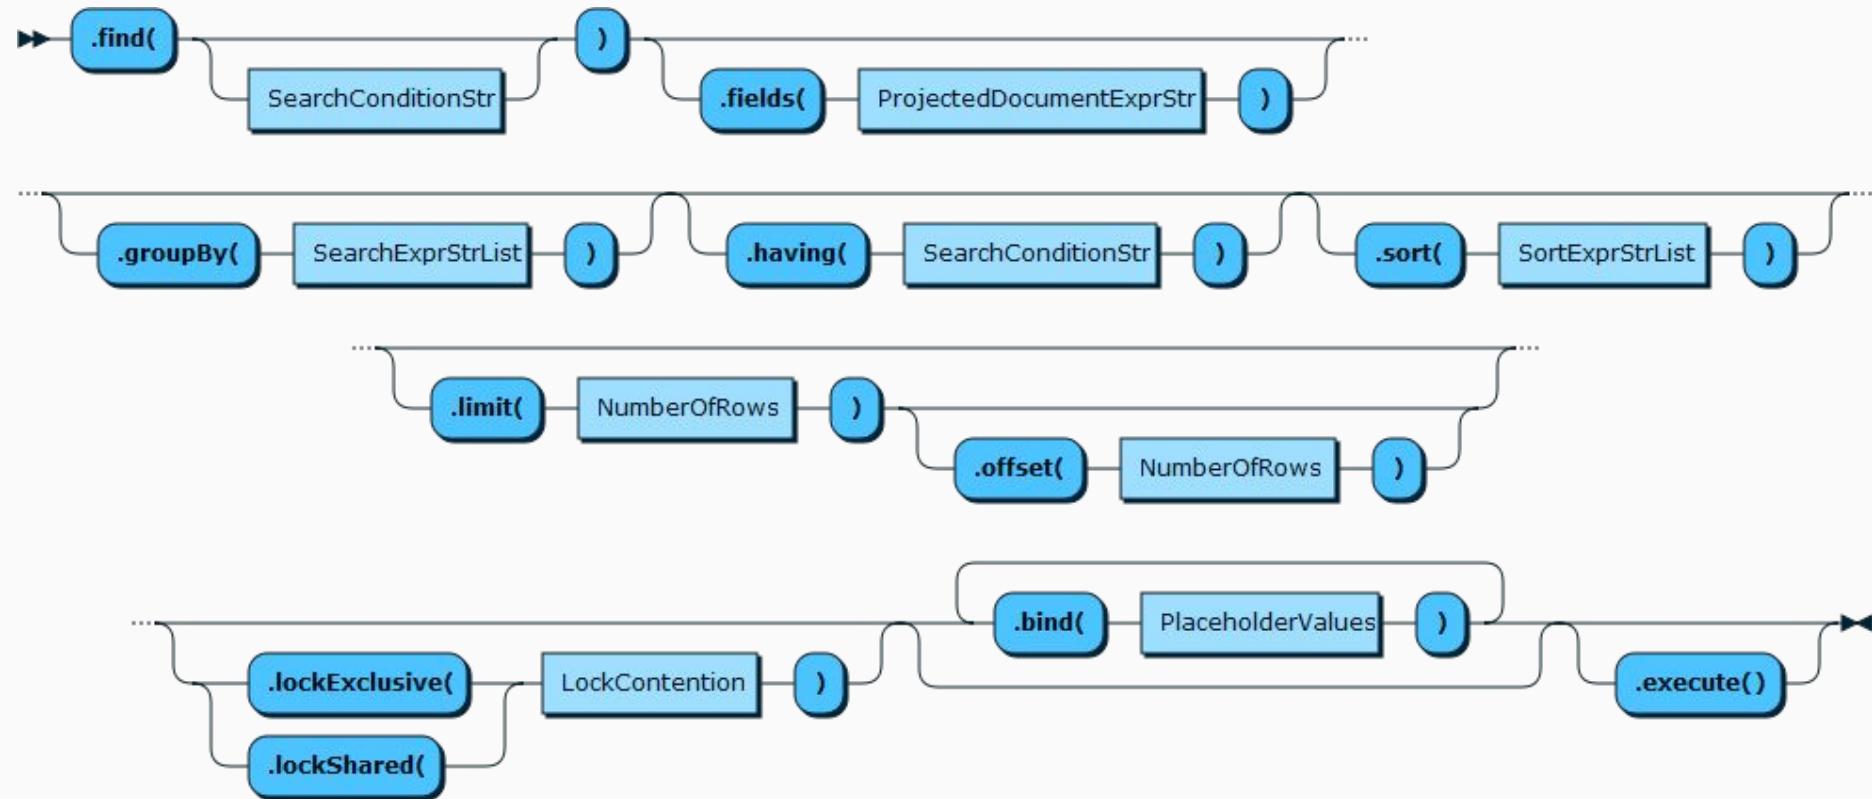
# Modify a Document



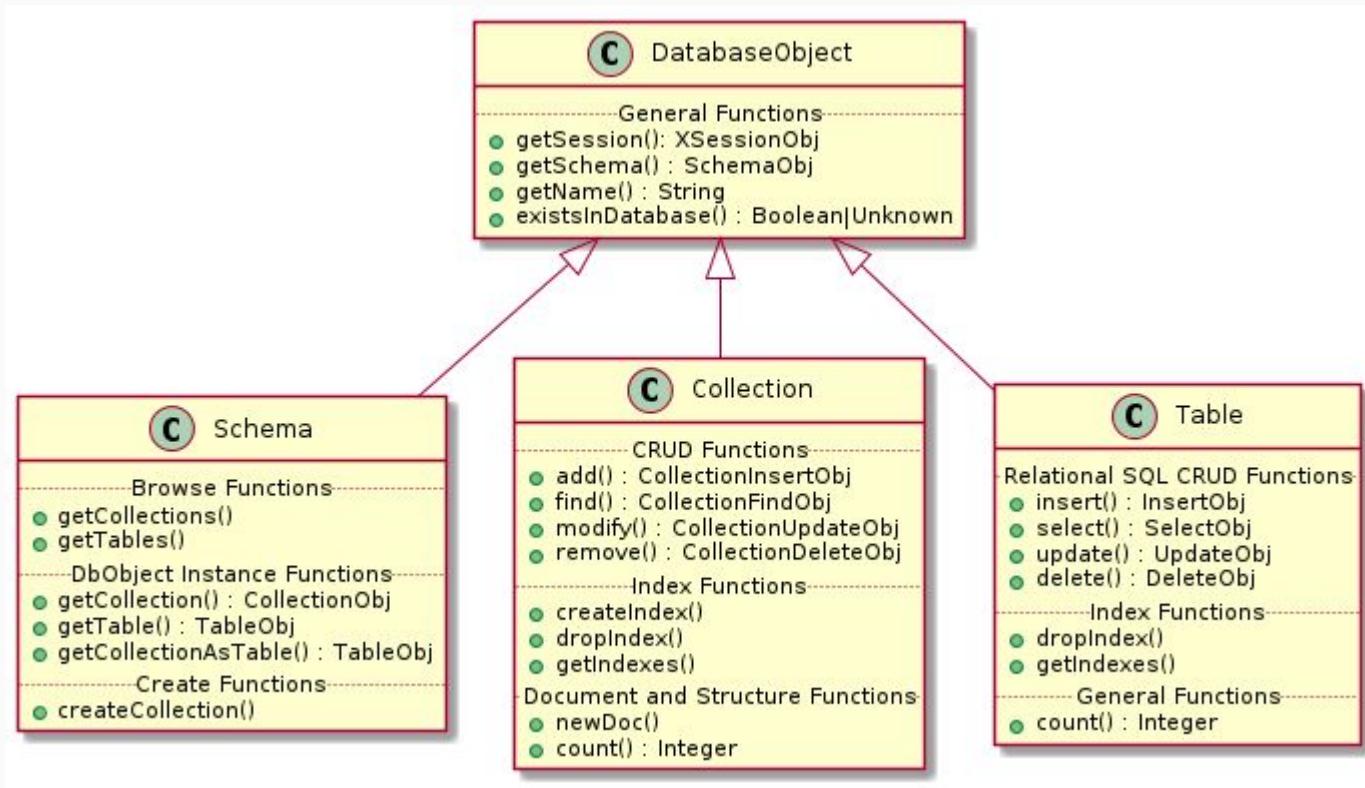
# Remove a Document



# Find a Document



# MySQL Document Store Objects Summary



MySQL → localhost:33060+ fred → JS → session.startTransaction()  
Query OK, 0 rows affected (0.0006 sec)

MySQL → localhost:33060+ fred → JS → test.add({name: 'the René'})  
Query OK, 1 item affected (0.1661 sec)

MySQL → localhost:33060+ fred → JS → test.find()

```
[  
  {  
    "_id": "00005ade551000000000000000000001",  
    "name": "fred"  
  },  
  {  
    "_id": "00005ade551000000000000000000002",  
    "name": "the René"  
  }]  
2 documents in set (0.0019 sec)
```

```
MySQL ➔ localhost:33060+ fred ➔ JS ➔ session.rollback()
Query OK, 0 rows affected (0.0992 sec)
```

```
MySQL ➔ localhost:33060+ fred ➔ JS ➔ test.find()
```

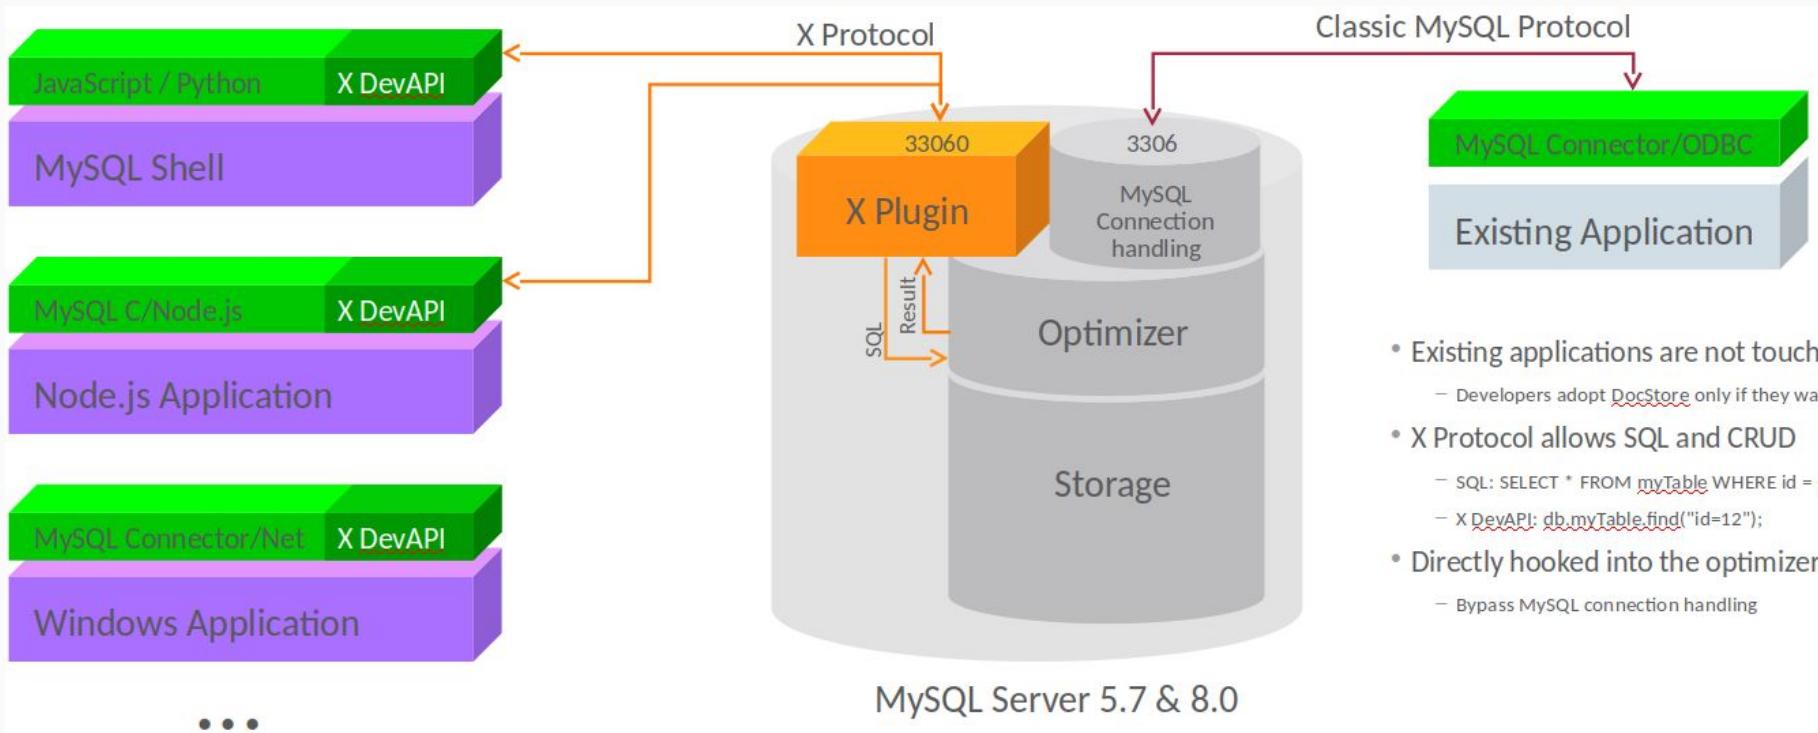
```
[  
  {  
    "_id": "00005ade5510000000000000000001",  
    "name": "fred"  
  }  
]  
1 document in set (0.0130 sec)
```



What about old SQL? The Hidden Part of the Iceberg

# JSON datatype is behind the scene

- ★ Native datatype (since 5.7.8)
- ★ JSON values are stored in MySQL tables using UTF8MB4
- ★ Conversion from "native" SQL types to JSON values
- ★ JSON manipulation functions (JSON\_EXTRACT,  
JSON\_KEYS, JSON\_SEARCH, JSON\_TABLES, ...)
- ★ Generated/virtual columns
  - Indexing JSON data
  - Foreign Keys to JSON data
  - SQL Views to JSON data



- Existing applications are not touched
  - Developers adopt DocStore only if they want to
- X Protocol allows SQL and CRUD
  - SQL: `SELECT * FROM myTable WHERE id = 12;`
  - X DevAPI: `db.myTable.find("id=12");`
- Directly hooked into the optimizer
  - Bypass MySQL connection handling

```
MySQL ➔ localhost:33060+ 🔒 ➔ docstore ➔ SQL ➔ DESC restaurants;  
+-----+-----+-----+-----+-----+  
| Field | Type          | Null | Key | Default | Extra           |  
+-----+-----+-----+-----+-----+  
| doc   | json          | YES  |      | NULL    |                 |  
| _id   | varbinary(32) | NO   | PRI  | NULL    | STORED GENERATED |  
+-----+-----+-----+-----+-----+  
2 rows in set (0.0668 sec)
```

```
MySQL ➔ localhost:33060+ 🔒 ➔ docstore ➔ SQL ➔ show create table restaurants\G  
***** 1. row *****  
    Table: restaurants  
Create Table: CREATE TABLE `restaurants` (  
  `doc` json DEFAULT NULL,  
  `_id` varbinary(32) GENERATED ALWAYS AS (json_unquote(json_extract(`doc`,_utf8mb4'$.id')))) STORED NOT NULL,  
  PRIMARY KEY (`_id`)  
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci
```

What does a collection look like on the server ?

# \_id

Every document has a unique identifier called the *document ID*, which can be thought of as the equivalent of a table's primary key. The document ID value can be manually assigned when adding a document.

If no value is assigned, a document ID is generated and assigned to the document automatically !

Use `getDocumentId()` or `getDocumentIds()` to get \_ids(s)

# Mapping to SQL Examples

```
createCollection('mycollection')

CREATE TABLE `mycoll` (
    doc JSON,
    PRIMARY KEY(`id`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4;

INSERT INTO `mycoll` (`doc`) VALUES (
    {
        "id": "5a76a807fb367e6d",
        "name": "test",
        "age": 123
    }
);
```

# More Mapping to SQL Examples

```
mycollection.find("test > 100")
```

```
SELECT doc  
  FROM `test`.`mycoll`  
 WHERE (JSON_EXTRACT(doc,'$.test') >100);
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ ALTER TABLE restaurants ADD COLUMN borough  
VARCHAR( 20) GENERATED ALWAYS AS (json_unquote(json_extract(`doc` , '$.borough'))) VIRTUAL;
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ SELECT _id, borough FROM restaurants LIMIT 5;
```

_id	borough
59ca58986b977d7c822812b7	Bronx
59ca58986b977d7c822812b8	Brooklyn
59ca58986b977d7c822812b9	Manhattan
59ca58986b977d7c822812ba	Brooklyn
59ca58986b977d7c822812bb	Queens

```
5 rows in set (0.00 sec)
```

```
MySQL ➔ localhost:33060+ 🔒 ➔ docstore ➔ JS ➔ db.restaurants.createIndex('cuisine_idx', {fields:{field: "$.cuisine", required: false, type: "text(20)"}})
```

```
MySQL ➔ localhost:33060+ 🔒 ➔ docstore ➔ SQL ➔ show create table restaurants\G
***** 1. row *****
Table: restaurants
Create Table: CREATE TABLE `restaurants` (
  `doc` json DEFAULT NULL,
  `_id` varbinary(32) GENERATED ALWAYS AS (json_unquote(json_extract(`doc`,_utf8mb4'$.id')))
) STORED NOT NULL,
  `$ix_t20_BC26D4DF1273E3F7412529AEE9E95A0CC8475CEB` text GENERATED ALWAYS AS (json_unquote(
  json_extract(`doc`,_utf8mb4'$.cuisine')))) VIRTUAL,
  PRIMARY KEY (`_id`),
  KEY `cuisine_idx` (`$ix_t20_BC26D4DF1273E3F7412529AEE9E95A0CC8475CEB`(20))
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci
```

**It's also possible to create indexes without using SQL syntax**

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ ALTER TABLE restaurants ADD COLUMN cuisine  
VARCHAR(20) GENERATED ALWAYS AS (doc->>"$.cuisine") VIRTUAL, WITH VALIDATION;  
ERROR: 1406: Data too long for column 'cuisine' at row 10
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ ALTER TABLE restaurants ADD COLUMN cuisine  
VARCHAR(20) GENERATED ALWAYS AS (LEFT(doc->>"$.cuisine",20)) VIRTUAL, WITH VALIDATION;  
Query OK, 25359 rows affected (0.00 sec)
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ SELECT _id, borough, cuisine FROM restaurants  
LIMIT 5;
```

_id	borough	cuisine
59ca58986b977d7c822812b7	Bronx	Bakery
59ca58986b977d7c822812b8	Brooklyn	Hamburgers
59ca58986b977d7c822812b9	Manhattan	Irish
59ca58986b977d7c822812ba	Brooklyn	American
59ca58986b977d7c822812bb	Queens	Jewish/Kosher

MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ EXPLAIN SELECT doc->("\$.name" AS name, cuisine, borough FROM restaurants WHERE cuisine='Italian' AND borough='Brooklyn' LIMIT 2\G

\*\*\*\*\* 1. row \*\*\*\*\*

```
    id: 1
select_type: SIMPLE
    table: restaurants
  partitions: NULL
      type: ALL
possible_keys: NULL
        key: NULL
    key_len: NULL
        ref: NULL
      rows: 24890
  filtered: 1.0000001192092896
    Extra: Using where
1 row in set, 1 warning (0.00 sec)
Note (code 1003): /* select#1 */ select json_unquote(json_extract(`docstore`.`restaurants`.`doc`, '$.name')) AS `name`, `docstore`.`restaurants`.`cuisine` AS `cuisine`, `docstore`.`restaurants`.`borough` AS `borough` from `docstore`.`restaurants` where ((`docstore`.`restaurants`.`borough` = 'Brooklyn') and (`docstore`.`restaurants`.`cuisine` = 'Italian')) limit 2
```

MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ EXPLAIN SELECT doc->("\$.name" AS name, cuisine, borough FROM restaurants WHERE cuisine='Italian' AND borough='Brooklyn' LIMIT 2\G

\*\*\*\*\* 1. row \*\*\*\*\*

```
    id: 1
select_type: SIMPLE
    table: restaurants
  partitions: NULL
      type: ALL ←
possible_keys: NULL
        key: NULL
    key_len: NULL
        ref: NULL
      rows: 24890 ←
filtered: 1.0000001192092896
    Extra: Using where
1 row in set, 1 warning (0.00 sec)
Note (code 1003): /* select#1 */ select json_unquote(json_extract(`docstore`.`restaurants`.`doc`, '$.name')) AS `name`, `docstore`.`restaurants`.`cuisine` AS `cuisine`, `docstore`.`restaurants`.`borough` AS `borough` from `docstore`.`restaurants` where ((`docstore`.`restaurants`.`borough` = 'Brooklyn') and (`docstore`.`restaurants`.`cuisine` = 'Italian')) limit 2
```

## SQL and JSON Example (3): explain

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ ALTER TABLE restaurants ADD INDEX cuisine_borought_idx(cuisine,borough);
Query OK, 0 rows affected (0.00 sec)
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ EXPLAIN SELECT doc->"$.name" AS name, cuisine, borough
FROM restaurants WHERE cuisine='Italian' AND borough='Brooklyn' LIMIT 2\G
***** 1. row *****
      id: 1
select_type: SIMPLE
      table: restaurants
     partitions: NULL
        type: ref
possible_keys: cuisine_borought_idx
          key: cuisine_borought_idx
      key_len: 166
        ref: const,const
       rows: 192
     filtered: 100
       Extra: NULL
1 row in set, 1 warning (0.00 sec)
Note (code 1003): /* select#1 */ select json_unquote(json_extract(`docstore`.`restaurants`.`doc`,'$.name')) AS `name`, `docstore`.`restaurants`.`cuisine` AS `cuisine`, `docstore`.`restaurants`.`borough` AS `borough` from `docstore`.`restaurants` where ((`docstore`.`restaurants`.`borough` = 'Brooklyn') and (`docstore`.`restaurants`.`cuisine` = 'Italian')) limit 2
```

## SQL and JSON Example (4): add index

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ ALTER TABLE restaurants ADD INDEX cuisine_borought_idx(cuisine,borough);
Query OK, 0 rows affected (0.00 sec)
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ EXPLAIN SELECT doc->"$.name" AS name, cuisine, borough
FROM restaurants WHERE cuisine='Italian' AND borough='Brooklyn' LIMIT 2\G
***** 1. row *****
      id: 1
select_type: SIMPLE
      table: restaurants
     partitions: NULL
        type: ref
possible_keys: cuisine_borought_idx
          key: cuisine_borought_idx
      key_len: 166
        ref: const,const
       rows: 192
     filtered: 100
       Extra: NULL
1 row in set, 1 warning (0.00 sec)
Note (code 1003): /* select#1 */ select json_unquote(json_extract(`docstore`.`restaurants`.`doc`, '$.name')) AS `name`, `docstore`.`restaurants`.`cuisine` AS `cuisine`, `docstore`.`restaurants`.`borough` AS `borough` from `docstore`.`restaurants` where ((`docstore`.`restaurants`.`borough` = 'Brooklyn') and (`docstore`.`restaurants`.`cuisine` = 'Italian')) limit 2
```

## SQL and JSON Example (4): add index

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ select doc->>'$.grades' from restaurants limit 1\G
***** 1. row *****
doc->'$.grades': [{"date": "2014-03-03T00:00:00Z", "grade": "A", "score": 2}, {"date": "2013-09-11T0
0:00:00Z", "grade": "A", "score": 6}, {"date": "2013-01-24T00:00:00Z", "grade": "A", "score": 10}, {"date": "2011-11-23T00:00:00Z", "grade": "A", "score": 9}, {"date": "2011-03-10T00:00:00Z", "grade": "B", "score": 14}]
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ select doc->'$.grades[0]' from restaurants limit 1\G
***** 1. row *****
doc->'$.grades[0]': {"date": "2014-03-03T00:00:00Z", "grade": "A", "score": 2}
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ select doc->'$.grades[last]' from restaurants limit 1\G
***** 1. row *****
doc->'$.grades[last]': {"date": "2011-03-10T00:00:00Z", "grade": "B", "score": 14}
```

```
MySQL ➔ localhost:33060+ ➔ docstore ➔ SQL ➔ select doc->'$.grades[1 to 2]' from restaurants limit 1
\G
***** 1. row *****
doc->'$.grades[1 to 2]': [{"date": "2013-09-11T00:00:00Z", "grade": "A", "score": 6}, {"date": "2013
-01-24T00:00:00Z", "grade": "A", "score": 10}]
```

```
SELECT _id, borough, cuisine, name  
FROM restaurants, JSON_TABLE(doc, "$"  
COLUMNS(name CHAR(40) PATH ".$.name"))  
t2 LIMIT 2;
```

_id	borough	cuisine	name
59c3a6946273b7d975cf3b39	Bronx	Bakery	Morris Park Bake Shop
59c3a6946273b7d975cf3b3a	Brooklyn	Hamburgers	Wendy'S

**JSON\_TABLE turns your  
un-structured JSON data into a  
temporary structured table!**

```
127.0.0.1:33060+ docstore SQL
SELECT name, borough, cuisine, street, zipcode
FROM restaurants, JSON_TABLE(doc, "$" COLUMNS(
    name CHAR(40) PATH ".$name",
    street CHAR(20) PATH ".$address.street",
    zipcode CHAR(10) PATH ".$address.zipcode"))
LIMIT 2;
```

name	borough	cuisine	street	zipcode
Morris Park Bake Shop	Bronx	Bakery	Morris Park Ave	10462
Wendy's	Brooklyn	Hamburgers	Flatbush Avenue	11225

This temporary structured table can  
be treated like any other table --  
**LIMIT, WHERE, GROUP BY ...**

# More Sophisticated Analysis

Dig deeper into  
your data for  
results

```
WITH cte1 AS (SELECT doc->>"$.name" AS 'name',
doc->>"$.cuisine" AS 'cuisine',
(SELECT AVG(score) FROM
JSON_TABLE(doc,("$.grades[*]")
COLUMNS (score INT PATH ".$score")) as r ) AS avg_score
FROM restaurants)
SELECT *, rank() OVER
(PARTITION BY cuisine ORDER BY avg_score) AS `rank`
FROM cte1
ORDER by `rank`, avg_score DESC limit 10;
```

This query uses a **Common Table Expression (CTE)** and a **Windowing Function** to rank the average scores of each restaurant, by each cuisine with unstructured JSON data

Find the top 10 restaurants by grade for each cuisine

# This is the best of the two worlds in one product !

- **Data integrity**
- **ACID Compliant**
- **Transactions**
- **SQL**
- **Schemaless**
- **flexible data structure**
- **easy to start (CRUD)**

# 16. Locking Changes

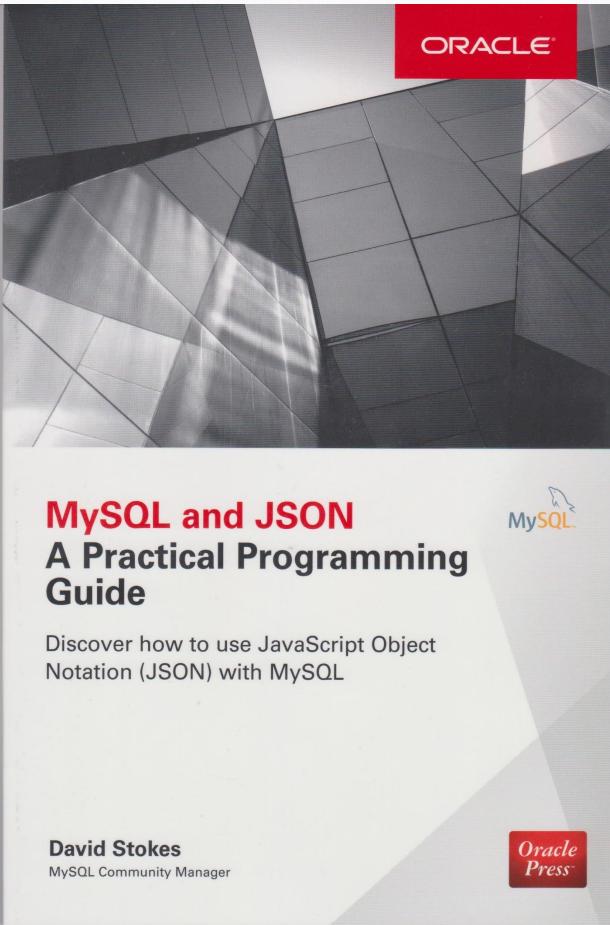
# SKIP LOCKED and NOWAIT

```
START TRANSACTION;  
SELECT * FROM seats WHERE seat_no BETWEEN 2 AND 3  
AND booked = 'NO'  
FOR UPDATE SKIP LOCKED;  
-----  
SELECT seat_no  
FROM seats JOIN seat_rows USING ( row_no )  
WHERE seat_no IN (3,4) AND seat_rows.row_no IN (12)  
AND booked = 'NO'  
FOR UPDATE OF seats SKIP LOCKED  
FOR SHARE OF seat_rows NOWAIT;
```

# Conclusion

# **Big Changes**

- 1. Constant Integration**
- 2. Smarter about environment**
- 3. More powerful SQL**
- 4. Data Dictionary**
- 5. NoSQL and SQL -- Best of both worlds**
- 6. Better Command and Control**



Please Buy My Book!!!

# Thanks!

Contact info:

Dave Stokes

[David.Stokes@Oracle.com](mailto:David.Stokes@Oracle.com)

@Stoker

[slideshare.net/davidmstokes](http://slideshare.net/davidmstokes)

[speakerdeck.com/davidmstokes](http://speakerdeck.com/davidmstokes)

[Elephantdolphin.blogger.com](http://Elephantdolphin.blogger.com)

[opensourcedba.Wordpress.com](http://opensourcedba.Wordpress.com)

