



# MySQL 8.0 Optimizer Guide

Morgan Tocker  
MySQL Product Manager (Server)

ORACLE®

# Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

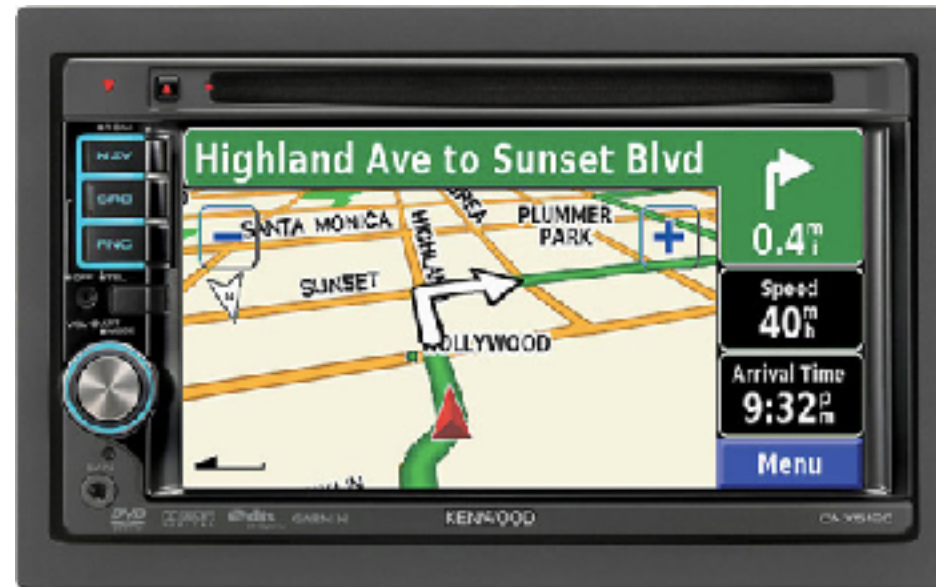
# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets


# Introduction

- SQL is declarative
- You state “*what you want*” not “*how you want*”
- Can’t usually sight check queries to understand execution efficiency
- Database management system is like a GPS navigation system. It finds the “best” route.

# GPS...



# MySQL Optimizer



# Diagnostic Commands

- EXPLAIN (all versions)
- EXPLAIN FORMAT=JSON (5.6+)
  - Supported by Workbench in Visual format
- Optimizer Trace (5.6+)



# Examples from “The World Schema”

- Contains Cities, Countries, Language statistics
- Download from:
  - <https://dev.mysql.com/doc/index-other.html>
- Very small data set
  - Good for learning
  - Not good for explaining performance differences

# Primary Table we are using

```
CREATE TABLE `Country` (  
  `Code` char(3) NOT NULL DEFAULT '',  
  `Name` char(52) NOT NULL DEFAULT '',  
  `Continent` enum('Asia','Europe','North America','Africa','Oceania','Antarctica','South America') NOT NULL DEFAULT  
  'Asia',  
  `Region` char(26) NOT NULL DEFAULT '',  
  `SurfaceArea` float(10,2) NOT NULL DEFAULT '0.00',  
  `IndepYear` smallint(6) DEFAULT NULL,  
  `Population` int(11) NOT NULL DEFAULT '0',  
  `LifeExpectancy` float(3,1) DEFAULT NULL,  
  `GNP` float(10,2) DEFAULT NULL,  
  `GNPold` float(10,2) DEFAULT NULL,  
  `LocalName` char(45) NOT NULL DEFAULT '',  
  `GovernmentForm` char(45) NOT NULL DEFAULT '',  
  `HeadOfState` char(60) DEFAULT NULL,  
  `Capital` int(11) DEFAULT NULL,  
  `Code2` char(2) NOT NULL DEFAULT '',  
  PRIMARY KEY (`Code`)  
) ENGINE=InnoDB DEFAULT CHARSET=latin1  
1 row in set (0.00 sec)
```

# Companion Website

- Content from “The Unofficial MySQL 8.0 Optimizer Guide”
  - <http://www.unofficialmysqlguide.com/>
- More detailed text for many of the examples here...
- Most still applies to 5.6+
  - EXPLAIN FORMAT=JSON in 5.6 does not show cost
  - Costs will be different
  - Output from Optimizer Trace may differ
  - Some features will be missing

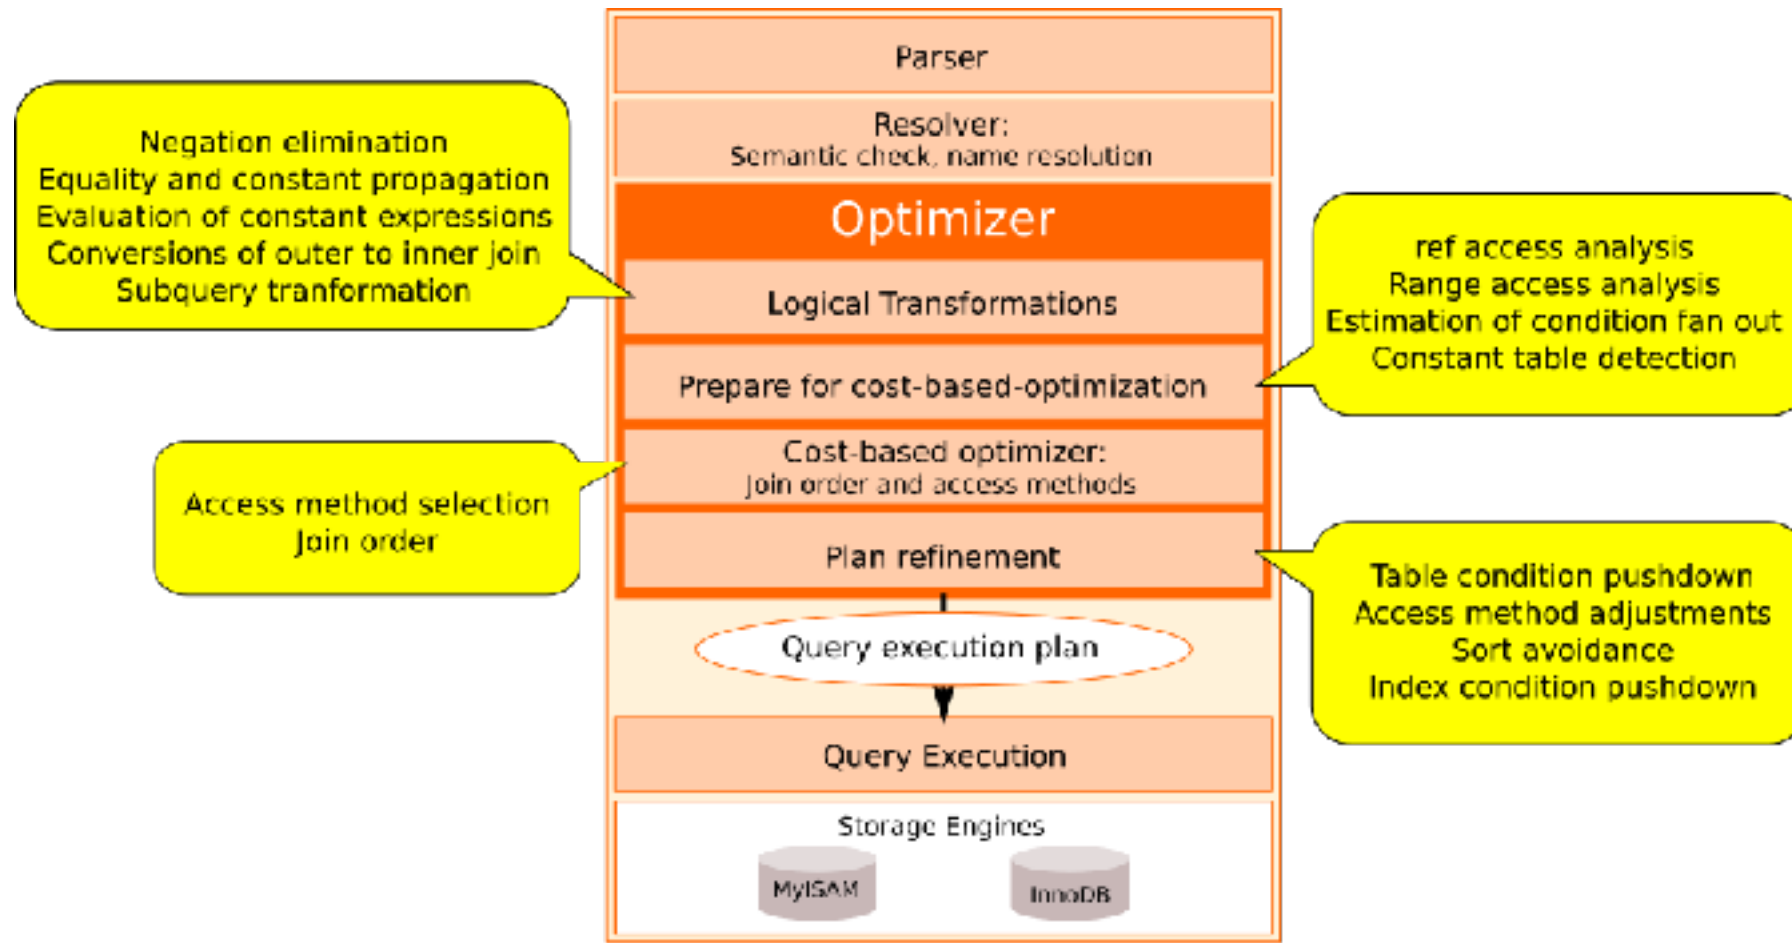
# Danger: Code on slides!

- Some examples may appear small
- Please feel free to download this deck from:
  - <https://www.slideshare.net/morgo/mysql-80-optimizer-guide>
- Follow along on your laptop

# Agenda

1. Introduction
2. **Server Architecture**
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Server Architecture



# Just the Important Parts

- Comprised of the Server and Storage Engines
- Query Optimization happens at the Server Level
- Semantically there are four stages of Query Optimization
- Followed by Query Execution

# Agenda

1. Introduction
2. Server Architecture
3. **B+trees**
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets




# B+trees

- When we mean “*add an index*” we usually mean “*add a B+tree index*”:
  - Includes PRIMARY, UNIQUE, INDEX type indexes.
- Understanding the basic structure of B+trees helps with optimization

# Binary Tree

- Not the same as a B+tree
- Understand Binary Tree first then compare and contrast



Locate 829813 in a  
(balanced) binary tree of  
1MM  $\approx$  20 hops.

**is this good?**


# B+tree

- Amortizes disk accesses by clustering into pages:
- Can achieve same outcome in **two hops**:

```
CREATE TABLE users (  
  id INT NOT NULL auto_increment,  
  username VARCHAR(32) NOT NULL,  
  payload TEXT,  
  PRIMARY KEY (id)  
);
```

# B+tree

- Amortizes disk accesses by clustering into pages
- Can achieve same outcome in **two hops**:



# B-trees are wide not deep

- From the root page: values  $\geq 800788$  but  $< 829908$  are on page 16386.
- From page 16386: values  $\geq 829804$  but  $< 829830$  are on **leaf page** 32012.
- Large fan out factor; 1000+ keys/page which point to another index page with 1000+ keys/page

# InnoDB uses a Clustered Index

- In InnoDB the data rows are also stored in a B+tree, organized by the primary key
- Secondary key indexes always include the value of the primary key

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. **EXPLAIN**
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# EXPLAIN

- Pre-execution view of how *MySQL intends* to execute a query
- Prints what *MySQL considers the best plan* after a process of considering potentially thousands of choices



```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM Country WHERE continent='Asia' and population > 5000000;
```

```
{
```

```
  "query_block": {
```

```
    "select_id": 1,
```

```
    "cost_info": {
```

```
      "query_cost": "25.40"
```

```
    },
```

```
    "table": {
```

```
      "table_name": "country",
```

```
      "access_type": "ALL",
```

```
      "rows_examined_per_scan": 239,
```

```
      "rows_produced_per_join": 11,
```

```
      "filtered": "6.46",
```


```
..
```

```
and ("attached_condition": "((`world`.`country`.`Continent` = 'Asia')  
and (`world`.`country`.`Population` > 5000000))")
```

```
  }
```

```
}
```

```
}
```



# What indexes will make this query faster?

- Some Suggestions:
  - Index on p (population)
  - Index on c (continent)
  - Index on p\_c (population, continent)
  - Index on c\_p (continent, population)

```
ALTER TABLE Country ADD INDEX p (population);
```

```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM Country WHERE continent='Asia' and population > 5000000;
```

```
{
```

```
  "query_block": {
```

```
    "select_id": 1,
```

```
    "cost_info": {
```

```
      "query_cost": "25.40"
```

```
    },
```

```
    "table": {
```

```
      "table_name": "Country",
```

```
      "access_type": "ALL",
```

```
      "possible_keys": [
```

```
        "p"
```

```
      ],
```

```
      "rows_examined_per_scan": 239,
```


```
      "rows_produced_per_join": 15,
```

```
      "filtered": "6.46",
```

```
..
```

```
    "attached_condition": "((`world`.`country`.`Continent` = 'Asia') and  
(`world`.`country`.`Population` > 5000000))"
```

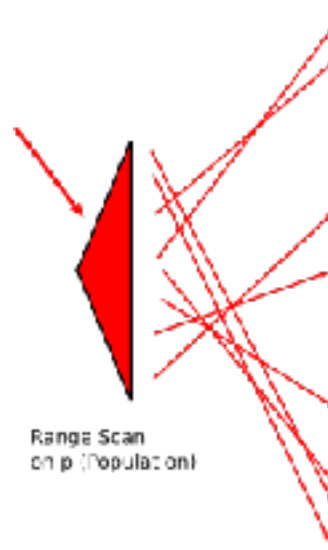
```
..
```



# Why would an index not be used?

```
SELECT * FROM Country
WHERE continent='Asia' AND population > 5000000;
```

Country



ABW	Aruba, North America, Caribbean, .. Beatrix, 129, AW
AFG	Afghanistan, Asia, Southern and Central Asia, .. 1, AF
AGO	Angola, Africa, Central Africa, .. 55, AO
ALA	Anguilla, North America, Caribbean, .. Elisabeth II, 62, AI
ALB	Albania, Europe, Southern Europe, .. Rexhep Mejdani, 34, AL
AND	Andorra, Europe, Southern Europe, .. 55, AD
..	..
CHN	China, Asia, Eastern Asia, 9572900.00, .. Jiang Zemin, 1391, CN
..	..
IND	India, Asia, Southern and Central Asia, .. 1109, IN
..	..
IO	British Indian Ocean Territory, Africa, .. Elisabeth II, NULL, IO
IRL	Ireland, Europe, British Islands, .. Mary McAldese, 1447, IE
IRN	Iran, Asia, Southern and Central Asia, .. 1380, IR
..	..
ZAF	South Africa, Africa, Southern Africa, .. Thabo Mbeki, 716, ZA
ZMB	Zambia, Africa, Eastern Africa, .. Frederick Chiluba, 3162, ZM
ZWE	Zimbabwe, Africa, Eastern Africa, .. Robert G. Mugabe, 4068, ZW

VS

Country

ABW	Aruba, North America, Caribbean, .. Beatrix, 129, AW
AFG	Afghanistan, Asia, Southern and Central Asia, .. 1, AF
AGO	Angola, Africa, Central Africa, .. 55, AO
ALA	Anguilla, North America, Caribbean, .. Elisabeth II, 62, AI
ALB	Albania, Europe, Southern Europe, .. Rexhep Mejdani, 34, AL
AND	Andorra, Europe, Southern Europe, .. 55, AD
..	..
CHN	China, Asia, Eastern Asia, 9572900.00, .. Jiang Zemin, 1391, CN
..	..
IND	India, Asia, Southern and Central Asia, .. 1109, IN
..	..
IO	British Indian Ocean Territory, Africa, .. Elisabeth II, NULL, IO
IRL	Ireland, Europe, British Islands, .. Mary McAldese, 1447, IE
IRN	Iran, Asia, Southern and Central Asia, .. 1380, IR
..	..
ZAF	South Africa, Africa, Southern Africa, .. Thabo Mbeki, 716, ZA
ZMB	Zambia, Africa, Eastern Africa, .. Frederick Chiluba, 3162, ZM
ZWE	Zimbabwe, Africa, Eastern Africa, .. Robert G. Mugabe, 4068, ZW

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. **Optimizer Trace**
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets


# Optimizer Trace

- What other choices did EXPLAIN not show?
- Why was that choice made?
- Output is quite verbose

```
ALTER TABLE Country ADD INDEX p (population);
EXPLAIN FORMAT=JSON
SELECT * FROM Country WHERE continent='Asia' and population > 5000000;
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "25.40"
    },
    "table": {
      "table_name": "Country",
      "access_type": "ALL",
      "possible_keys": [
        "p"
      ],
      "rows_examined_per_scan": 239,
      "rows_produced_per_join": 15,
      "filtered": "6.46",
      "cost_info": {
        "read_cost": "23.86",
        "eval_cost": "1.54",
        "prefix_cost": "25.40",
        "data_read_per_join": "3K"
      },
      ..
      "attached_condition": "((`world`.`country`.`Continent` = 'Asia') and
(`world`.`country`.`Population` > 5000000))"
      ..
    }
  }
}
```

It's available but not used. Why?



```

SET optimizer_trace="enabled=on";
SELECT * FROM Country WHERE continent='Asia' and population > 5000000;
SELECT * FROM information_schema.optimizer_trace;
{
  "steps": [
    {
      "join_preparation": {
        "select#": 1,
        "steps": [
          {
            "expanded_query": "/* select#1 */ select `country`.`Code` AS `Code`,`country`.`Name` AS `Name`,`country`.`Continent` AS `Continent`,`country`.`Region` AS `Region`,`country`.`SurfaceArea` AS `SurfaceArea`,`country`.`IndepYear` AS `IndepYear`,`country`.`Population` AS `Population`,`country`.`LifeExpectancy` AS `LifeExpectancy`,`country`.`GNP` AS `GNP`,`country`.`GNPOld` AS `GNPOld`,`country`.`LocalName` AS `LocalName`,`country`.`GovernmentForm` AS `GovernmentForm`,`country`.`HeadOfState` AS `HeadOfState`,`country`.`Capital` AS `Capital`,`country`.`Code2` AS `Code2` from `country` where ((`country`.`Continent` = 'Asia') and (`country`.`Population` > 5000000))"
          }
        ]
      }
    },
    {
      "join_optimization": {
        "select#": 1,
        "steps": [
          {
            "condition_processing": {
              "condition": "WHERE",

```



```

    "original_condition": "((`country`.`Continent` = 'Asia') and (`country`.`Population` > 5000000))",
    "steps": [
      {
        "transformation": "equality_propagation",
        "resulting_condition": "((`country`.`Population` > 5000000) and multiple equal('Asia',
`country`.`Continent`))"
      },
      {
        "transformation": "constant_propagation",
        "resulting_condition": "((`country`.`Population` > 5000000) and multiple equal('Asia',
`country`.`Continent`))"
      },
      {
        "transformation": "trivial_condition_removal",
        "resulting_condition": "((`country`.`Population` > 5000000) and multiple equal('Asia',
`country`.`Continent`))"
      }
    ]
  },
  {
    "substitute_generated_columns": {
    },
  },
  {
    "table_dependencies": [
      {
        "table": "`country`",
        "row_may_be_null": false,
        "map_bit": 0,

```

```

        "depends_on_map_bits": [
        ]
    }
]
},
{
    "ref_optimizer_key_uses": [
    ]
},
{
    "rows_estimation": [
    {
        "table": "`country`",
        "range_analysis": {
            "table_scan": {
                "rows": 239,
                "cost": 27.5
            },
            "potential_range_indexes": [
                {
                    "index": "PRIMARY",
                    "usable": false,
                    "cause": "not_applicable"
                },
                {
                    "index": "p",
                    "usable": true,
                    "key_parts": [
                        "Population",
                        "Code"
                    ]
                }
            ]
        }
    ]
}

```

```

    }
  ],
  "setup_range_conditions": [
  ],
  "group_index_range": {
    "chosen": false,
    "cause": "not_group_by_or_distinct"
  },
  "analyzing_range_alternatives": {
    "range_scan_alternatives": [
      {
        "index": "p",
        "ranges": [
          "5000000 < Population"
        ],
        "index_dives_for_eq_ranges": true,
        "rowid_ordered": false,
        "using_mrr": false,
        "index_only": false,
        "rows": 108,
        "cost": 38.06,
        "chosen": false,
        "cause": "cost"
      }
    ],
    "analyzing_roworder_intersect": {
      "usable": false,
      "cause": "too_few_roworder_scans"
    }
  }
}

```

Aha! It was too expensive.

```

    }
  ]
},
{
  "considered_execution_plans": [
    {
      "plan_prefix": [
      ],
      "table": "`country`",
      "best_access_path": {
        "considered_access_paths": [
          {
            "rows_to_scan": 239,
            "access_type": "scan",
            "resulting_rows": 239,
            "cost": 25.4,
            "chosen": true
          }
        ]
      },
      "condition_filtering_pct": 100,
      "rows_for_plan": 239,
      "cost_for_plan": 25.4,
      "chosen": true
    }
  ],
},
{
  "attaching_conditions_to_tables": {
    "original_condition": "((`country`.`Continent` =
'Asia') and (`country`.`Population` > 5000000))",

```

Prefer to table scan  
instead

```

    "attached_conditions_computation": [
    ],
    "attached_conditions_summary": [
      {
        "table": "`country`",
        "attached": "((`country`.`Continent` = 'Asia') and (`country`.`Population` > 5000000))"
      }
    ]
  },
  {
    "refine_plan": [
      {
        "table": "`country`"
      }
    ]
  }
],
},
{
  "join_execution": {
    "select#": 1,
    "steps": [
    ]
  }
}
]
}

```

# Why would an index not be used?

## OPTIMIZER TRACE:

```
"analyzing_range_alternatives": {  
  "range_scan_alternatives": [  
    {  
      "index": "p",  
      "ranges": [  
        "5000000 < Population"  
      ],  
      "index_dives_for_eq_ranges": true,  
      "rowid_ordered": false,  
      "using_mrr": false,  
      "index_only": false,  
      "rows": 108,  
      "cost": 38.06,  
      "chosen": false,  
      "cause": "cost"  
    }  
  ],  
}
```

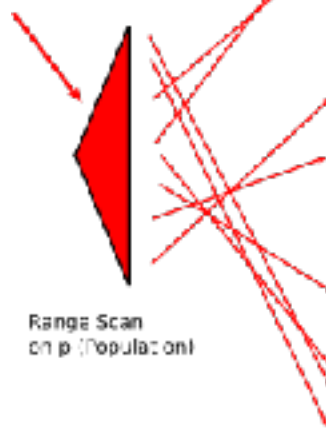
## FORCE INDEX (p):

```
..  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "48.86"  
    },  
    "table": {  
      "table_name": "Country",  
      "access_type": "range",  
      "possible_keys": [  
        "p"  
      ],  
      "key": "p",  
    },  
  },  
..
```

# Reason again...

```
SELECT * FROM Country
WHERE continent='Asia' AND population > 5000000;
```

Country



Range Scan on p (Population)

ABW	Aruba, North America, Caribbean, .. Beatrix, 129, AW
AFG	Afghanistan, Asia, Southern and Central Asia, .. 1, AF
AGO	Angola, Africa, Central Africa, .. 55, AO
AIA	Anguilla, North America, Caribbean, .. Elisabeth II, 62, AI
ALB	Albania, Europe, Southern Europe, .. Rexhep Mejdani, 34, AL
AND	Andorra, Europe, Southern Europe, .. 55, AD
..	..
CHN	China, Asia, Eastern Asia, 9572900.00, .. Jiang Zemin, 1891, CN
..	..
IND	India, Asia, Southern and Central Asia, .. 1109, IN
..	..
IO	British Indian Ocean Territory, Africa, .. Elisabeth II, NULL, IO
IRL	Ireland, Europe, British Islands, .. Mary McAleese, 1447, IE
IRN	Iran, Asia, Southern and Central Asia, .. 1380, IR
..	..
ZAF	South Africa, Africa, Southern Africa, .. Thabo Mbeki, 716, ZA
ZMB	Zambia, Africa, Eastern Africa, .. Frederick Chiluba, 3162, ZM
ZWE	Zimbabwe, Africa, Eastern Africa, .. Robert G. Mugabe, 4068, ZW

VS

Country

ABW	Aruba, North America, Caribbean, .. Beatrix, 129, AW
AFG	Afghanistan, Asia, Southern and Central Asia, .. 1, AF
AGO	Angola, Africa, Central Africa, .. 55, AO
AIA	Anguilla, North America, Caribbean, .. Elisabeth II, 62, AI
ALB	Albania, Europe, Southern Europe, .. Rexhep Mejdani, 34, AL
AND	Andorra, Europe, Southern Europe, .. 55, AD
..	..
CHN	China, Asia, Eastern Asia, 9572900.00, .. Jiang Zemin, 1891, CN
..	..
IND	India, Asia, Southern and Central Asia, .. 1109, IN
..	..
IO	British Indian Ocean Territory, Africa, .. Elisabeth II, NULL, IO
IRL	Ireland, Europe, British Islands, .. Mary McAleese, 1447, IE
IRN	Iran, Asia, Southern and Central Asia, .. 1380, IR
..	..
ZAF	South Africa, Africa, Southern Africa, .. Thabo Mbeki, 716, ZA
ZMB	Zambia, Africa, Eastern Africa, .. Frederick Chiluba, 3162, ZM
ZWE	Zimbabwe, Africa, Eastern Africa, .. Robert G. Mugabe, 4068, ZW

# Agenda

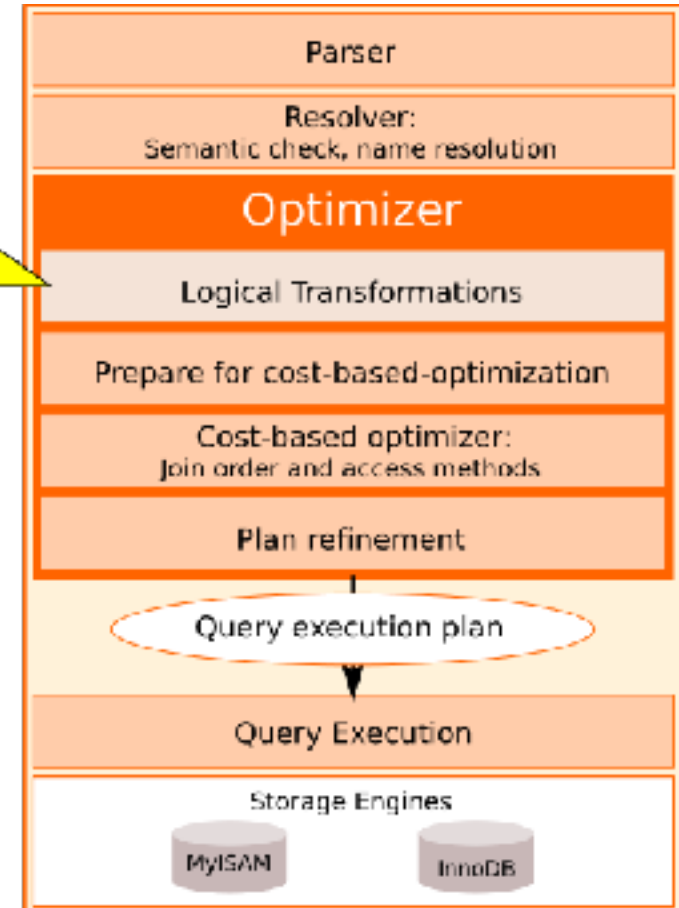
1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. **Logical Transformations**
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets



# Logical Transformations

- First part of optimization is eliminating unnecessary work

Negation elimination  
Equality and constant propagation  
Evaluation of constant expressions  
Conversions of outer to inner join  
Subquery transformation



# Why eliminate unnecessary work?

- Short-cut/reduce number of execution plans that need to be evaluated
- Transform parts of queries to take advantage of better execution strategies
- Think of a how a compiler transforms code to be more efficient
  - MySQL does similar at runtime

## Example:

```
SELECT * FROM Country  
WHERE population > 5000000 AND continent='Asia'  
AND 1=1;
```

# SHOW WARNINGS says:

```
EXPLAIN FORMAT=JSON SELECT * FROM Country WHERE population > 5000000 AND 1=1;
```

**SHOW WARNINGS;**

```
/* select#1 */ select
`world`.`Country`.`Code` AS `Code`,
`world`.`Country`.`Name` AS `Name`,
`world`.`Country`.`Continent` AS `Continent`,
`world`.`Country`.`Region` AS `Region`,
`world`.`Country`.`SurfaceArea` AS `SurfaceArea`,
`world`.`Country`.`IndepYear` AS `IndepYear`,
`world`.`Country`.`Population` AS `Population`,
`world`.`Country`.`LifeExpectancy` AS `LifeExpectancy`,
`world`.`Country`.`GNP` AS `GNP`,
`world`.`Country`.`GNPOld` AS `GNPOld`,
`world`.`Country`.`LocalName` AS `LocalName`,
`world`.`Country`.`GovernmentForm` AS `GovernmentForm`,
`world`.`Country`.`HeadOfState` AS `HeadOfState`,
`world`.`Country`.`Capital` AS `Capital`,
`world`.`Country`.`Code2` AS `Code2`
from `world`.`Country`
where (`world`.`Country`.`Population` > 5000000)
```

# OPTIMIZER TRACE says:

..

```
"steps": [  
  {  
    "condition_processing": {  
      "condition": "WHERE",  
      "original_condition": "((`Country`.`Population` > 5000000) and (1 = 1))",  
      "steps": [  
        {  
          "transformation": "equality_propagation",  
          "resulting_condition": "((`Country`.`Population` > 5000000) and (1 = 1))"  
        },  
        {  
          "transformation": "constant_propagation",  
          "resulting_condition": "((`Country`.`Population` > 5000000) and (1 = 1))"  
        },  
        {  
          "transformation": "trivial_condition_removal",  
          "resulting_condition": "(`Country`.`Population` > 5000000)"  
        }  
      ]  
    }  
  ]  
..
```

# What sort of transformations can occur?

- Merging views back with definition of base tables
- Derived table in FROM clause merged back into base tables
- Unique subqueries converted directly to INNER JOIN statements
- Primary key lookup converted to constant values.
  - Shortcut plans that will need to be evaluated.

# Primary Key Lookup

```
SELECT * FROM Country WHERE code='CAN'  
/* select#1 */ select  
'CAN' AS `Code`,  
'Canada' AS `Name`,  
'North America' AS `Continent`,  
'North America' AS `Region`,  
'9970610.00' AS `SurfaceArea`,  
'1867' AS `IndepYear`,  
'31147000' AS `Population`,  
'79.4' AS `LifeExpectancy`,  
'598862.00' AS `GNP`,  
'625626.00' AS `GNPOld`,  
'Canada' AS `LocalName`,  
'Constitutional Monarchy, Federation' AS `GovernmentForm`,  
'Elisabeth II' AS `HeadOfState`,  
'1822' AS `Capital`,  
'CA' AS `Code2`  
from `world`.`Country` where 1
```

# Primary key does not exist

```
SELECT * FROM Country WHERE code='XYZ'
```

```
/* select#1 */ select NULL AS `Code`,NULL AS  
`Name`,NULL AS `Continent`,NULL AS `Region`, NULL AS  
`SurfaceArea`,NULL AS `IndepYear`,NULL AS  
`Population`,NULL AS `LifeExpectancy`,NULL AS `GNP`,  
NULL AS `GNPOld`,NULL AS `LocalName`,NULL AS  
`GovernmentForm`,NULL AS `HeadOfState`,NULL AS  
`Capital`, NULL AS `Code2` from `world`.`Country`  
where multiple equal('XYZ', NULL)
```



# Impossible WHERE

```
SELECT * FROM Country WHERE code='CAN' AND 1=0

/* select#1 */ select `world`.`Country`.`Code` AS
`Code`,`world`.`Country`.`Name` AS `Name`,
`world`.`Country`.`Continent` AS `Continent`,`world`.`Country`.`Region`
AS `Region`,`world`.`Country`.`SurfaceArea` AS
`SurfaceArea`,`world`.`Country`.`IndepYear` AS `IndepYear`,
`world`.`Country`.`Population` AS
`Population`,`world`.`Country`.`LifeExpectancy` AS `LifeExpectancy`,
`world`.`Country`.`GNP` AS `GNP`,`world`.`Country`.`GNPOld` AS
`GNPOld`,`world`.`Country`.`LocalName` AS
`LocalName`,`world`.`Country`.`GovernmentForm` AS `GovernmentForm`,
`world`.`Country`.`HeadOfState` AS
`HeadOfState`,`world`.`Country`.`Capital` AS `Capital`,
`world`.`Country`.`Code2` AS `Code2` from `world`.`Country` where 0
```

# Are transformations always safe?

- Yes they should be
- New transformations (and execution strategies) may return *non deterministic* queries in a different order
- Some illegal statements as a result of derived\_merge transformation

MySQL 5.7.6-m16 default opt: x

Morgan

Securehttps://github.com/rails/rails/issues/19281#issuecomment-78260683

rails / rails

Watch2,416Star85,262Fork14,392


CodeIssues460Pull requests714Projects0PulseGraphs

# MySQL 5.7.6-m16 default optimizer\_switch derived\_merge=on causes Error: You can't specify target table 'pets' for update in FROM clause: #19281

New Issue

Closed

yahonda opened this issue on Mar 10, 2015 · 8 comments



yahonda commented on Mar 10, 2015

Contributor

MySQL 5.7.6-m16 has been released.  
Since upgrading my environments to this version, some of mysql and mysql2 test cases get errors.

These errors can be resolved by setting `optimizer_switch=derived_merge=off` in `/etc/my.cnf`.  
Then I'm not sure these errors can be addressed in Rails yet.

- Testcases

```
for i in mysql mysql2
do
  echo $i
  ActiveRecord ruby -Itest test/cases/scoping/relation_scoping_test.rb -n test_delete_all_default
  ActiveRecord ruby -Itest test/cases/persistence_test.rb -n /test_delete_all_with_join_and_where
done
```

- Environments

```
Server version: 5.7.6-m16 MySQL Community Server (GPL)
$ ruby -v
ruby 2.2.1p85 (2015-02-26 revision 48766) [x86_64-linux]
```

Assignees

No one assigned

Labels

activerecord

Projects

None yet

Milestone

No milestone

Notifications

Unsubscribe

You're receiving notifications because you were mentioned.

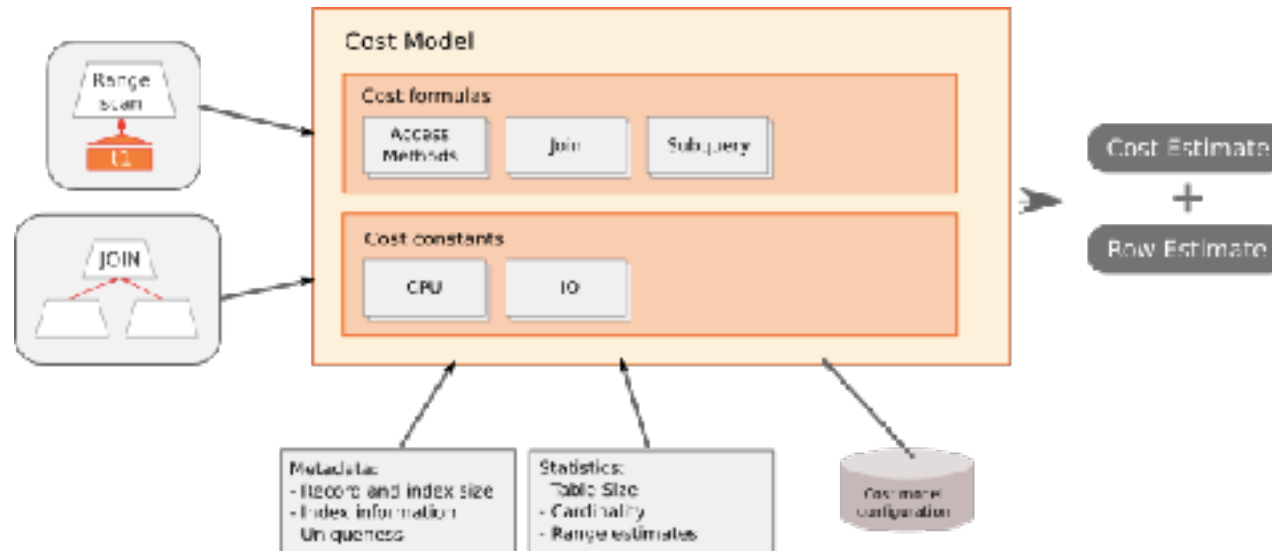
3 participants

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Query Optimizer Strategy

- Model each of the possible execution plans (using support from statistics and meta data)
- Pick the plan with the lowest cost



# Model you say?

1. Assign a cost to each operation
2. Evaluate how many operations each possible plan would take
3. Sum up the total
4. Choose the plan with the lowest overall cost

# How are statistics calculated?

- Dictionary Information
- Cardinality Statistics
- Records In Range Dynamic Sampling
- Table Size

# Example Model: Table Scan

*SELECT \* FROM Country WHERE continent='Asia' and population > 5000000;*

*IO Cost:*

# pages in table \* (IO\_BLOCK\_READ\_COST | MEMORY\_BLOCK\_READ\_COST)

*CPU Cost:*

# records \* ROW\_EVALUATE\_COST

*Defaults:*

IO\_BLOCK\_READ\_COST = 1

MEMORY\_BLOCK\_READ\_COST = 0.25

ROW\_EVALUATE\_COST=0.1

*Values:*

# pages in table = 6

# records = 239

100% on Disk:

$= (6 * 1) + (0.1 * 239)$

**= 29.9**

100% in Memory:

$= (6 * 0.25) + (0.1 * 239)$

**= 25.4**

*SELECT clust\_index\_size from  
INNODB\_SYS\_TABLESTATS WHERE  
name='world/country'*

**New!** MySQL 8.0 estimates how many of the pages will be in memory.

EXPLAIN said  
cost was 25.40



# Example Model: Range Scan

*SELECT \* FROM Country WHERE continent='Asia' and **population > 5000000;***

*IO Cost:*

# records\_in\_range \* (IO\_BLOCK\_READ\_COST | MEMORY\_BLOCK\_READ\_COST)

*CPU Cost:*

# records\_in\_range \* ROW\_EVALUATE\_COST  
+ # records\_in\_range \* ROW\_EVALUATE\_COST

Evaluate range condition

Evaluate WHERE condition

= (108 \* 0.25) + ( (108 \* 0.1) + (108 \* 0.1) )  
**= 48.6**

Compares to "query\_cost":  
"48.86" in EXPLAIN.

100% in memory.  
On disk = 129.6

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "25.40"
    },
    "table": {
      "table_name": "country",
      "access_type": "ALL",
      "possible_keys": [
        "p"
      ],
      ..
    },
    "cost_info": {
      "read_cost": "23.86",
      "eval_cost": "1.54",
      "prefix_cost": "25.40",
      "data_read_per_join": "3K"
    },
    ..
  },
  ..
}
```

CPU Cost

IO Cost

Total Cost

# Cost Constant Refinement

```
select * from mysql.server_cost;
```

cost_name	cost_value	last_update	comment	default_value
disk_temptable_create_cost	NULL	2017-04-14 16:01:42	NULL	20
disk_temptable_row_cost	NULL	2017-04-14 16:01:42	NULL	0.5
key_compare_cost	NULL	2017-04-14 16:01:42	NULL	0.05
memory_temptable_create_cost	NULL	2017-04-14 16:01:42	NULL	1
memory_temptable_row_cost	NULL	2017-04-14 16:01:42	NULL	0.1
row_evaluate_cost	NULL	2017-04-14 16:01:42	NULL	0.1

```
6 rows in set (0.00 sec)
```

```
select * from mysql.engine_cost\G
```

```
***** 1. row *****
```

```
engine_name: default
```

```
device_type: 0
```

```
cost_name: io_block_read_cost
```

```
cost_value: NULL
```

```
last_update: 2017-04-14 16:01:42
```

```
comment: NULL
```

```
default_value: 1
```

# Cost Constant Refinement

```
UPDATE mysql.server_cost SET cost_value=1 WHERE cost_name='row_evaluate_cost';  
UPDATE mysql.engine_cost set cost_value = 1;  
FLUSH OPTIMIZER_COSTS;  
EXPLAIN FORMAT=JSON SELECT * FROM Country WHERE continent='Asia' and population > 5000000;
```

```
{  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "245.00"  
    },  
    "table": {  
      "table_name": "Country",  
      "access_type": "ALL",  
      ..  
    }  
  }  
}
```

Increase row evaluate cost from 0.1 to 1. Make memory and IO block read cost the same.

New Table Scan Cost:  
$$= (6 * 1) + (1 * 239)$$
$$= 245$$

# Are plans exhaustively evaluated?

- Short cuts are taken to not spend too much time in planning:
  - Some parts of queries may be transformed to limit plans evaluated
  - The optimizer will by default limit the search depth of bad plans:  
`optimizer_search_depth=64`  
`optimizer_prune_level=1`

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
- 8. Hints and Switches**
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# How often is the query optimizer wrong?

- Yes it happens
- Similar to GPS; you may not have traffic data available for all streets
- The model may be incomplete or imperfect
- There exist method(s) to overwrite it

# Hints and Switches

- Typically a better level of override to modifying cost constants
- Come in three varieties:
  - Old Style Hints
  - New Comment-Style Hints
  - Switches



# Old Style Hints

- Have SQL and Hint intermingled
- Cause errors when indexes don't exist

```
SELECT * FROM Country FORCE INDEX (p) WHERE population > 5000000;
```

```
SELECT * FROM Country IGNORE INDEX (p) WHERE population > 5000000;
```

```
SELECT * FROM Country USE INDEX (p) WHERE population > 5000000;
```

```
SELECT STRAIGHT_JOIN ..;
```

```
SELECT * FROM Country STRAIGHT_JOIN ..;
```

# New Comment-Style Hints

- Can be added by a system that doesn't understand SQL
- Clearer defined semantics as a hint not a directive
- Fine granularity

```
SELECT
```

```
/*+ NO_RANGE_OPTIMIZATION (Country) */
```

```
* FROM Country
```

```
WHERE Population > 1000000000 AND Continent='Asia';
```

# Switches

- As new optimizations are added, some cause regressions
- Allow the specific optimization to be disabled (SESSION or GLOBAL)

```
SELECT @@optimizer_switch;
```

```
index_merge=on,index_merge_union=on,index_merge_sort_union=on,  
index_merge_intersection=on,engine_condition_pushdown=on,  
index_condition_pushdown=on,mrr=on,mrr_cost_based=on,  
block_nested_loop=on,batched_key_access=off,materialization=on,  
semijoin=on,loosescan=on,firstmatch=on,duplicateweedout=on,  
subquery_materialization_cost_based=on,use_index_extensions=on,  
condition_fanout_filter=on,derived_merge=on
```

# How to consider hints and switches

- They provide immediate pain relief to production problems at the cost of maintenance
- They add technical debt to your applications

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
- 9. Comparing Plans**
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Our simple query with $n$ candidate indexes

- Indexes exist on `p(population)` and `c(continent)`:

```
SELECT * FROM Country  
WHERE population > 50000000 AND continent='Asia';
```

>50M, how many  
are less?

How many countries in  
Asia vs total world?

Does order of  
predicates matter? **No.**


# Role of the Optimizer

- Given these many choices, which is the best choice?
- A good GPS navigator finds the fastest route!
- We can expect a good query optimizer to do similar

```
ALTER TABLE Country ADD INDEX c (continent);
EXPLAIN FORMAT=JSON # 50M
SELECT * FROM Country WHERE population > 50000000 AND continent='Asia';
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "9.60"
    },
    "table": {
      "table_name": "Country",
      "access_type": "ref",
      "possible_keys": [
        "p",
        "c"
      ],
      "key": "c",
      "used_key_parts": [
        "Continent"
      ],
      "key_length": "1",
      "ref": [
        "const"
      ],
      ..
      "attached_condition": "(`world`.`country`.`Population` > 50000000)"
      ..
    }
  }
}
```

Continent is determined to be lower cost.





```
EXPLAIN FORMAT=JSON # 500M
```

```
SELECT * FROM Country WHERE continent='Asia' and population > 500000000;
```

```
{
```

```
  "query_block": {
```

```
    "select_id": 1,
```

```
    "cost_info": {
```

```
      "query_cost": "1.16"
```

```
    },
```

```
    "table": {
```

```
      "table_name": "Country",
```

```
      "access_type": "range",
```

```
      "possible_keys": [
```

```
        "p",
```

```
        "c"
```

```
      ],
```

```
      "key": "p",
```

```
      "used_key_parts": [
```

```
        "Population"
```

```
      ],
```


```
      "key_length": "4",
```

```
..
```

```
  "attached_condition": "(`world`.`country`.`Continent` = 'Asia')"
```

```
..
```

Change the predicate,  
the query plan changes.



# Query Plan Evaluation

- Evaluated for each query, and thus each set of predicates
- Currently not cached\*
- For prepared statements, permanent transformations are cached

\* Cardinality statistics are cached. Don't get confused.

# Cost Estimates

	p>5M c='Asia'	p>50M, c='Asia'	p>500M, c='Asia'
p	48.86	11.06	1.16
c	9.60	9.60	9.60
ALL	25.40	25.40	25.40

$p$  

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
- 10. Composite Indexes**
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# The role of composite indexes

- Useful when two or more predicates combined improves filtering effect. i.e.

Not all countries with a population > 5M are in Asia


# Composite Indexes

- p\_c (population, continent)
- c\_p (continent, population)

```
ALTER TABLE Country ADD INDEX p_c (Population, Continent);  
EXPLAIN FORMAT=JSON  
SELECT * FROM Country FORCE INDEX (p_c) WHERE continent='Asia' and population >  
5000000;
```

```
{  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "48.86"  
    },  
    "table": {  
      "table_name": "Country",  
      "access_type": "range",  
      "possible_keys": [  
        "p_c"  
      ],  
      "key": "p_c",  
      "used_key_parts": [  
        "Population"  
      ],  
      "key_length": "4",  
      ..
```

Only part of the key is  
used!



# Rule of Thumb

- Index on (const, range) instead of (range, const)
- Applies to all databases



```

ALTER TABLE Country ADD INDEX c_p (Continent, Population);
EXPLAIN FORMAT=JSON
SELECT * FROM Country WHERE continent='Asia' and population > 5000000;
{


```

```

  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "7.91"
    },
    "table": {
      "table_name": "Country",
      "access_type": "range",
      "possible_keys": [
        "p",
        "c",
        "p_c",
        "c_p"
      ],
      "key": "c_p",
      "used_key_parts": [
        "Continent",
        "Population"
      ],
      "key_length": "5",

```

All of the key is used



# Composite Left-most Rule

- An index on (Continent, Population) can also be used as an index on (Continent)
- It can not be used as an index on (Population)

```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM Country FORCE INDEX (c_p) WHERE population >  
500000000;
```

```
{
```

```
  "query_block": {
```

```
    "select_id": 1,
```

```
    "cost_info": {
```

```
      "query_cost": "83.90"
```

```
    },
```

```
    "table": {
```

```
      "table_name": "Country",
```

```
      "access_type": "ALL",
```

```
      "rows_examined_per_scan": 239,
```

```
      "rows_produced_per_join": 79,
```

```
      "filtered": "33.33",
```

```
..
```

```
      "attached_condition": "(`world`.`country`.`Population` >  
5000000000)"
```

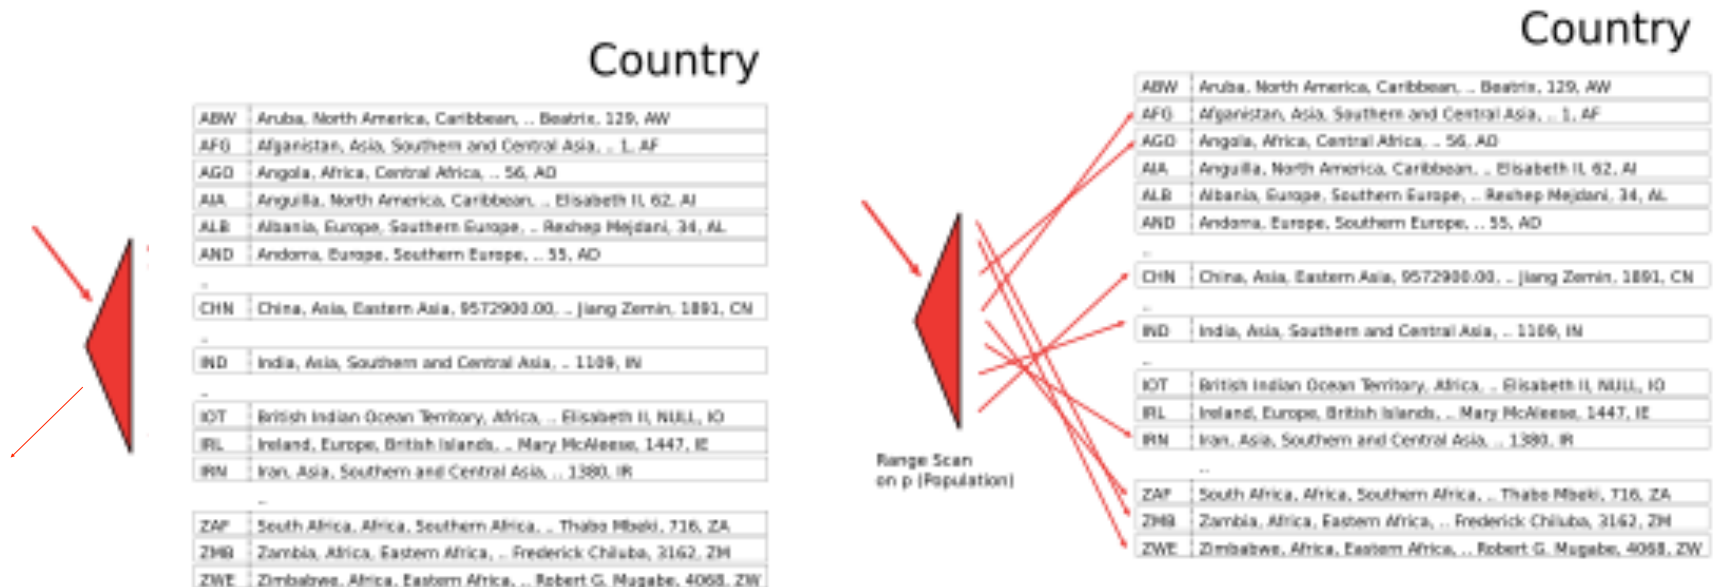
```
..
```

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
- 11. Covering Indexes**
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Covering Indexes

- A special kind of composite index
- All information returned just by accessing the index




```
ALTER TABLE Country ADD INDEX c_p_n (Continent,Population,Name);
EXPLAIN FORMAT=JSON
SELECT Name FROM Country WHERE continent='Asia' and population > 5000000;
{
```

```
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "3.72"
    },
    "table": {
      "table_name": "Country",
      "access_type": "range",
      "possible_keys": [
        ..
        "c_p_n"
      ],
      "key": "c_p_n",
      "used_key_parts": [
        "Continent",
        "Population"
      ],
      "key_length": "5",
      ..
      "filtered": "100.00",
      "using_index": true,
      ..
    }
  }
```

Cost is reduced by  
53%

Using index means  
"covering index"



# Use cases

- Can be used as in this example
- Also beneficial in join conditions (join through covering index on intermediate table)
- Useful in aggregate queries

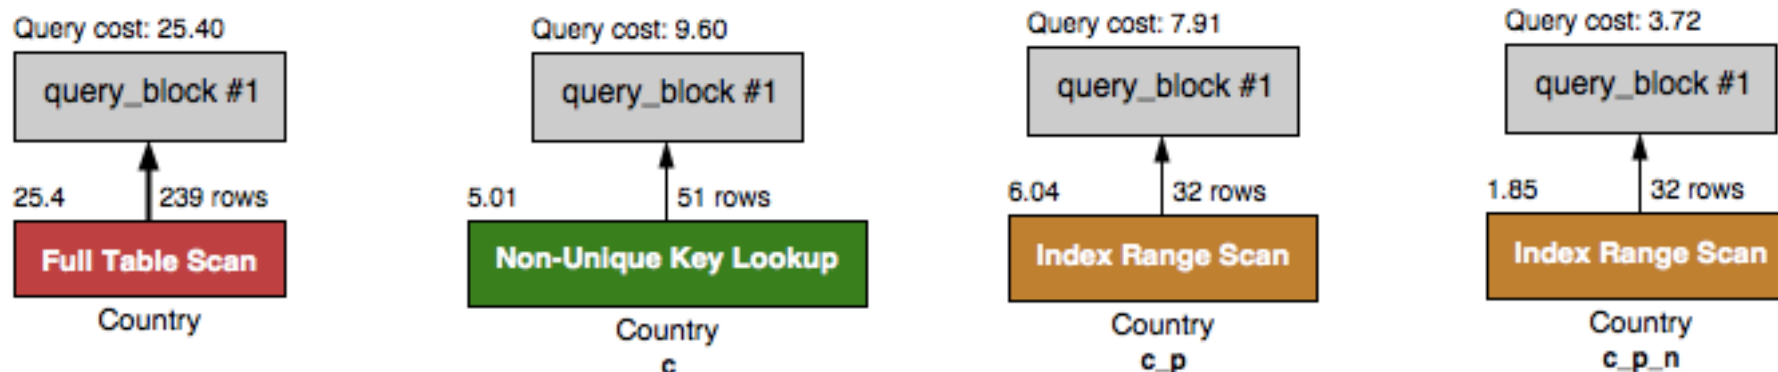
# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
- 12. Visual Explain**
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets



# Visual Explain

- For complex queries, it is useful to see visual representation
- Visualizations in this deck are produced by MySQL Workbench.



# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
- 13. Transient Plans**
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets


## A quick recap:

- So far we've talked about 4 candidate indexes:
  - p (population)
  - c (continent)
  - p\_c (population, continent)
  - c\_p (continent, population)
- We've always used c='Asia' and p > 5M

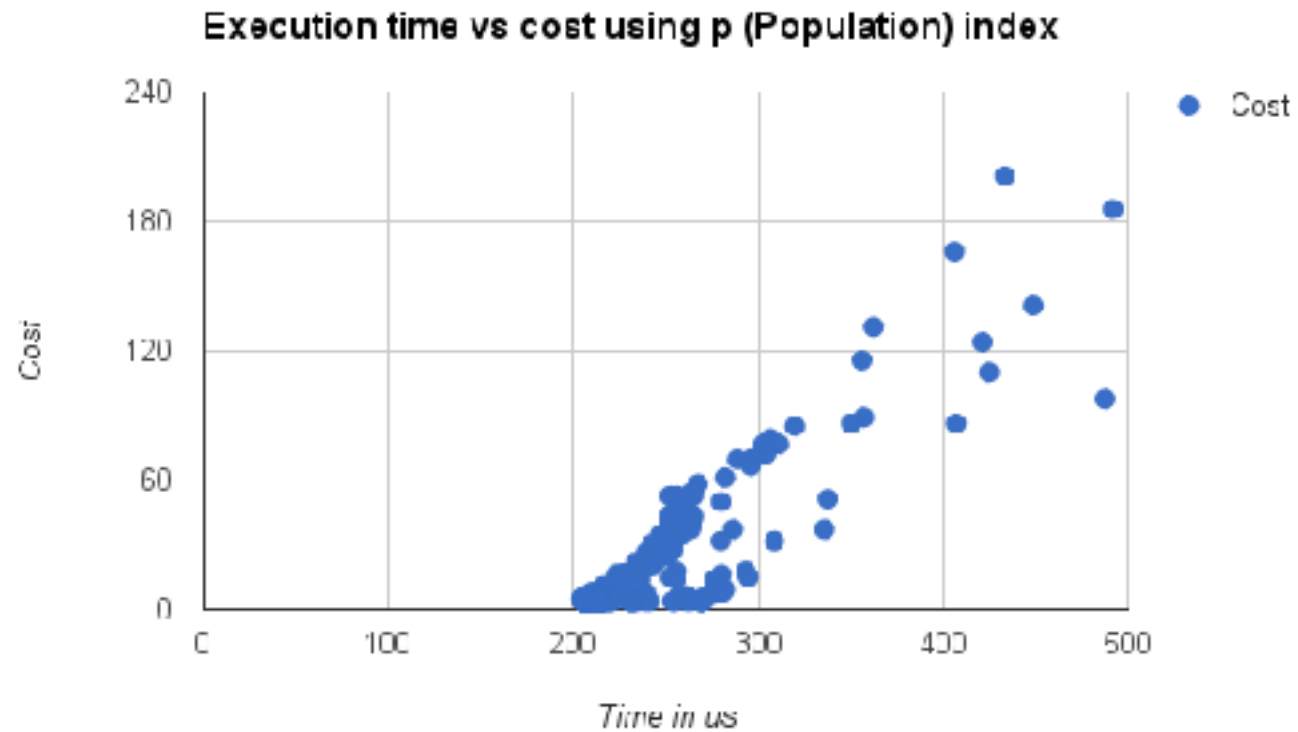
# Cost Estimates

	p>5M c='Asia'	p>5M c='Antarctica'	p>50M, c='Asia'	p>50M c='Antarctica'	p>500M, c='Asia'	p>500M c='Antarctica'
p	48.86	48.86	11.06	11.06	1.16	1.16
c	9.60	1.75	9.60	1.75	9.60	1.75
c_p	7.91	0.71	5.21	0.71	1.16	0.71
p_c	48.86	48.86	11.06	11.06	1.16	1.16
ALL	25.40	25.40	25.40	25.40	25.40	25.40

# Cost Estimates



# Actual Execution Time



# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
- 14. Subqueries**
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Subquery (Scalar)

- Can optimize away the inner part first and then cache it.
- This avoids re-executing the inner part for-each-row

```
SELECT * FROM Country WHERE  
Code = (SELECT CountryCode FROM City WHERE  
name='Toronto' );
```



EXPLAIN FORMAT=JSON

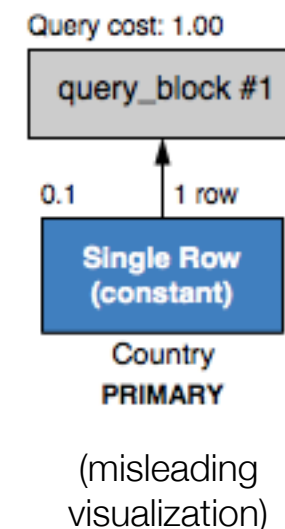
```
SELECT * FROM Country WHERE Code = (SELECT CountryCode FROM City WHERE name='Toronto');
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "1.00"
    },
    "table": {
      "table_name": "Country",
      "access_type": "const",
      ..
      "key": "PRIMARY",
      ..
    },
    "optimized_away_subqueries": [
      {
        "dependent": false,
        "cacheable": true,
        "query_block": {
          "select_id": 2,
          "cost_info": {
            "query_cost": "425.05"
          },
          "table": {
            "table_name": "City",
            "access_type": "ALL",
            ..

```

First query + its cost

Second query + its cost



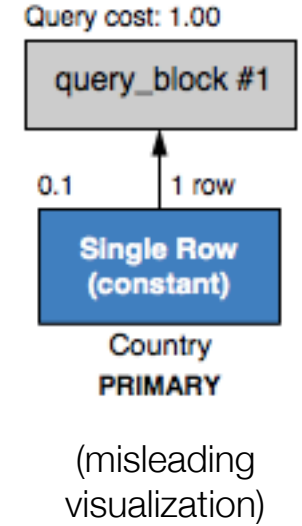
```

ALTER TABLE city ADD INDEX n (name);
EXPLAIN FORMAT=JSON
SELECT * FROM Country WHERE Code = (SELECT CountryCode FROM City WHERE name='Toronto');
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "1.00"
    },
    "table": {
      "table_name": "Country",
      "access_type": "const",
      ..
      "key": "PRIMARY",
      ..
    },
    "optimized_away_subqueries": [
      {
        "dependent": false,
        "cacheable": true,
        "query_block": {
          "select_id": 2,
          "cost_info": {
            "query_cost": "0.35"
          },
          "table": {
            "table_name": "City",
            "access_type": "ref",
            "possible_keys": [
              "n"
            ],
            "key": "n",
            ..

```

First query + its cost

Second query + its cost



# Subquery (IN list)

- When the result inner subquery returns unique results it can safely be transformed to an inner join:

```
EXPLAIN FORMAT=JSON SELECT * FROM City WHERE CountryCode IN  
(SELECT Code FROM Country WHERE Continent = 'Asia');
```

```
show warnings;
```

```
/* select#1 */ select `world`.`city`.`ID` AS `ID`,`world`.`city`.`Name` AS  
`Name`,`world`.`city`.`CountryCode` AS `CountryCode`,`world`.`city`.`District`  
AS `District`,`world`.`city`.`Population` AS `Population` from  
`world`.`country` join `world`.`city` where ((`world`.`city`.`CountryCode` =  
`world`.`country`.`Code`) and (`world`.`country`.`Continent` = 'Asia'))  
1 row in set (0.00 sec)
```

EXPLAIN FORMAT=JSON

```
SELECT * FROM City WHERE CountryCode IN  
(SELECT Code FROM Country WHERE Continent = 'Asia');
```

```
{
```

```
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "327.58"
```

```
    },
```

```
    "nested_loop": [  
      {
```

```
        {
```

```
          "table": {  
            "table_name": "Country",  
            "access_type": "ref",
```

```
..
```

```
            "key": "c",
```

```
..
```

```
            "using_index": true,
```

```
..
```

```
..
```

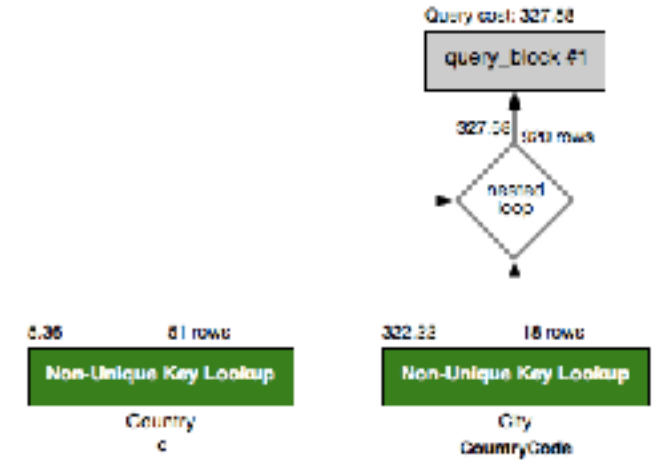
```
            "used_columns": [  
              "Code",  
              "Continent"
```

```
..
```

```
..
```

```
{
```

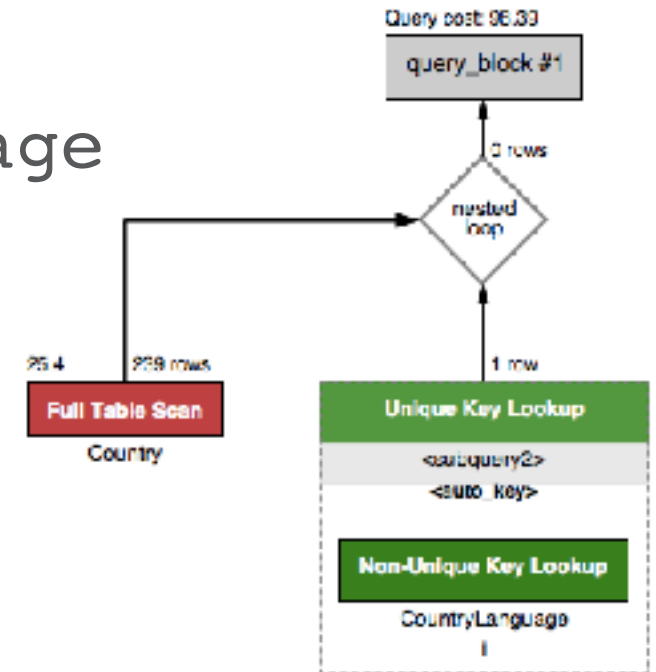
```
  "table": {  
    "table_name": "City",  
    "access_type": "ref",  
    "possible_keys": [  
      "CountryCode"  
    ],  
    "key": "CountryCode",  
  
    "ref": [  
      "world.Country.Code"  
    ],
```



## Subquery (cont.)

- When non-unique the optimizer needs to pick a semi-join strategy
- Multiple options: `FirstMatch`, `MaterializeLookup`, `DuplicatesWeedout`

```
SELECT * FROM Country WHERE Code IN
(SELECT CountryCode FROM CountryLanguage
WHERE isOfficial=1);
```



```
ALTER TABLE CountryLanguage ADD INDEX i (isOfficial);
```

```
EXPLAIN FORMAT=JSON SELECT * FROM Country WHERE Code IN  
(SELECT CountryCode FROM CountryLanguage WHERE isOfficial=1);
```

```
{
```

```
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "98.39"  
    },
```

```
    "nested_loop": [  
      {
```

```
        "table": {  
          "table_name": "Country",  
          "access_type": "ALL",  
          "possible_keys": [  
            "PRIMARY"  
          ],
```

```
        ..
```

```
    ..
```

```
      "filtered": "100.00",
```

```
      ..
```

```
    ..
```

```
    ..
```

```
  "table": {  
    "table_name": "<subquery2>",  
    "access_type": "eq_ref",  
    "key": "<auto_key>",  
    "key_length": "3",  
    "ref": [  
      "world.Country.Code"
```

```
    ],
```

```
    "rows_examined_per_scan": 1,
```

```
    "materialized_from_subquery": {
```

```
      "using_temporary_table": true,
```

```
      "query_block": {
```

```
        "table": {
```

```
          "table_name": "CountryLanguage",
```

```
          "access_type": "ref",
```

```
          "key": "i",
```

```
          "using_index": true,
```

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
- 15. CTEs and Views**
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Views

- A way of saving a SELECT statement as a table
- Allows for simplified queries
- Processed using one of two methods internally:
  - **Merge** - transform the view to be combined with the query.
  - **Materialize** - save the contents of the view in a temporary table, then begin querying




```
ALTER TABLE country ADD INDEX c_n (continent, name);
CREATE VIEW vCountry_Asia AS SELECT * FROM Country WHERE Continent='Asia';
EXPLAIN FORMAT=JSON
SELECT * FROM vCountry_Asia WHERE Name='China';
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "0.35"
    },
    "table": {
      "table_name": "country",
      "access_type": "ref",
      "possible_keys": [
        ..
        "c_n"
      ],
      "key": "c_n",
      "used_key_parts": [
        "Continent",
        "Name"
      ],
      "key_length": "53",
      "ref": [
        "const",
        "const"
      ],
      ..
    }
  }
}
```

This is the base table

Predicates from the view definition and query combined



```
SHOW WARNINGS;
/* select#1 */ select
`world`.`Country`.`Code` AS `Code`,
`world`.`Country`.`Name` AS `Name`,
`world`.`Country`.`Continent` AS `Continent`,
`world`.`Country`.`Region` AS `Region`,
`world`.`Country`.`SurfaceArea` AS `SurfaceArea`,
`world`.`Country`.`IndepYear` AS `IndepYear`,
`world`.`Country`.`Population` AS `Population`,
`world`.`Country`.`LifeExpectancy` AS `LifeExpectancy`,
`world`.`Country`.`GNP` AS `GNP`,
`world`.`Country`.`GNPOld` AS `GNPOld`,
`world`.`Country`.`LocalName` AS `LocalName`,
`world`.`Country`.`GovernmentForm` AS `GovernmentForm`,
`world`.`Country`.`HeadOfState` AS `HeadOfState`,
`world`.`Country`.`Capital` AS `Capital`,
`world`.`Country`.`Code2` AS `Code2`
from `world`.`Country`
where
((`world`.`Country`.`Continent` = 'Asia')
and (`world`.`Country`.`Name` = 'China'))
```

```
CREATE VIEW vCountrys_Per_Continent AS
SELECT Continent, COUNT(*) as Count FROM Country
GROUP BY Continent;
EXPLAIN FORMAT=JSON
SELECT * FROM vCountrys_Per_Continent WHERE Continent='Asia';
```

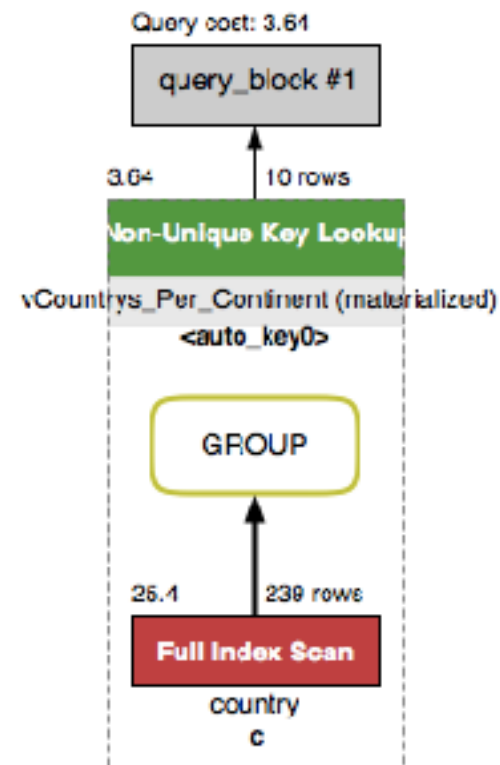
```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "3.64"
    },
  },
  "table": {
    "table_name": "vCountrys_Per_Continent",
    "access_type": "ref",
    "possible_keys": [
      "<auto_key0>"
    ],
    "key": "<auto_key0>",
    "used_key_parts": [
      "Continent"
    ],
    "key_length": "1",
    "ref": [
      "const"
    ],
    ..
    "used_columns": [
      "Continent",
      "Count"
    ],
    ..
  },
  ..
}
```

This is only the cost of  
accessing the materialized table

This is the view name

This step happens first.

```
..
"materialized_from_subquery": {
  "using_temporary_table": true,
  "dependent": false,
  "cacheable": true,
  "query_block": {
    "select_id": 2,
    "cost_info": {
      "query_cost": "25.40"
    },
  },
}
```



```
SHOW WARNINGS;  
/* select#1 */ select  
`vCountrys_Per_Continent`.`Continent` AS `Continent`,  
`vCountrys_Per_Continent`.`Count` AS `Count`  
from `world`.`vCountrys_Per_Continent`  
where (`vCountrys_Per_Continent`.`Continent` = 'Asia')
```

# WITH (CTE)

- A view for query-only duration
- Same optimizations available as views:
  - **Merge** - transform the CTE to be combined with the query.
  - **Materialize** - save the contents of the CTE in a temporary table, then begin querying

## # Identical Queries - CTE and VIEW

```
WITH vCountry_Asia AS (SELECT * FROM Country WHERE  
Continent='Asia')
```

```
SELECT * FROM vCountry_Asia WHERE Name='China';
```

```
CREATE VIEW vCountry_Asia AS SELECT * FROM Country WHERE  
Continent='Asia';
```

```
SELECT * FROM vCountry_Asia WHERE Name='China';
```

# CTEs are new!

- May provide performance enhancements over legacy code using temporary tables - which never merge.
- Derived tables may need to materialize more than once. A CTE does not! i.e.

```
SELECT * FROM my_table, (SELECT ... ) as t1 ...
```

```
UNION ALL
```

```
SELECT * FROM my_table, (SELECT ... ) as t1 ...
```

# WITH RECURSIVE - new!

```
WITH RECURSIVE my_cte AS (  
  SELECT 1 AS n  
  UNION ALL  
  SELECT 1+n FROM my_cte WHERE n<10  
)  
SELECT * FROM my_cte;
```

```
+-----+
```

```
|  n  |
```

```
+-----+
```

```
|    1 |
```

```
|    2 |
```

```
..
```

```
|    9 |
```

```
|   10 |
```

```
+-----+
```

```
10 rows in set (0.01 sec)
```



```

{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "2.84"
    },
    "table": {
      "table_name": "my_cte",
      "access_type": "ALL",
      ..
    },
    "used_columns": [
      "n"
    ],
    "materialized_from_subquery": {
      "using_temporary_table": true,
      "dependent": false,
      "cacheable": true,
      "query_block": {
        "union_result": {
          "using_temporary_table": false,
          ..
        }
      }
    },
    "dependent": false,
    "cacheable": true,
    "query_block": {
      "select_id": 3,
      "recursive": true,
      "cost_info": {
        "query_cost": "2.72"
      },
      "table": {
        "table_name": "my_cte",
        "access_type": "ALL",
        "used_columns": [
          "n"
        ],
        "attached_condition":
          "(`my_cte`.`n` < 10)"
      },
      ..
    }
  }
}

```

Cost per iteration

Requires a temporary table for intermediate results

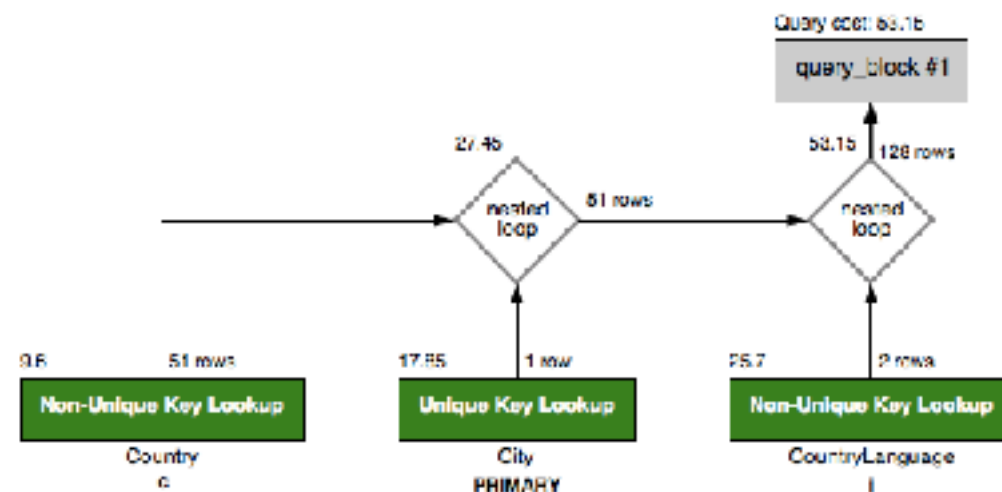
# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
- 16. Joins**
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

```
SELECT
    Country.Name as Country, City.Name as Capital,
    Language
FROM
    City
    INNER JOIN Country ON Country.Capital=City.id
    INNER JOIN CountryLanguage ON
CountryLanguage.CountryCode=Country.code
WHERE
    Country.Continent='Asia' and
CountryLanguage.IsOfficial='T';
```

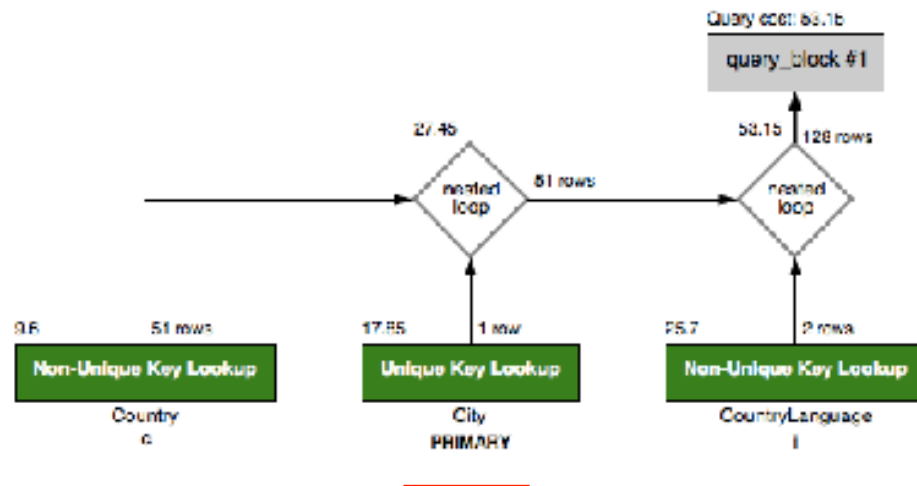
# Join Strategy (Nested Loop Join)

1. Pick Driving Table (Country)
2. For each row in Country step through to City table
3. For each row in City table step through to CountryLanguage table
4. Repeat



# Join efficiency

- Important to eliminate work before accessing other tables (WHERE clause should have lots of predicates that filter driving table)
- Indexes are required on the columns that connect between driving table, and subsequent tables:



ON Country.Capital=**City.id**

# INNER JOIN vs LEFT JOIN

- LEFT JOIN semantically says “right row is optional”.
  - Forces JOIN order to be left side first.
  - Reduces possible ways to join tables

# Join Order Hints

- One of the most frequent types of hints to apply
- New join order hints in 8.0:
  - `JOIN_FIXED_ORDER`
  - `JOIN_ORDER`
  - `JOIN_PREFIX`
  - `JOIN_SUFFIX`

# Agenda

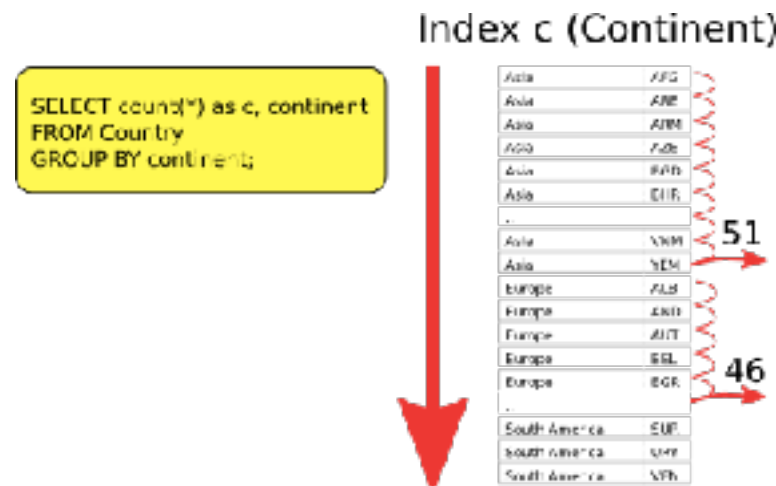
1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
- 17. Aggregation**
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets



# Group By - Loose Index Scan

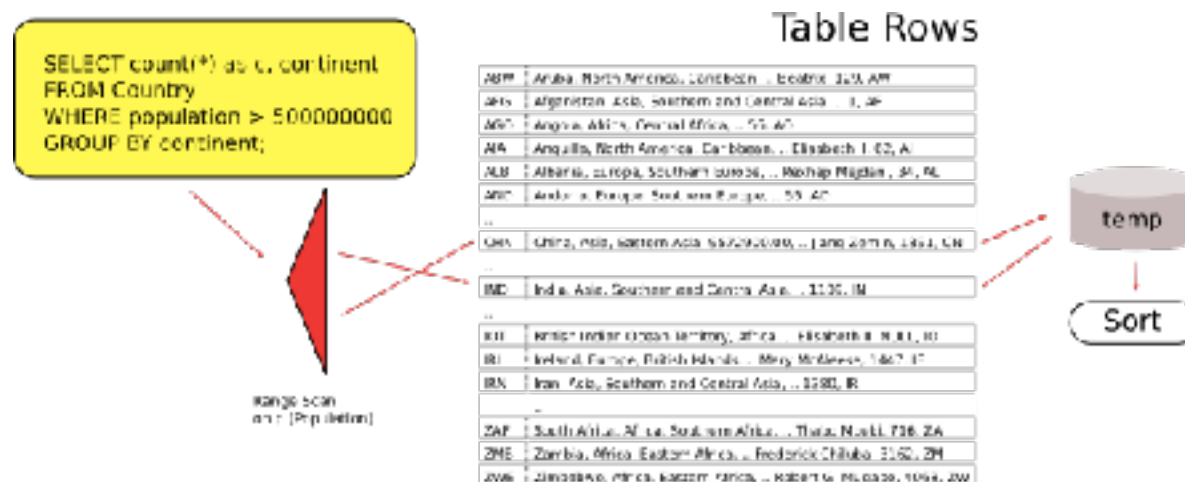
- Scan the index from start to finish without buffering. Results are pipelined to client:

```
SELECT count(*) as c, continent FROM Country  
GROUP BY continent;
```



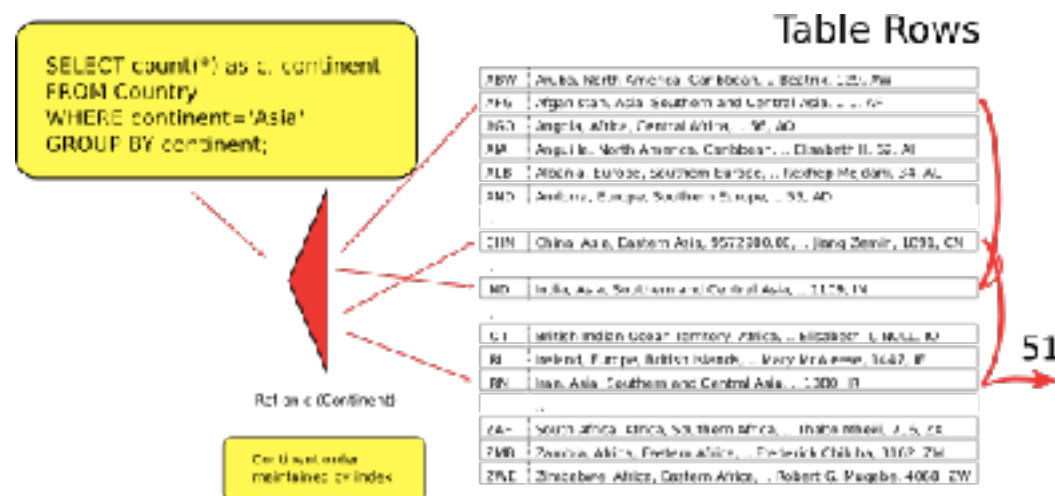
# Group By - Index Filtering Rows

- Use the index to eliminate as much work as possible
- Store rows in intermediate temporary file and then sort



# Group By - Index Filtering + Guaranteed Order

- Use the index to eliminate as much work as possible
- The index also maintains order



# UNION

- Requires an intermediate temporary table to weed out duplicate rows
- The optimizer does not really have any optimizations for UNION (such as a merge with views)

```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM City WHERE CountryCode = 'CAN'
```

```
UNION
```

```
SELECT * FROM City WHERE CountryCode = 'USA'
```

```
{
```

```
  "union_result": {
```

```
    "using_temporary_table": true,
```

```
    "table_name": "<union1,2>",
```

```
    "access_type": "ALL",
```

```
    "query_specifications": [
```

```
      {
```

```
        "dependent": false,
```

```
        "cacheable": true,
```

```
        "query_block": {
```

```
          "select_id": 1,
```

```
          "cost_info": {
```

```
            "query_cost": "17.15"
```

```
          },
```

```
          "table": {
```

```
            "table_name": "City",
```

```
            "access_type": "ref",
```

```
..
```

```
            "key": "CountryCode",
```

```
..
```

Temporary table to  
de-duplicate

```
{
```

```
  "dependent": false,
```

```
  "cacheable": true,
```

```
  "query_block": {
```

```
    "select_id": 2,
```

```
    "cost_info": {
```

```
      "query_cost": "46.15"
```

```
    },
```

```
    "table": {
```

```
      "table_name": "City",
```

```
      "access_type": "ref",
```

```
      "possible_keys": [
```

```
        "CountryCode"
```

```
      ],
```

```
      "key": "CountryCode",
```

```
      "used_key_parts": [
```

```
        "CountryCode"
```

```
      ],
```

```
      "key_length": "3",
```

```
      "ref": [
```

```
        "const"
```

```
      ],
```

```
..
```

# UNION ALL

- Results may contain duplicate rows
- Does not require an intermediate temporary table in simple use cases. i.e. no result ordering.
- Otherwise similar to UNION

```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM City WHERE CountryCode = 'CAN'
```

```
UNION ALL
```

```
SELECT * FROM City WHERE CountryCode = 'USA'
```

```
{
```

```
  "query_block": {
```

```
    "union_result": {
```

```
      "using_temporary_table": false,
```

```
      "query_specifications": [
```

```
        {
```

```
          "dependent": false,
```

```
          "cacheable": true,
```

```
          "query_block": {
```

```
            "select_id": 1,
```

```
            "cost_info": {
```

```
              "query_cost": "17.15"
```

```
            },
```

```
            "table": {
```

```
              "table_name": "City",
```

```
              "access_type": "ref",
```

```
..
```

```
              "key": "CountryCode",
```

```
..
```

No temporary table

```
{
```

```
  "dependent": false,
```

```
  "cacheable": true,
```

```
  "query_block": {
```

```
    "select_id": 2,
```

```
    "cost_info": {
```

```
      "query_cost": "46.15"
```

```
    },
```

```
    "table": {
```

```
      "table_name": "City",
```

```
      "access_type": "ref",
```

```
      "possible_keys": [
```

```
        "CountryCode"
```

```
      ],
```

```
      "key": "CountryCode",
```

```
      "used_key_parts": [
```

```
        "CountryCode"
```

```
      ],
```

```
      "key_length": "3",
```

```
      "ref": [
```

```
        "const"
```

```
      ],
```

```
..
```

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
- 18. Descending Indexes**
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets



# Descending Indexes

- B+tree indexes are ordered
- In 8.0 you can specify the order
- Use cases:
  - Faster to scan in order
  - Can't change direction in a composite index

```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM Country WHERE continent='Asia' AND population > 5000000  
ORDER BY population DESC;
```

```
{  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "7.91"  
    },  
    "ordering_operation": {  
      "using_filesort": false,  
      "table": {  
        "table_name": "Country",  
        "access_type": "range",  
        ..  
        "key": "c_p",  
        ..  
        "backward_index_scan": true,  
        ..  
      }  
    }  
  }  
}
```



Still uses the index, but  
about 15% slower

EXPLAIN FORMAT=JSON

SELECT \* FROM Country WHERE continent IN ('Asia', 'Oceania') AND population > 5000000  
ORDER BY **continent ASC, population DESC**

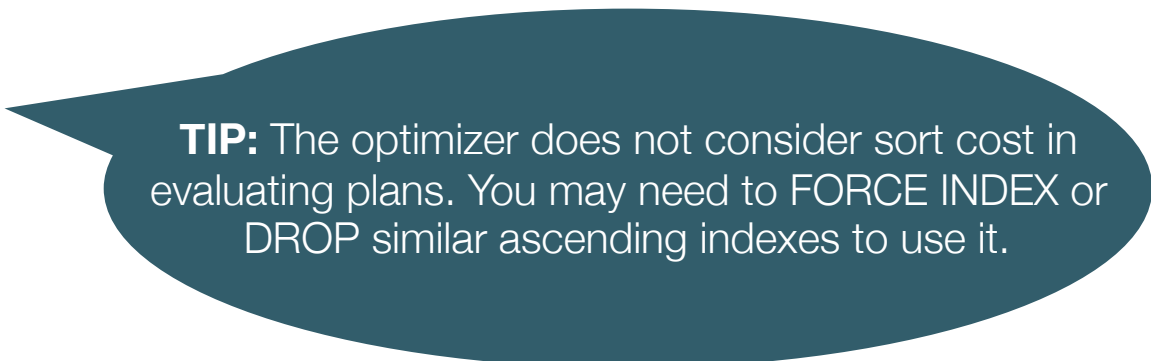
```
{  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "48.36"  
    },  
    "ordering_operation": {  
      "using_filesort": true,  
      "cost_info": {  
        "sort_cost": "33.00"  
      },  
      "table": {  
        "table_name": "Country",  
        "access_type": "range",  
        "key": "c_p",  
        ..  
        "rows_examined_per_scan": 33,  
        "rows_produced_per_join": 33,  
        "filtered": "100.00",  
        ..  
      }  
    }  
  }  
}
```



Must sort values of  
population in reverse

```
ALTER TABLE Country DROP INDEX c_p, DROP INDEX c_p_n,  
ADD INDEX c_p_desc (continent ASC, population DESC);  
EXPLAIN FORMAT=JSON  
SELECT * FROM Country WHERE continent IN ('Asia', 'Oceania') AND population > 5000000  
ORDER BY continent ASC, population DESC;
```

```
{  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "15.36"  
    },  
    "ordering_operation": {  
      "using_filesort": false,  
      "table": {  
        "table_name": "Country",  
        "access_type": "range",  
        ..  
        "key": "c_p_desc",  
        "used_key_parts": [  
          "Continent",  
          "Population"  
        ],  
        "key_length": "5",  
        ..  
      }  
    }  
  }  
}
```



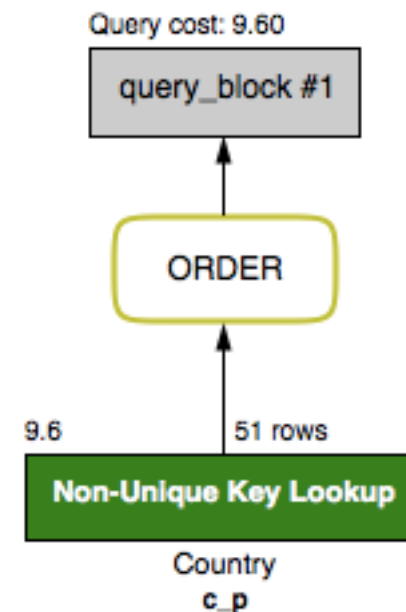
**TIP:** The optimizer does not consider sort cost in evaluating plans. You may need to **FORCE INDEX** or **DROP** similar ascending indexes to use it.

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
- 19. Sorting**
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# How is ORDER BY optimized?

1. Via an Index
2. Top N Buffer (“priority queue”)
3. Using temporary files



# Via an Index

- B+tree indexes are ordered
- Some ORDER BY queries do not require sorting at all

```
EXPLAIN FORMAT=JSON
SELECT * FROM Country WHERE continent='Asia' ORDER BY population;
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "9.60"
    },
    "ordering_operation": {
      "using_filesort": false,
      ..
      "key": "c_p",
```

The order is provided by  
c\_p

# Via a Priority Queue

- Special ORDER BY + small limit optimization
- Keeps top N records in an in memory buffer
- Usage is **NOT** shown in EXPLAIN

```
SELECT * FROM Country IGNORE INDEX (p, p_c)  
ORDER BY population LIMIT 10;
```



```

"select#": 1,
"steps": [
  {
    "filesort_information": [
      {
        "direction": "asc",
        "table": "`country` IGNORE INDEX (`p_c`) IGNORE INDEX (`p`)",
        "field": "Population"
      }
    ],
    "filesort_priority_queue_optimization": {
      "limit": 10,
      "chosen": true
    },
    "filesort_execution": [
    ],
    "filesort_summary": {
      "memory_available": 262144,
      "key_size": 4,
      "row_size": 272,
      "max_rows_per_buffer": 11,
      "num_rows_estimate": 587,
      "num_rows_found": 11,
      "num_examined_rows": 239,
      "num_tmp_files": 0,
      "sort_buffer_size": 3080,
      "sort_algorithm": "std::sort",
      "unpacked_addon_fields": "using_priority_queue",
      "sort_mode": "<fixed_sort_key, additional_fields>"
    }
  }
]

```



OPTIMIZER TRACE  
showing Priority Queue for  
sort

# Using Temporary Files

- Either “Alternative Sort Algorithm” (no blobs present) or “Original Sort Algorithm”

```
SELECT * FROM Country IGNORE INDEX (p, p_c)  
ORDER BY population;
```

```
"select#": 1,
"steps": [
  {
    "filesort_information": [
      {
        "direction": "asc",
        "table": "`country` IGNORE INDEX (`p_c`) IGNORE INDEX (`p`)",
        "field": "Population"
      }
    ],
    "filesort_priority_queue_optimization": {
      "usable": false,
      "cause": "not applicable (no LIMIT)"
    },
    "filesort_execution": [
    ],
    "filesort_summary": {
      "memory_available": 262144,
      "key_size": 4,
      "row_size": 274,
      "max_rows_per_buffer": 587,
      "num_rows_estimate": 587,
      "num_rows_found": 239,
      "num_examined_rows": 239,
      "num_tmp_files": 0,
      "sort_buffer_size": 165536,
      "sort_algorithm": "std::stable_sort",
      "sort_mode": "<fixed_sort_key, packed_additional_fields>"
    }
  }
]
```

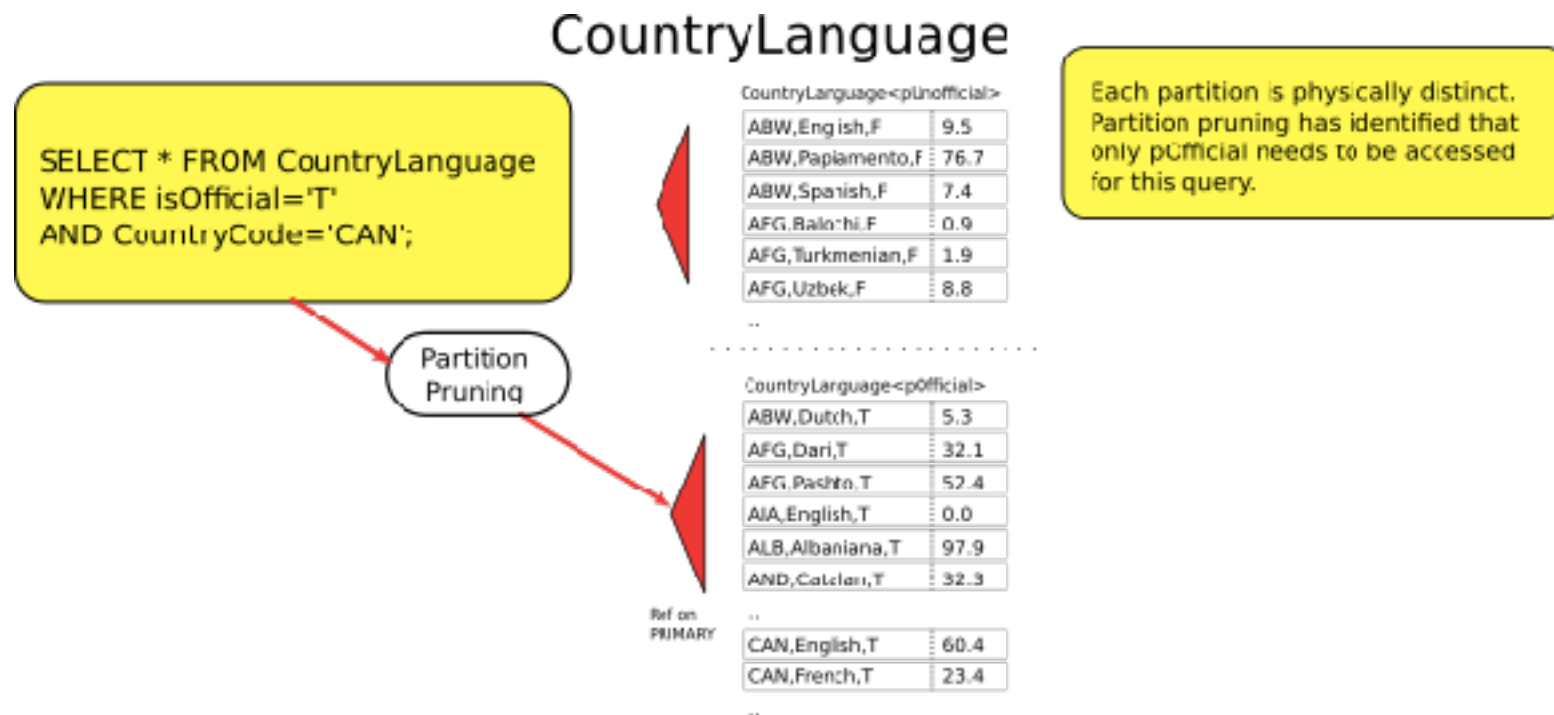
Not Using Priority Sort

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
- 20. Partitioning**
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Partitioning

- Split a table physically into smaller tables
- At the user-level make it still appear as one table



# Use Cases

- Can be a better fit low cardinality columns than indexing
- Useful for time series data with retention scheme
  - i.e. drop data older than 3 months
- Data where queries always have some locality
  - i.e. store\_id, region

# Partition Pruning

- Optimizer looks at query and identifies which partitions need to be accessed

```
ALTER TABLE CountryLanguage MODIFY IsOfficial CHAR(1) NOT NULL DEFAULT 'F', DROP  
PRIMARY KEY, ADD PRIMARY KEY(CountryCode, Language, IsOfficial);
```

```
ALTER TABLE CountryLanguage PARTITION BY LIST COLUMNS (IsOfficial) (  
    PARTITION pUnofficial VALUES IN ('F'),  
    PARTITION pOfficial VALUES IN ('T')  
);
```

```
EXPLAIN FORMAT=JSON
```

```
SELECT * FROM CountryLanguage WHERE isOfficial='T' AND  
CountryCode='CAN';
```

```
{  
  "query_block": {  
    "select_id": 1,  
    "cost_info": {  
      "query_cost": "2.40"  
    },  
    "table": {  
      "table_name": "CountryLanguage",  
      "partitions": [  
        "p0fficial"  
      ],  
      "access_type": "ref",  
      ..  
      "key": "PRIMARY",  
      ..  
    }  
  }  
}
```



Only accesses one  
partition



# Explicit Partition Selection

- Also possible to “target” a partition
- Consider this similar to query hints

```
SELECT * FROM CountryLanguage PARTITION (pOfficial)  
WHERE CountryCode='CAN';
```

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
- 21. Query Rewrite**
22. Invisible Indexes
23. Profiling
24. JSON
25. Character Sets

# Query Rewrite

- MySQL allows you to change queries before they are executed
- Insert a hint, or remove a join that is not required

```
mysql -u root -p < install_rewriter.sql
```

```
INSERT INTO query_rewrite.rewrite_rules(pattern_database, pattern,  
replacement) VALUES (  
"world",  
"SELECT * FROM Country WHERE population > ? AND continent=?",  
"SELECT * FROM Country WHERE population > ? AND continent=? LIMIT 1"  
);  
CALL query_rewrite.flush_rewrite_rules();
```

```
SELECT * FROM Country WHERE population > 5000000 AND  
continent='Asia';
```

```
SHOW WARNINGS;
```

```
***** 1. row *****
```

```
Level: Note
```

```
Code: 1105
```

```
Message: Query 'SELECT * FROM Country WHERE population >  
5000000 AND continent='Asia'' rewritten to 'SELECT *  
FROM Country WHERE population > 5000000 AND  
continent='Asia' LIMIT 1' by a query rewrite plugin
```

```
1 row in set (0.00 sec)
```

```
SELECT * FROM query_rewrite.rewrite_rules\G
```

```
***** 1. row *****
```

```
      id: 1
```

```
      pattern: SELECT * FROM Country WHERE  
population > ? AND continent=?
```

```
      pattern_database: world
```

```
      replacement: SELECT * FROM Country WHERE  
population > ? AND continent=? LIMIT 1
```

```
      enabled: YES
```

```
      message: NULL
```

```
      pattern_digest: 88876bbb502cef6efddcc661cce77deb
```

```
normalized_pattern: select `*` from `world`.`country`  
where ((`population` > ?) and (`continent` = ?))
```

```
1 row in set (0.00 sec)
```

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
- 22. Invisible Indexes**
23. Profiling
24. JSON
25. Character Sets

# Changing Indexes is a *Destructive Operation*

- Removing an index can make some queries much slower
- Adding can cause some existing query plans to change
- Old-style hints will generate errors if indexes are removed

# Invisible Indexes, the “Recycle Bin”

- Hide the indexes from the optimizer
- Will no longer be considered as part of query execution plans
- Still kept up to date and are maintained by insert/update/delete statements



# Invisible Indexes: Soft Delete

```
ALTER TABLE Country ALTER INDEX c INVISIBLE;
SELECT * FROM information_schema.statistics WHERE is_visible='NO';
***** 1. row *****

TABLE_CATALOG: def
TABLE_SCHEMA: world
TABLE_NAME: Country
NON_UNIQUE: 1
INDEX_SCHEMA: world
INDEX_NAME: c
SEQ_IN_INDEX: 1
COLUMN_NAME: Continent
COLLATION: A
CARDINALITY: 7
SUB_PART: NULL
PACKED: NULL
NULLABLE:
INDEX_TYPE: BTREE
COMMENT: disabled
INDEX_COMMENT:
IS_VISIBLE: NO
```

# Invisible Indexes: Staged Rollout

```
ALTER TABLE Country ADD INDEX c (Continent)  
INVISIBLE;
```

# after some time

```
ALTER TABLE Country ALTER INDEX c VISIBLE;
```

# Finding Unused Indexes

```
SELECT * FROM sys.schema_unused_indexes;
```

object_schema	object_name	index_name
world	Country	p
world	Country	p_c

```
2 rows in set (0.01 sec)
```

# Do indexes hurt reads or writes?

- They can have some impact on both:
  - On writes, indexes need to space, and to be maintained
  - On reads, lets use an example...

# Indexes Hurting Reads

```
CREATE TABLE t1 (  
  id INT NOT NULL primary key auto_increment,  
  a VARCHAR(255) NOT NULL,  
  b VARCHAR(255) NOT NULL,  
  c TEXT,  
  d TEXT,  
  INDEX a (a),  
  INDEX ab (a,b));
```

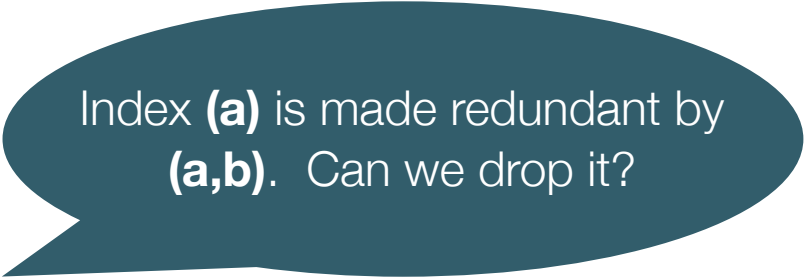
**Both** indexes are candidates.  
**Both** will be examined.

# Sample Query

```
SELECT * FROM t1 WHERE a = 'abc' AND b = 'bcd';
```

# A use case for invisible indexes!

```
CREATE TABLE t1 (  
  id INT NOT NULL primary key auto_increment,  
  a VARCHAR(255) NOT NULL,  
  b VARCHAR(255) NOT NULL,  
  c TEXT,  
  d TEXT,  
INDEX a (a),  
INDEX ab (a,b));
```



Index **(a)** is made redundant by **(a,b)**. Can we drop it?

# Consider:

```
SELECT count(*) FROM t1 FORCE INDEX (a)  
WHERE a='1234' AND id=1234;
```

# No, due to clustered Index!

```
FORCE INDEX (a) WHERE a='1234' AND id=1234;
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "0.35"
    },
    "table": {
      "table_name": "t1",
      "access_type": "const",
      "possible_keys": [
        "a"
      ],
      "key": "a",
      "used_key_parts": [
        "a",
        "id"
      ],
    },
  },
}
```

```
FORCE INDEX (ab) WHERE a='1234' AND id=1234;
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "11.80"
    },
    "table": {
      "table_name": "t1",
      "access_type": "ref",
      "possible_keys": [
        "ab"
      ],
      "key": "ab",
      "used_key_parts": [
        "a"
      ],
    },
  },
}
```

..

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
- 23. Profiling**
24. JSON
25. Character Sets



# Profiling

- Optimizer only shows estimates from pre-execution view
- Can be useful to know actual time spent
- Support for profiling is only very basic

```
wget http://www.tocker.ca/files/ps-show-profiles.sql  
mysql -u root -p < ps-show-profiles.sql
```

```
CALL sys.enable_profiling();
CALL sys.show_profiles;
***** 1. row *****
Event_ID: 22
Duration: 495.02 us
    Query: SELECT * FROM Country WHERE co ... Asia' and population > 5000000
1 row in set (0.00 sec)
```

```
CALL sys.show_profile_for_event_id(22);
```

Status	Duration
starting	64.82 us
checking permissions	4.10 us
Opening tables	11.87 us
init	29.74 us
System lock	5.63 us
optimizing	8.74 us
statistics	139.38 us
preparing	11.94 us
executing	348.00 ns
Sending data	192.59 us
end	1.17 us
query end	4.60 us
closing tables	4.07 us
freeing items	13.60 us
cleaning up	734.00 ns

```
15 rows in set (0.00 sec)
```

```
SELECT * FROM Country WHERE Continent='Antarctica' and SLEEP(5);
CALL sys.show_profiles();
CALL sys.show_profile_for_event_id(<event_id>);
```

Status	Duration
starting	103.89 us
checking permissions	4.48 us
Opening tables	17.78 us
init	45.75 us
System lock	8.37 us
optimizing	11.98 us
statistics	144.78 us
preparing	15.78 us
executing	634.00 ns
Sending data	116.15 us
<b>User sleep</b>	<b>5.00 s</b>
<b>User sleep</b>	<b>5.00 s</b>
<b>User sleep</b>	<b>5.00 s</b>
<b>User sleep</b>	<b>5.00 s</b>
<b>User sleep</b>	<b>5.00 s</b>
end	2.05 us
query end	5.63 us
closing tables	7.30 us
freeing items	20.19 us
cleaning up	1.20 us

20 rows in set (0.01 sec)

Sleeps for each row after  
index used on (c)

```
SELECT region, count(*) as c FROM Country GROUP BY region;
CALL sys.show_profiles();
CALL sys.show_profile_for_event_id(<event_id>);
```

Status	Duration
starting	87.43 us
checking permissions	4.93 us
Opening tables	17.35 us
init	25.81 us
System lock	9.04 us
optimizing	3.37 us
statistics	18.31 us
preparing	10.94 us
<b>Creating tmp table</b>	<b>35.57 us</b>
<b>Sorting result</b>	<b>2.38 us</b>
executing	741.00 ns
Sending data	446.03 us
<b>Creating sort index</b>	<b>49.45 us</b>
end	1.71 us
query end	4.85 us
removing tmp table	4.71 us
closing tables	6.12 us
freeing items	17.17 us
cleaning up	1.00 us

19 rows in set (0.01 sec)

```
SELECT * FROM performance_schema.events_statements_history_long
```

```
WHERE event_id=<event_id>\G
```

```
***** 1. row *****
```

```
THREAD_ID: 3062
```

```
EVENT_ID: 1566
```

```
END_EVENT_ID: 1585
```

```
EVENT_NAME: statement/sql/select
```

```
SOURCE: init_net_server_extension.cc:80
```

```
TIMER_START: 588883869566277000
```

```
TIMER_END: 588883870317683000
```

```
TIMER_WAIT: 751406000
```

```
LOCK_TIME: 132000000
```

```
count(*) as c FROM Country GROUP BY region
```

```
DIGEST: d3a04b346fe48da4f1f5c2e06628a245
```

```
DIGEST_TEXT: SELECT `region` ,  
COUNT ( * ) AS `c` FROM `Country`  
GROUP BY `region`
```

```
CURRENT_SCHEMA: world
```

```
OBJECT_TYPE: NULL
```

```
OBJECT_SCHEMA: NULL
```

```
OBJECT_NAME: NULL
```

```
OBJECT_INSTANCE_BEGIN: NULL
```

```
MYSQL_ERRNO: 0
```

```
RETURNED_SQLSTATE: NULL
```

```
MESSAGE_TEXT: NULL
```

```
ERRORS: 0
```

```
WARNINGS: 0
```

```
..
```

For non-aggregate queries rows sent  
vs. rows examined helps indicate index  
effectiveness.

```
..
```

```
ROWS_AFFECTED: 0
```

```
ROWS_SENT: 25
```

```
ROWS_EXAMINED: 289
```

```
CREATED_TMP_DISK_TABLES: 0
```

```
CREATED_TMP_TABLES: 1
```

```
SELECT_FULL_JOIN: 0
```

```
SELECT_FULL_RANGE_JOIN: 0
```

```
SELECT_RANGE: 0
```

```
SELECT_RANGE_CHECK: 0
```

```
SELECT_SCAN: 1
```

```
SORT_MERGE_PASSES: 0
```

```
SORT_RANGE: 0
```

```
SORT_ROWS: 25
```

```
SORT_SCAN: 1
```

```
NO_INDEX_USED: 1
```

```
NO_GOOD_INDEX_USED: 0
```

```
NESTING_EVENT_ID: NULL
```

```
NESTING_EVENT_TYPE: NULL
```

```
NESTING_EVENT_LEVEL: 0
```

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
- 24. JSON**
25. Character Sets

# JSON

- Optimizer has native support for JSON with indexes on generated columns used for matching JSON path expressions

```
CREATE TABLE CountryJson (Code char(3) not null primary key, doc JSON NOT NULL);
INSERT INTO CountryJSON SELECT code,
  JSON_OBJECT(
    'Name', Name,
    'Continent', Continent,
    ..
    'HeadOfState', HeadOfState,
    'Capital', Capital,
    'Code2', Code2
  ) FROM Country;
```

EXPLAIN FORMAT=JSON

SELECT \* FROM CountryJSON where **doc->>"\$.Name"** = 'Canada' ;

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "48.80"
    },
    "table": {
      "table_name": "CountryJSON",
      "access_type": "ALL",
      "rows_examined_per_scan": 239,
      "rows_produced_per_join": 239,
      "filtered": "100.00",
      "cost_info": {
        "read_cost": "1.00",
        "eval_cost": "47.80",
        "prefix_cost": "48.80",
        "data_read_per_join": "3K"
      },
    },
  },

```

..



```
ALTER TABLE CountryJSON ADD Name char(52) AS (doc->>"$.Name"),
ADD INDEX n (Name);
EXPLAIN FORMAT=JSON
SELECT * FROM CountryJSON where doc->>"$.Name" = 'Canada';
```

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "1.20"
    },
    "table": {
      "table_name": "CountryJSON",
      "access_type": "ref"
    },
    ..
    "key": "n",
    ..
    "key_length": "53",
    "ref": [
      "const"
    ],
    ..
  }
```

Matches expression from  
indexed virtual column

Key from virtual column

# JSON Comparator

- JSON types compare to MySQL types

```
SELECT CountryJSON.* FROM CountryJSON
INNER JOIN Country ON CountryJSON.doc->>"$.Name" = Country.Name WHERE
Country.Name='Canada';
```

```
***** 1. row *****
```

Code: CAN

```
doc: {"GNP": 598862, "Name": "Canada", "Code2": "CA", "GNPOld": 625626, "Region":
"North America", "Capital": 1822, "Continent": "North America", "IndepYear": 1867,
"LocalName": "Canada", "Population": 31147000, "HeadOfState": "Elisabeth II",
"SurfaceArea": 9970610, "GovernmentForm": "Constitutional Monarchy, Federation",
"LifeExpectancy": 79.4000015258789}
```

Name: Canada

# Agenda

1. Introduction
2. Server Architecture
3. B+trees
4. EXPLAIN
5. Optimizer Trace
6. Logical Transformations
7. Cost Based Optimization
8. Hints and Switches
9. Comparing Plans
10. Composite Indexes
11. Covering Indexes
12. Visual Explain
13. Transient Plans
14. Subqueries
15. CTEs and Views
16. Joins
17. Aggregation
18. Descending Indexes
19. Sorting
20. Partitioning
21. Query Rewrite
22. Invisible Indexes
23. Profiling
24. JSON
- 25. Character Sets**

# Character Sets

- The default character set in MySQL 8.0 is utf8mb4
- Utf8mb4 is variable length (1-4 bytes)
- InnoDB will always store as variable size for both CHAR and VARCHAR
- Some buffers inside MySQL may require the fixed length (4 bytes)

## Character Sets (cont.)

- CHAR(n) or VARCHAR(n) refers to n characters - x4 for maximum length
- EXPLAIN will always show the maximum length
- Mysqldump will preserve character set

```
ALTER TABLE City DROP FOREIGN KEY city_ibfk_1;
```

```
ALTER TABLE CountryLanguage DROP FOREIGN KEY  
countryLanguage_ibfk_1;
```

```
ALTER TABLE Country CONVERT TO CHARACTER SET  
utf8mb4;
```

```

{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "0.35"
    },
    "table": {
      "table_name": "Country",
      "access_type": "ref",
      "possible_keys": [
        "n"
      ],
      "key": "n",
      "used_key_parts": [
        "Name"
      ],
      "key_length": "52",
      ..
      "rows_examined_per_scan": 1,
      "rows_produced_per_join": 1,
      "filtered": "100.00",
      "cost_info": {
        "read_cost": "0.25",
        "eval_cost": "0.10",
        "prefix_cost": "0.35",
        "data_read_per_join": "264"
      },
      ..

```

Key length as latin1

```

{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "0.35"
    },
    "table": {
      "table_name": "Country",
      "access_type": "ref",
      "possible_keys": [
        "n"
      ],
      "key": "n",
      "used_key_parts": [
        "Name"
      ],
      "key_length": "208",
      ..
      "rows_examined_per_scan": 1,
      "rows_produced_per_join": 1,
      "filtered": "100.00",
      "cost_info": {
        "read_cost": "0.25",
        "eval_cost": "0.10",
        "prefix_cost": "0.35",
        "data_read_per_join": "968"
      },
      ..

```

Key length as utf8

# Conclusion

- Thank you for coming!
- This presentation is available as a website:  
[www.unofficialmysqlguide.com](http://www.unofficialmysqlguide.com)

ORACLE®