



# MySQL Group Replication

## An Overview

# Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

# Program Agenda

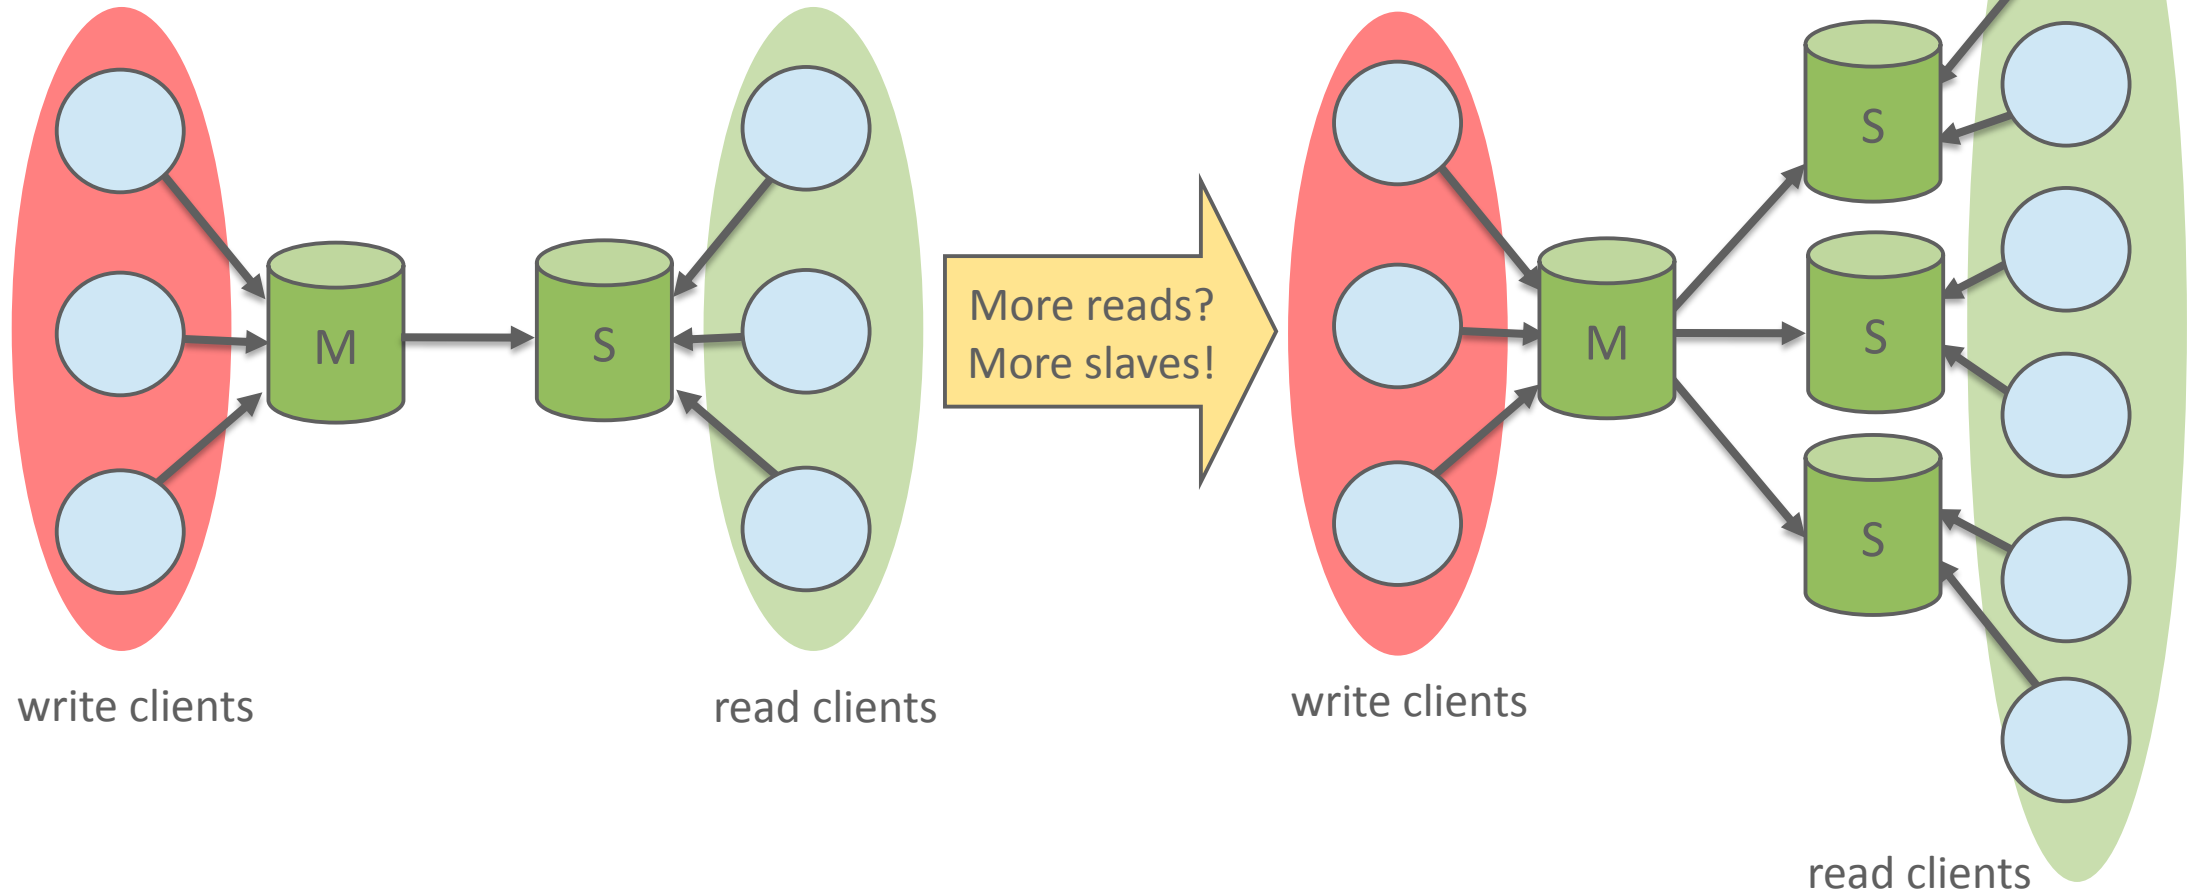
# Program Agenda

- 1 ➤ The Theory
- 2 ➤ How It Works
- 3 ➤ How to Use It
- 4 ➤ Conclusion

# 1 The Theory


# Background: What is Replication Used For?

## Read scale-out




# Background: What is Replication Used For?

**Redundancy:** If master crashes, **promote** slave to master



# MySQL Group Replication





# MySQL Group Replication

- **What is MySQL Group Replication?**

“Multi-master **update anywhere** replication plugin for MySQL with built-in **automatic distributed recovery, conflict detection, and group membership.**”

- **What does the MySQL Group Replication plugin do for the user?**

- Removes the need for manually handling server fail-over
- Provides distributed fault tolerance
- Enables Active/Active update anywhere setups
- Automates group reconfiguration (handling of crashes, failures, re-connects)
- Provides a highly available distributed database service

# Use Cases

- **Elastic Replication**

- Environments that require a very fluid replication infrastructure, where the number of servers has to grow or shrink dynamically and with as little manual intervention as possible.

- **Highly Available Shards**

- Sharding is a popular approach to achieve write scale-out. Users can use MySQL Group Replication to implement highly available shards. Each shard can map to an individual Replication Group.

- **Alternative to Master-Slave Replication**

- It may be that a single master server makes it a single point of contention. Writing to an entire group may prove more scalable under certain circumstances.

# The Theory Behind It

- Implementation is based on “Replicated Database State Machines”
  - Group Communication primitives resemble general properties of Databases
  - Distributed systems meet Databases: [Pedone, Guerraoui, and Schiper paper](#)
- Deferred update replication: before committing locally we certify it on all nodes
  - In order to implement it one needs Atomic Broadcast – the change occurs everywhere or nowhere
  - This is necessary to ensure data consistency across all nodes
- Membership Service
  - Group Membership: it allows one to know that at a given moment in time all the members that are participating in the protocol are associated with the same logical identifier (view id)
  - View Synchrony: ensure that messages from past views are all delivered before a new view is installed

## 2 How It Works

# Example: Adding a Node to an Existing Group

- Create a user for automated distributed recovery:

```
./bin/mysql -uroot -h 127.0.0.1 -P 13001 -p --prompt='server1>'  
  
server1>  
CREATE USER 'rpl_user'@'%' IDENTIFIED BY 'rpl_pass';  
GRANT REPLICATION SLAVE ON *.* TO rpl_user@'%';
```

- Server needs to be started with a valid configuration:

```
./bin/mysqld --no-defaults --basedir=. --datadir=<DATADIR_LOCATION> -P <PORT> \  
--socket=mysqld<ID>.sock --log-bin=master-bin --server-id=<ID> \  
--gtid-mode=on --enforce-gtid-consistency --log-slave-updates \  
--binlog-checksum=NONE --binlog-format=row \  
--master-info-repository=TABLE --relay-log-info-repository=TABLE \  
--transaction-write-set-extraction=MURMUR32 \  
--plugin-dir=lib/plugin --plugin-load=group_replication.so
```

## Example: Adding a Node to an Existing Group (2)

- Now lets configure the replication user used for recovery:


```
SET GLOBAL group_replication_recovery_user='rpl_user';  
SET GLOBAL group_replication_recovery_password='rpl_pass';
```

- Now the group communication backbone:

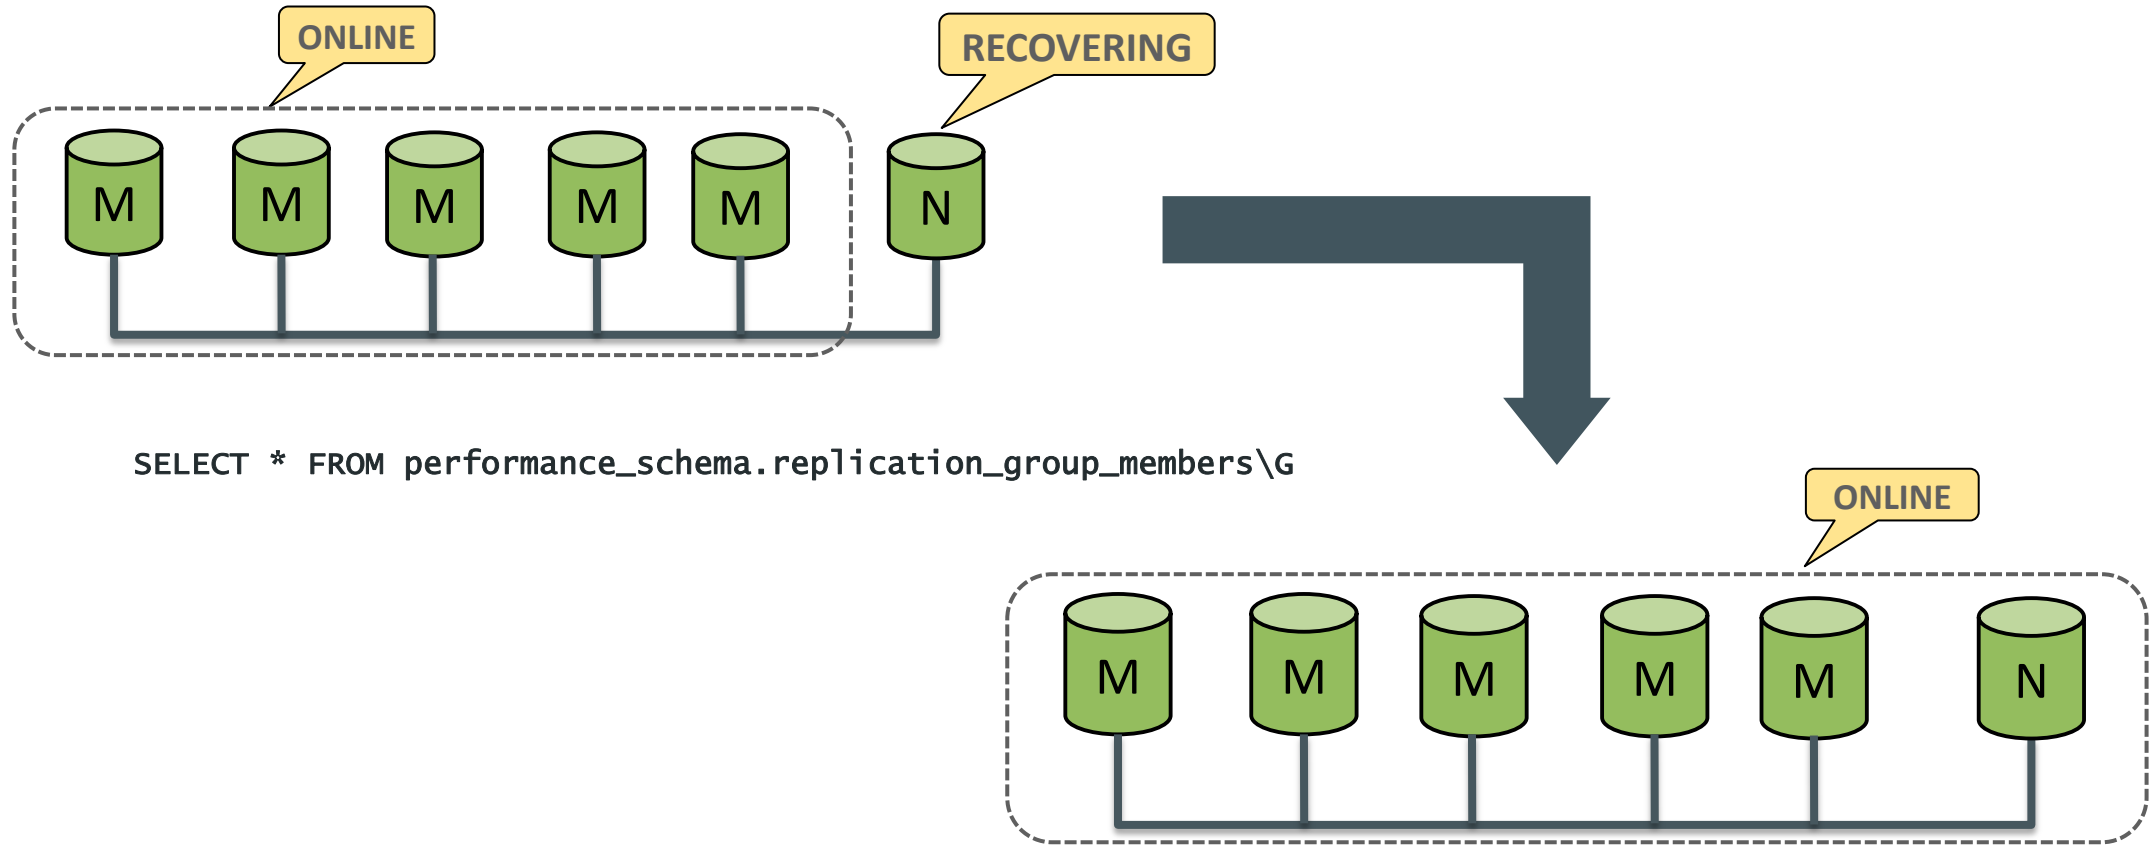
```
SET GLOBAL group_replication_group_name= <valid UUID>;  
SET GLOBAL group_replication_local_address=<this node address:port for the  
communication backbone>;  
SET GLOBAL group_replication_peer_addresses= <comma-separated list of all other  
nodes in the group>;  
SET GLOBAL group_replication_bootstrap_group= 0;
```

## Example: Adding a Node to an Existing Group (3)

- Server that joins the group will automatically synchronize with the others
  - It will retrieve the diff between its data and the rest of the group
  - Hint: provision the new node with base data (i.e. restore a backup) before joining an existing group



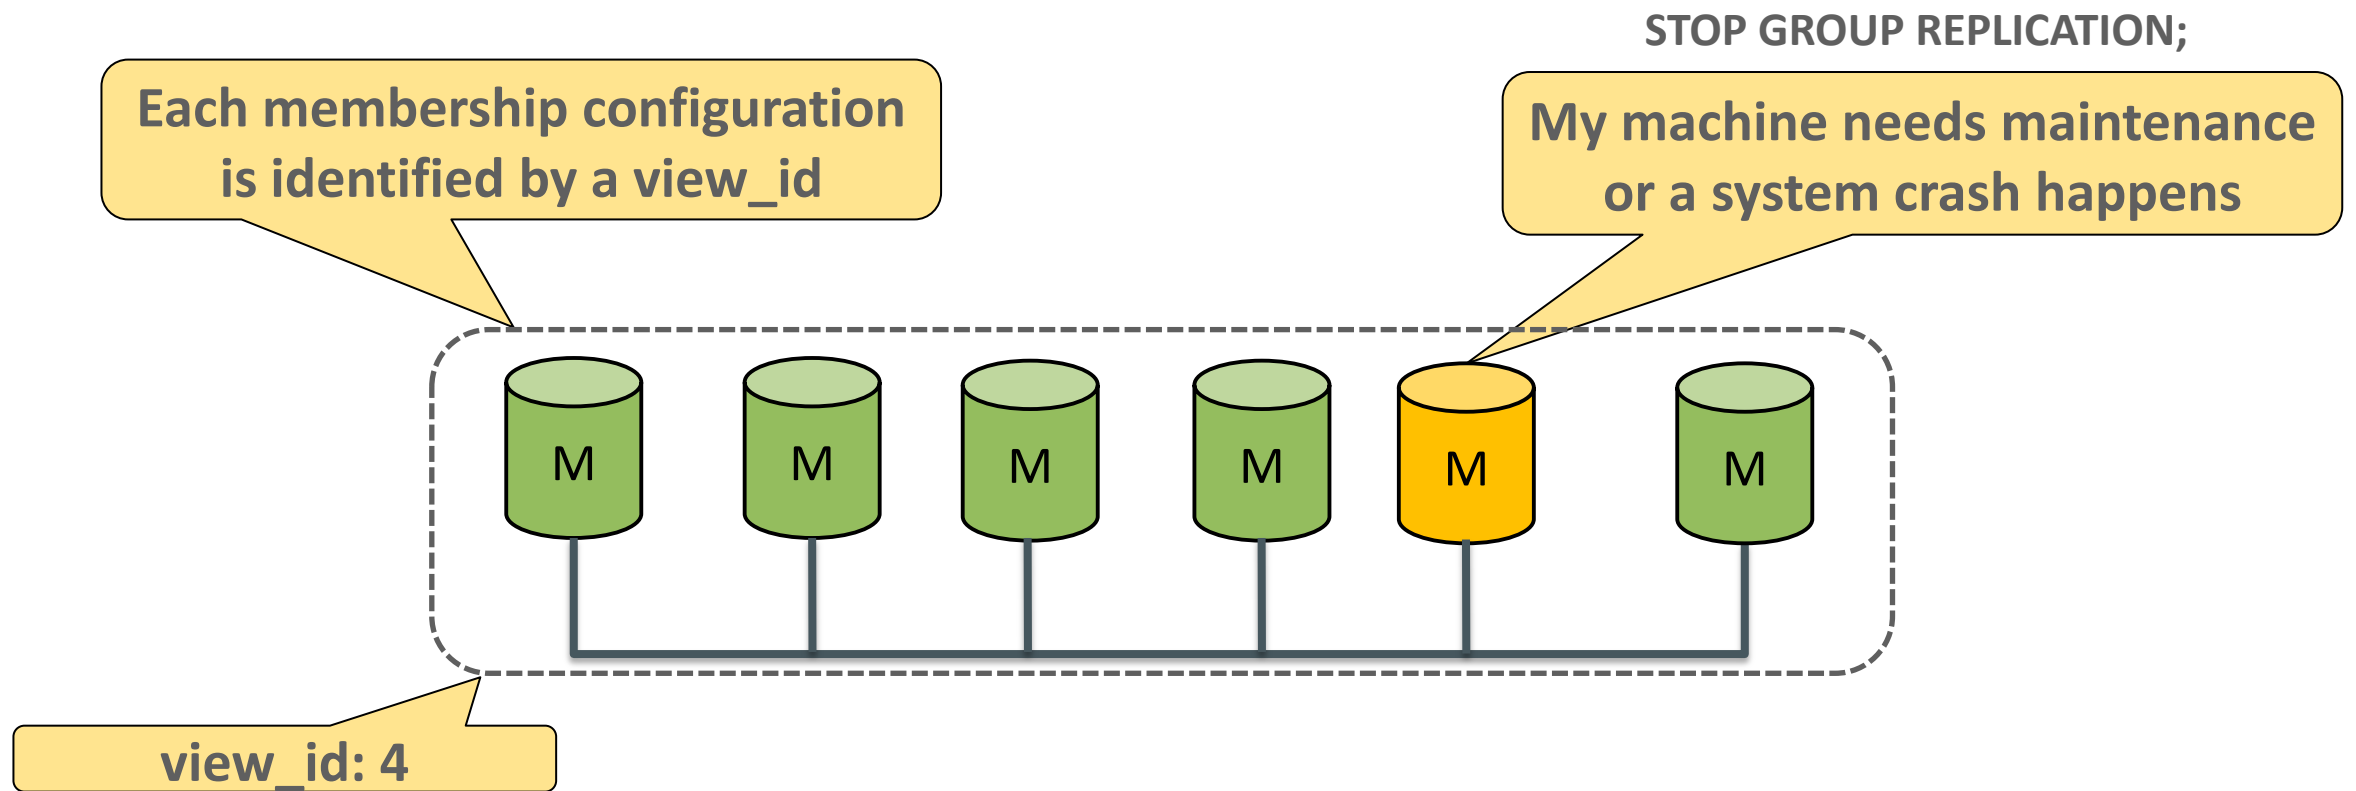
## Example: Adding a Node to an Existing Group (4)






# Example: Automated Group Membership

- If a server leaves the group, the others will automatically be informed




## Example: Automated Group Membership (2)

- If a server leaves the group, the others will automatically be informed





- Server that (re)joins the group will automatically sync with the others



## How to use


# Multi-Master Update Anywhere!

- Any two transactions on different servers can write to the same row
- Conflicts will automatically be detected and handled
  - *First committer wins rule*



# Full GTID Support

- All group members share the same UUID, which is the group name



# Monitoring the Replication Group

- Monitor group health and stats using Performance Schema tables

```
mysql> SELECT * FROM
performance_schema.replication_connection_status\G
***** 1. row *****
CHANNEL_NAME: group_replication_applier
GROUP_NAME: 8a94f357-aab4-11df-86ab-c80aa9429563
SOURCE_UUID: 8a94f357-aab4-11df-86ab-c80aa9429563
THREAD_ID: NULL
SERVICE_STATE: ON
...
```

```
mysql> SELECT * FROM performance_schema.replication_group_member_stats\G
***** 1. row *****
CHANNEL_NAME: group_replication_applier
VIEW_ID: 1428497631:3
MEMBER_ID: e38fdea8-dded-11e4-b211-e8b1fc3848de
COUNT_TRANSACTIONS_IN_QUEUE: 0
COUNT_TRANSACTIONS_CHECKED: 12
COUNT_CONFLICTS_DETECTED: 5
COUNT_TRANSACTIONS_VALIDATING: 6
TRANSACTIONS_COMMITTED_ALL_MEMBERS: 8a84f397-aaa4-18df-89ab-
c70aa9823561:1-7
LAST_CONFLICT_FREE_TRANSACTION: 8a84f397-aaa4-18df-89ab-c70aa9823561:7
```

```
mysql> SELECT * FROM performance_schema.replication_group_members\G
***** 1. row *****
CHANNEL_NAME: group_replication_applier
MEMBER_ID: 597dbb72-3e2c-11e4-9d9d-ecf4bb227f3b
MEMBER_HOST: nightfury
MEMBER_PORT: 13000
MEMBER_STATE: ONLINE
***** 2. row *****
...
```

## 4 Conclusion

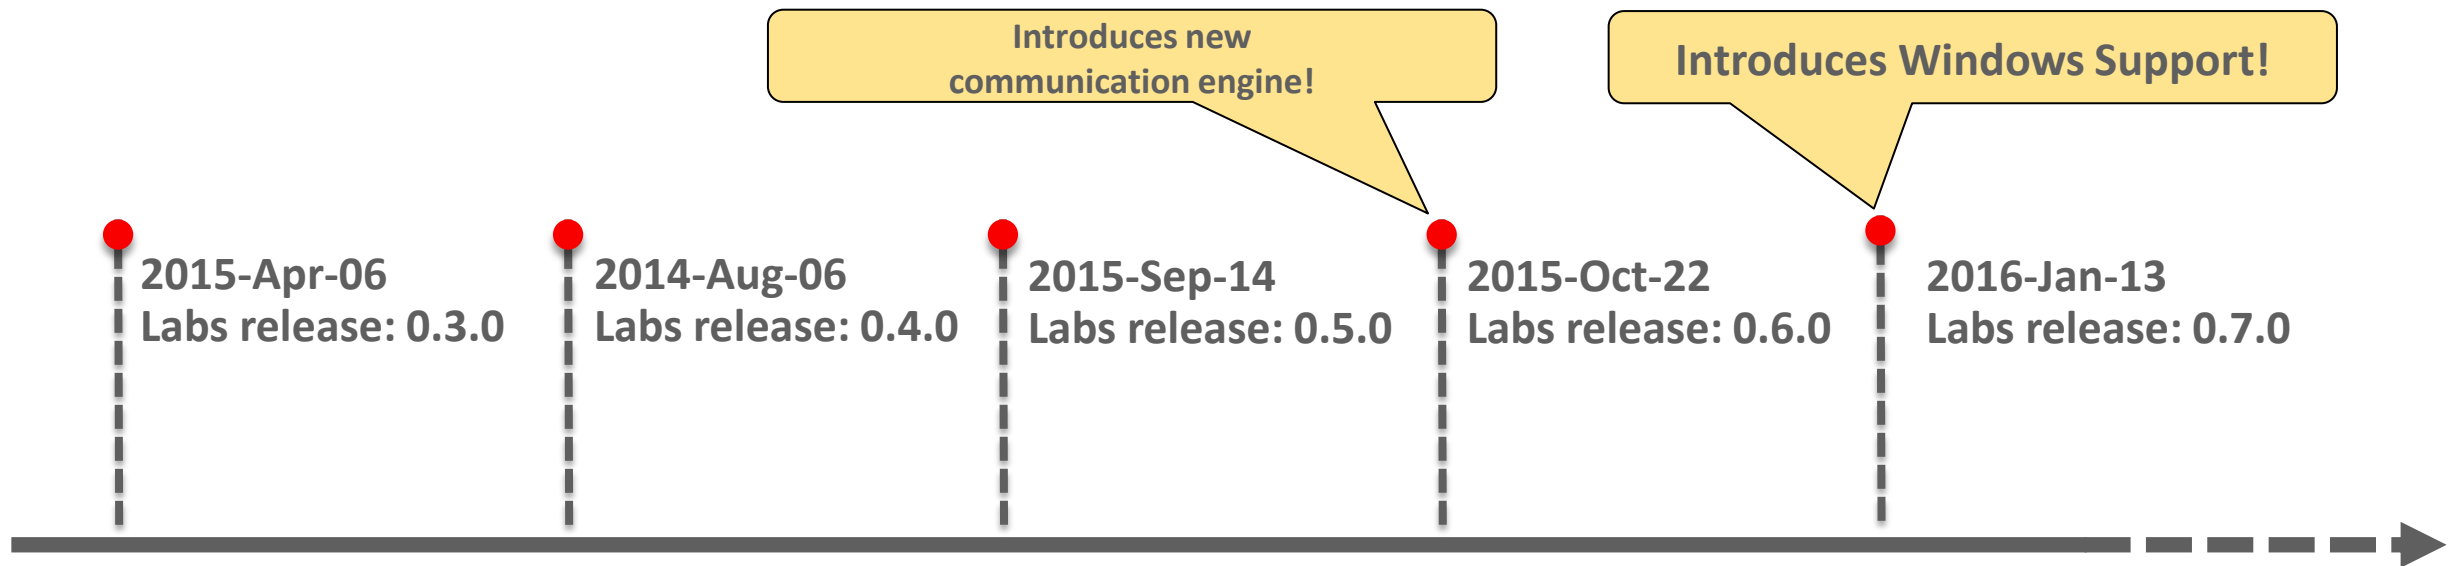
# Summary

- **Cloud Friendly**
  - Great technology for deployments where high availability and elasticity is a requirement, such as Cloud based infrastructures
- **Integrated**
  - With standard MySQL servers through a well defined API
  - With standard GTIDs, ROW based replication, and Performance Schema tables
- **Autonomic and Operations Friendly**
  - It is self-healing: no admin overhead for handling server fail-overs
  - Provides fault-tolerance: enables multi-master update anywhere and a resilient distributed MySQL service
- Lab releases provide a sneak peek at what is coming -- a new replication plugin and exciting new infrastructure: **MySQL Group Replication and MySQL Router**



# Releases

- Strong development cycles and continuous community engagement through regular lab releases



# Where to go from here?

- Packages
  - <http://labs.mysql.com>
- Blogs from the Engineers (news, technical information, and much more)
  - <http://mysqlhighavailability.com/tag/mysql-group-replication/>

ORACLE®