

# MYSQL PERFORMANCE TUNING

---

## I. MySQL Configuration (incl. MySQL 5.7)





# AURIMAS MIKALAUSKAS

---

- MySQL hacker since 1999 (yep, v3.23)
- In love with performance & scalability since 2004 Danga Interactive talk on LiveJournal's Backend at LISA04
- Percona Senior Performance Consultant & Architect 2006-2015
- Performance Engineer at EstanteVirtual.com.br
- Independent data performance consultant, instructor at Speedemy.com

# ON THE AGENDA FOR TODAY

.....

*How to Choose The Right MySQL Distribution*


*Essentials of MySQL Configuration Tuning*

*17 Key MySQL Settings for Best Performance*



# DISTROS

.....  
*MySQL? Percona Server?  
MariaDB? WebScaleSQL?*



# COMMUNITY MYSQL

---



- Community MySQL 5.1
  - innodb plugin not enabled by default
  - very poor InnoDB scalability (by default)
- Community MySQL 5.5
  - still missing a number of performance improvements
  - stability issues with high amounts of writes
- Community MySQL 5.6
  - not many reasons to use alternative distro
  - still not perfectly stable with high end hardware in a write-intensive environment

# COMMUNITY MYSQL 5.7

---



- Current GA release
- Rock solid
- New features:
  - multi-source replication
  - JSON support
  - proper multi-threaded replication
  - online buffer pool resize
  - spatial data types for InnoDB
  - sys schema
  - etc. (see full list)

# PERCONA SERVER

---



- Launched July 2008
  - as Percona patches for MySQL
  - in reality earlier than that
- Always up to date with upstream
  - i.e. Percona Server 5.1.73 = MySQL 5.1.73 + Percona code
  - also always backwards compatible
    - meaning no problem to keep switching between Percona Server and MySQL. No upgrade/downgrade scripts need to be used.
  - some exceptions to backwards compatibility in the past
    - none of these options were enabled by default
    - very clear when it's the case
- Makes it easy to try it out and switch back at any point

# PERCONA SERVER VERSION

---



- Using MySQL 5.1?
  - definitely switch to Percona Server 5.1
- Using MySQL 5.5? Switch to Percona Server 5.5 if...
  - adaptive hash index is a bottle-neck
  - you want faster checksums (hardware accelerated)
  - you have write-intensive workload
  - you need PAM authentication, audit log, thread pool
- Using MySQL 5.6? Switch to Percona Server 5.6 if...
  - you have very write-intensive workload and high end hardware
  - you need some of the features mentioned with 5.5
- MySQL 5.7? Use MySQL, Percona Server 5.7 coming soon.



# PERCONA SERVER KEYSTONE FEATURE: EXTENDED SLOW QUERY LOG

.....



```
# Time: 150811 12:15:06.959000 ←————— microsecond TS
# User@Host: root[root] @ localhost [localhost] Id: 62591948
# Schema: Core_DB Last_errno: 0 Killed: 0
# Query_time: 2.411376 Lock_time: 0.002444 Rows_sent: 0 Rows_examined: 10811959 Rows_affected: 0
# Bytes_sent: 0 Tmp_tables: 1 Tmp_disk_tables: 0 Tmp_table_sizes: 0
# QC_Hit: No Full_scan: Yes Full_join: No Tmp_table: Yes Tmp_table_on_disk: No
# Filesort: No Filesort_on_disk: No Merge_passes: 0
# InnoDB_trx_id: 20A40488F
# InnoDB_IO_r_ops: 37 InnoDB_IO_r_bytes: 606208 InnoDB_IO_r_wait: 0.000225
# InnoDB_rec_lock_wait: 0.000000 InnoDB_queue_wait: 0.000000
# InnoDB_pages_distinct: 1747
SELECT ...
```

-VS-

```
# Time: 150811 12:15:06
# User@Host: root[root] @ localhost [localhost] Id: 62591948
# Schema: Core_DB Last_errno: 1051 Killed: 0
# Query_time: 2.411376 Lock_time: 0.002444 Rows_sent: 0 Rows_examined: 10811959
SELECT ...
```

# EXTENDED SLOW QUERY LOG (2)



.....

```
# Time: 150811 12:15:06.959000
# User@Host: root[root] @ localhost [localhost] Id: 62591948
# Schema: Core_DB Last_errno: 0 Killed: 0
# Query_time: 2.411376 Lock_time: 0.002444 Rows_sent: 0 Rows_examined: 10811959 Rows_affected: 0
# Bytes_sent: 0 Tmp_tables: 1 Tmp_disk_tables: 0 Tmp_table_sizes: 0
# QC_Hit: No Full_scan: Yes Full_join: No Tmp_table: Yes Tmp_table_on_disk: No
# Filesort: No Filesort_on_disk: No Merge_passes: 0
# InnoDB_trx_id: 20A40488F
# InnoDB_IO_r_ops: 37 InnoDB_IO_r_bytes: 606208 InnoDB_IO_r_wait: 0.000225
# InnoDB_rec_lock_wait: 0.000000 InnoDB_queue_wait: 0.000000
# InnoDB_pages_distinct: 1747
SELECT ...
```

← query plan

-VS-

```
# Time: 150811 12:15:06
# User@Host: root[root] @ localhost [localhost] Id: 62591948
# Schema: Core_DB Last_errno: 1051 Killed: 0
# Query_time: 2.411376 Lock_time: 0.002444 Rows_sent: 0 Rows_examined: 10811959
SELECT ...
```

# EXTENDED SLOW QUERY LOG (3)



.....

```
# Time: 150811 12:15:06.959000
# User@Host: root[root] @ localhost [localhost] Id: 62591948
# Schema: Core_DB Last_errno: 0 Killed: 0
# Query_time: 2.411376 Lock_time: 0.002444 Rows_sent: 0 Rows_examined: 10811959 Rows_affected: 0
# Bytes_sent: 0 Tmp_tables: 1 Tmp_disk_tables: 0 Tmp_table_sizes: 0
# QC_Hit: No Full_scan: Yes Full_join: No Tmp_table: Yes Tmp_table_on_disk: No
# Filesort: No Filesort_on_disk: No Merge_passes: 0
# InnoDB_trx_id: 20A40488F
# InnoDB_IO_r_ops: 37 InnoDB_IO_r_bytes: 606208 InnoDB_IO_r_wait: 0.000225
# InnoDB_rec_lock_wait: 0.000000 InnoDB_queue_wait: 0.000000
# InnoDB_pages_distinct: 1747
SELECT ...
```

← extra innodb info

-VS-

```
# Time: 150811 12:15:06
# User@Host: root[root] @ localhost [localhost] Id: 62591948
# Schema: Core_DB Last_errno: 1051 Killed: 0
# Query_time: 2.411376 Lock_time: 0.002444 Rows_sent: 0 Rows_examined: 10811959
SELECT ...
```



# MARIADB

---



- Made by Monty Widenius
  - creator of MySQL and MyISAM
- **Advertised** as the *good* MySQL
  - as opposed to *bad* MySQL that's being “killed” by ~~Evil Corp~~ Oracle
- MariaDB 5.5 = MySQL 5.5 + XtraDB + MariaDB 5.3:
  - query optimizer improvements
  - multi-master replication
  - group commit fix
- MariaDB 10
  - real *fork* of MySQL 5.6
  - may become “backwards” incompatible

# MARIADB 10 FEATURES

---



- Parallel replication
- Multi-source replication
- Cassandra, Spider, TokuDB storage engines
- and few others.

# SWITCH TO MARIADB?

---




- Think twice before switching
- No great advantages over MySQL 5.7 or Percona Server
  - And yet a number of performance improvements from 5.7 will be missed in MariaDB 10
- Some features may be appealing, if that's your case - go for it
- Backwards incompatibility is a bit worrying
  - full mysqldump->import should still be possible



# WEBSALESQL


---



- Collaboration by a few heavy-duty MySQL users:
  - Alibaba, Facebook, Google, Linkedin & Twitter
- MySQL 5.6 at its core
  - good stuff being back-ported from MySQL 5.7
- Has very special mission
  - is not meant to be general purpose MySQL Server
  - addresses needs of running MySQL at **SCALE**

# FEW WEBSCALESQL FEATURES

---




- Ability to specify millisecond timeouts
- Super read-only mode
- Ability to disable deadlock detection
- Prefix index query optimization
- Performance Schema not compiled in
- InnoDB flushing performance fixes

read [this article by Laurynas @ Percona](#) to learn more.

# SWITCH TO WEBSCALESQL?

---



- Think thrice
- Only switch if you REALLY understand what it is
- Great performance features that don't *hurt* functionality will be ported to other variants (some already are)



# THE ESSENTIALS

---

*Of MySQL Configuration Tuning*

# MYSQL DEFAULTS ARE POOR

---

- Okay for development, not for busy production db
- Good news:
  - MySQL 5.6 defaults are better than versions < 5.6
  - MySQL 5.7 defaults are better than 5.6
  - MySQL is super easy to configure

# MYSQL CONFIGURATION FILE

---

- One configuration file (unless it contains *include* directives)
- Location:
  - /etc/my.cnf - most systems
  - /etc/mysql/my.cnf - Debian style
  - On Windows - best use data directory (create the file)



# COMMON MISTAKES

---

- Using Trial and Error approach
  - change something, see if it *feels* better
  - don't do it
- Asking Google for performance advice
  - answers often lack context
  - use case could be benchmarks - often non-production suitable
  - settings are hardware (or else) dependant
- Obsessing about fine-tuning the my.cnf
  - 10-15 variables is often all you need to change
  - fine-tuning won't give you significant wins

# COMMON MISTAKES (2)

---

- Changing many things at once
  - makes it very hard to figure out what caused which effect
  - instead, change one thing, then *measure* impact
- Not keeping my.cnf in sync with the changes you make
  - changing settings online is convenient, but...
  - don't forget to update my.cnf
- Redundant entries in my.cnf
  - MySQL won't mind them
  - Last value will be used
  - “-” and “\_” can be used interchangeably

# COMMON MISTAKES (3)

---

- Multiplying buffer sizes
  - don't do it
  - some buffers are local, some server-wide
  - few variables need to be increased after hardware upgrades
- Using the wrong my.cnf section
  - [mysql], [client], [mysqld\_safe] — all are incorrect choices
  - [mysqld] — put ALL of the server configuration here

# CHANGING CONFIGURATION ONLINE

---

- Many things can be changed online (*dynamically*)
- Even *innodb\_buffer\_pool\_size* starting with MySQL 5.7
- Example of changing *innodb\_thread\_concurrency* online:

```
mysql> show global variables like 'innodb%';
```

Variable_name	Value
...	
innodb_thread_concurrency	0
...	

```
140 rows in set (0.02 sec)
```

```
mysql> set global innodb_thread_concurrency = 8;  
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> select @@global.innodb_thread_concurrency;
```

@@global.innodb_thread_concurrency
8

```
1 row in set (0.00 sec)
```



# GLOBAL -VS- LOCAL SCOPE

---

- In many cases, you only want to change local session buffers and leave global configuration as is
  - *(sort|join|read|read\_rnd)\_buffer\_size* are all good examples
- So:
  - For a query that needs to sort a lot of data, before you run it:

```
set sort_buffer_size = 64 * 1024 * 1024;
```

- Some queries would find index merge intersect optimization harmful, so for such queries you can just:

```
set optimizer_switch = 'index_merge_intersection=off';
```

# 17

---

## *Key MySQL Settings*

# READY TO USE MY.CNF WITH ALL SETTINGS

---

I have prepared a my.cnf with short handy descriptions near each variable and appropriate links to learn more for your convenience. Download it here:

<http://www.speedemy.com/17>

# 1. DEFAULT\_STORAGE\_ENGINE — CHOOSE THE RIGHT ENGINE FIRST

---

- Already on InnoDB? Skip to next variable
- Otherwise, bear with me.

# STORAGE ENGINE WHAT?

---

- MySQL uses pluggable storage engines since the beginning
- Different storage engines store data in different ways
  - some may even store it in a remote server (e.g. Federated)
  - others may only pretend they are storing data (e.g. Blackhole)
- Two most commonly used ones:
  - MyISAM
  - InnoDB
- Few other famous SEs:
  - Blackhole
  - Archive
  - CSV
  - TokuDB



# WHAT'S WRONG WITH MYISAM

---

- Non transactional:
  - no **A**tomicity
  - no **C**onsistency
  - no **I**solation
  - no **D**urability
- Table level locks only
- Not scalable when it comes to write intensive workload
  - with an exception of INSERT-only workload

# MEET INNODB

---

- Shipped with MySQL since 2001
- Fully ACID transactional storage engine, designed to handle highly concurrent workload and scale.
- Road to multi-core multi-disk systems was bumpy
  - handles high concurrency pretty well since MySQL 5.0.30
  - is hard to beat in many ways nowadays
- Default MySQL storage engine since version 5.5.5

# HERE'S STORAGE ENGINES YOU ARE USING

---

- Run this on your server:

```
mysql> SELECT engine,
count(*) as TABLES,
concat(round(sum(table_rows)/1000000,2),'M') rows,
concat(round(sum(data_length)/(1024*1024*1024),2),'G') DATA,
concat(round(sum(index_length)/(1024*1024*1024),2),'G') idx,
concat(round(sum(data_length+index_length)/(1024*1024*1024),2),'G') total_size,
round(sum(index_length)/sum(data_length),2) idxfrac
FROM information_schema.TABLES
WHERE table_schema not in ('mysql', 'performance_schema', 'information_schema')
GROUP BY engine
ORDER BY sum(data_length+index_length) DESC LIMIT 10;
```

engine	TABLES	rows	DATA	idx	total_size	idxfrac
InnoDB	181	457.58M	92.34G	54.58G	146.92G	0.59
MyISAM	13	22.91M	7.85G	2.12G	9.97G	0.27

2 rows in set (0.22 sec)

\* no need to rewrite this, it will be included in the ebook you will get at [speedemy.com/17](http://speedemy.com/17)

# HERE'S MYISAM TABLES ON YOUR SERVER

---

- Run this:

```
mysql> SELECT
    concat(table_schema, '.', table_name) tbl,
    engine,
    concat(round(table_rows/1000000,2),'M') rows,
    concat(round(data_length/(1024*1024*1024),2),'G') DATA,
    concat(round(index_length/(1024*1024*1024),2),'G') idx,
    concat(round((data_length+index_length)/(1024*1024*1024),2),'G') total_size,
    round(index_length/data_length,2) idxfrac
FROM information_schema.TABLES
WHERE table_schema not in ('mysql', 'performance_schema', 'information_schema')
    AND engine = 'MyISAM'
ORDER BY data_length+index_length DESC;
```

# CONVERTING TO INNODB

---

- It's not enough to change *default-storage-engine*
- Tables need to be converted. One. By. One.
- Simply converting largest tables is not enough:
  - if at least one table in a join is MyISAM, the entire query is using table level locks. So having even a small MyISAM table in a large join can be very bad for concurrency
- Convert only when the server is configured for InnoDB properly



# CONVERTING EACH TABLE TO INNODB

---

- This will give you a list of commands to run:

```
SET @DB_NAME = 'your_database';
```

```
SELECT CONCAT('ALTER TABLE `', TABLE_NAME, '` ENGINE=InnoDB;') AS sql_statements  
FROM INFORMATION_SCHEMA.TABLES AS tb  
WHERE TABLE_SCHEMA = @DB_NAME  
AND ENGINE = 'MyISAM'  
AND TABLE_TYPE = 'BASE TABLE'  
ORDER BY TABLE_NAME DESC;
```

- To use InnoDB even when MyISAM is specified as a storage engine, in Percona Server you can set this in my.cnf:

```
enforce_storage_engine = InnoDB
```

## 2. INNODB\_BUFFER\_POOL\_SIZE — GET THE BEST OUT OF YOUR MEMORY


---

- Most important variable for InnoDB
- InnoDB Buffer Pool is:
  - a cache for read data (these are stored in pages of 16kb)
  - also cache for indexes, modified (dirty) data
  - and place for many internal InnoDB structures

# QUITE A TYPICAL BUFFER POOL DISTRIBUTION

.....

InnoDB Buffer Pool



Innodb_buffer_pool_pages_data	Min: 3791688	Max: 3796394	Avg: 3792557	StdDev: 494	Upper75: 3792852
Innodb_buffer_pool_pages_dirty	Min: 0	Max: 34501	Avg: 7080	StdDev: 7177	Upper75: 11186
Innodb_buffer_pool_pages_free	Min: 7596	Max: 8608	Avg: 8172	StdDev: 62	Upper75: 8191
Innodb_buffer_pool_pages_misc	Min: 258955	Max: 263115	Avg: 262496	StdDev: 490	Upper75: 262789

# INNODB\_BUFFER\_POOL\_SIZE = ?

---

- On a dedicated server, set to 80% of total memory
  - leaves room for other internal structures outside the buffer pool
  - query execution (session buffers)
  - OS cache (binary logs, relay logs, innodb transaction log, etc.)
  - OS memory structures
  - even on servers with 256-512GB of RAM
- Avoid swapping at all costs!
  - swapping is not the same as reading data from disk
  - it's much much worse

# INNODB\_BUFFER\_POOL\_SIZE = ?

---

- On a shared server:
  - check total size of db, maybe it all fits in, say, 50% of RAM?
  - don't strive to make it perfect
  - again, avoid swapping at all costs (use *vmstat 1* to check)
- Number of random reads from disk into BP can be checked by:

```
$ mysqladmin ext -ri1 | grep Innodb_buffer_pool_reads
| Innodb_buffer_pool_reads      | 1832098003 |
| Innodb_buffer_pool_reads      | 595        |
| Innodb_buffer_pool_reads      | 915        |
| Innodb_buffer_pool_reads      | 734        |
| Innodb_buffer_pool_reads      | 622        |
| Innodb_buffer_pool_reads      | 710        |
| Innodb_buffer_pool_reads      | 664        |
...
```

# CHANGING INNODB\_BUFFER\_POOL\_SIZE

---

- On MySQL 5.7, can be done online:

```
mysql> set global innodb_buffer_pool_size = size_in_bytes;
```

- MySQL log will go like:

```
[Note] InnoDB: Resizing buffer pool from 134217728 to 21474836480. (unit=134217728)
[Note] InnoDB: disabled adaptive hash index.
[Note] InnoDB: buffer pool 0 : 159 chunks (1302369 blocks) were added.
[Note] InnoDB: buffer pool 0 : hash tables were resized.
[Note] InnoDB: Resized hash tables at lock_sys, adaptive hash index, dictionary.
[Note] InnoDB: Completed to resize buffer pool from 134217728 to 21474836480.
[Note] InnoDB: Re-enabled adaptive hash index.
```

- Don't forget to update my.cnf after this
- On any previous version of MySQL:
  - set an appropriate *innodb\_buffer\_pool\_size* in *my.cnf*
  - restart MySQL server



### 3. INNODB\_LOG\_FILE\_SIZE — ROOM FOR MYSQL'S REDO LOG

---

- Size for redo (a.k.a. transaction) logs
- 5MB up until MySQL 5.6.8 — way too small for just about any workload
  - Current default of 48MB still too small for most
- What's innodb log file?
  - next slide, please

# WHAT IS INNODB REDO LOG?

---

- Transactional storage engines (including InnoDB) need both undo and redo functions, similar to image editor
  - well not exactly; let me explain
- Undo happens when rollback occurs
  - Implicitly or explicitly
- Redo happens after a database crash
  - Also when preparing a backup and other similar cases

# UNDO

---

- When record is changed (but before COMMIT)
  - changes are not written to data files right away
  - first they are written to transaction log
  - and data is modified in memory
  - original *unmodified* copy is moved to *rollback segment*
- if Rollback occurs:
  - Undo needs to happen, to restore data to previous state
  - InnoDB removes the copy from the rollback segment (and copies it back if need be), removes the dirty page
  - and marks in transaction log that change was rolled back
- That's Undo.

# REDO

---

- After COMMIT, changes are ready to be written to data files
  - but they aren't actually written (yet). This would be inefficient.
- Instead, changes are written to Redo log (and modified pages are only stored in memory)
- If MySQL crashes (before changes are written to data files):
  - Redo needs to happen
  - else, these changes would be lost forever
- After restart, InnoDB finds the last checkpoint position
  - and re-applies the same changes it had in memory before restart
- It is a gross oversimplification, but yeah, that's Redo.

# REDO LOG SIZE

---

- Size matters:
  - Small log files make writes slower and crash recovery faster
  - Large log files make writes faster and crash recovery slower
- Log file is like a buffer, so:
  - small log files = small buffer, so flushing needs to happen often
  - big log files = big buffer, hence flushing is more streamlined
- But:
  - with large log files, in case of crash, more work needs to be done to restore the database to the consistent state (i.e. the Redo operation)

# SIZING REDO LOG

---

- Rule of Thumb:
  - Check that total size of your Redo logs fits in 1-2h worth of writes during your busy period

```
mysql> pager grep seq
mysql> show engine innodb status\G select sleep(60); show engine innodb status\G
Log sequence number 1777308180429
...
Log sequence number 1777354541591

mysql> nopager
mysql> select (1777354541591-1777308180429)*60/1024/1024;
+-----+
| (1777354541591-1777308180429)*60/1024/1024 |
+-----+
| 2652.80696869 |
+-----+
1 row in set (0.00 sec)
```

- Based on this 60s sample, InnoDB *could be* writing around 2.6GB per hour, so *innodb\_log\_file\_size=2560M* sounds like a good start to get 5GB worth of redo logs total

# CHANGING THE REDO LOG SIZE

---

- MySQL  $\geq$  5.6 - change my.cnf and restart the server.
- MySQL  $<$  5.6, my.cnf change is **not enough**:
  1. change innodb\_log\_file\_size in my.cnf
  2. stop MySQL server
  3. ensure MySQL had a clean shutdown (mysql log is your friend)
  4. remove old log files, *usually* by running the following command: `rm -f /var/lib/mysql/ib_logfile*`
  5. start MySQL server – it should take a bit longer to start because it is going to be creating new transaction log files
- MySQL  $<$  5.6.2 only supports 4GB total!



## 4. INNODB\_FLUSH\_LOG\_AT\_TRX\_COMMIT – DURABLE OR NOT? THAT IS THE QUESTION!

.....

- *innodb\_flush\_log\_at\_trx\_commit=1* by default:
  - FLUSH and SYNC after EVERY transaction commit - full durability
  - INSERT/UPDATE/DELETE is a *transaction* if *autocommit=1* (default behaviour)
- SYNC is often expensive - it's synchronisation to disks
  - Exception is if you have non-volatile cache
    - e.g. battery-backup unit (BBU) protected write-back cache
    - or Super-capacitor with a flash memory chip
- Alternatives values for *innodb\_flush\_log\_at\_trx\_commit*:
  - 0 means FLUSH the buffers to OS, but DO NOT SYNC (no actual IO is performed on commit)
  - 2 means DON'T FLUSH and DON'T SYNC (again no actual IO is performed on commit)

# INNODB\_FLUSH\_LOG\_AT\_TRX\_COMMIT = ?

---

- When 0 or 2 is chosen, SYNC is performed once per second:
  - Means you may lose up to 1s worth of committed data
  - Once per second is not guaranteed\*
- *innodb\_flush\_log\_at\_trx\_commit=1* is full durability
  - Required for bank transactions and similar financial operations
  - Many websites use 0 or 2 instead
    - After all, MyISAM would lose up to 30s worth of data in case of crash and it's been default for many many years
- So then, 0 or 2?
  - Small difference, because neither causes SYNC to disk
  - 0 is good in that no data is lost if MySQL crashes (but the machine stays ON)

\* I've seen SYNC delays due to mutex contention issues

## 5. SYNC\_BINLOG – THAT'S FOR DURABLE BINLOG

---

- *sync\_binlog=1* makes binary logs durable
- Therefore:
  - if you have no slaves & no backups - use *sync\_binlog=0* (default)
  - if you do have replication and/or backups, but you don't mind losing few events (on slave or otherwise) in case of server power loss in order to gain better performance, *sync\_binlog=0*
  - if consistency is really important and you also use *innodb\_flush\_log\_at\_trx\_commit=1* anyway, do use *sync\_binlog=1* and make sure you run MySQL 5.6+, Percona Server 5.6+ or MariaDB 5.5+ as these versions have a binlog group commit fix (meaning that sync calls are grouped together)

## 6. INNODB\_FLUSH\_METHOD – AVOID DOUBLE BUFFERING

---

- Set *innodb\_flush\_method=O\_DIRECT*
  - Only supported in Linux
  - Overcomes OS cache for reads and for writes
- If not used, double buffering occurs
  - Bad because memory is wasted on storing the same data in memory twice (InnoDB buffer pool and OS cache)
- Very few exceptions to this
  - If you're not sure, you're not an exception.

## 7. INNODB\_BUFFER\_POOL\_INSTANCES – REDUCE MUTEX CONTENTION

---

- Introduced in MySQL 5.5
- Reduces global buffer pool mutex contention
  - splits buffer pool into multiple buffer pools
- On MySQL 5.5 (including variants) be more conservative:
  - *innodb\_buffer\_pool\_instances=4* should be good enough
- On MySQL 5.6+ (including variants):
  - *innodb\_buffer\_pool\_instances=8* or even *16* is a safe bet

## 8. INNODB\_THREAD\_CONCURRENCY — CONTROL YOUR THREADS

---

- *innodb\_thread\_concurrency=0* is default and often used in benchmarks, but
  - with high workloads, setting a limit *may* work much better
- How does it work? I wish I could show you, but let me at least describe it.
  - next slide, please

# INNODB CONCURRENCY MECHANISM

---

- *innodb\_thread\_concurrency* controls how many threads can be executing in parallel
  - if 0, all requests will be served immediately
    - that's fine if you have 32 CPU cores and 4 requests
    - not so if you have 32 CPU intensive requests and 4 CPU cores
    - can become a mess when all are executing at the same time and new requests keep coming in
- number > 0 caps the number of threads that are executing at the same time
  - But it's not a simple FIFO queue
  - Each request is given a certain number of tickets
    - 500 by default on MySQL 5.5 and earlier
    - 5000 by default on MySQL 5.6 and newer

# INNODB CONCURRENCY MECHANISM (2)

---

- A thread waits in the queue for a slot to become available
  - once it starts executing, it starts using the tickets
  - one ticket is used for every row read, insert, update, etc.
  - when all tickets are used, a thread is sent to the back of the queue and the cycle repeats until operation is complete
- The advantage is that long running queries don't prevent quick queries from ever getting into the queue, even if that does prolong their execution.
- And of course that there's no fighting for resources



# INNODB CONCURRENCY CONTROLS

---

- *innodb\_concurrency\_tickets* determines the number of tickets given
  - increase for long queries to run longer before letting others in
- *innodb\_thread\_sleep\_delay* sets amount of sleep time before joining the innodb queue (in microseconds)
- *innodb\_thread\_concurrency* sets how many slots are available for execution
  - use if you have spikes or hardware is often saturated
  - set to 8 and go up until you see good hardware utilization
  - it is a dynamic variable:
    - set global innodb\_thread\_concurrency=8;

# MONITORING THE QUEUE

---

- You can see the number of requests queued up and executing in the “*show engine innodb status\G*” output.
- Look for something like:

22 queries inside InnoDB, 104 queries in queue

## 9. SKIP\_NAME\_RESOLVE — DO SKIP THAT REVERSE IP LOOKUP

---

- Add skip\_name\_resolve to avoid DNS resolution on connect
- No impact when all is working fine
- When DNS server fails, takes a long time to figure things out
  - slow connections due to DNS failure don't help to solve this faster
- Exception: local hosts file based names



# HOW ABOUT A DEEP BREATH

---

*and a quick stretch ?*

**BETTER?**

**OK, LET'S CONTINUE**

## 10. INNODB\_IO\_CAPACITY(\_MAX)? – CAP INNODB IO USAGE

---

- *innodb\_io\_capacity* controls how many write IO requests per second (IOPS) will MySQL issue when flushing the dirty data
- *innodb\_io\_capacity\_max* controls how many write IOPS will MySQL issue flushing the dirty data when it's under stress
- IO activity related to background writes only
- Under stress means MySQL is behind with flushing activity and needs to shift gears or things may go bad
- Set *innodb\_io\_capacity* to 50-75% of write capacity
  - and *innodb\_io\_capacity\_max* to 100% or close
  - write capacity = number of random write iops server can handle ([more on it here](#))

# 11. INNODB\_STATS\_ON\_METADATA – TURN THEM OFF!

---

- OFF by default on MySQL 5.6 and 5.7
- Safe to switch OFF on MySQL 5.5 and 5.1 too, so in my.cnf:
  - `innodb_stats_on_metadata = 0`
  - can be changed online too
- Makes “show table status” and some queries against INFORMATION\_SCHEMA faster.
- InnoDB stats are still going to be updated, don't worry!

## 12. INNODB\_BUFFER\_POOL\_(DUMP\_AT\_SHUTDOWN|LOAD\_AT\_STARTUP)

---

- Makes warm-up much faster
  - works even with SSDs
  - by loading contents of buffer pool on server startup
- Only page reference numbers are stored in a file
- Asynchronous activity, no direct performance impact
- *innodb\_buffer\_pool\_dump\_pct* in MySQL 5.7 to control how much of the buffer pool to dump (25 by default, I recommend 75-100)
- Supported in MySQL since 5.6, in Percona Server - since 5.1



## 13. INNODB\_ADAPTIVE\_HASH\_INDEX\_PARTS – SPLIT THE AHI MUTEX

---

- Adaptive Hash Index is ON by default
- Dynamic Hash index maintained by InnoDB to improve certain query patterns. Usually very helpful
  - except when requests for mutex start backfiring
- Starting with MySQL 5.7 mutex can be split (and is by default split into 8 partitions), i.e.
  - *innodb\_adaptive\_hash\_index\_parts=8*
- Using MySQL 5.6 and earlier?
  - Switch to an appropriate Percona Server or MariaDB version and use *innodb\_adaptive\_hash\_index\_partitions*

## 14. QUERY\_CACHE\_TYPE – ON? OFF? ON DEMAND?

---

- Before it becomes a bottle-neck, query cache is great with:
  - small databases with few updates
  - low concurrency workload
  - read-only databases
- Even if it's helpful, don't exceed *query\_cache\_size=256M*
  - wait time on invalidation increases significantly
  - innocent queries get blocked
- With high concurrency workload, often a bottle-neck
- Domas Mituzas suggests using this query cache tuner (next slide, please)

## Query cache tuner

Optimal size for your query cache: **0**

# QUERY CACHE TUNING

---

- More seriously though, mutex is still locked even if *query\_cache\_size=0*
- Use the following configuration:
  - *query\_cache\_size=0*
  - *query\_cache\_type=OFF*
- Requires a MySQL restart to disable mutex
- Works with MySQL 5.5 or newer
  - If using MySQL 5.1, switch to Percona Server 5.1 to get the same effect

## 15. INNODB\_CHECKSUM\_ALGORITHM – THE SECRET HARDWARE ACCELERATION TRICK

---

- Old checksum algorithm is expensive (CPU overhead)
- `crc32` can use native CPU instructions. That's a YUUGE win.
- Use `innodb_checksum_algorithm=crc32` with MySQL 5.6
  - used by default in MySQL 5.7
- Safe to change, no need to reload data
- Can be changed online in fact

## 16. TABLE\_OPEN\_CACHE\_INSTANCES – IT'S THERE FOR A REASON

---

- Introduced in MySQL 5.6.6 to split open cache instances
- Avoid server lock-up when opening many tables or when opening a table is slow
- Starting with MySQL 5.7.8, *table\_open\_cache\_instances=16*
  - Set this manually in MySQL 5.6

## 17. INNODB\_(READ|WRITE)\_IO\_THREADS – LAST AND, YES, LEAST

---

- MySQL supports Asynchronous IO on Linux since MySQL 5.5, so this is not as important as it may seem.
- These threads are used for background activities only
- Set it to match number of bearing read/write disks
  - e.g. *innodb\_read\_io\_threads=8* and *innodb\_write\_io\_threads=4* on RAID10 with 8 disks.
  - on SSDs, set at 32/16 (or similar) respectively

# NOT THAT I DIDN'T MENTION IT, BUT

---

I have prepared a my.cnf with short handy descriptions near each variable and appropriate links to learn more for your convenience. Download it here:

<http://www.speedemy.com/17>



# FINAL THOUGHTS

---

- Configuration file is only part of the deal
- More often than not, real the devil is in queries
- To learn about query optimization, register for a free webinar on query optimization:

<http://www.speedemy.com/webinars>

# END



Speedemy



Speedemy\_com



Speedemy.com

*P.S. I have prepared a video where I have discussed MySQL configuration in great detail. It also contains an additional section about reading MySQL Status Counters, in case you're interested. [Register here](#) to get the free video once post production is finished.*