

MySQL Server 8.0

COSCUP 2017 in Taiwan

created: 2017/08/02

MySQL Global Business Unit

Master Principal Sales Consultant/Shinya Sugiyama

Safe Harbor Statement

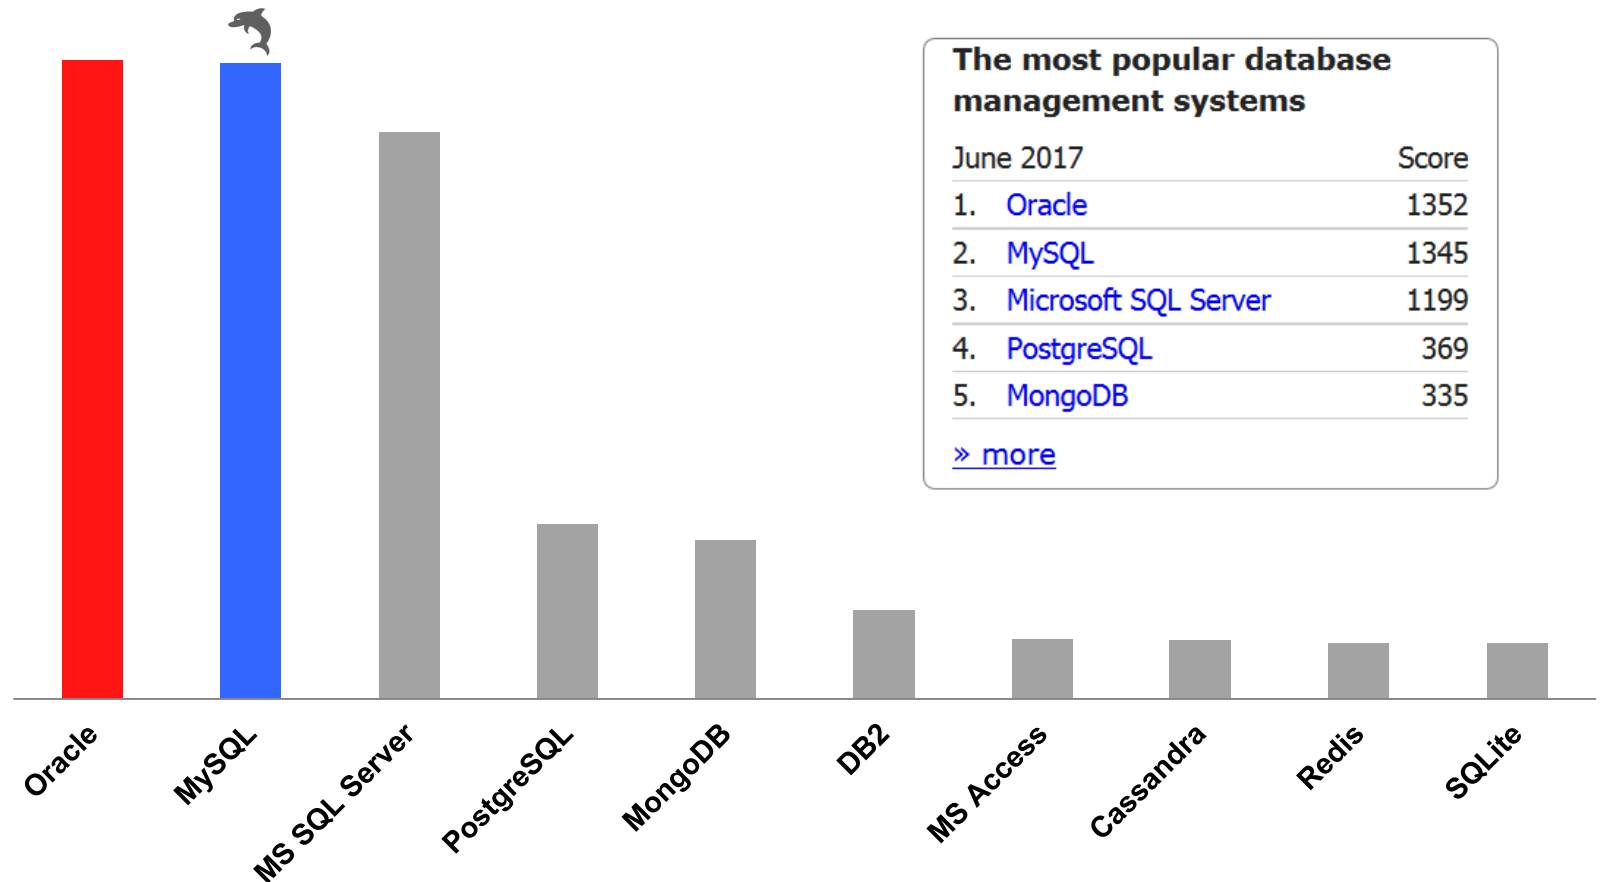
The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

The world's one of the most popular open source database

<http://db-engines.com/en/>

The DB-Engines Ranking ranks database management systems according to their popularity. The ranking is updated monthly.

https://db-engines.com/en/ranking_definition



In Global

Many Web Site Use MySQL as their default database.



In Japan

Many Web Companies also use MySQL for their backed Database.

<http://www.alexa.com/topsites/countries/JP>

Company	URL
Square Enix	http://www.jp.square-enix.com/
Oricon ME	http://www.oricon.co.jp/
forTravel	http://4travel.jp/
Sony	https://www.sony.net/

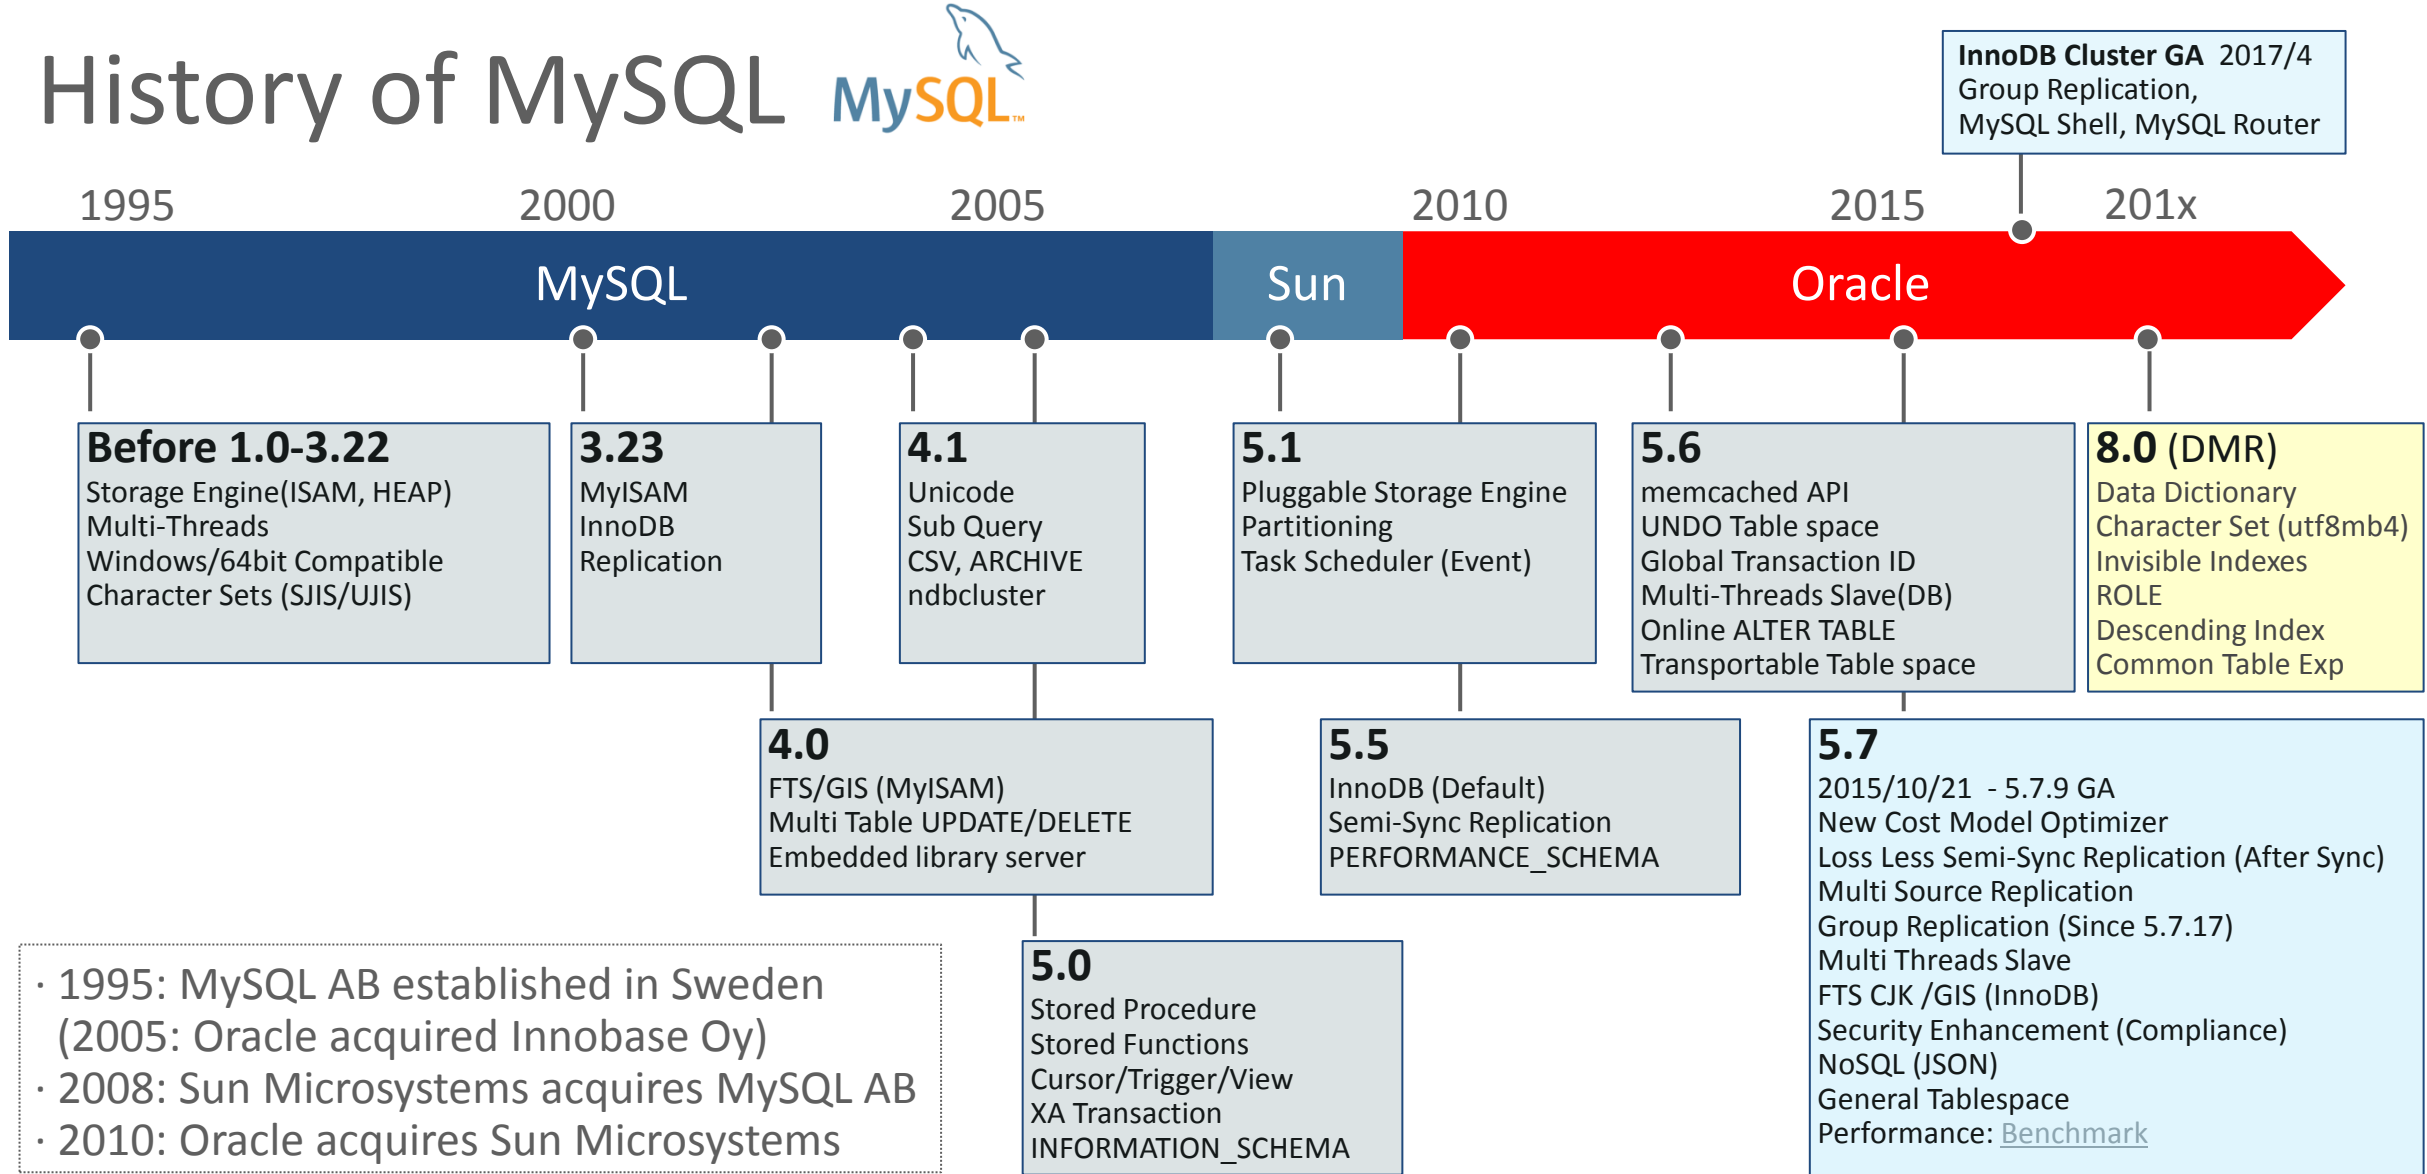
分野	OSS	使用率
OS	Linux 系	67.3%
	BSD 系	12.9%
RDBMS	MySQL	53.1%
	PostgreSQL	35.0%
アプリケーションサーバー	Tomcat	35.6%
	JBoss	12.0%
システム運用管理	Zabbix	16.2%
	Nagios	7.1%
	Chef	3.9%
	Hinemos	1.9%
システムソフトウェア	Samba	21.4%
	BIND	13.6%
ハイパーバイザー	Xen	16.2%
	KVM	10.7%
クラウド基盤	OpenStack	6.1%
	Docker	4.5%
	CloudStack	3.6%
	Cloud Foundry	2.9%
データ分散処理	Hadoop	6.8%
	Spark	1.3%
NoSQL	MongoDB	4.5%
	Scalaris	4.2%
	Cassandra	2.6%
	Hypertable	2.6%

n=309

OSS usage ratio in Japanese company that use OSS in their business (Source [IDC Japan](#), 2/2016)



History of MySQL




Looking Back MySQL 5.7 – Improvements across the board

- Replication
- InnoDB
- Optimizer
- Security
- Performance Schema
- GIS
- FTS (Chinese, Korean, Japanese)

- Triggers
- Partitioning
- **New! SYS Schema**
- **New! JSON**
- Performance

**200+ new
features
in total!**



Continuing the JSON Developer Experience

MySQL 8.0

- MySQL Document Store
 - Document collections
 - Relational tables
 - Combine them as you like
- Administer MySQL via the MySQL Shell
 - One stop DevOps tool for MySQL
 - Out of the Box HA
 - Use your preferred language: SQL, JavaScript, Python, ...
 - Relational or Document
- Additional JSON Functions
 - [JSON_ARRAYAGG\(\)](#), [JSON_OBJECTAGG\(\)](#), [JSON_PRETTY\(\)](#),
 - [JSON_STORAGE_FREE\(\)](#), [JSON_STORAGE_SIZE\(\)](#)

```
select json_pretty(body) from T_JSON_DOC where id = 1;G
***** 1. row *****

json_pretty(body): {
  "id": 1,
  "name": "",
  "price": 10000,
  "Conditions": [
    "NEW",
    2015,
    "Excellent"
  ]
}
```

WL#7987: JSON aggregation functions
<https://dev.mysql.com/worklog/task/?id=7987>
WL#9191: Add JSON_PRETTY function
<https://dev.mysql.com/worklog/task/?id=9191>
WL#9192: Add JSON_STORAGE_SIZE / JSON_STORAGE_FREE functions
<https://dev.mysql.com/worklog/task/?id=9192>

GIS

- Geography support
 - st_distance()
- Spatial Reference Systems (SRS) Support
- SQL/MM Information Schema views
- Standard compliant axis ordering in import/export functions
- Helper functions to manipulate and convert data:
 - st_x(geom, x)
 - st_y(geom, y)
 - st_srid(geom, srid)

```
mysql> SET @g1 = ST_GeomFromText('POINT(1 1)', 4326);  
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> SET @g2 = ST_GeomFromText('POINT(2 2)', 4326);  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> SELECT ST_Distance(@g1, @g2);
```

```
+-----+  
| ST_Distance(@g1, @g2) |  
+-----+  
|      156874.3859490455 |  
+-----+  
1 row in set (0.06 sec)
```



WL#9347: Ellipsoidal ST_Distance for point and multipoint

<https://dev.mysql.com/worklog/task/?id=9347>

WL#8606: Mutator ST_X and ST_Y

<https://dev.mysql.com/worklog/task/?id=8606>

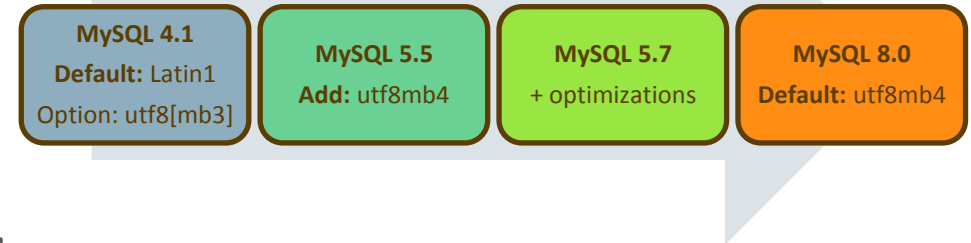
WL#8543: Mutator ST_SRID

<https://dev.mysql.com/worklog/task/?id=8543>

utf8mb4 as default character set

MySQL 8.0

- Support for the latest Unicode 9.0
- Accent (ai, as) and case sensitive (ci,cs) collations
- Language specific collations (including Japanese!)



```
mysql> select *,@@version from information_schema.COLLATIONS where CHARACTER_SET_NAME = 'utf8mb4' and IS_DEFAULT = 'YES';
```

COLLATION_NAME	CHARACTER_SET_NAME	ID	IS_DEFAULT	IS_COMPILED	SORTLEN	@@version
utf8mb4_general_ci	utf8mb4	45	Yes	Yes	1	5.7.18-enterprise-commercial

```
mysql> select *,@@version from information_schema.COLLATIONS where CHARACTER_SET_NAME = 'utf8mb4' and IS_DEFAULT = 'YES';
```

COLLATION_NAME	CHARACTER_SET_NAME	ID	IS_DEFAULT	IS_COMPILED	SORTLEN	PAD_ATTRIBUTE	@@version
utf8mb4_0900_ai_ci	utf8mb4	255	Yes	Yes	0	NO PAD	8.0.2-dmr

<http://www.unicode.org/>


WL#10818: Add utf8mb4 accent sensitive and case insensitive collation
<https://dev.mysql.com/worklog/task/?id=10818>
WL#7554: Switch to new default character set and change mtr test cases
<https://dev.mysql.com/worklog/task/?id=7554>

UTF-8

The character set for the Web

- UTF-8 is the dominating character set in today's applications
- Requires 1-4 bytes for storing characters
- Historically a performance problem

1 byte	Basic Latin letters, digits, and punctuation signs use.
2 byte	Most European and Middle East script letters, extended Latin letters
3 byte ~ 4 byte	Korean, Chinese, and Japanese



<https://en.wikipedia.org/wiki/UTF-8>

	UTF8(UTF8MB3)	UTF8MB4
🐬	X	○
鯨	X	○

<http://mysqlserverteam.com/mysql-8-0-collations-migrating-from-older-collations/>

MySQL 8.0 vs MySQL 5.7 utf8mb4


TPS

+300-350% in OLTP RO


+176-233% in OLTP RW

+1500-1800% in SELECT DISTINCT_RANGES


OLTP RO



OLTP RW



SELECT DISTINCT_RANGES




MySQL 8.0 utf8mb4 vs MySQL 5.7 utf8mb3


TPS

+270-340% in OLTP RO
+150-200% in OLTP RW
+1300-1600% in SELECT DISTINCT_RANGES


OLTP RO



OLTP RW



SELECT DISTINCT_RANGES



UUID and Bit-wise Improvements

- Functions to convert UUID to and from binary:
 - UUID_TO_BIN(), BIN_TO_UUID(), IS_UUID()
- Bit-wise operations on binary data types
 - Designed with IPv6 in mind: INET6_ATON(address) & INET6_ATON(network)

UUIDs are a good alternative to AUTO_INCREMENT PRIMARY KEY. But they also come with some disadvantages:

- increased storage: 36 characters
- more difficult to debug
- performance issues
- Size and not being ordered

```
mysql> select INET6_NTOA(INET6_ATON('59b0:c4d6:48b4:3717:f031:d05b:705d:6c65'));
+-----+
| INET6_NTOA(INET6_ATON('59b0:c4d6:48b4:3717:f031:d05b:705d:6c65')) |
+-----+
| 59b0:c4d6:48b4:3717:f031:d05b:705d:6c65 |
+-----+
```

```
mysql> SELECT INET6_NTOA(subnet(ip_address, net_len)) as range_from,
-> INET6_NTOA(subnet(ip_address, net_len) | host_mask(net_len)) as range_to FROM cidr LIMIT 2;
+-----+-----+
| range_from | range_to |
+-----+-----+
| d1b:7977:5915:437:6830:: | d1b:7977:5915:437:683f:ffff:ffff:ffff |
| 8337:1c78:f38f:2805:4648:38ab:2ec2:: | 8337:1c78:f38f:2805:4648:38ab:2ec3:ffff |
+-----+-----+
```

WL#8920: Improve usability of UUID manipulations
<https://dev.mysql.com/worklog/task/?id=8920>
WL#8699: Bit-wise operations on binary data types
<https://dev.mysql.com/worklog/task/?id=8699>

UUID_TO_BIN Optimization

- Binary format is now smaller and insert-order efficient:

From **VARCHAR(36)** 53303f87-78fe-11e6-a477-8c89a52c4f3b


To **VARBINARY(16)** 11e678fe53303f87a4778c89a52c4f3b

```
mysql> select bin_to_uuid(id_bin),id_var from T_UUID limit 1;
```


bin_to_uuid(id_bin)	id_var
lacebc1d-4f6d-11e7-a11b-0242ac110002	lacebc6d-4f6d-11e7-a11b-0242ac110002

```
mysql> select hex(id_bin),hex(id_var) from T_UUID limit 1;
```

hex(id_bin)	hex(id_var)
1ACEBC1D4F6D11E7A11B0242AC110002	31616365626336642D346636642D313165372D613131622D303234326163313130303032



MySQL 8.0: Better Handling of Hot Rows



```
SELECT seat_no FROM seats
JOIN seat_rows USING ( row_no )
WHERE seat_no IN (3,4)
AND seat_rows.row_no IN (12)
AND booked = 'NO'
FOR UPDATE OF seats SKIP LOCKED
FOR SHARE OF seat_rows NOWAIT;
```

Non deterministically skip over locked rows
Ex) skip orders which are pending

Error immediately if a row is already locked

<http://mysqlserverteam.com/mysql-8-0-1-using-skip-locked-and-nowait-to-handle-hot-rows/>

WL#3597: Implement NOWAIT and SKIP LOCKED
<https://dev.mysql.com/worklog/task/?id=3597>
WL#8919: InnoDB: Implement NOWAIT and SKIP LOCKED
<https://dev.mysql.com/worklog/task/?id=8919>

Common Table Expressions

- “With queries”
- Both Recursive and Non-Recursive Forms
- Simplifies writing complex SQL:

```
WITH t1 AS (SELECT * FROM city WHERE CountryCode = 'TWN')  
SELECT * FROM t1 limit 5;
```

ID	Name	CountryCode	District	Population
3263	Taipei	TWN	Taipei	2641312
3264	Kaohsiung	TWN	Kaohsiung	1475505
3265	Taichung	TWN	Taichung	940589
3266	Tainan	TWN	Tainan	728060
3267	Panchiao	TWN	Taipei	523850

5 rows in set (0.01 sec)

WL#3634: Recursive WITH (CTE)
<https://dev.mysql.com/worklog/task/?id=3634>
WL#883: Non-recursive WITH clause (CTE)
<https://dev.mysql.com/worklog/task/?id=883>

DBT3 Query 15 Top Supplier Query

Using view

```
CREATE VIEW revenue0 (supplier_no, total_revenue) AS
SELECT l_suppkey, SUM(l_extendedprice * (1-l_discount))
FROM lineitem
WHERE l_shipdate >= '1996-07-01'
      AND l_shipdate < DATE_ADD('1996-07-01',
                                INTERVAL '90' day)
GROUP BY l_suppkey;
```

rewrite

Using CTE

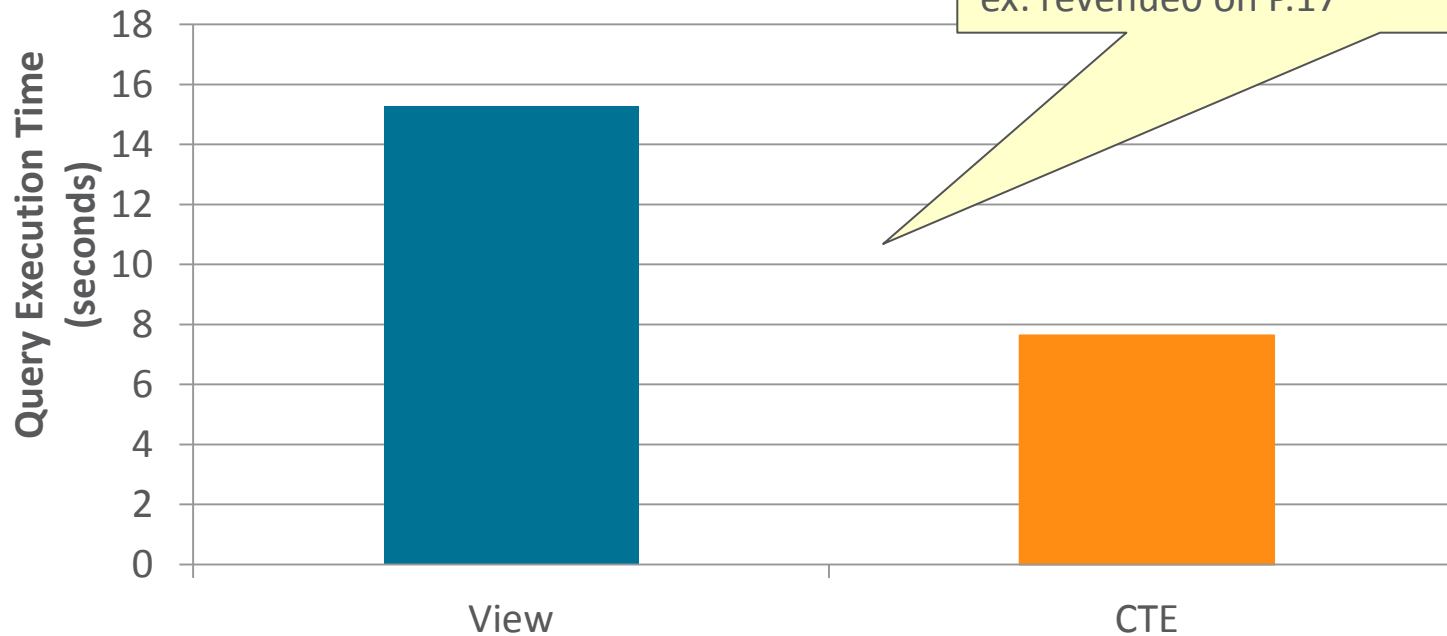
```
WITH revenue0 (supplier_no, total_revenue) AS
(SELECT l_suppkey, SUM(l_extendedprice * (1-l_discount))
FROM lineitem
WHERE l_shipdate >= '1996-07-01'
      AND l_shipdate < DATE_ADD('1996-07-01',
                                INTERVAL '90' day)
GROUP BY l_suppkey)
```

```
SELECT s_suppkey, s_name, s_address, s_phone, total_revenue
FROM supplier, revenue0
WHERE s_suppkey = supplier_no AND total_revenue = (SELECT MAX(total_revenue) FROM revenue0)
ORDER BY s_suppkey;
```

<http://mysqlserverteam.com/mysql-8-0-labs-recursive-common-table-expressions-in-mysql-ctes/>

DBT-3 Query 15

Query Performance



Another Example:

Derived table can not be referenced twice:

```
SELECT ...FROM (SELECT a, b, SUM(c) s FROM t1 GROUP BY a, b) AS d1 JOIN (SELECT a, b, SUM(c) s FROM t1 GROUP BY a, b) AS d2 ON d1.b = d2.a;
```

CTE can:

```
WITH d AS (SELECT a, b, SUM(c) s FROM t1 GROUP BY a, b) SELECT ... FROM d AS d1 JOIN d AS d2 ON d1.b = d2.a;
```


Recursive CTE

```
WITH RECURSIVE cte AS  
( SELECT ... FROM table_name /* "seed" SELECT */  
  UNION ALL  
  SELECT ... FROM cte, table_name) /* "recursive" SELECT */  
SELECT ... FROM cte;
```



Recursion

- A recursive CTE refers to itself in a subquery
- The “seed” SELECT is executed once to create the initial data subset, the recursive SELECT is repeatedly executed to return subsets of data until the complete result set is obtained.
- Useful to dig in hierarchies (parent/child, part/subpart)
- Similar to Oracle's CONNECT BY

Example: Recursive CTE

Print 1 to 10

```
WITH RECURSIVE qn AS  
(SELECT 1 AS a UNION ALL SELECT 1+a FROM qn WHERE a<10)  
SELECT * FROM qn;
```

a
1
2
3
4
5
6
7
8
9
10

Example: Recursive CTE

Hierarchy Traversal (List reporting chain)

```
mysql> select * from employees;
```

id	name	manager_id
29	Pedro	198
72	Pierre	29
123	Adil	692
198	John	333
333	Yasmina	NULL
692	Tarek	333
4610	Sarah	29

7 rows in set (0.00 sec)

```
mysql> WITH RECURSIVE
```

```
-> emp_ext (id, name, path) AS  
-> (SELECT id, name, CAST(id AS CHAR(200))  
-> FROM employees WHERE manager_id IS NULL  
-> UNION ALL SELECT s.id, s.name, CONCAT(m.path, ",", s.id)  
-> FROM emp_ext m JOIN employees s ON m.id=s.manager_id )  
-> SELECT * FROM emp_ext ORDER BY path;
```

id	name	path
333	Yasmina	333
198	John	333,198
29	Pedro	333,198,29
4610	Sarah	333,198,29,4610
72	Pierre	333,198,29,72
692	Tarek	333,692
123	Adil	333,692,123

7 rows in set (0.00 sec)

Window Functions

```
SELECT name, department_id, salary, SUM(salary)
OVER (PARTITION BY department_id) AS department_total
FROM employee ORDER BY department_id, name;
```

The **OVER** keyword signals a window function.
PARTITION == disjoint set of rows in result set.

name	department_id	salary	department_total
Ed	10	100000	370000
Fred	10	60000	370000
Lebedev	20	65000	130000
Pete	20	65000	130000
Jeff	30	300000	370000
Will	30	70000	370000

Ranking Windows Functions:

RANK, DENSE_RANK, PERCENT_RANK, CUME_DIST, ROW_NUMBER

Analytical Windows Functions:

NTILE, LEAD, LAG, FIRST_VALUE, LAST_VALUE, NTH_VALUE

Aggregates Windows Function:

COUNT, SUM, AVG, MAX, MIN

<https://dev.mysql.com/doc/refman/8.0/en/window-functions-usage.html>

<https://www.slideshare.net/DagHWanvik/sql-window-functions-for-mysql>

WL#9236: Add first batch of SQL window functions

<https://dev.mysql.com/worklog/task/?id=9236>

WL#9727: Additional aggregate window functions

<https://dev.mysql.com/worklog/task/?id=9727>

WL#9603: Add remaining non-aggregate window functions

<https://dev.mysql.com/worklog/task/?id=9603>

Window Functions : OVER

over_clause: {OVER (window_spec) | OVER window_name}

```
mysql> select * from t;
```

i
1
2
3
4

```
mysql> SELECT i,(SELECT SUM(i) FROM t) FROM t;
```

i	(SELECT SUM(i) FROM t)
1	10
2	10
3	10
4	10

The **OVER** keyword signals a window function.
If OVER() is empty, the window consists of all query rows and the window function computes a result using all rows.

```
mysql> SELECT i, SUM(i) OVER() AS sum FROM t;
```

i	sum
1	10
2	10
3	10
4	10

<http://mysqlserverteam.com/mysql-8-0-2-introducing-window-functions/>

Window Functions: PARTITION

partition_clause: A PARTITION BY clause indicates how to divide the query rows into groups.

```
mysql> select * from sales;
```

employee	date	sale
odin	2017-03-01	200
odin	2017-04-01	300
odin	2017-05-01	400
thor	2017-03-01	400
thor	2017-04-01	300
thor	2017-05-01	500

PARTITION == disjoint set of rows in result set.
employee (Odin, Thor)

```
mysql> select employee, SUM(sale)
-> FROM sales GROUP BY employee;
```

employee	SUM(sale)
odin	900
thor	1200

```
mysql> select employee,date,sale,SUM(sale)
-> OVER (PARTITION BY employee) AS sum FROM sales;
```

employee	date	sale	sum
odin	2017-03-01	200	900
odin	2017-04-01	300	900
odin	2017-05-01	400	900
thor	2017-03-01	400	1200
thor	2017-04-01	300	1200
thor	2017-05-01	500	1200

Window Functions: RANK

```
SELECT name, dept_id AS dept, salary,  
       RANK() OVER w AS `rank`  
FROM employee  
   WINDOW w AS (PARTITION BY dept_id  
                ORDER BY salary DESC);
```

name	dept_id	salary	rank
Newt	NULL	75000	1
Ed	10	100000	1
Newt	10	80000	2
Fred	10	70000	3
Michael	10	70000	3
Jon	10	60000	5
Dag	10	NULL	6
Pete	20	65000	1
Lebedev	20	65000	1
Jeff	30	300000	1
Will	30	70000	2

Window Functions: ROW_NUMBER

```
SELECT name, dept_id AS dept, salary,  
       RANK() OVER w AS `rank`,  
       DENSE_RANK() OVER w AS dense,  
       ROW_NUMBER() OVER w AS `rowno`  
FROM employee  
   WINDOW w AS (PARTITION BY dept_id  
                 ORDER BY salary DESC);
```

name	dept_id	salary	rank	dense	rowno
Newt	NULL	75000	1	1	1
Ed	10	100000	1	1	1
Newt	10	80000	2	2	2
Fred	10	70000	3	3	3
Michael	10	70000	3	3	4
Jon	10	60000	5	4	5
Dag	10	NULL	6	5	6
Pete	20	65000	1	1	1
Lebedev	20	65000	1	1	2
Jeff	30	300000	1	1	1
Will	30	70000	2	2	2

Optimizer Cost Model

Improved to consider buffer pool fit

```
SELECT * FROM Country WHERE population > 200000000;
```

Model for a table scan:

pages in table *
(IO_BLOCK_READ_COST |
MEMORY_BLOCK_READ_COST)

records * ROW_EVALUATE_COST

= 25.4 100% in memory

= 29.9 100% on disk

Model for a range scan:

records_in_range *
(IO_BLOCK_READ_COST |
MEMORY_BLOCK_READ_COST)

records_in_range *
ROW_EVALUATE_COST + #
records_in_range *
ROW_EVALUATE_COST

= **22.5 100% in memory**

= 60 100% on disk

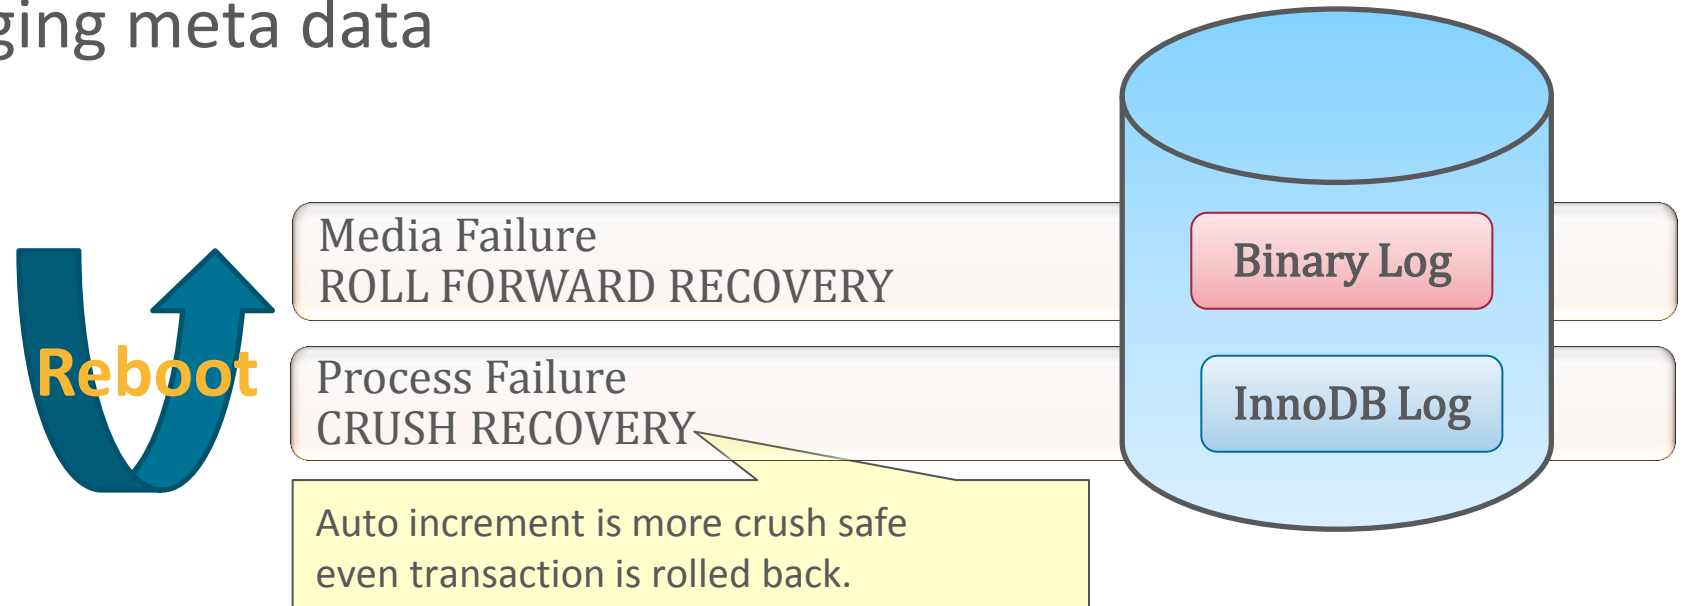
Model accounts for memory fit. For data on disk an IO block read defaults to 1.0. In memory defaults to 0.25.

Much larger performance difference for range scan not in memory (good)

WL#7093: Optimizer provides InnoDB with a bigger buffer
<https://dev.mysql.com/worklog/task/?id=7093>

InnoDB Auto Increment Persists

- First reported as BUG [#199](#)
- Auto increment counters are now written to the REDO log
- Allows for fast changing meta data



WL#6204: InnoDB persistent max value for autoinc columns
<https://dev.mysql.com/worklog/task/?id=6204>

Descending Indexes

For B+tree indexes

```
CREATE TABLE t1 (  
  a INT,b INT,  
  INDEX a_b (a DESC, b ASC)  
);
```

- In 5.7: Index in ascending order is created, server scans it backwards
- In 8.0: Index in descending order is created, server scans it forwards

Benefits:

- Forward index scan is faster than backward index scan
- Use indexes instead of filesort for ORDER BY clause with ASC/DESC sort key

WL#1074: Add Descending indexes support
<https://dev.mysql.com/worklog/task/?id=1074>

EX: Descending Indexes

```
mysql> explain select * from city2 order by city_id asc limit 3;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	city2	NULL	index	NULL	idx_asc_city_id	2	NULL	3	100.00	NULL

MySQL 5.5 ~ 5.7: can create ASC index Only.

```
mysql> explain select * from city2 order by city_id desc limit 3;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	city2	NULL	index	NULL	idx_asc_city_id	2	NULL	3	100.00	Backward index scan

```
mysql> explain select * from city2 order by city_id asc limit 3;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	city2	NULL	index	NULL	idx_asc_city_id	2	NULL	3	100.00	NULL

MySQL8.0 ~: can create both ASC and DESC index

```
mysql> explain select * from city2 order by city_id desc limit 3;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	city2	NULL	index	NULL	idx_desc_city_id	2	NULL	3	100.00	NULL

Invisible Indexes

- Indexes are “hidden” to the MySQL Optimizer
 - Not the same as “disabled indexes”
 - Contents are fully up to date and maintained by DML
- Two use cases:
 - Soft Delete (*Recycle Bin*)
 - Staged Rollout

WL#8697: Support for INVISIBLE indexes
<https://dev.mysql.com/worklog/task/?id=8697>

Soft Delete

Example Usage

- I don't think this index is used any more:
`ALTER TABLE Country ALTER INDEX c INVISIBLE;`
- I need to revert:
`ALTER TABLE Country ALTER INDEX c VISIBLE;`
- It is now safe to drop:
`ALTER TABLE Country DROP INDEX c;`

Misc01 Perform

You can confirm un-used Index by using sys schema...
But if not sure....

Report

- Top I/O by File by Time
- Top I/O by Event Category
- Top I/O in Time by Event Categories
- Top I/O Time by User/Thread
- ▼ High Cost SQL Statements
- Statement Analysis
- Statements in Highest 5 Percent by P
- Using Temp Tables
- With Sorting
- Full Table Scans
- Errors or Warnings
- Database Schema Statistics
- Schema Object Overview (High Over
- Schema Index Statistics
- Schema Table Statistics
- Schema Table Statistics (with InnoDe
- Tables with Full Table Scans
- Unused Indexes
- ▼ Wait Event Times (Expert)
- Global Waits by Time
- Waits by User by Time
- Wait Classes by Time
- Waits Classes by Average Time

List of indexes that were never used since the server started or since P_S data collection started.

Schema	Object	Index
NEWS7	Innovation_Day...	_id
NEWS7	jbooks	isbn
NEWS7	OC	_id
NEWS7	tbl_partition	idx_tbl_partition
NEWS7	tbl_partition	idx_tbl_partition_id
NEWS7	T_Generated_Co...	full_name
NEWS7	T_GIS	feature_type
NEWS7	T_GIS	feature_street
NEWS7	T_JSON	idx_feature_street
NEWS7	T_JSON	idx_feature_type
NEWS7	T_JSON_DOC	idx_json_price_v
NEWS7	T_ONLINE	idx_feature_street
NEWS7	T_SHORT_JSON	idx_feature_street
NEWS7	X_JSON	_id
NEWS7	x_posts	_id
NEWS7	X_PYTHON	_id
NEWS7	X_PYTHON	idx_name
NEWS7	X_PYTHON	ft_idx_text
ngram	articles	ngram_idx

Export... Copy Selected Copy Query Refresh

<https://github.com/mysql/mysql-sys>

Staged Rollout

- Adding any new index can change existing execution plans.
- All change introduces risk of regression
- Invisible indexes allows you to stage all changes
 - i.e. put the database in a “prepared” state
- Turn on changes at an opportune time

```
ALTER TABLE Country ADD INDEX c (Continent) INVISIBLE;  
# after some time  
ALTER TABLE Country ALTER INDEX c VISIBLE;
```

```

SELECT * FROM information_schema.statistics WHERE is_visible='NO';
***** 1. row *****
TABLE_CATALOG: def
TABLE_SCHEMA: world
TABLE_NAME: Country
NON_UNIQUE: 1
INDEX_SCHEMA: world
INDEX_NAME: c
SEQ_IN_INDEX: 1
COLUMN_NAME: Continent
COLLATION: A
CARDINALITY: 7
SUB_PART: NULL
PACKED: NULL
NULLABLE:
INDEX_TYPE: BTREE
COMMENT: disabled
INDEX_COMMENT:
IS_VISIBLE: NO

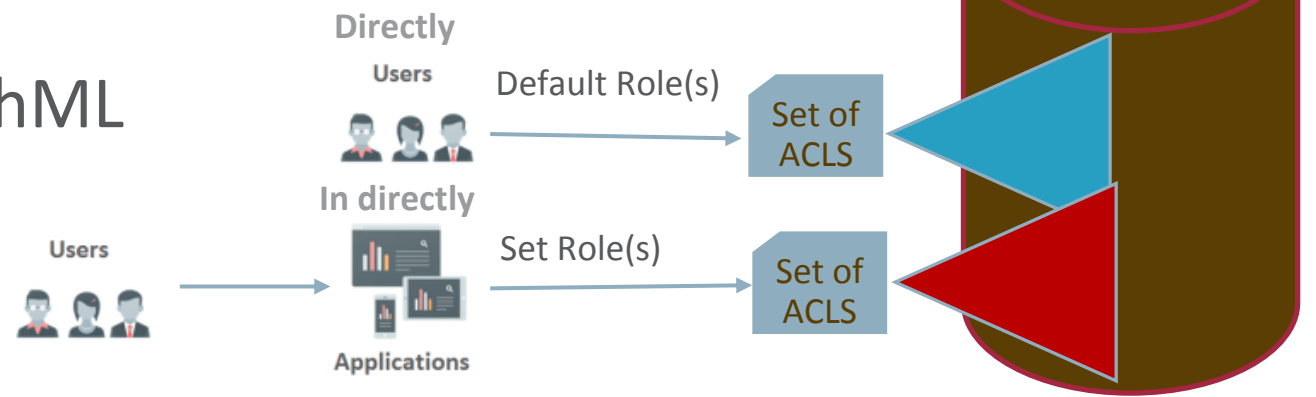
```

MySQL Roles

Improving MySQL Access Controls

- Introduced in the 8.0.0 DMR
- Easier to manage user and applications rights
- As standards compliant as practically possible
- Multiple default roles
- Can export the role graph in GraphML

```
mysql> select user(),current_role();
+-----+-----+
| user()          | current_role() |
+-----+-----+
| user01@localhost | `role80`@`%`   |
+-----+-----+
```



WL#988: Roles
<https://dev.mysql.com/worklog/task/?id=988>

Atomic ACL Statements

- Long standing MySQL issue!
 - For Replication, HA, Backups, etc.
- Possible now - ACL tables reside in InnoDB Data Dictionary
- Not just a table operation: memory caches need update too
- Applies to statements performing multiple logical operations, e.g.
 - CREATE USER u1, u2
 - GRANT SELECT ON *.* TO u1, u2
- Uses a custom MDL lock to block ACL related activity
 - While altering the ACL caches and tables

+	-----	+	-----	+	-----	+
	TABLE_SCHEMA		TABLE_NAME		ENGINE	
+	-----	+	-----	+	-----	+
	mysql		user		InnoDB	
+	-----	+	-----	+	-----	+

WL#9045: Make user management DDLs atomic
<https://dev.mysql.com/worklog/task/?id=9045>

InnoDB Redo and Undo Encryption

- AES 256 encryption
- *Encrypted when redo/undo log data is written to disk*
- Decryption occurs when redo/undo log data is read from disk
- Once redo/undo log data is read into memory, it is in unencrypted form.
- Two tiered encryption – like Innodb tablespace encryption
 - Fast key rotation, high performance
- *Easy to use*
 - Enabled using [innodb redo log encrypt](#) and [innodb undo log encrypt](#)

WL#9289: InnoDB: Support Transparent Data Encryption for Undo Tablespaces

<https://dev.mysql.com/worklog/task/?id=9289>

WL#9290: InnoDB: Support Transparent Data Encryption for Redo Log


<https://dev.mysql.com/worklog/task/?id=9290>

Performance Schema Indexes

Over 30x faster!

- Allows for more efficient access to Performance Schema tables
- A total of **90 indexes** across **89 tables**
- Adds zero overhead
 - A physical index is not maintained internally
 - Implementation of indexes *tricks* the optimizer into better execution plan

SELECT * FROM sys.session 1000 active sessions



WL#6616: PERFORMANCE_SCHEMA, INDEXES
<https://dev.mysql.com/worklog/task/?id=6616>

EX: Performance Schema Indexes

```
mysql> SELECT * FROM variables_by_thread IGNORE INDEX (primary)
WHERE thread_id = 34 AND variable_name = 'time_zone';
```

THREAD_ID	VARIABLE_NAME	VARIABLE_VALUE
34	time_zone	SYSTEM

1 row in set (0.00 sec)

```
mysql> show status like 'Handler_read_%';
```

Variable_name	Value
Handler_read_first	0
Handler_read_key	0
Handler_read_last	0
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	136

7 rows in set (0.01 sec)

Without Index

“Handler_read_rnd_next”
Is increased.

```
mysql> SELECT * FROM variables_by_thread
WHERE thread_id = 34 AND variable_name = 'time_zone';
```

THREAD_ID	VARIABLE_NAME	VARIABLE_VALUE
34	time_zone	SYSTEM

1 row in set (0.00 sec)

```
mysql> show status like 'Handler_read_%';
```

Variable_name	Value
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	0

7 rows in set (0.01 sec)

With Index

“Handler_read_rnd_next”
is maintain “0” in this case.

Handler_read_rnd_next: The number of requests to read the next row in the data file. This value is high if you are doing a lot of table scans.

Performance Schema Instrumenting SQL Errors

Aggregation	Table Name
By Account	events_errors_summary_by_account_by_error
By Host	events_errors_summary_by_host_by_error
By Thread	events_errors_summary_by_thread_by_error
By User	events_errors_summary_by_user_by_error
Global	events_errors_summary_global_by_error

```
mysql> select * from performance_schema.events_errors_summary_by_host_by_error where FIRST_SEEN is not NULL;
```

HOST	ERROR_NUMBER	ERROR_NAME	SQL_STATE	SUM_ERROR_RAISED	SUM_ERROR_HANDLED	FIRST_SEEN	LAST_SEEN
NULL	1045	ER_ACCESS_DENIED_ERROR	28000	1	0	2016-10-27 15:57:16	2016-10-27 15:57:16
localhost	1046	ER_NO_DB_ERROR	3D000	1	0	2016-10-27 16:00:37	2016-10-27 16:00:37
localhost	1049	ER_BAD_DB_ERROR	42000	1	0	2016-10-27 18:21:09	2016-10-27 18:21:09
localhost	1064	ER_PARSE_ERROR	42000	15	0	2016-10-27 15:58:01	2016-10-27 18:24:06
localhost	1146	ER_NO_SUCH_TABLE	42S02	2	0	2016-10-27 16:08:03	2016-10-27 18:14:41
localhost	1287	ER_WARN_DEPRECATED_SYNTAX	HY000	24	0	2016-10-27 16:07:10	2016-10-27 16:07:10
localhost	3554	ER_NO_SYSTEM_TABLE_ACCESS	HY000	140	0	2016-10-27 15:57:30	2016-10-27 18:38:09

```
mysql> select * from performance_schema.events_errors_summary_by_user_by_error where FIRST_SEEN is not NULL;
```

USER	ERROR_NUMBER	ERROR_NAME	SQL_STATE	SUM_ERROR_RAISED	SUM_ERROR_HANDLED	FIRST_SEEN	LAST_SEEN
NULL	1045	ER_ACCESS_DENIED_ERROR	28000	1	0	2016-10-27 15:57:16	2016-10-27 15:57:16
root	1046	ER_NO_DB_ERROR	3D000	1	0	2016-10-27 16:00:37	2016-10-27 16:00:37
root	1049	ER_BAD_DB_ERROR	42000	1	0	2016-10-27 18:21:09	2016-10-27 18:21:09
root	1064	ER_PARSE_ERROR	42000	15	0	2016-10-27 15:58:01	2016-10-27 18:24:06
root	1146	ER_NO_SUCH_TABLE	42S02	2	0	2016-10-27 16:08:03	2016-10-27 18:14:41
root	1287	ER_WARN_DEPRECATED_SYNTAX	HY000	24	0	2016-10-27 16:07:10	2016-10-27 16:07:10
root	3554	ER_NO_SYSTEM_TABLE_ACCESS	HY000	140	0	2016-10-27 15:57:30	2016-10-27 18:38:09

```
mysql> SELECT * FROM performance_schema.events_errors_summary_global_by_error WHERE sum_error_handled > 0 OR SUM_ERROR_RAISED > 0;
```

ERROR_NUMBER	ERROR_NAME	SQL_STATE	SUM_ERROR_RAISED	SUM_ERROR_HANDLED	FIRST_SEEN	LAST_SEEN
1049	ER_BAD_DB_ERROR	42000	3	0	2017-07-07 14:35:47	2017-07-07 14:36:49
1054	ER_BAD_FIELD_ERROR	42S22	1	0	2017-07-07 08:20:04	2017-07-07 08:20:04
1062	ER_DUP_ENTRY	23000	1	0	2017-07-07 13:30:58	2017-07-07 13:30:58
1064	ER_PARSE_ERROR	42000	6	0	2017-07-07 07:49:59	2017-07-07 14:36:08
1146	ER_NO_SUCH_TABLE	42S02	3	0	2017-07-07 13:30:38	2017-07-07 14:37:33
1287	ER_WARN_DEPRECATED_SYNTAX	HY000	7	0	2017-07-07 11:41:03	2017-07-07 13:39:42
1305	ER_SP_DOES_NOT_EXIST	42000	4	0	2017-07-07 12:44:54	2017-07-07 13:30:11
1411	ER_WRONG_VALUE_FOR_TYPE	HY000	9	0	2017-07-07 12:45:00	2017-07-07 12:45:16
3568	ER_UNRESOLVED_TABLE_LOCK	HY000	4	0	2017-07-07 13:44:45	2017-07-07 13:46:28

Histograms

- More consistent query execution for cases when data is skewed
- Lower cost to maintain than an index
- `ANALYZE TABLE t UPDATE HISTOGRAM ON c1 WITH 10 BUCKETS;`

```
mysql> SELECT * FROM events_statements_histogram_by_digest WHERE SCHEMA_NAME = 'sakila'  
-> AND DIGEST = 'a5980f0634db05c87a7aeb17e1344f84' AND COUNT_BUCKET > 0 limit 1
```

```
***** 1. row *****
```

```
SCHEMA_NAME: sakila  
DIGEST: a5980f0634db05c87a7aeb17e1344f84  
BUCKET_NUMBER: 153  
BUCKET_TIMER_LOW: 10964781961  
BUCKET_TIMER_HIGH: 11481536214  
COUNT_BUCKET: 1  
COUNT_BUCKET_AND_LOWER: 1  
BUCKET_QUANTILE: 0.500000
```

Statement Histogram Summary Tables
these values indicate that 50% of queries
run in under 11.48 microseconds:

WL#8706: Persistent storage of Histogram data
<https://dev.mysql.com/worklog/task/?id=8706>
WL#8707: Classes/structures for Histograms
<https://dev.mysql.com/worklog/task/?id=8707>
WL#8943: Extend ANALYZE TABLE with histogram support
<https://dev.mysql.com/worklog/task/?id=8943>

Performance Schema Histograms

Showing distribution of query time from a run of mysqlslap

```
mysql> WITH total as
-> (SELECT SUM(count_bucket) as t FROM performance_schema.events_statements_histogram_global)
-> SELECT sys.decimal_bucket_name(sys.decimal_bucket(MIN(bucket_timer_low))) as bucket,
-> sys.visualization(SUM(count_bucket) / (SELECT t FROM total), 50) as visualization,
-> SUM(count_bucket) as count
-> FROM performance_schema.events_statements_histogram_global
-> GROUP BY sys.decimal_bucket(bucket_timer_low);
```

bucket	visualization	count
0us+	#####	439148
10us+	#####	163434
100us+		1517
1ms+	#####	99453
10ms+		369
100ms+		29
1s+		2
10s+		1

8 rows in set (0.04 sec)

Generated with a quick CTE over
events_statements_histogram_global

WL#5384: PERFORMANCE_SCHEMA Histograms
<https://dev.mysql.com/worklog/task/?id=5384>



Performance Schema Histograms (cont.)

```
query: INSERT INTO `t1` VALUES (...)
```

```
db: mysqlslap
```

```
total_latency: 54.43 s
```

```
exec_count: 58377
```

```
lock_latency: 1.70 s
```

```
..
```

```
digest: 4e0c5b796c4052b0da4548fd7cb694be
```

```
first_seen: 2017-04-16 20:59:16
```

```
last_seen: 2017-04-16 21:00:34
```

```
latency_distribution:
```

```
0us+
```

```
10us+ #####
```

```
100us+ #####
```

```
1ms+ #
```

```
10ms+
```

```
100ms+
```

```
1s+
```

```
10s+
```

Available on a per statement digest level.
Can quickly aggregate top-N statements
with latency distribution.

Performance Schema Data Locks

```
SELECT thread_id, object_name, index_name, lock_type, lock_mode, lock_data
FROM performance_schema.data_locks WHERE object_name = 'seats';
```

thread_id	object_name	index_name	lock_type	lock_mode	lock_data
33	seats	NULL	TABLE	IX	NULL
33	seats	PRIMARY	RECORD	X	3, 5
33	seats	PRIMARY	RECORD	X	3, 6
33	seats	PRIMARY	RECORD	X	4, 5
33	seats	PRIMARY	RECORD	X	4, 6

5 rows in set (0.00 sec)

- which data is locked,
- who owns the lock,
- who waits for the data.

WL#6657: PERFORMANCE_SCHEMA, DATA LOCKS

<https://dev.mysql.com/worklog/task/?id=6657>

WL#9275: DEPRECATE INFORMATION_SCHEMA.INNODB_LOCKS IN 5.7

<https://dev.mysql.com/worklog/task/?id=9275>

Persist Configuration

- Persist GLOBAL Server Variables
 - SET PERSIST max_connections = 500;
- Examples Include: Offline_mode, Read_Only
- Requires no direct filesystem access
- Includes timestamp and change user

```
mysql> SET PERSIST log_timestamps='SYSTEM';  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> select * from performance_schema.variables_info where variable_source='PERSISTED';
```

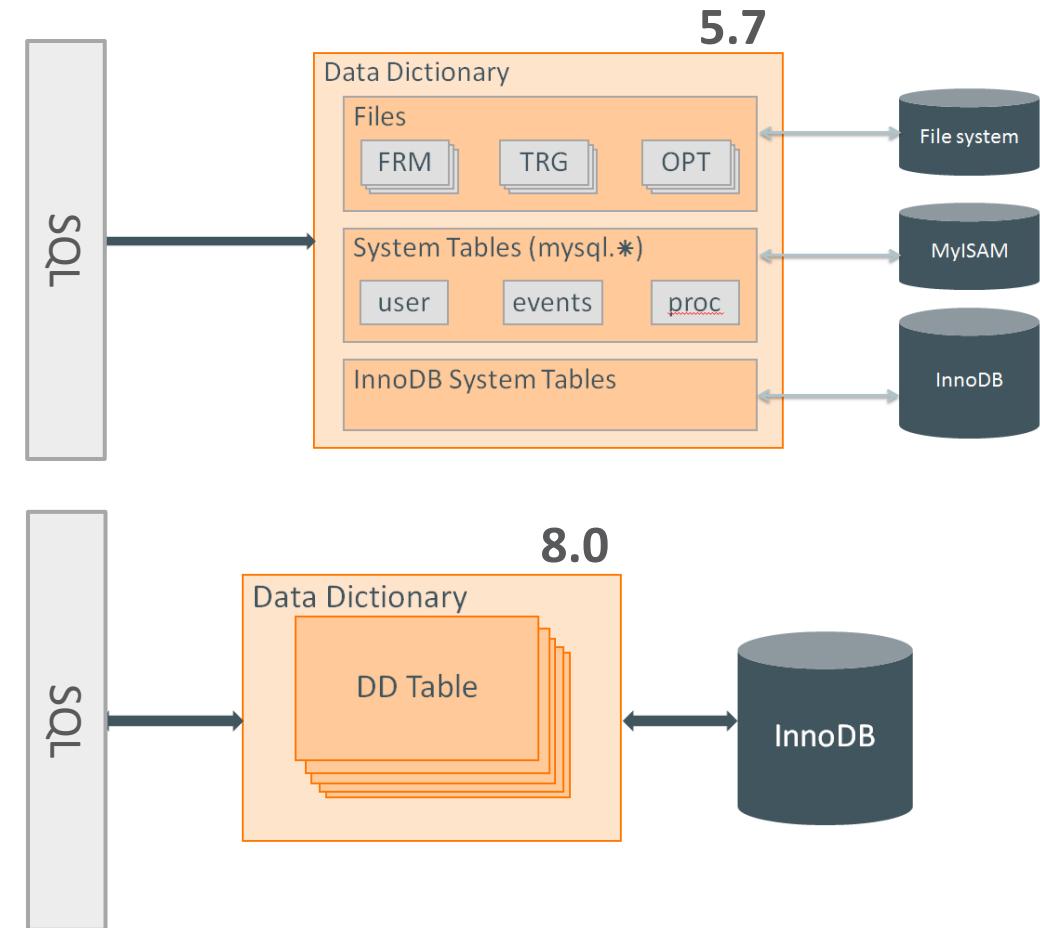
VARIABLE_NAME	VARIABLE_SOURCE	VARIABLE_PATH	MIN_VALUE	MAX_VALUE	SET_TIME	SET_USER	SET_HOST
log_timestamps	PERSISTED	/var/lib/mysql/mysqld-auto.cnf	0	0	2017-07-14 14:48:28		

```
#cat /var/lib/mysql/mysqld-auto.cnf  
{ "mysql_server": { "log_timestamps": "SYSTEM" } }
```

WL#8688: Support ability to persist SET GLOBAL settings
<https://dev.mysql.com/worklog/task/?id=8688>

Transactional Data Dictionary

- Increased Reliability
- Using InnoDB internally for data dictionary
 - No FRM files
 - No DB.OPT files
 - No TRG files
 - No TRN files
 - No PAR files
- No longer contain MyISAM tables



```
[root@DockerHost oracle]# ls -l /docker/docker802/world
total 1084
-rw-r----- 1 27 27 638976 Jul 18 01:25 city.ibd
-rw-r----- 1 27 27 196608 Jul 18 01:25 country.ibd
-rw-r----- 1 27 27 262144 Jul 18 01:25 countrylanguage.ibd
```

8.0

WL#6379: Schema definitions for new DD
<https://dev.mysql.com/worklog/task/?id=6379>
WL#6392: Upgrade to Transactional Data Dictionary
<https://dev.mysql.com/worklog/task/?id=6392>
WL#6394: Bootstrap code for new DD
<https://dev.mysql.com/worklog/task/?id=6394> and more

Transactional Data Dictionary


- Better cross-platform experience
 - No dependencies on filesystem semantics
- Atomic DDL
 - Better Replication, Simplifies server edge cases
- MDL(Metadata Lock) for Foreign Keys
- Flexible Metadata API
 - Easier path to adding new features

Ver	Delete Tables	Delete Stored Programs	Delete Schema	ATOMICITY
5.7	<ul style="list-style-type: none">- Metadata, TRN/TRG/FRM files- Data, InnoDB tables	<ul style="list-style-type: none">- Metadata, rows in MyISAM (non-transactional)	<ul style="list-style-type: none">- Metadata, DB.OPT file	Mix of filesystem, non-transactional/transactional storage and multiple commits
8.0	<ul style="list-style-type: none">- Metadata, rows in InnoDB- Data, InnoDB tables	<ul style="list-style-type: none">- Metadata, rows in InnoDB	<ul style="list-style-type: none">- Metadata, rows in InnoDB	Updates to transactional storage, one commit

Information Schema Performance

100 schemas times 50 tables (5000 tables)

Already faster at **7/10**
queries in our test suite!



30x Faster

```
mysql> SELECT TABLE_SCHEMA, TABLE_NAME, TABLE_TYPE, ENGINE, ROW_FORMAT  
-> FROM information_schema.tables WHERE TABLE_SCHEMA LIKE 'db%';
```

TABLE_SCHEMA	TABLE_NAME	TABLE_TYPE	ENGINE	ROW_FORMAT
db1000	T_PARTITION	BASE TABLE	InnoDB	Dynamic
db1000	actor	BASE TABLE	InnoDB	Dynamic
db1000	actor2	BASE TABLE	InnoDB	Dynamic
db1000	actor_info	VIEW	NULL	NULL
db1000	address	BASE TABLE	InnoDB	Dynamic

<SNIP>

Test Performed with 100 schemas, each with 50 tables.

Memcached Enhancement (Range Query)

Support of "multiple get" operation

- Improve the read performance.
- User can fetch multiple key value pairs in a single memcached query.
- Reduce communication traffic between client and server.

Add support for "range queries"

With Range queries, user can specify a particular range, and fetch all the qualified values in this range.

```
get A B C D
get @>B
get @<D
get @<=D
get @>B@<D
```

city_id	name	state	country
A	BANGALORE	BANGALORE	IN
B	CHENNAI	TAMIL NADU	IN
C	DELHI	DELHI	IN
D	HYDERABAD	TELANGANA	IN

WL#6650: InnoDB_Memcached: support multiple get and range search
<https://dev.mysql.com/worklog/task/?id=6650>

All these features plus...

Source code now documented with Doxygen

Plugin Infrastructure!

Expanded GIS Support

Expanded Query Hints Support

Improved Scan Query Performance

Improved BLOB Storage

Cost Model Improvements

Scalability Improvements

Atomicity in Privileges

Dynamic Privileges

Parser Refactoring

Improvements to Temporary Tables

C++11 and Toolchain Improvements

Improve undo tablespace management

- Replication Applier Thread Progress Reports
- GTID_PURGED always settable
- Improved Parallel Replication
- Multi-Source Replication Filters
- Multi-Source Replication Filters In Performance Schema
- Binary Log Expiration Period in Seconds
- SQL Grouping Function
- Optimizer Trace detailed sort statistics
- Smaller Package Downloads
- JSON performance improvements
- Expanded Query Hints
- Improved usability of cost constant configuration

<http://mysqlserverteam.com/the-mysql-8-0-0-milestone-release-is-available/>

<http://mysqlhighavailability.com/replication-features-in-mysql-8-0-2/>

<http://www.unofficialmysqlguide.com/index.html>

Thank you!!

“Please Enjoy Upcoming New Features. We hope it will support your business.”

Appendix: Just In case...

Oracle Lifetime Support for MySQL

Features	Premier (Years 1-5)	Extended (Years 6-8)	Sustain (Years 9+)
24x7 Support	•	•	•
Unlimited Support Incidents	•	•	•
Knowledge Base	•	•	•
Maintenance Releases, Bug Fixes, Patches, Updates	•	•	Pre-Existing only
MySQL Consultative Support	•	•	•

<https://www.mysql.com/support/>

<https://www.mysql.com/support/consultative.html>

Concern about security? Require useful tools? Please check our products.

<https://www.mysql.com/products/>

<https://www.mysql.com/products/enterprise/>

Health Check (Check your MySQL Environment)

<https://www.mysql.com/news-and-events/health-check/>

Do you want to
reduce TCO?
Need remote DBA?

Do you want to
save more time and
focus on service
development?

Do you require
Comprehensive
security solution?

Integrated Cloud

Applications & Platform Services

ORACLE®