

Secondary Index Search in MySQL

Mijin An

meeeeeejin@gmail.com



Contents

- Secondary/Primary Key Record Format
- Secondary Index Search
- Proposal for Improvement
- Improved Secondary Index Search

Secondary Key Record Format

	Variable field lengths (1-2 bytes per var. field)
	Nullable field bitmap (1 bit per nullable field)
N-5	Info Flags (4 bits)
N-4	Number of Records Owned (4 bits)
	Order (13 bits)
N-2	Record Type (3 bits)
N	Next Record Offset (2)
N+k	Secondary Key Fields (k)
N+k+j	Cluster Key Fields (j)

= Primary Key Value

Primary Key Record Format

	Variable field lengths (1-2 bytes per var. field)
	Nullable field bitmap (1 bit per nullable field)
N-5	Info Flags (4 bits)
	Number of Records Owned (4 bits)
N-4	Order (13 bits)
	Record Type (3 bits)
N-2	Next Record Offset (2)
N	Cluster Key Fields (k)
N+k	Transaction ID (6)
N+k+6	Roll Pointer (7)
N+k+13	Non-Key Fields (j)
N+k+13+j	

Secondary Index Search

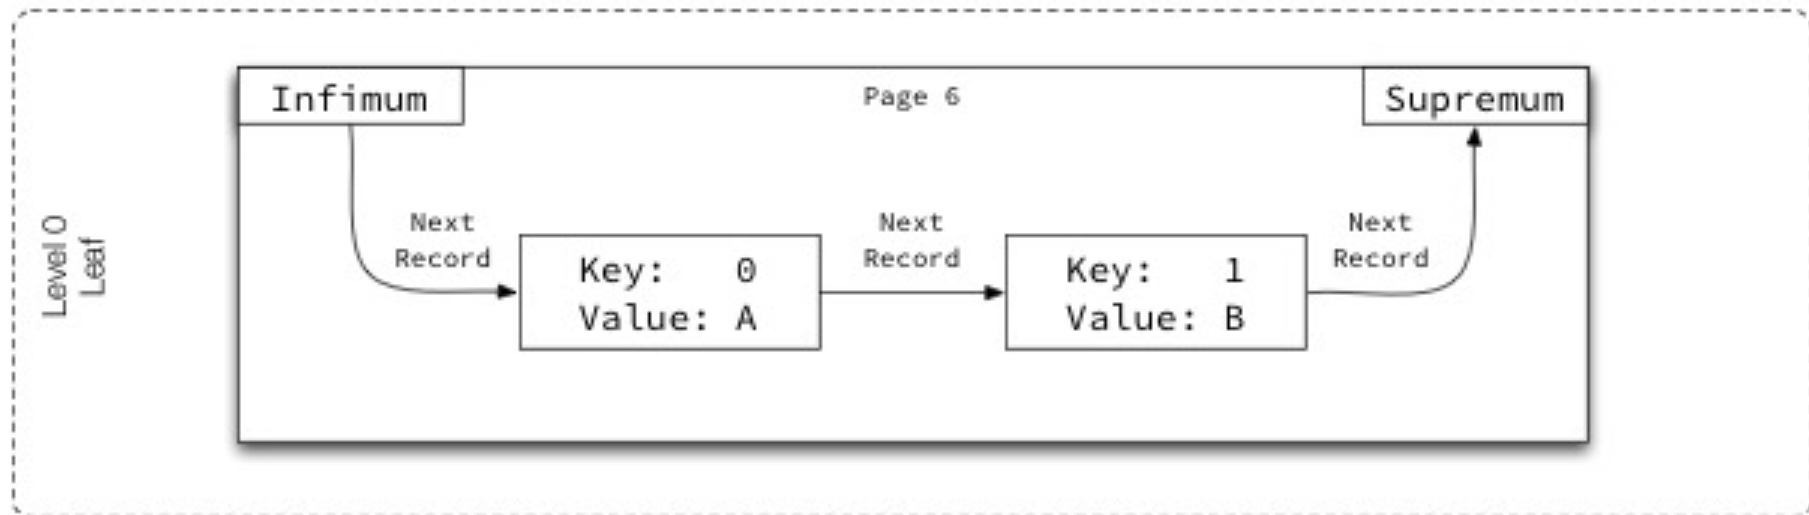
`row_search_for_mysql()`

- ① Start a mini-transaction
- ② Searches an index tree, then positions a tree cursor(*pcur*) on a record corresponding to the index
`(btr_pcur_open_with_no_init())`
- ③ Get the record that *pcur* indicates(*rec*)
- ④ Go to *rec_loop*:

rec_loop:

Look for matching records in a loop

- a. Get the record that *pcur* indicates(*rec*)
- b. If record is infimum or supremum, go to *next_rec*:



- c.
 - If record is infimum → move *pcur* to the next record
 - Else if record is supremum → move *pcur* to the first record of next page
 - Else → move *pcur* according to the search direction
- d. If we may need to process the record the cursor is now on, go to *rec_loop*:

rec_loop:

Look for matching records in a loop

- a. Get the record that *pcur* indicates(*rec*)
- b. If record is infimum or supremum, go to *next_rec:*
- c. Caculates the offsets to each field in the record (*offsets*)
- d. If *select_lock_type* == *LOCK_S* or *LOCK_X*,
 - try to place a lock on the index record
- e. Go to *locks_ok:*
 - If *index* != *clust_index* && *need_to_access_clustered*,
 - go to *requires_clust_rec:*

requires_clust_rec:

Get clustered record using secondary index

- a. Retrieves the clustered record(*clust_rec*) corresponding to a record in a secondary index `(row_sel_get_clust_rec_for_mysql())`

row_sel_get_clust_rec_for_mysql()

Get clustered record using secondary index

- a. Get the *clust_index* from *sec_index*
- b. Searches an index tree, then positions a tree cursor (*clust_pcur*) on a record corresponding to the index
(btr_pcur_open_with_no_init())
- c. Get the record that *clust_pcur* indicates (*clust_rec*)
- d. Calculate the offsets to each field in the record (*offsets*)
- e. If *select_lock_type* == *LOCK_S* or *LOCK_X*,
 - try to place a lock on the index record

row_sel_get_clust_rec_for_mysql()

Get clustered record using secondary index

- f. Go to *func_exit*:
- g. In *func_exit*, store the current *clust_pcur* position if *select_lock_type* is *LOCK_S* or *LOCK_X*

requires_clust_rec:

Get clustered record using secondary index

- a. Retrieves the clustered record(*clust_rec*) corresponding to a record in a secondary index (row_sel_get_clust_rec_for_mysql())
- b. If *clust_rec* is delete marked, skip it → *next_rec*:
- c. Store *clust_rec* into the *result_rec*
- d. Go to *normal_return*:


Secondary Index Search

row_search_for_mysql() → normal_return: → func_exit:

- ⑤ Commit mini-transaction
- ⑥ Go to *func_exit:*
- ⑦ In *func_exit:*, frees the space occupied by a memory heap then exit

Secondary Index Search

Secondary index tree




Primary index tree




Secondary Index Search

Secondary index tree




Primary index tree




Data records

Secondary Index Search

Secondary index tree





Primary index tree



Data records


Proposal: Prefetching in MySQL

Secondary index tree



Primary key list(unsorted)


Primary index tree



Data records

Proposal: Prefetching in MySQL

Secondary index tree




→ Primary key



Primary key list(unsorted)


Primary index tree



Data records

Proposal: Prefetching in MySQL

Secondary index tree




→ Primary key



Primary key list(unsorted)


Primary index tree



Data records

Proposal: Prefetching in MySQL

Secondary index tree




→ Primary key



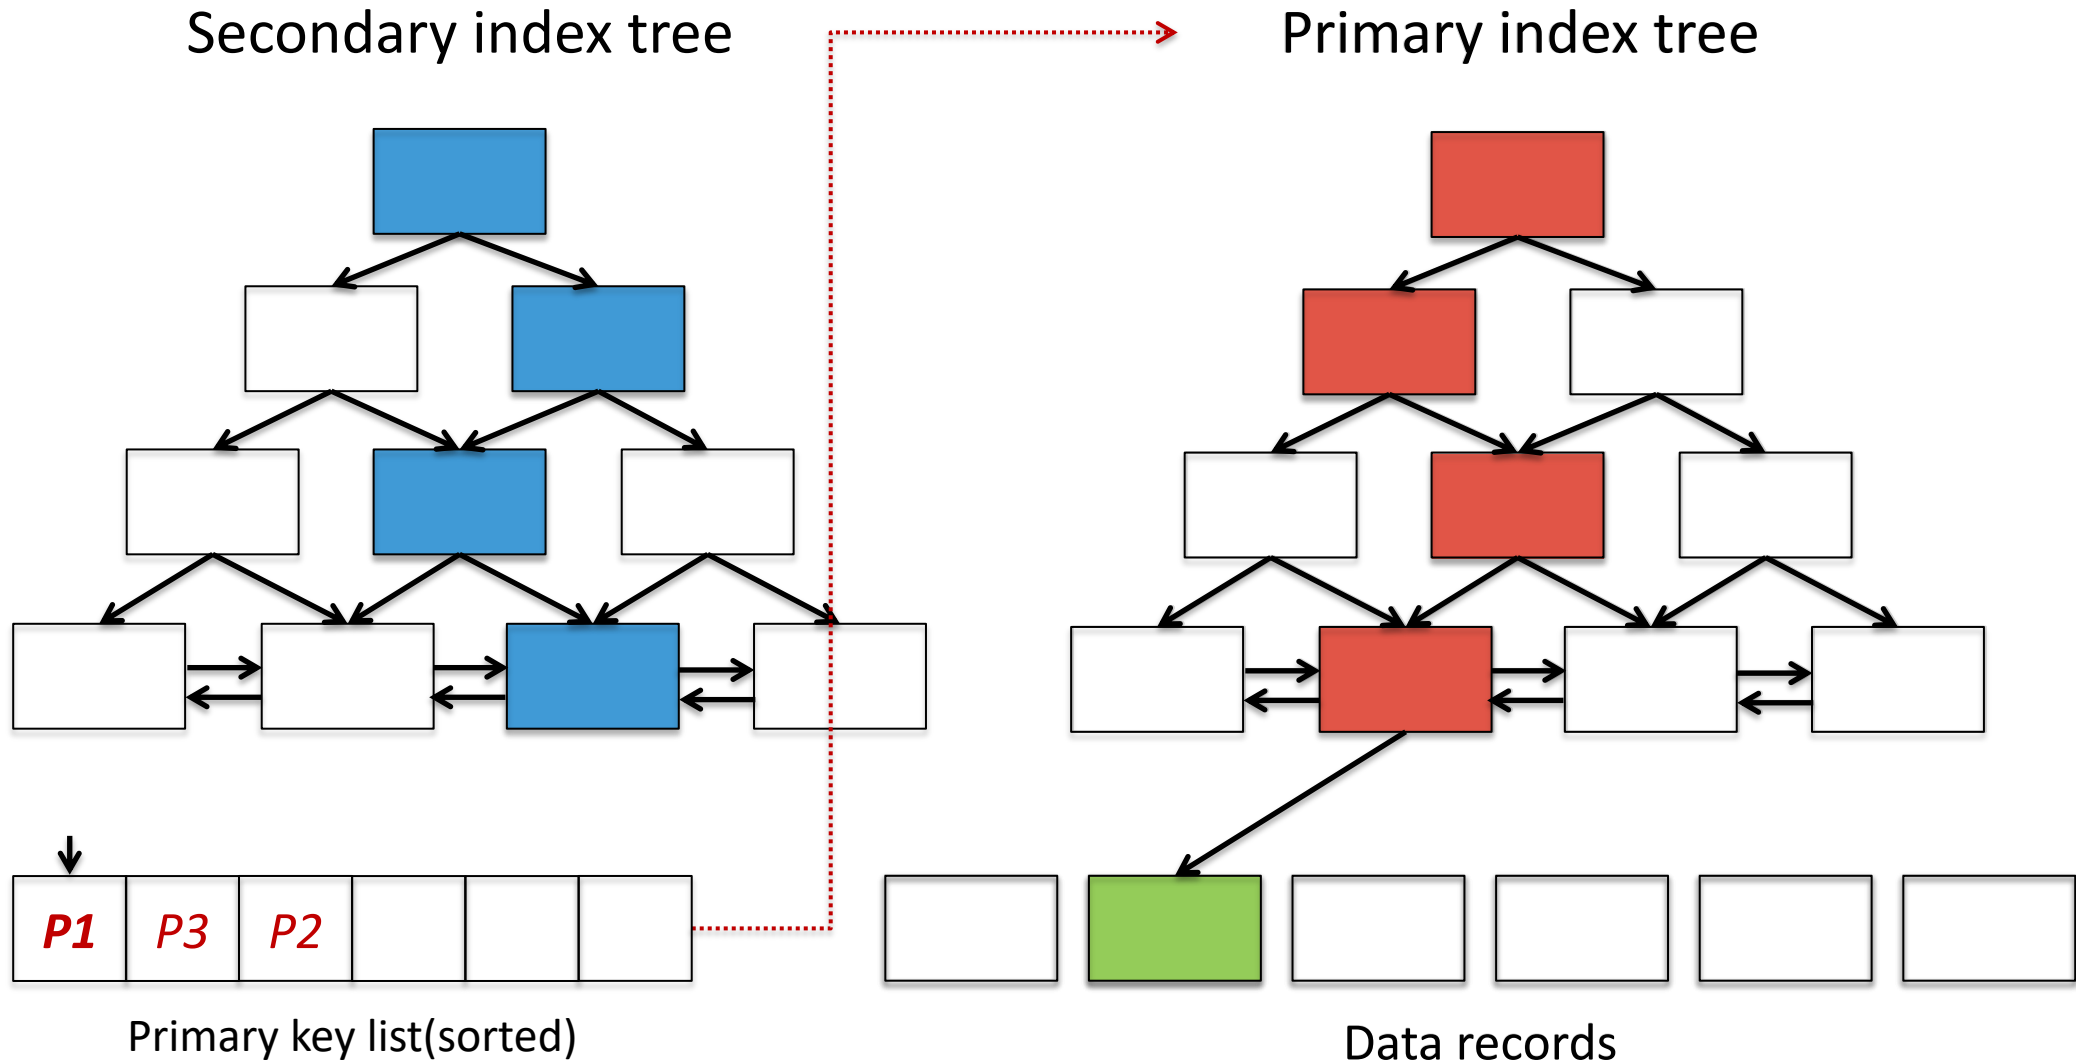
Primary key list (sorted)

Primary index tree

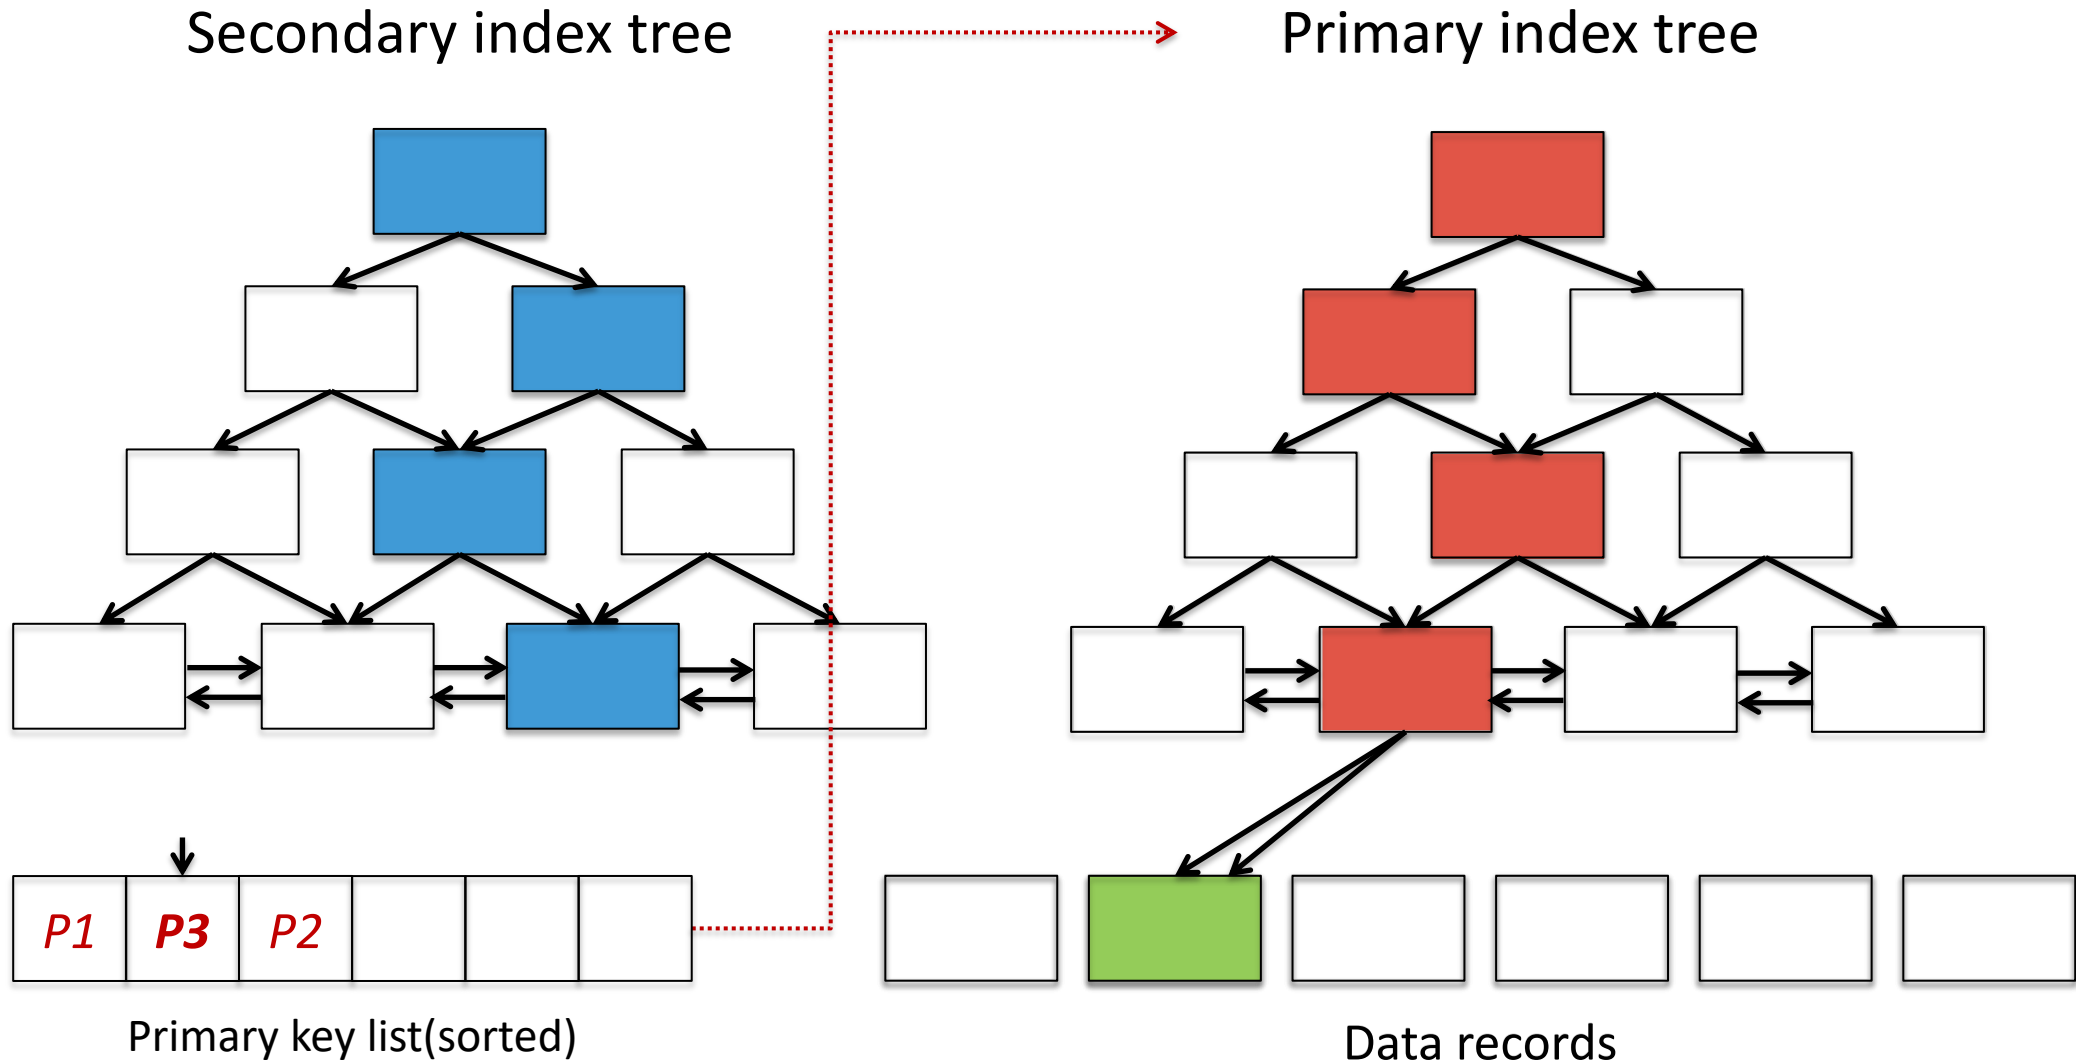


Data records

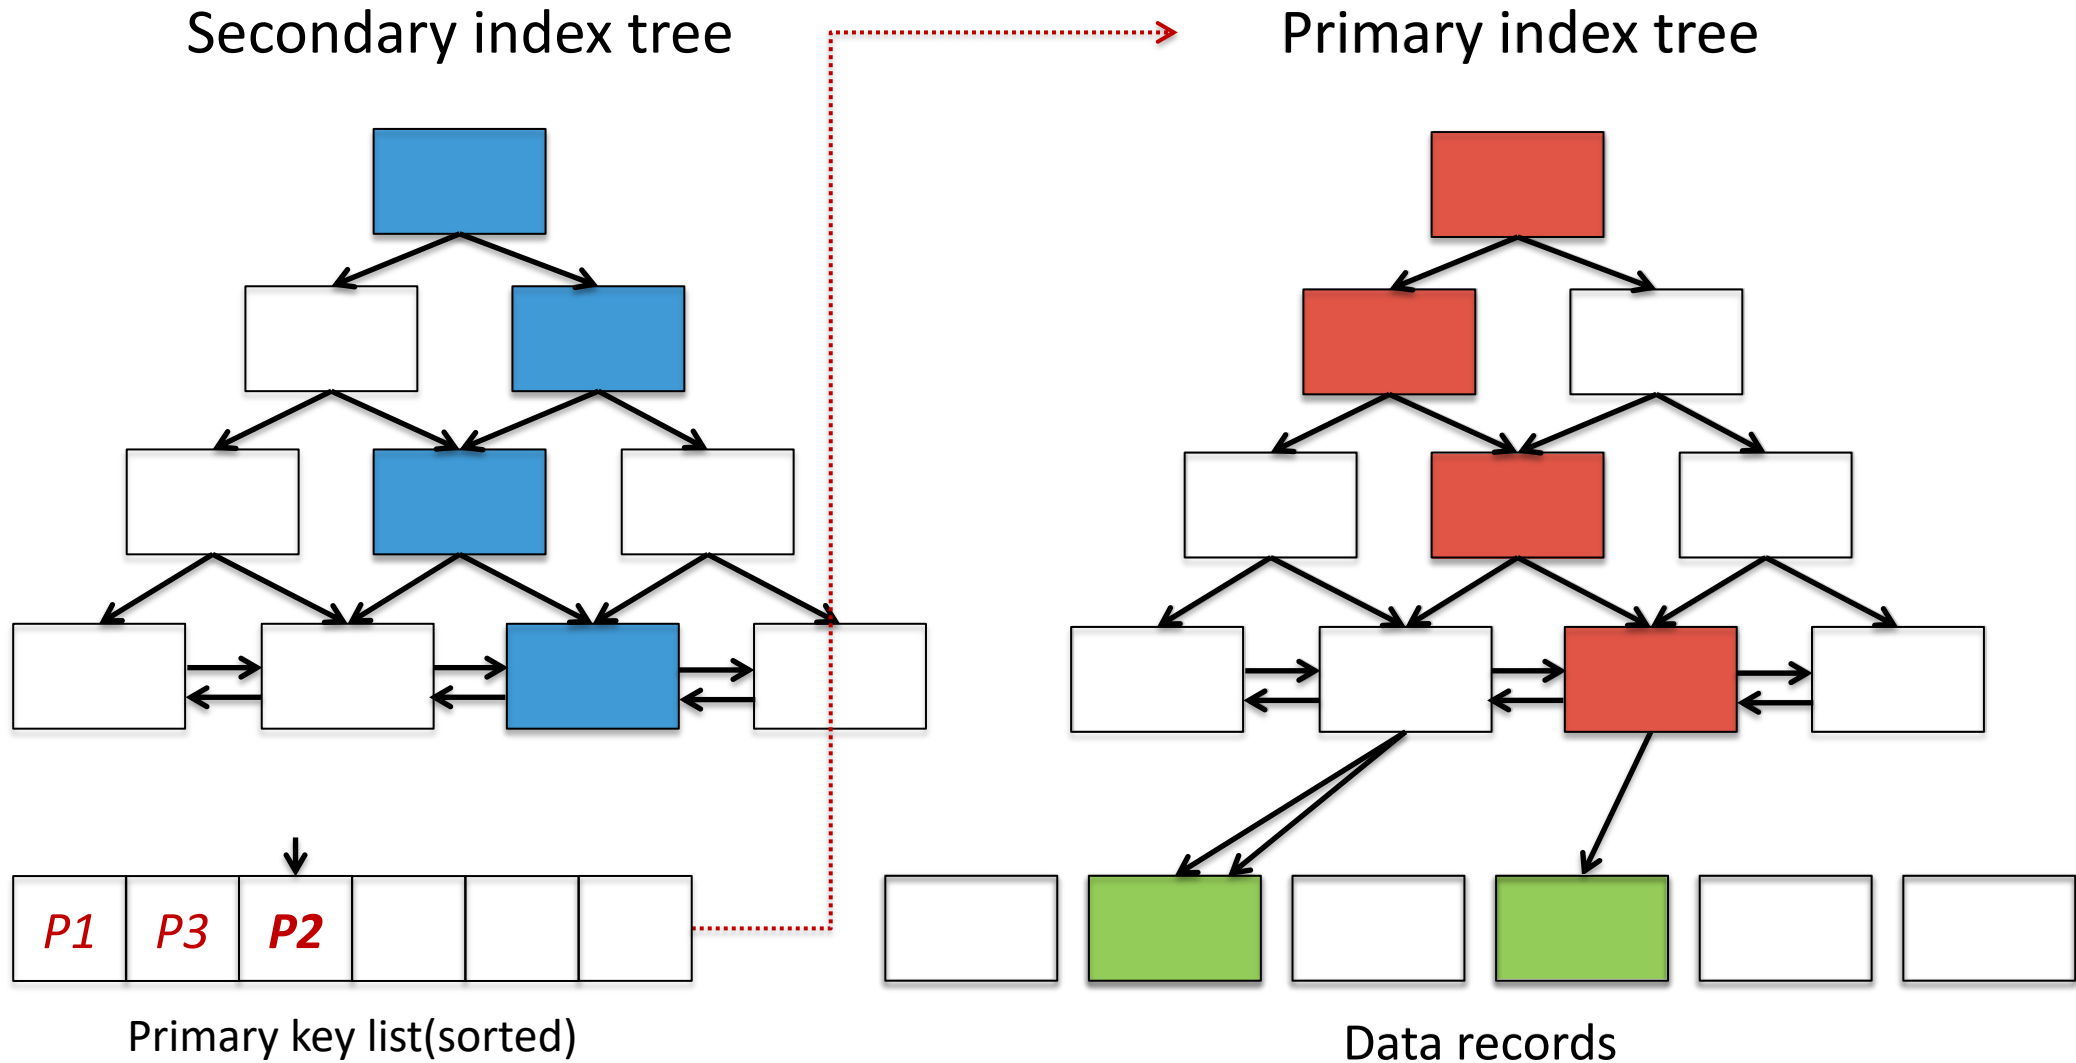
Proposal: Prefetching in MySQL



Proposal: Prefetching in MySQL



Proposal: Prefetching in MySQL



Improved Secondary Index Search

- ① Search a data record using secondary index
- ② Get primary key and insert it into primary key list, then repeat ①~② process until there is no record to search using secondary index
- ③ Sort the primary key list
- ④ Search primary index tree sequentially using sorted primary key list
- ⑤ Access data record sequentially

Reference

- [1] “MySQL 5.6 Reference Manual”, MySQL, <https://dev.mysql.com/doc/refman/5.6/en/>
- [2] Jeremy Cole, “InnoDB”, <https://blog.jcole.us/innodb/>