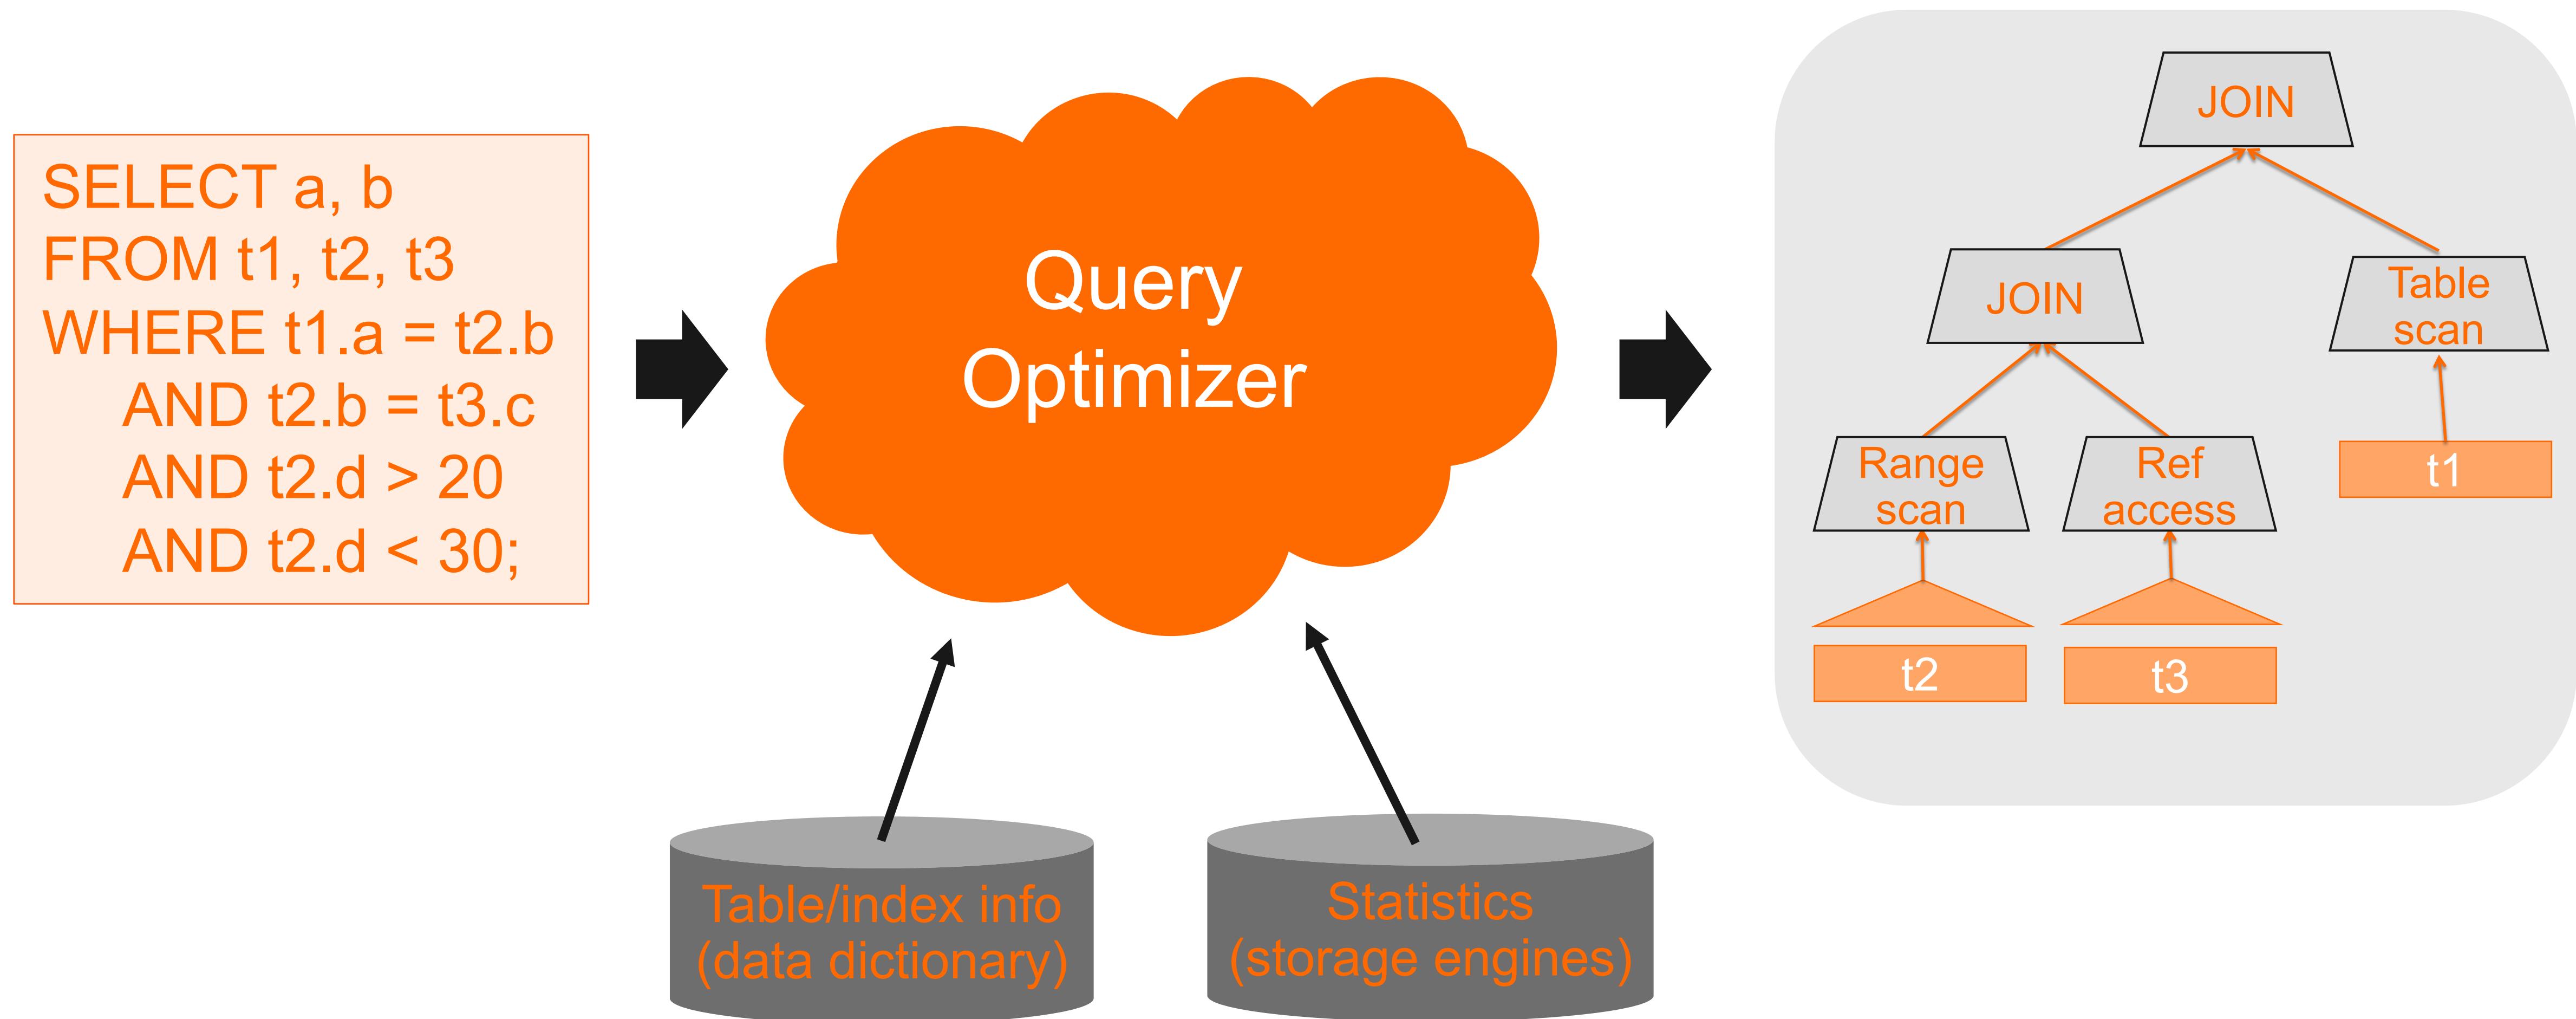


The MySQL Query Optimizer

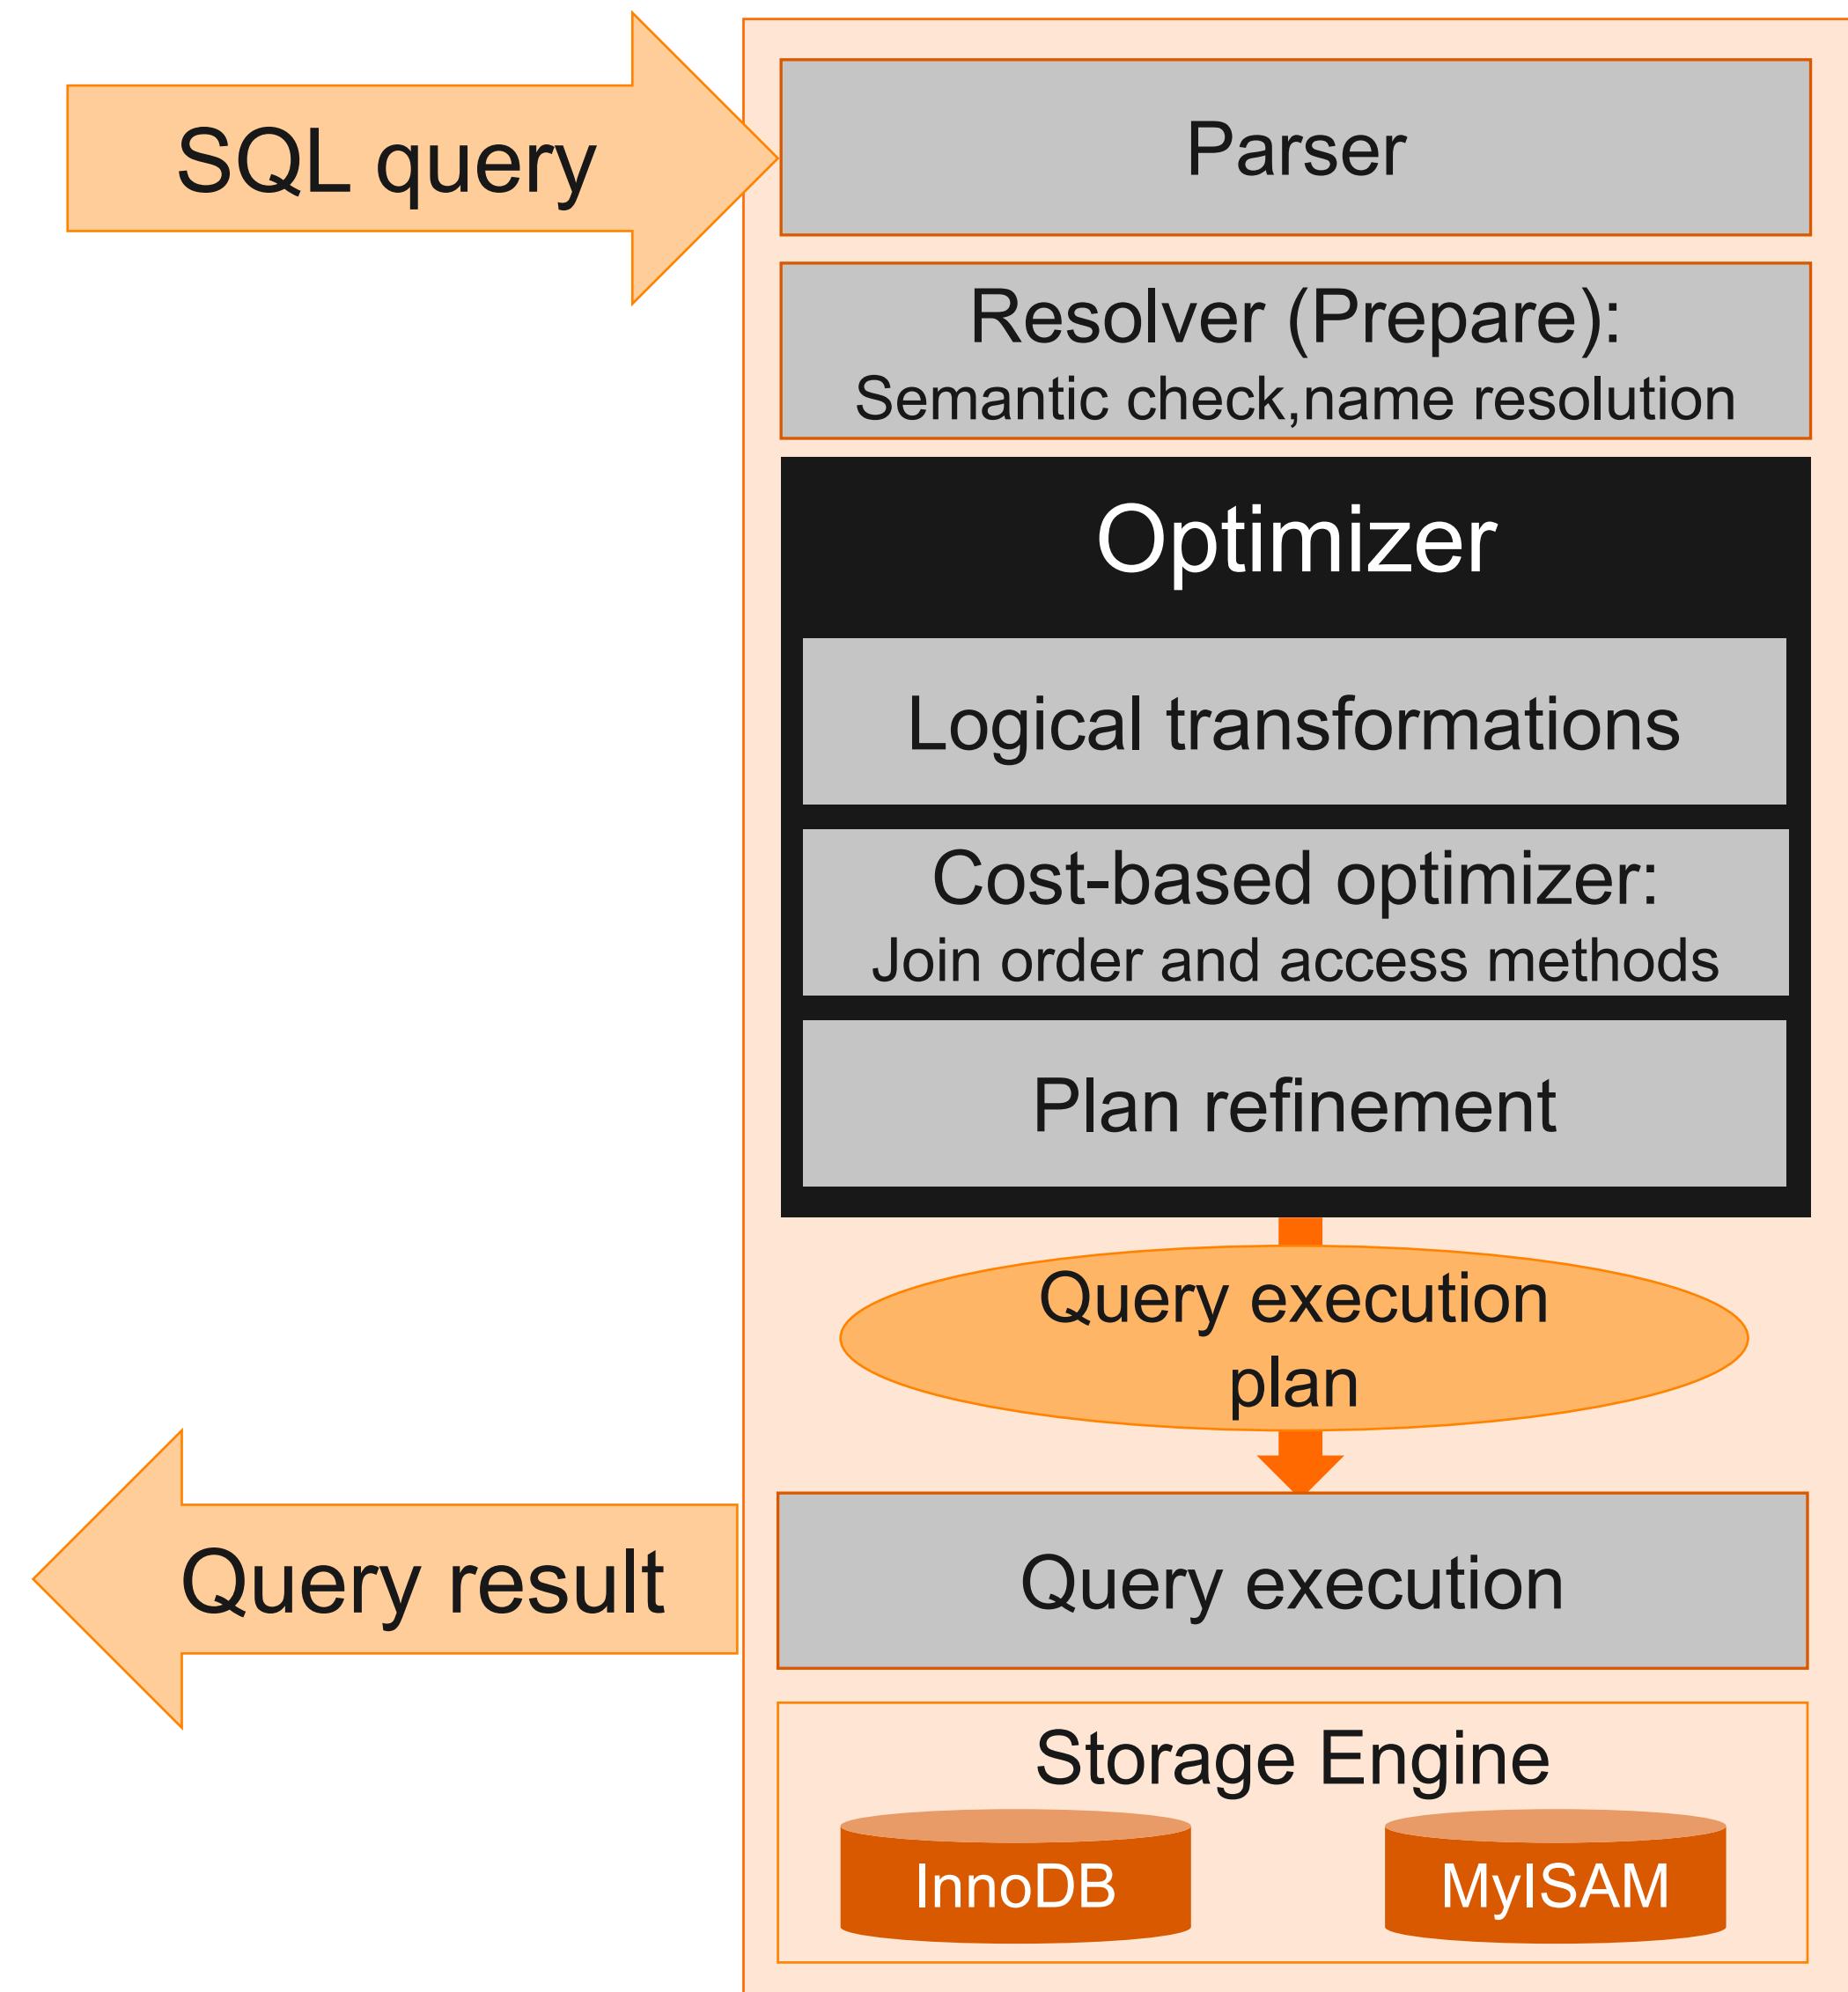
Explained Through Optimizer Trace

Øystein Grøvlen
Senior Staff Engineer
Alibaba Cloud

MySQL Query Optimizer



MySQL Architecture



Optimizer Trace

How to generate it

- EXPLAIN shows the selected plan
- Optimizer trace shows WHY the plan was selected

```
SET optimizer_trace= "enabled=on";
SELECT * FROM t1, t2 WHERE f1=1 AND f1=f2 AND f2>0;
SELECT trace FROM information_schema.optimizer_trace
INTO OUTFILE filename LINES TERMINATED BY '';
SET optimizer_trace="enabled=off";
```

QUERY	SELECT * FROM t1, t2 WHERE f1=1 AND f1=f2 AND f2>0;
TRACE	{ "steps": [{ "join_preparation": { "select#": 1,... } ... } ...] }
MISSING_BYTES_BEYOND_MAX_MEM_SIZE	0
INSUFFICIENT_PRIVILEGES	0

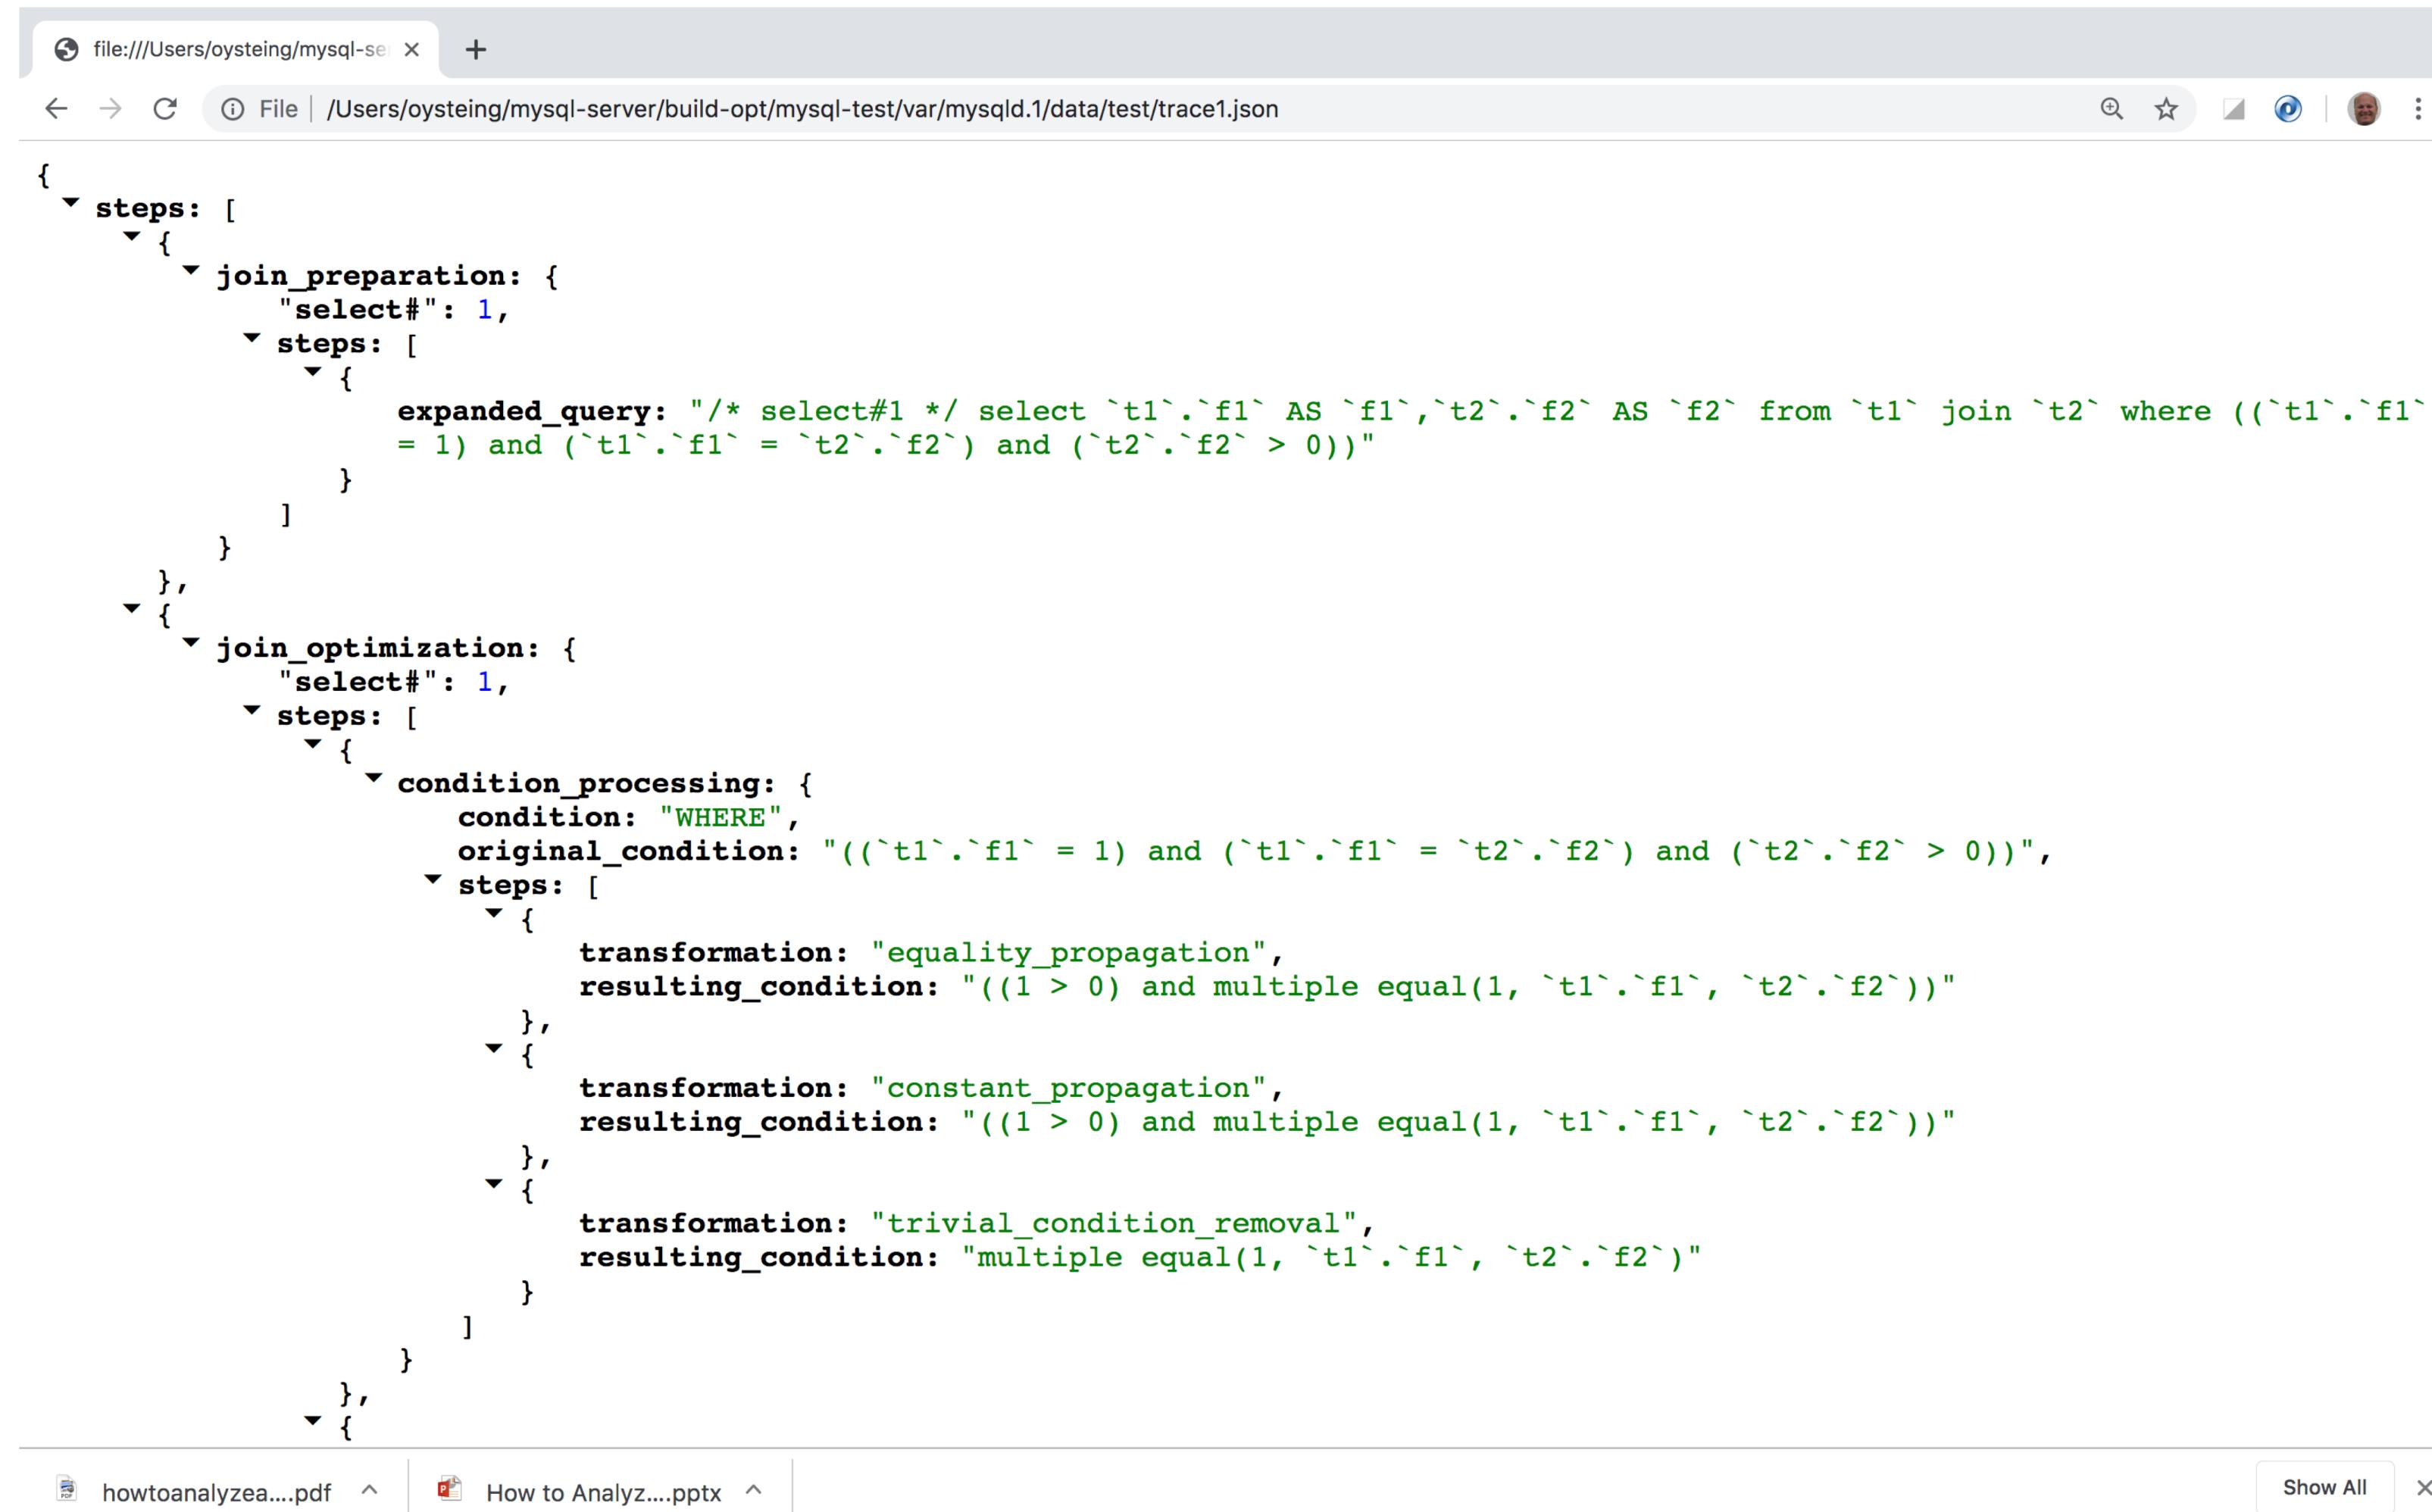
Optimizer Trace

Example

```
{  
  "steps": [  
    {  
      "join_preparation": {  
        "select#": 1,  
        "steps": [  
          {  
            "expanded_query": "/* select#1 */ select `t1`.`f1` AS `f1`, `t2`.`f2` AS `f2` from `t1` join `t2` where ((`t1`.`f1` = 1) and (`t1`.`f1` = `t2`.`f2`) and (`t2`.`f2` > 0))"  
          }  
        ]  
      },  
      {  
        "join_optimization": {  
          "select#": 1,  
          "steps": [  
            {  
              "condition_processing": {  
                "condition": "WHERE",  
                "original_condition": "((`t1`.`f1` = 1) and (`t1`.`f1` = `t2`.`f2`) and (`t2`.`f2` > 0))",  
                "steps": [  
                  {  
                    ...  
                  }  
                ]  
              }  
            ]  
          }  
        }  
      }  
    ]  
  ]  
}
```

Optimizer Trace

JSON browser plugin



```
{  
  "steps": [  
    {  
      "join_preparation": {  
        "select#": 1,  
        "steps": [  
          {  
            "expanded_query": "/* select#1 */ select `t1`.`f1` AS `f1`, `t2`.`f2` AS `f2` from `t1` join `t2` where ((`t1`.`f1` = 1) and (`t1`.`f1` = `t2`.`f2`) and (`t2`.`f2` > 0))"  
          }  
        ]  
      },  
      {  
        "join_optimization": {  
          "select#": 1,  
          "steps": [  
            {  
              "condition_processing": {  
                "condition": "WHERE",  
                "original_condition": "((`t1`.`f1` = 1) and (`t1`.`f1` = `t2`.`f2`) and (`t2`.`f2` > 0))",  
                "steps": [  
                  {  
                    "transformation": "equality_propagation",  
                    "resulting_condition": "((1 > 0) and multiple equal(1, `t1`.`f1`, `t2`.`f2"))"  
                  },  
                  {  
                    "transformation": "constant_propagation",  
                    "resulting_condition": "((1 > 0) and multiple equal(1, `t1`.`f1`, `t2`.`f2"))"  
                  },  
                  {  
                    "transformation": "trivial_condition_removal",  
                    "resulting_condition": "multiple equal(1, `t1`.`f1`, `t2`.`f2")"  
                  }  
                ]  
              }  
            }  
          ]  
        }  
      }  
    }  
  ]  
}
```

Browser Plugin

Collapse to see main trace objects/phases

```
{  
  ▼ steps: [  
    ▼ {  
      ▶ join_preparation: { ... }  
    },  
    ▼ {  
      ▶ join_optimization: { ... }  
    },  
    ▼ {  
      ▶ join_execution: { ... }  
    }  
  ]  
}
```

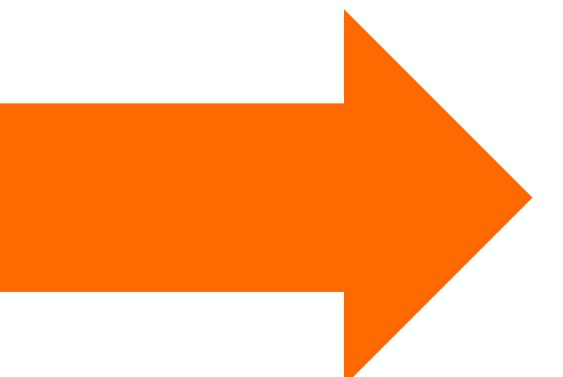
Browser Plugin

Expand JSON objects

```
{  
  ▼ steps: [  
    ▼ {  
      ▼ join_preparation: {  
        "select#": 1,  
        ▼ steps: [  
          ▼ {  
            expanded_query: "/* select#1 */ select `t1`.`f1` AS  
              `f1`, `t2`.`f2` AS `f2` from `t1` join `t2` where ((`t1`.`f1` =  
              1) and (`t1`.`f1` = `t2`.`f2`) and (`t2`.`f2` > 0))"  
          }  
        ]  
      }  
    },  
    ▼ {  
      ▶ join_optimization: { ... }  
    },  
    ▼ {  
      ▶ join_execution: { ... }  
    }  
  ]  
}
```

Prepare Phase

- Name resolving
 - Map names to database objects (tables, columns, ...)
- Semantic checks
- Permanent transformations:
 - Conversion of outer join to inner join
 - Merging of views and derived tables
 - Subquery transformations
 - IN to EXISTS
 - IN to Semijoin (5.6)
 - EXISTS to IN (8.0.16)
 - Etc.



Simpler query to
optimize and
execute



Prepare for later
optimizations

Conversion of outer join to inner join

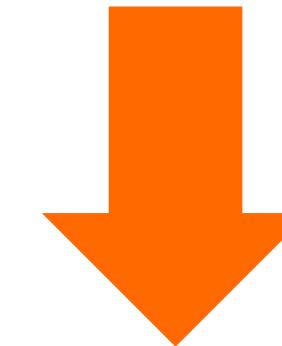
```
SELECT o_orderkey FROM orders LEFT JOIN lineitem ON o_orderkey = l_orderkey  
WHERE l_discount > 0.10;
```

```
"join_preparation": {  
  "select#": 1,  
  "steps": [  
    {  
      "expanded_query": "/* select#1 */ select `orders`.`o_orderkey` AS `o_orderkey` from (`orders` left join `lineitem` on  
((`orders`.`o_orderkey` = `lineitem`.`l_orderkey`))) where (`lineitem`.`l_discount` > 0.10)"  
    },  
    {  
      "transformations_to_nested_joins": {  
        "transformations": [  
          "outer_join_to_inner_join",  
          "JOIN_condition_to_WHERE",  
          "parenthesis_removal"  
        ],  
        "expanded_query": "/* select#1 */ select `orders`.`o_orderkey` AS `o_orderkey` from `orders` join `lineitem` where  
((`lineitem`.`l_discount` > 0.10) and(`orders`.`o_orderkey` = `lineitem`.`l_orderkey`))"  
      } ...  
    } ...  
  ]  
}
```

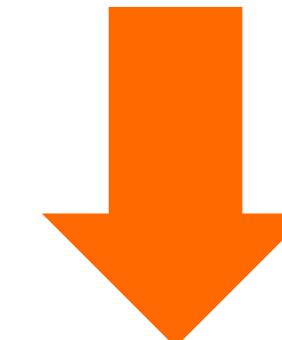
Conversion from EXISTS-subquery to IN-subquery

New in 8.0.16

```
SELECT o_orderpriority, COUNT(*) AS order_count FROM orders  
WHERE EXISTS  
  (SELECT * FROM lineitem WHERE l_orderkey = o_orderkey AND l_commitdate < l_receiptdate)  
GROUP BY o_orderpriority ORDER BY o_orderpriority;
```



```
SELECT o_orderpriority, COUNT(*) AS order_count FROM orders  
WHERE o_orderkey IN  
  (SELECT l_orderkey FROM lineitem WHERE l_commitdate < l_receiptdate)  
GROUP BY o_orderpriority ORDER BY o_orderpriority;
```



```
SELECT o_orderpriority, COUNT(*) AS order_count  
FROM orders SEMIJOIN lineitem ON l_orderkey = o_orderkey  
WHERE l_commitdate < l_receiptdate  
GROUP BY o_orderpriority ORDER BY o_orderpriority;
```

Conversion from EXISTS-subquery to IN-subquery

Optimizer trace

```
"join_preparation": {
    "select#": 1,
    "steps": [
        {
            "join_preparation": {
                "select#": 2,
                "steps": [
                    {
                        "expanded_query": "/* select#2 */ select 1 from `lineitem` where ((`lineitem`.`l_orderkey` = `orders`.`o_orderkey`)
and (`lineitem`.`l_commitDATE` < `lineitem`.`l_receiptDATE`))"
                },
                {
                    "transformation": {
                        "select#": 2,
                        "from": "EXISTS (SELECT)",
                        "to": "semijoin",
                        "chosen": true
                    }
                }
            }
        ]
    ...
}
```

Conversion from EXISTS-subquery to IN-subquery

Optimizer trace, cont.

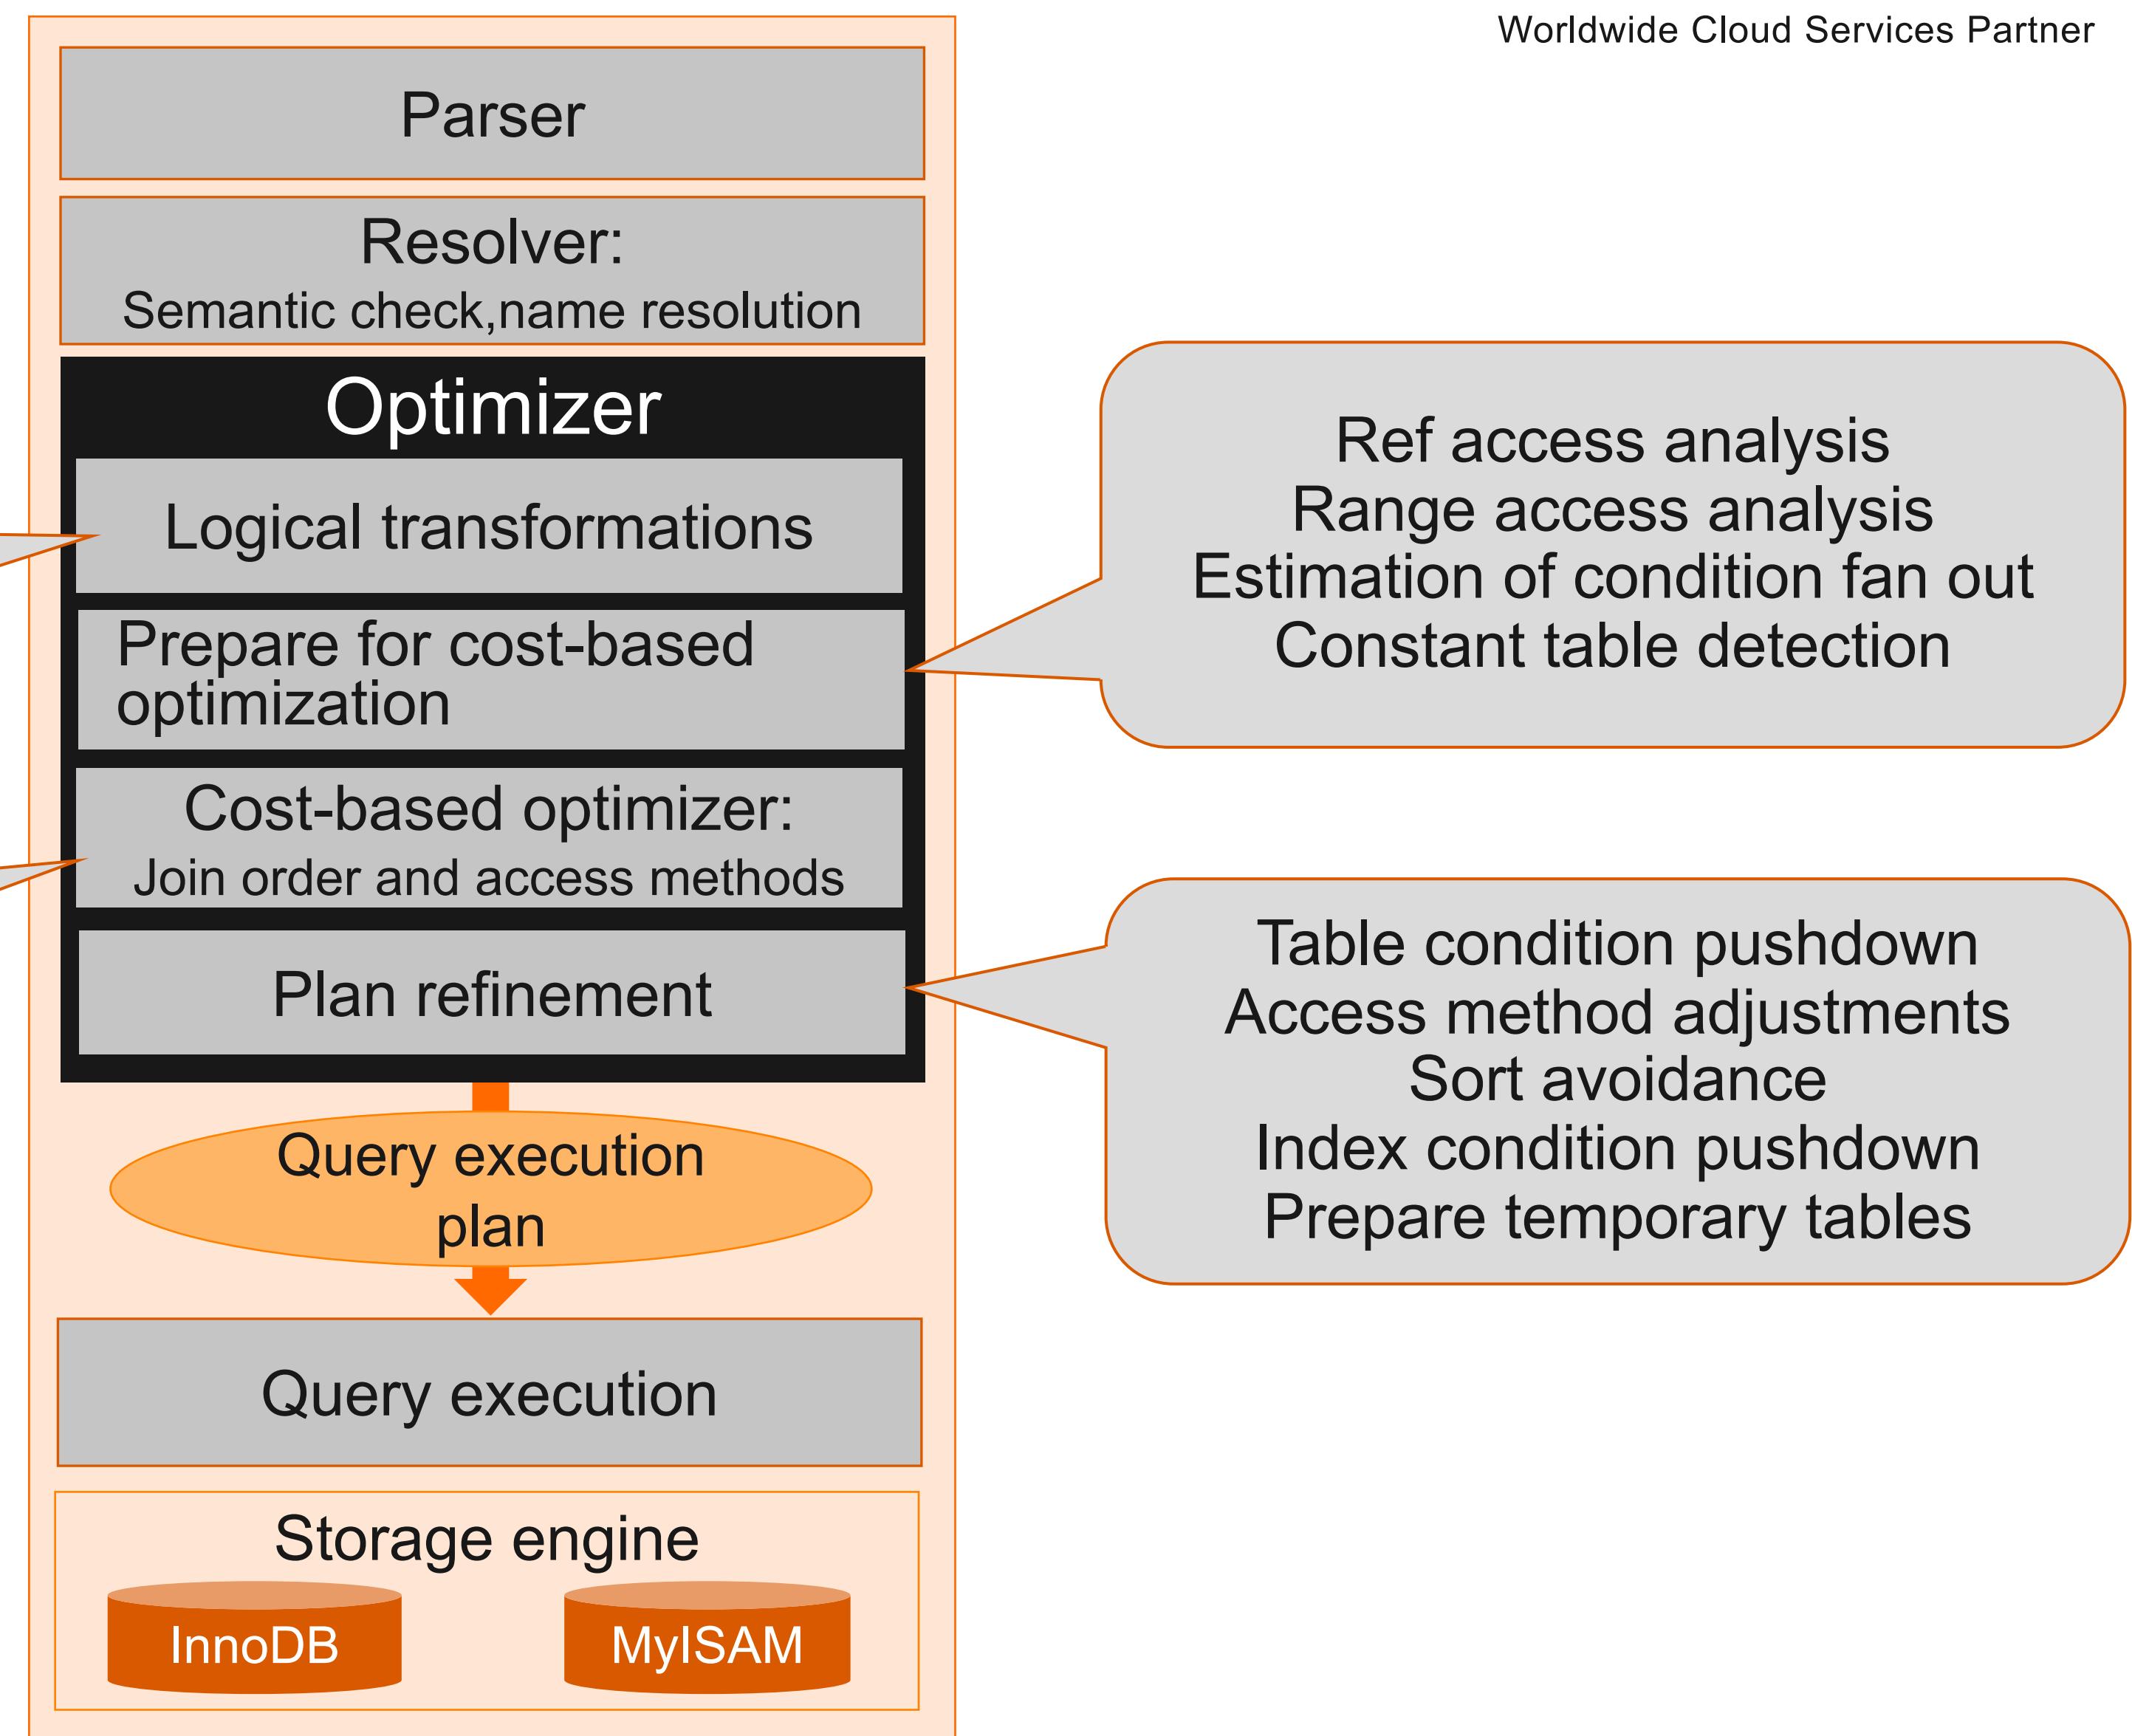
```
  ▼ join_preparation: {
    "select#": 1,
    ▼ steps: [
      ▼ {
        ▶ join_preparation: { ... }
      },
      ▼ {
        expanded_query: "/* select#1 */ select `orders`.`o_orderpriority` AS `o_orderpriority`,count(0) AS
          `order_count` from `orders` where exists(/* select#2 */ select 1 from `lineitem` where
          ((`lineitem`.`l_orderkey` = `orders`.`o_orderkey`) and (`lineitem`.`l_commitDATE` <
          `lineitem`.`l_receiptDATE`))) group by `orders`.`o_orderpriority` order by `orders`.`o_orderpriority`"
      },
      ▼ {
        ▼ transformation: {
          "select#": 2,
          from: "IN (SELECT)",
          to: "semijoin",
          chosen: true,
          evaluating_constant_semijoin_conditions: [ ]
        }
      },
      ▼ {
        ▼ transformations_to_nested_joins: {
          ▼ transformations: [
            "semijoin"
          ],
          expanded_query: "/* select#1 */ select `orders`.`o_orderpriority` AS `o_orderpriority`,count(0) AS
            `order_count` from `orders` semi join (`lineitem`) where ((`lineitem`.`l_commitDATE` <
            `lineitem`.`l_receiptDATE`) and (`orders`.`o_orderkey` = `lineitem`.`l_orderkey`)) group by
            `orders`.`o_orderpriority` order by `orders`.`o_orderpriority`"
        }
      }
    ]
  }
```

Query Optimization

Main phases

Negation elimination
Equality and constant propagation
Evaluation of constant expressions
Substitution of generated columns

Access method selection
Join order

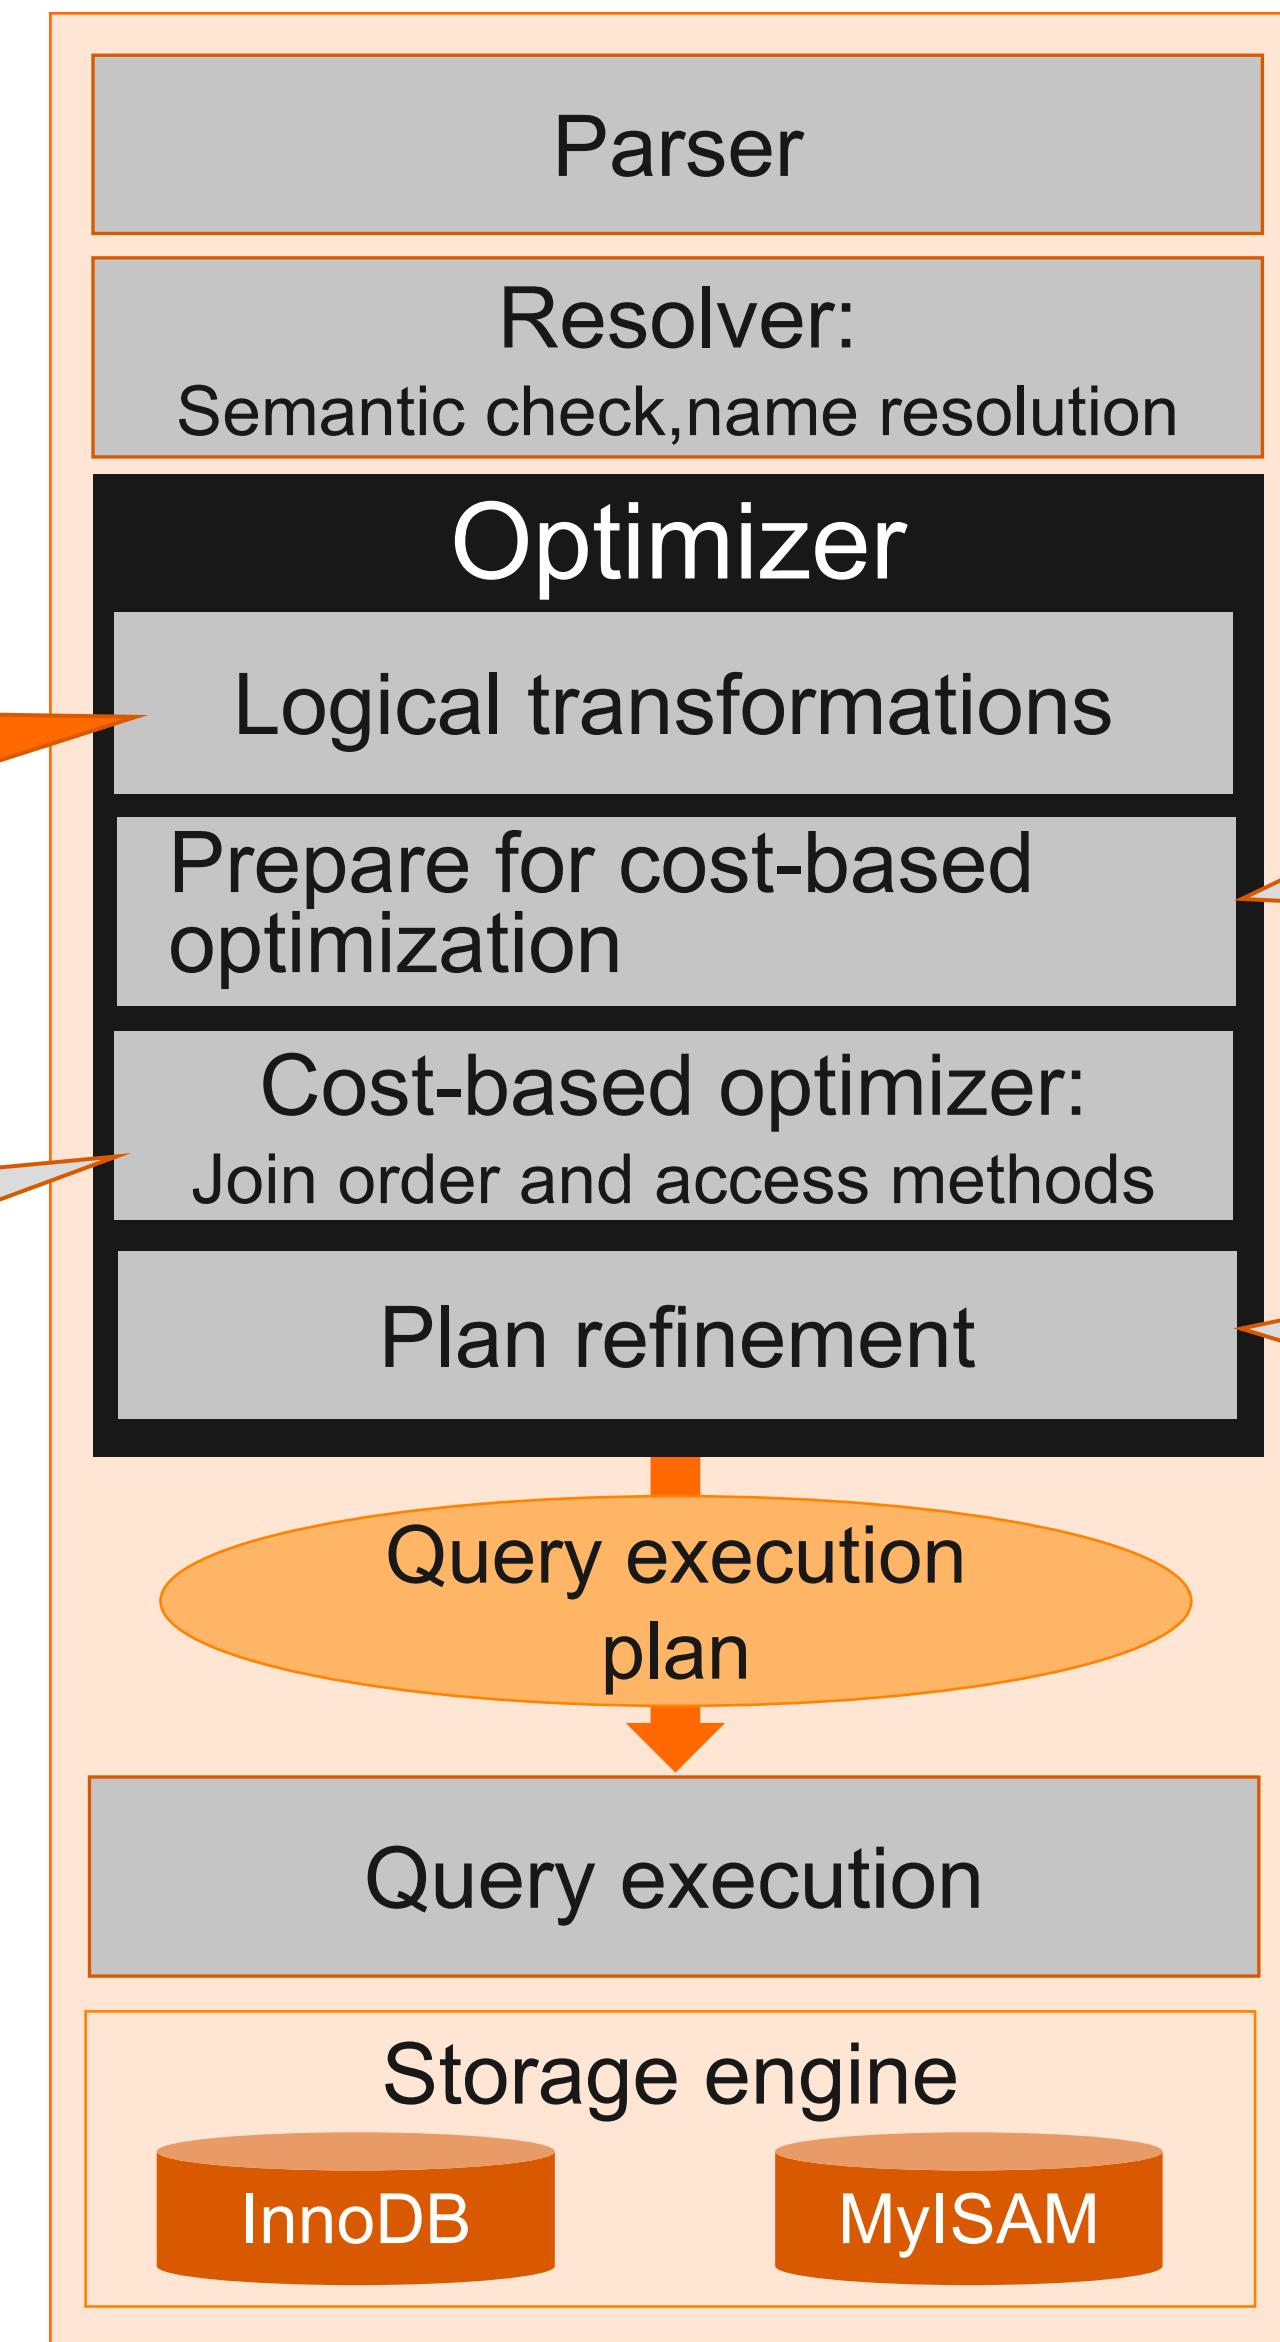


Query Optimization

Main phases

Negation elimination
Equality and constant propagation
Evaluation of constant expressions
Substitution of generated columns

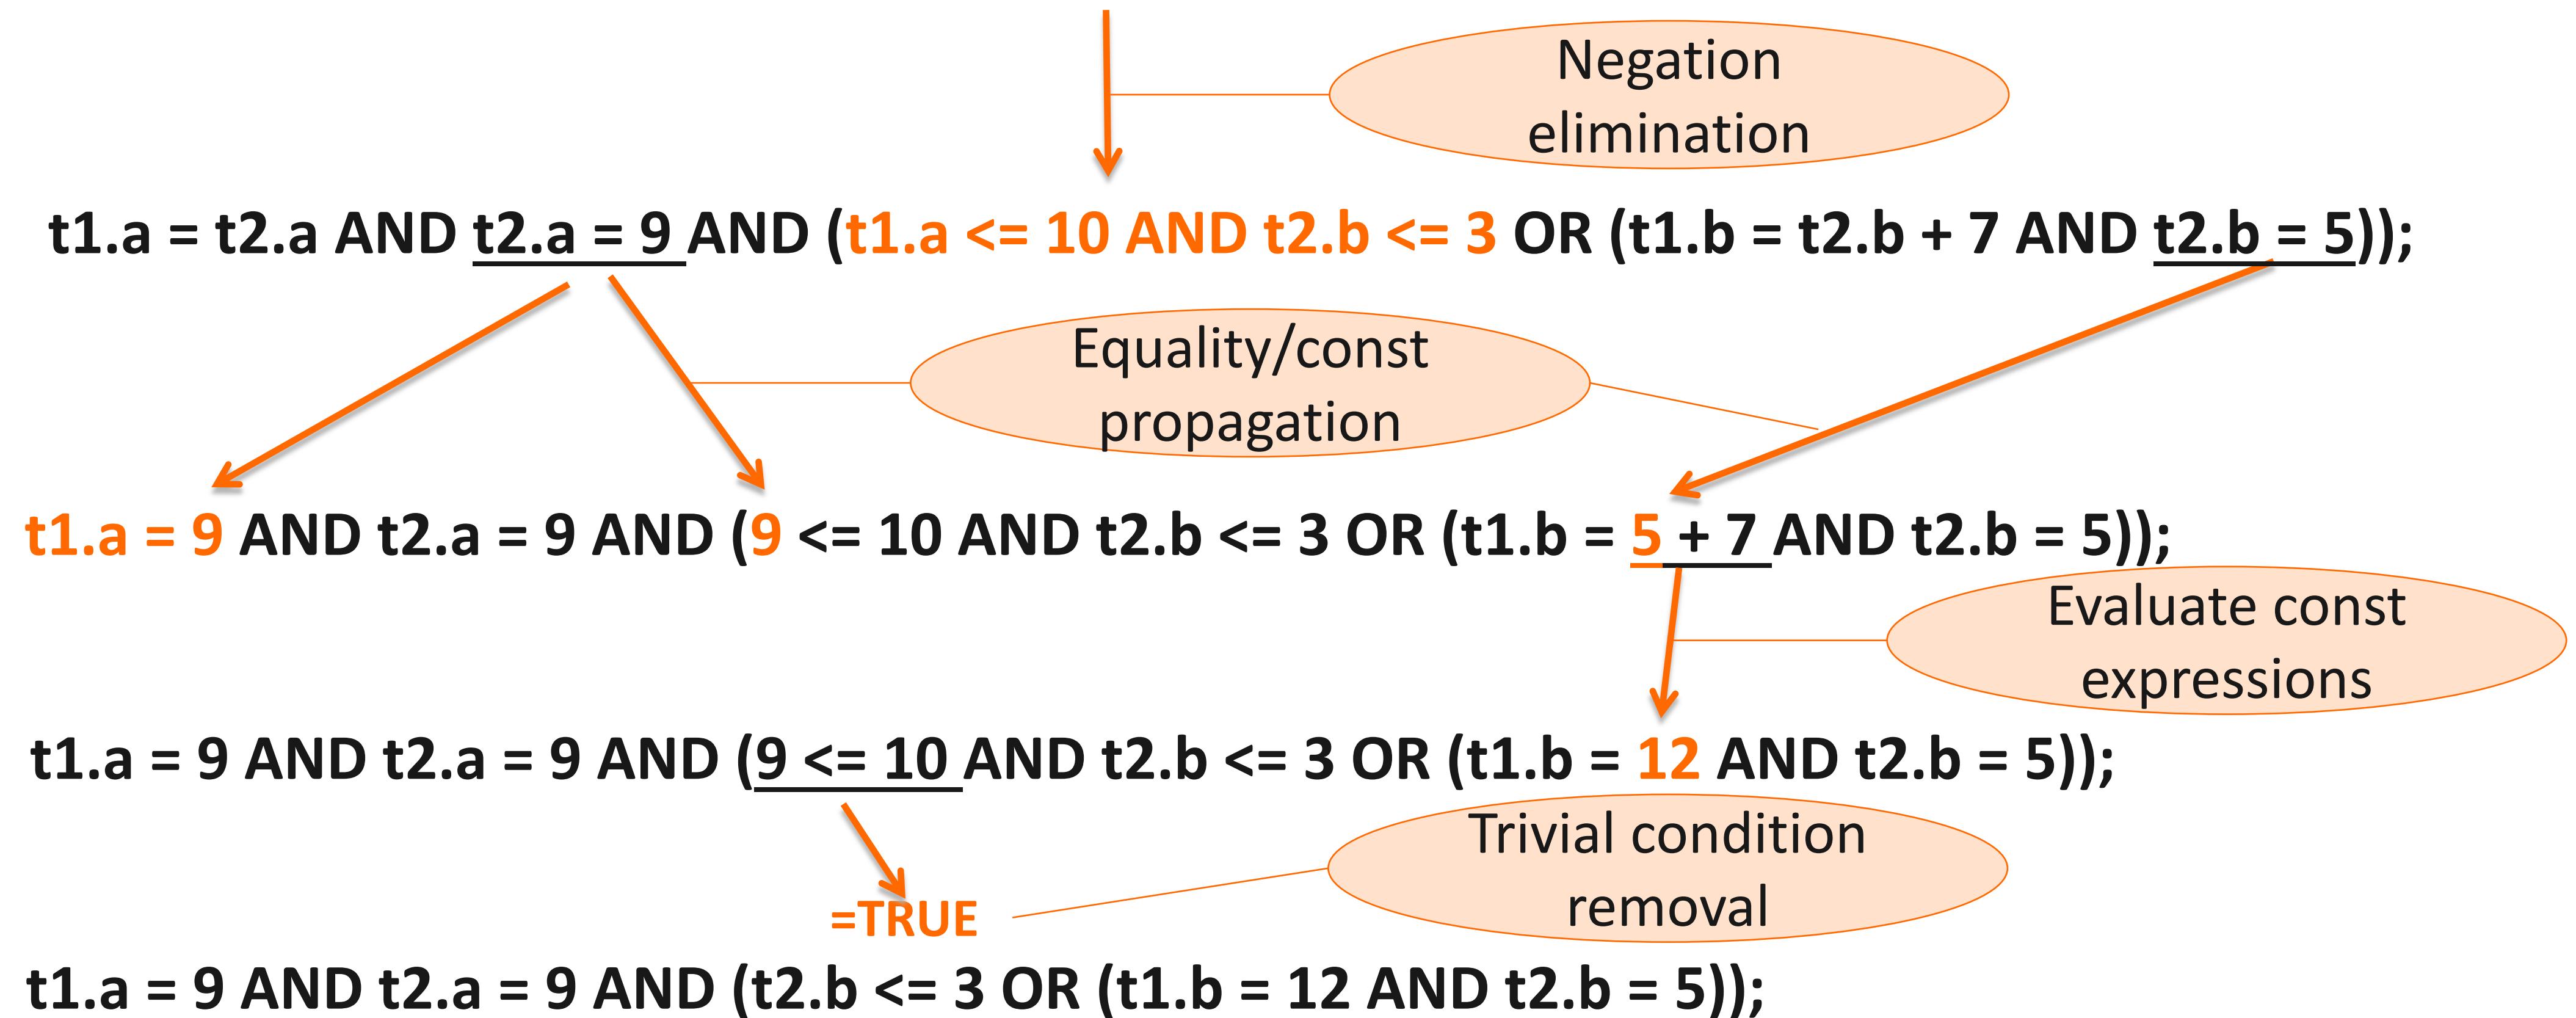
Access method selection
Join order



Logical Transformations

Condition processing

```
SELECT * FROM t1, t2 WHERE  
t1.a = t2.a AND t2.a = 9 AND (NOT (t1.a > 10 OR t2.b > 3) OR (t1.b = t2.b + 7 AND t2.b = 5));
```



Logical Transformations

Optimizer Trace

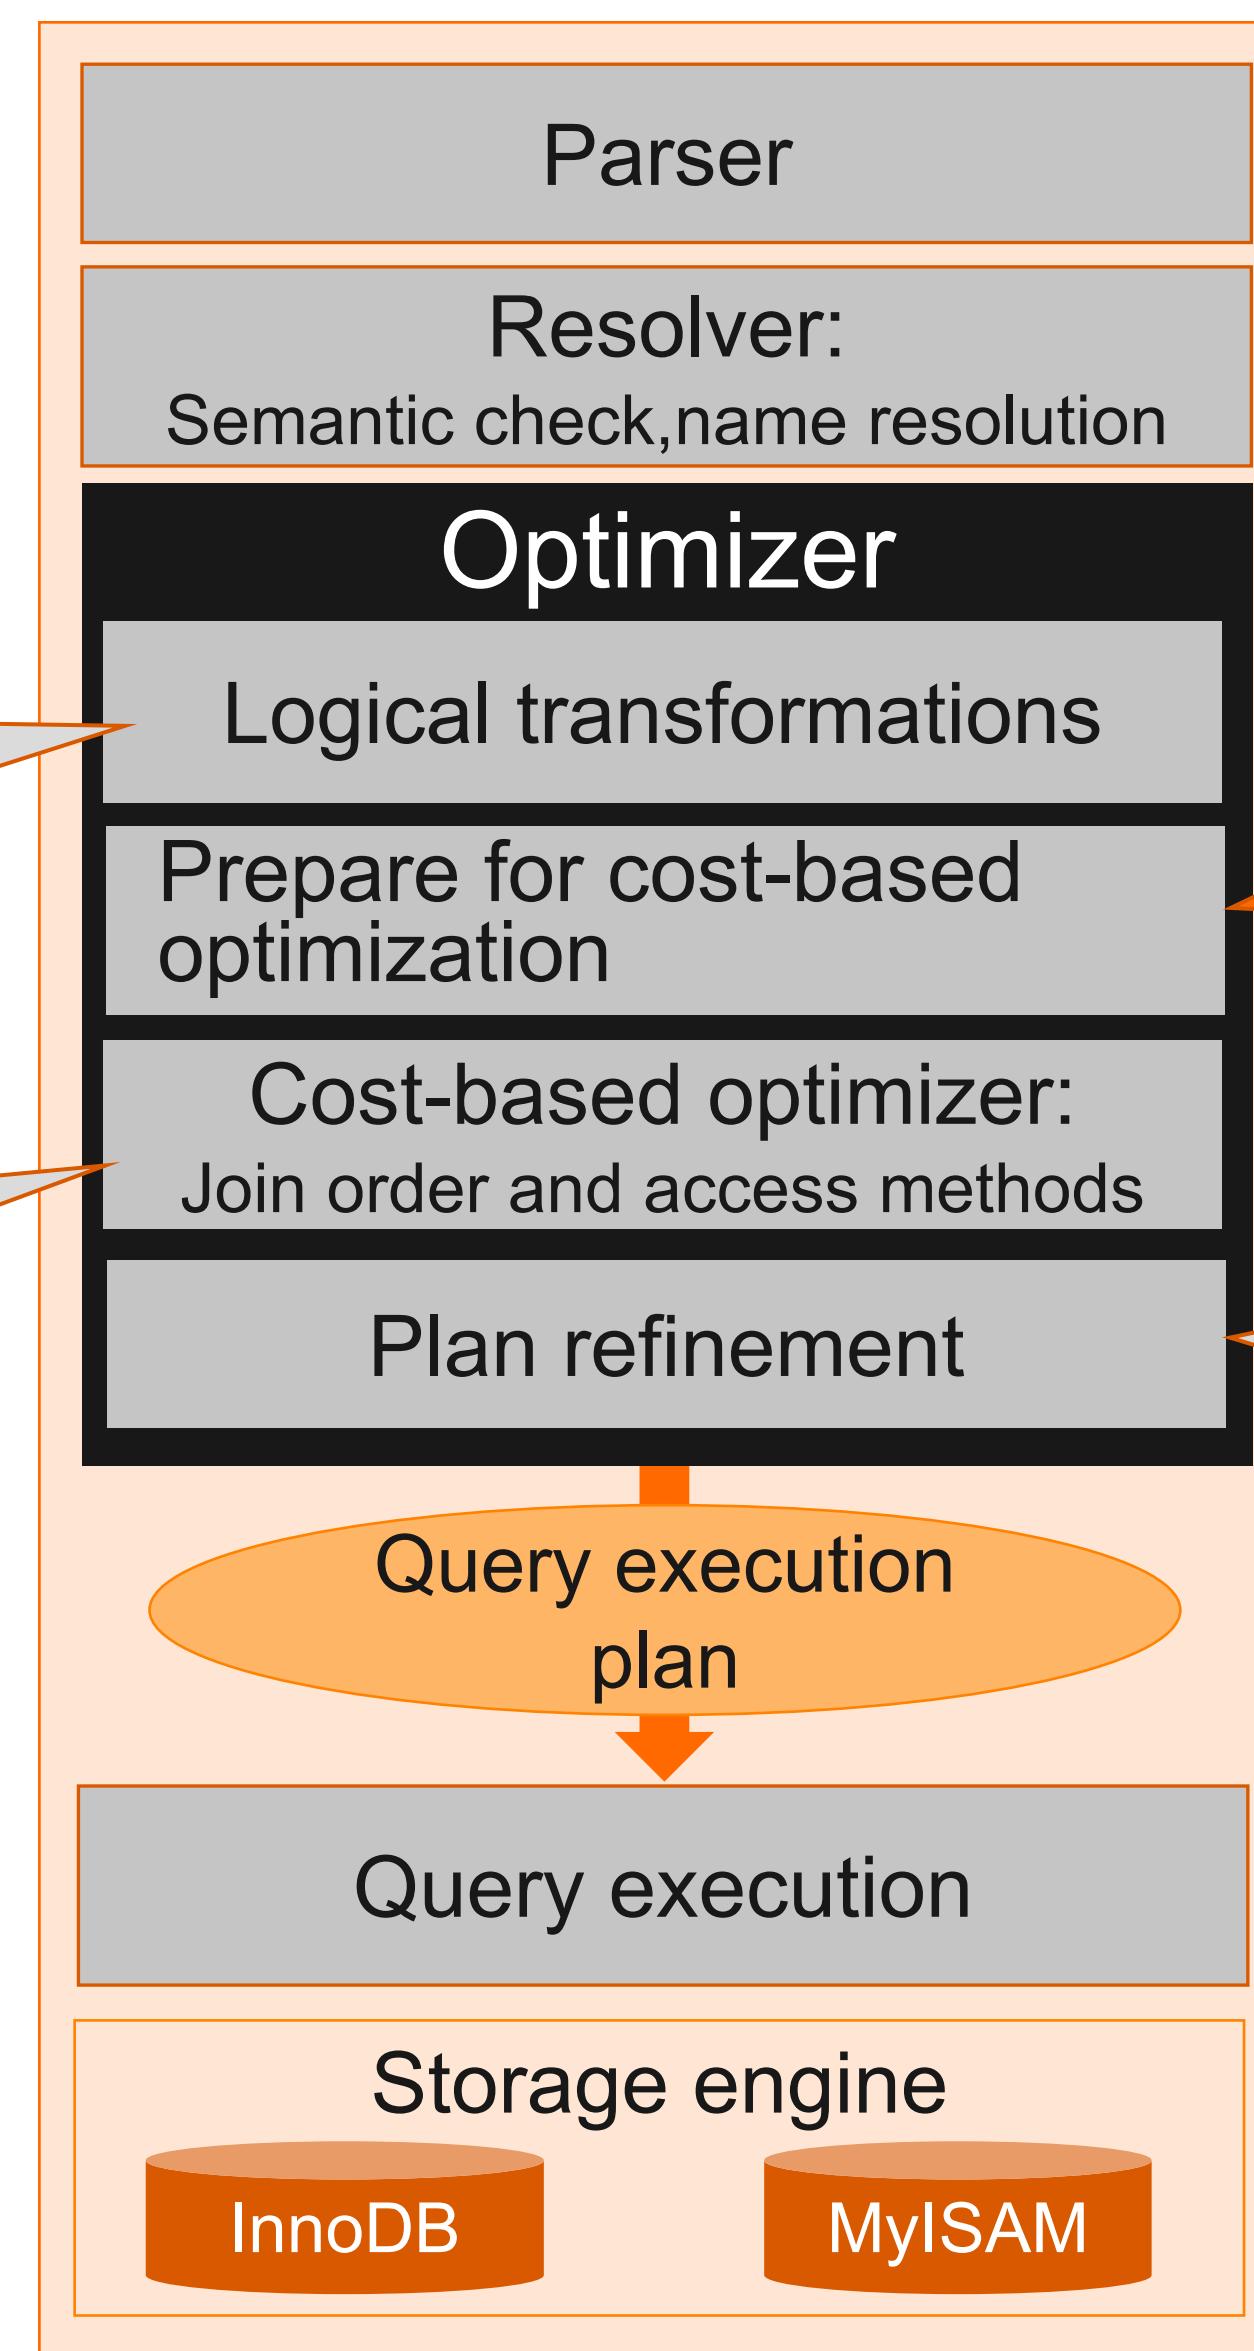
```
"join_optimization": {
    "select#": 1,
    "steps": [
        {
            "condition_processing": {
                "condition": "WHERE",
                "original_condition": "(`t1`.`a` = `t2`.`a`) and (`t2`.`a` = 9) and (((`t1`.`a` <= 10) and (`t2`.`b` <= 3)) or ((`t1`.`b` =(`t2`.`b` + 7)) and (`t2`.`b` = 5)))",
                "steps": [
                    {
                        "transformation": "equality_propagation",
                        "resulting_condition": "(((9 <= 10) and (`t2`.`b` <= 3)) or ((`t1`.`b` = (5 + 7)) and multiple_equal(5, `t2`.`b`))) and multiple_equal(9, `t1`.`a`, `t2`.`a`)"
                    },
                    {
                        "transformation": "constant_propagation",
                        "resulting_condition": "(((9 <= 10) and (`t2`.`b` <= 3)) or ((`t1`.`b` = 12) and multiple_equal(5, `t2`.`b`))) and multiple_equal(9, `t1`.`a`, `t2`.`a`)"
                    },
                    {
                        "transformation": "trivial_condition_removal",
                        "resulting_condition": "(((`t2`.`b` <= 3) or ((`t1`.`b` = 12) and multiple_equal(5, `t2`.`b`))) and multiple_equal(9, `t1`.`a`, `t2`.`a`))"
                    }
                ]
            }
        }
    ]
}
```

Query Optimization

Main phases

Negation elimination
Equality and constant propagation
Evaluation of constant expressions
Substitution of generated columns

Access method selection
Join order



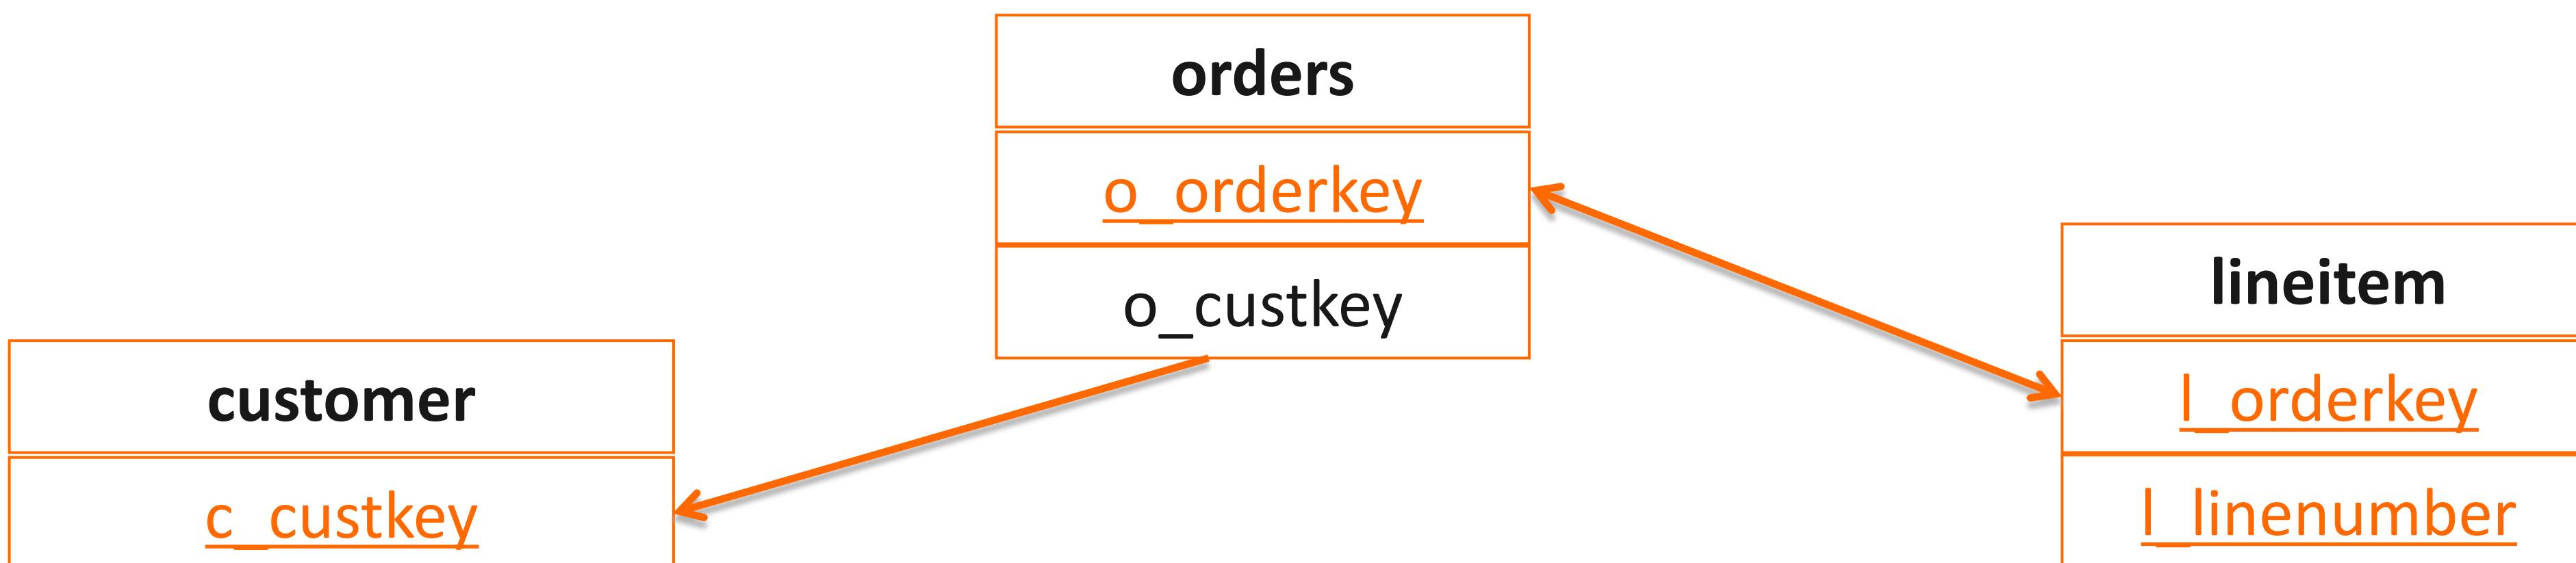
Ref access analysis
Range access analysis
Estimation of condition fan out
Constant table detection

Table condition pushdown
Access method adjustments
Sort avoidance
Index condition pushdown
Prepare temporary tables

Ref Access Analysis

Determine which indexes that can be used for index lookup in a join

```
SELECT l_orderkey, sum(l_extendedprice * (1 - l_discount)) AS revenue,  
       o_orderdate, o_shippriority  
FROM customer JOIN orders ON c_custkey = o_custkey  
                JOIN lineitem ON l_orderkey = o_orderkey  
WHERE c_mktsegment = 'FURNITURE'  
      AND o_orderdate < '1997-04-15' AND l_shipdate > '1997-04-15'  
GROUP by l_orderkey, o_orderdate, o_shippriority  
ORDER by revenue desc, o_orderdate  
LIMIT 10;
```



Ref Access Analysis

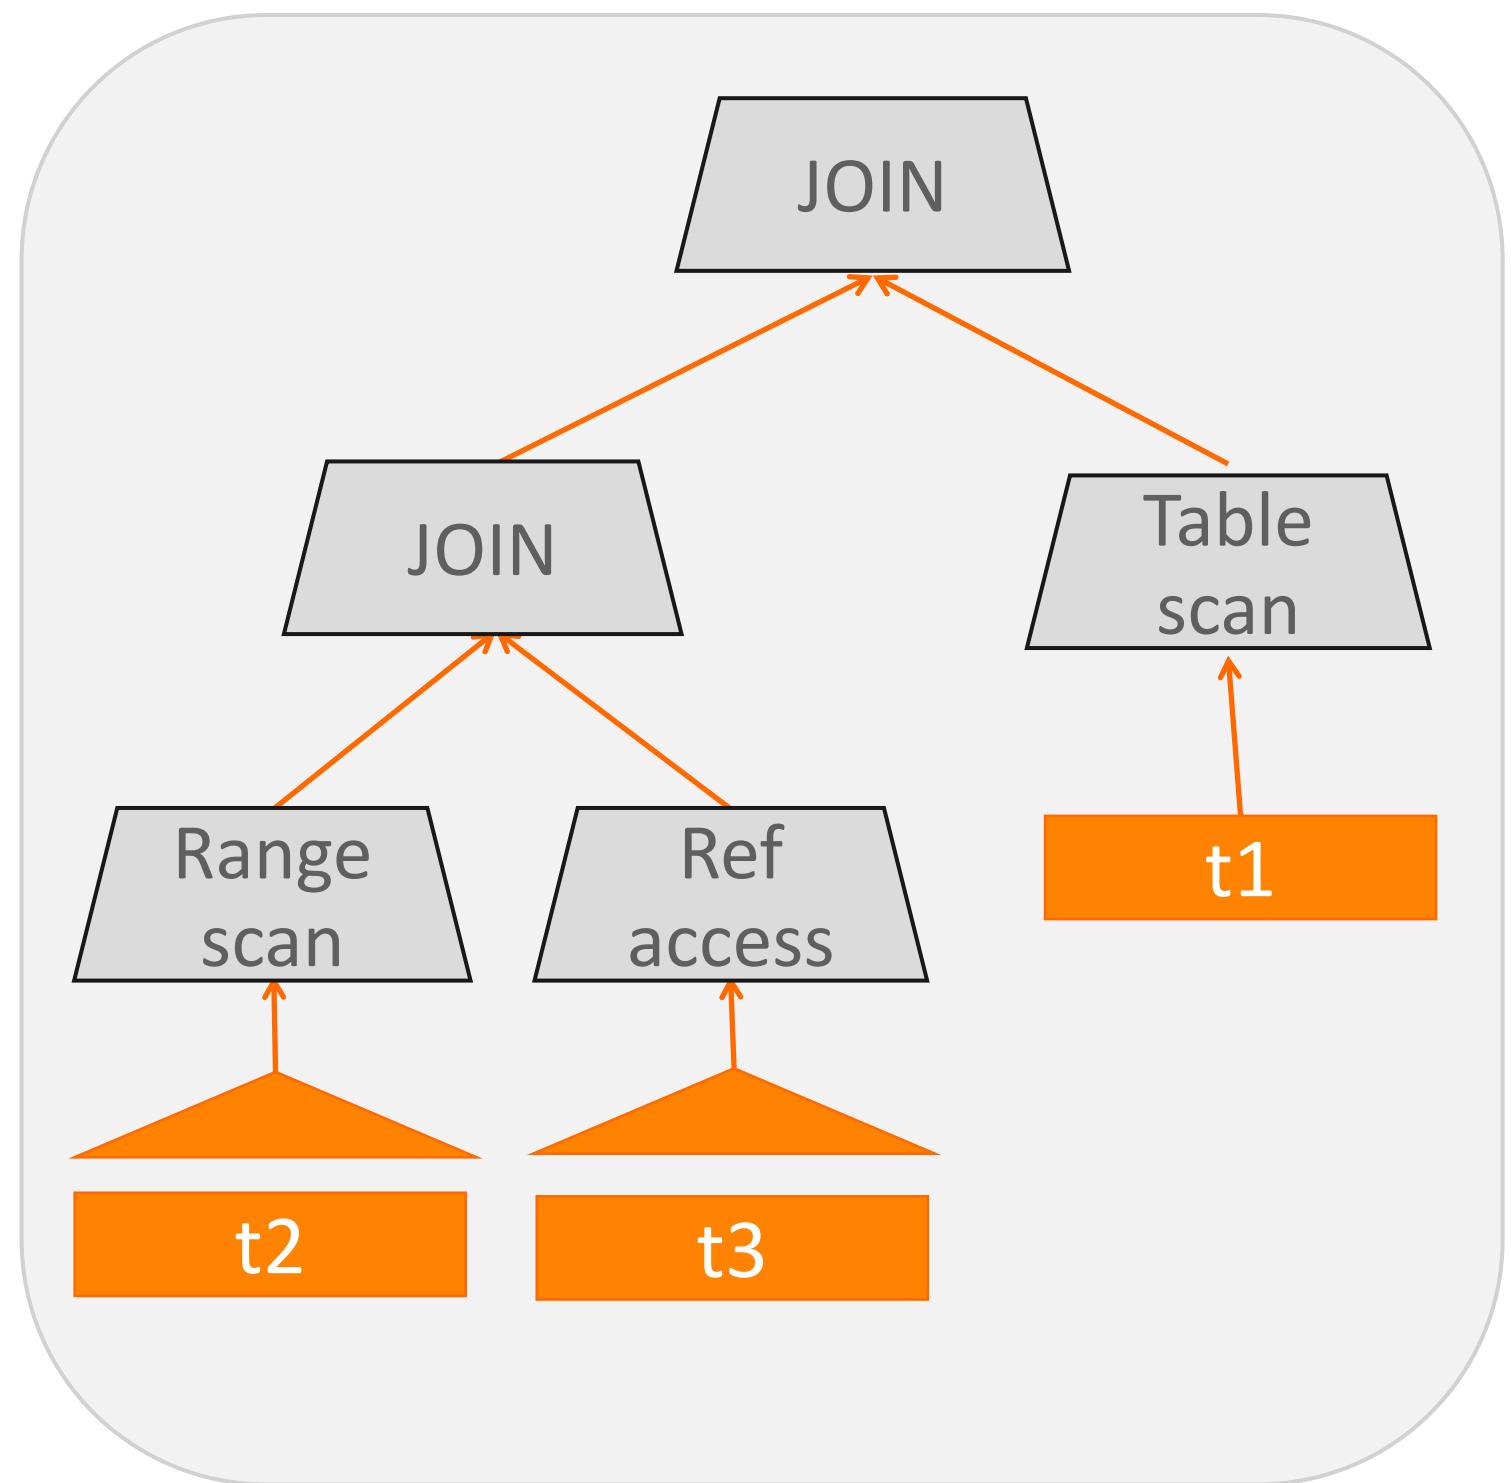
Optimizer trace

```
{  
  "ref_optimizer_key_uses": [  
    {  
      "table": "`customer`",  
      "field": "c_custkey",  
      "equals": "`orders`.`o_custkey`",  
      "null_rejecting": true  
    },  
    {  
      "table": "`orders`",  
      "field": "o_orderkey",  
      "equals": "`lineitem`.`l_orderkey`",  
      "null_rejecting": false  
    },  
    {  
      "table": "`lineitem`",  
      "field": "l_orderkey",  
      "equals": "`orders`.`o_orderkey`",  
      "null_rejecting": false  
    }  
  ]  
},
```

Cost-based Query Optimization

General idea:

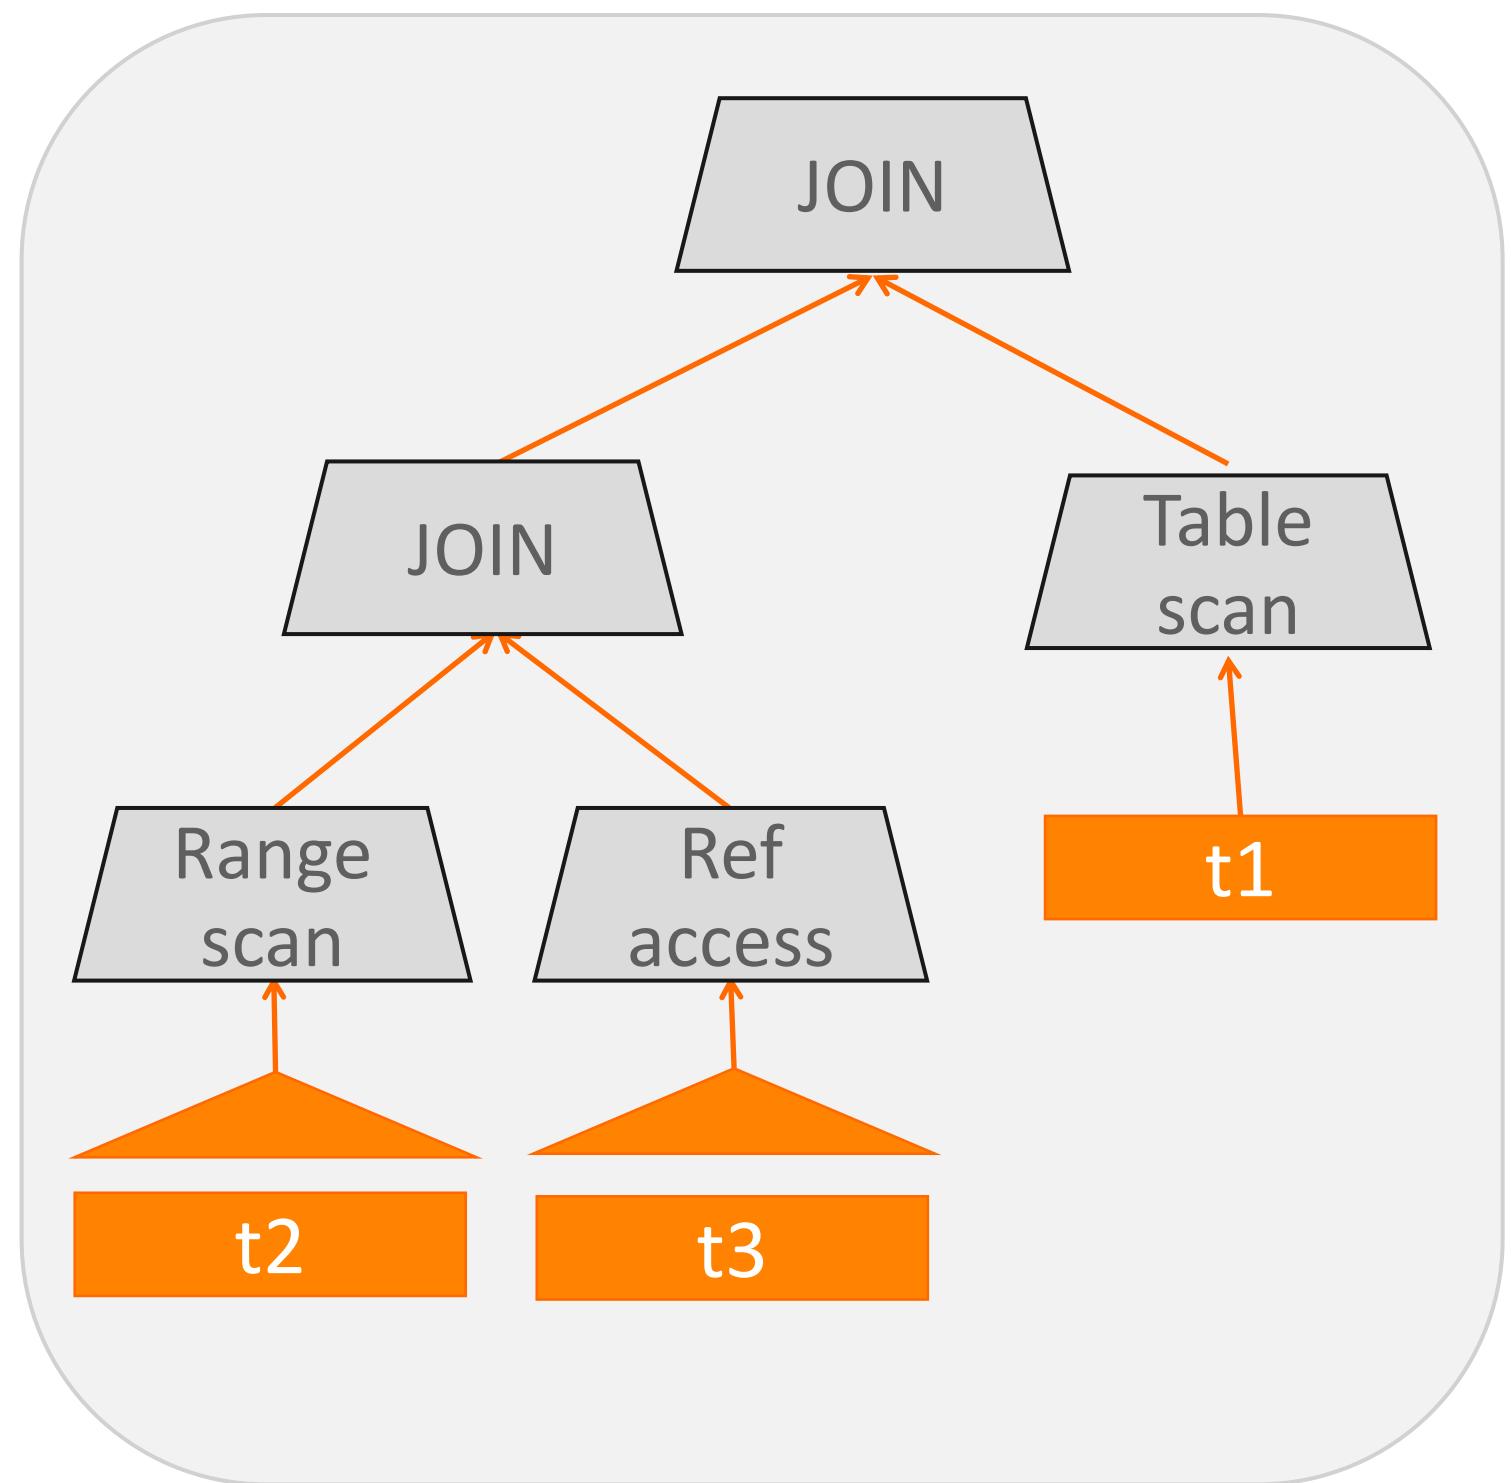
- Assign cost to operations
- Assign cost to partial or alternative plans
- Search for plan with lowest cost



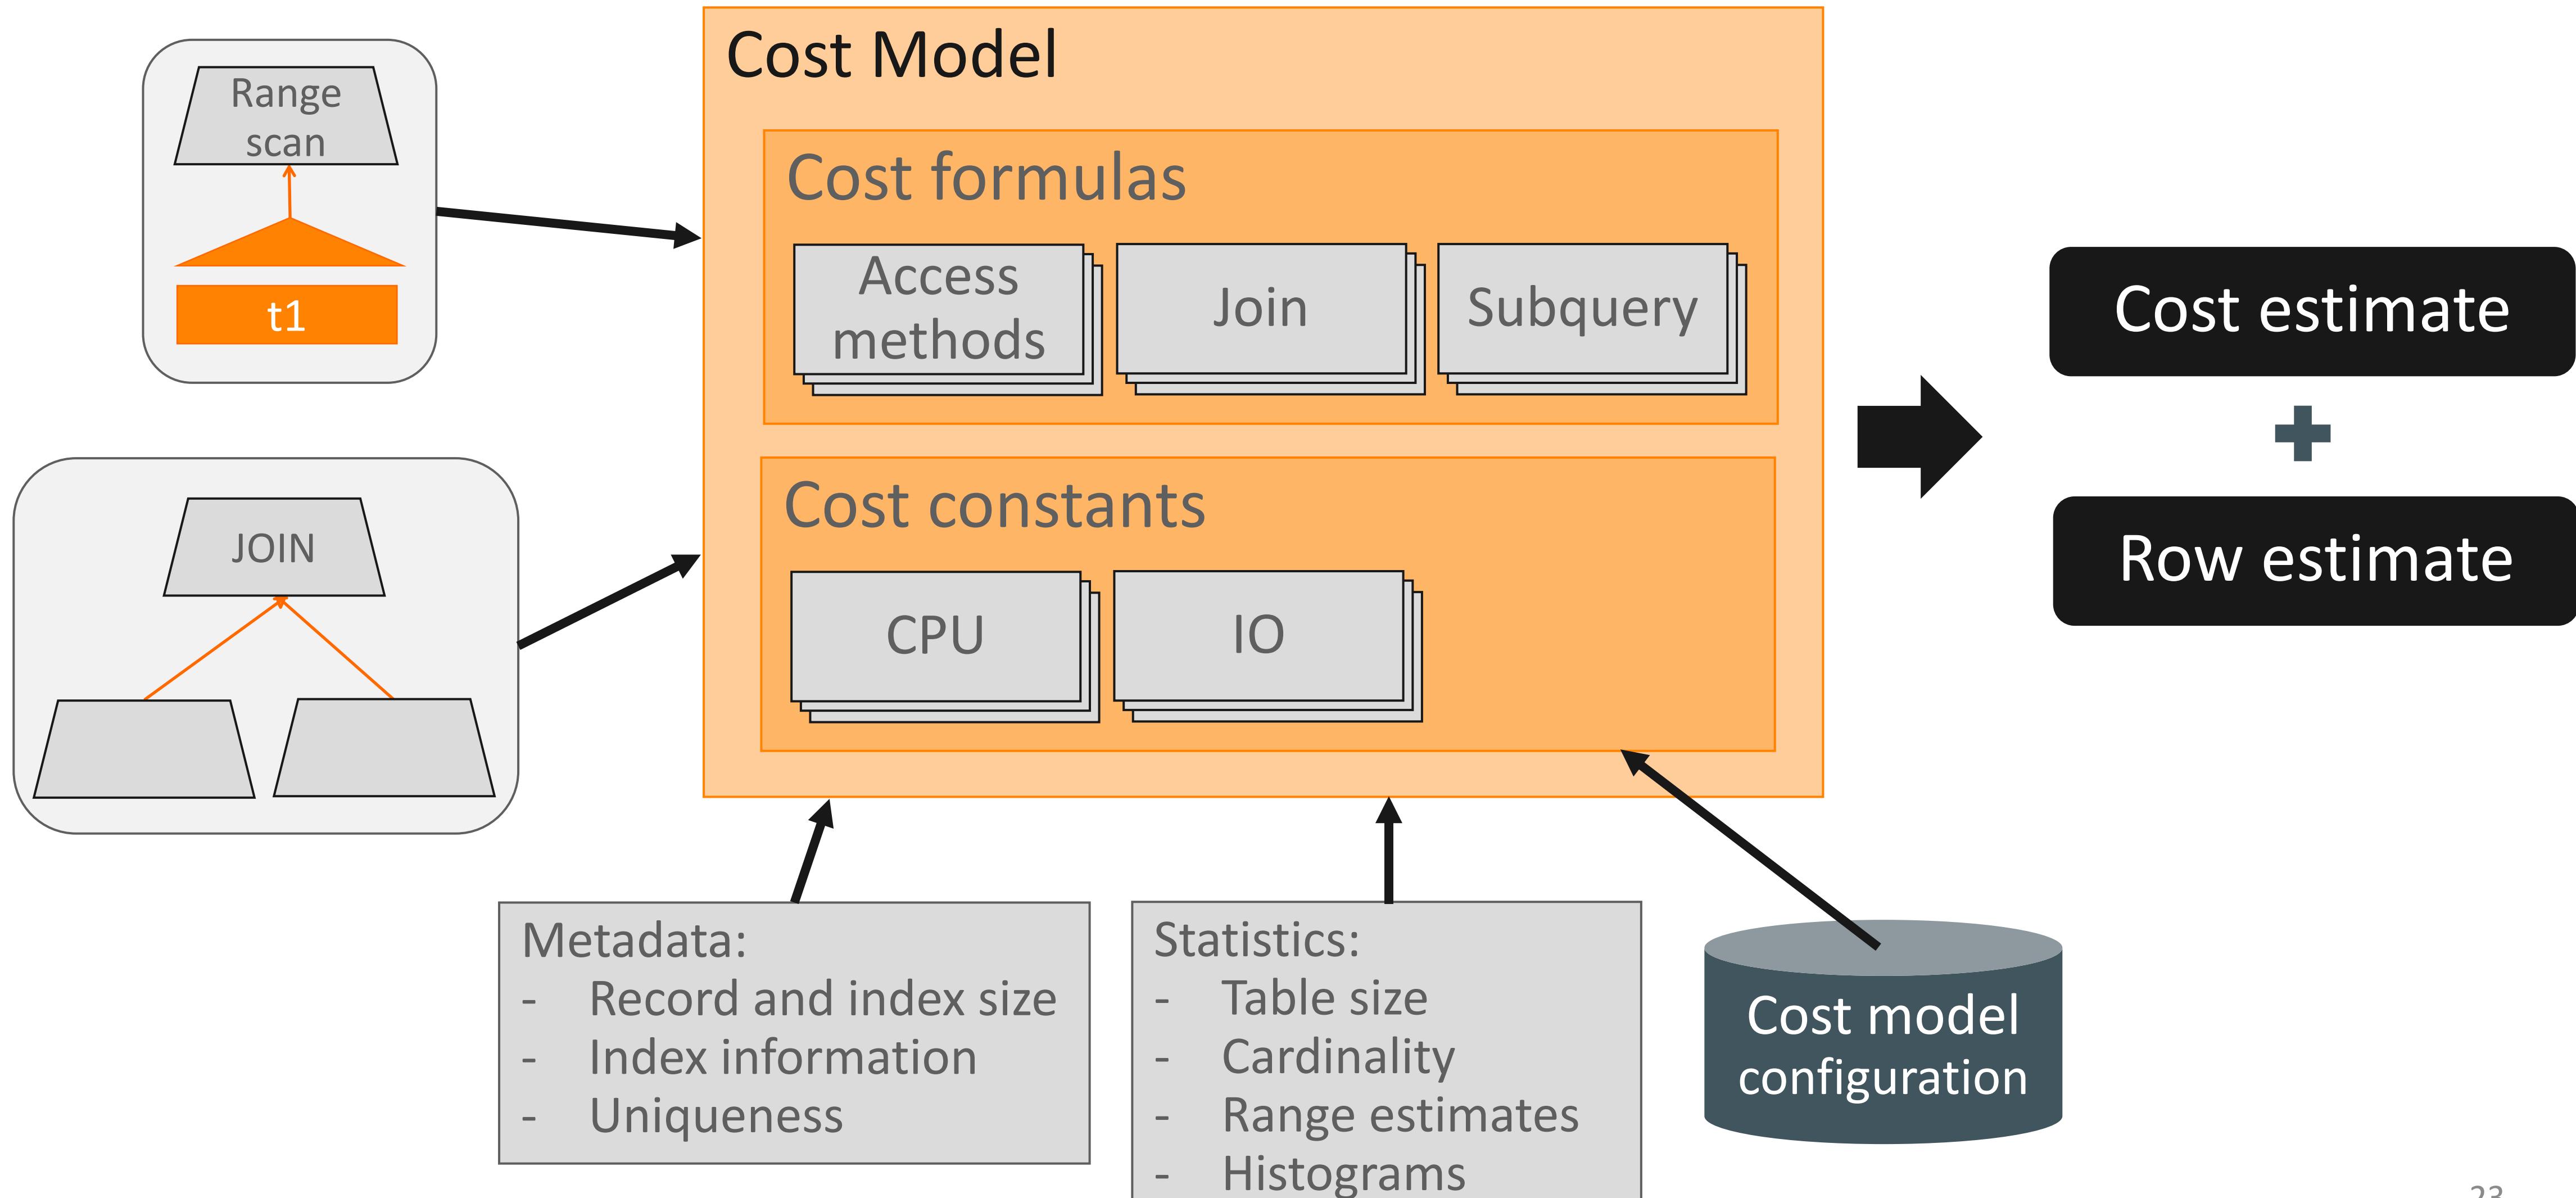
Cost-based Query Optimizations

The main cost-based optimizations:

- Index and access method:
 - Table scan
 - Index scan
 - Range scan
 - Index lookup (ref access)
- Join order
- Join buffering strategy
- Subquery strategy



Optimizer Cost Model



Cost Estimates

- The **cost** for **executing** a query
- **Cost unit:**
 - “read a random data page from disk”
- Main cost factors:
 - IO cost:
 - #pages read from table
 - #pages read from index
 - CPU cost:
 - Evaluating query conditions
 - Comparing keys/records
 - Sorting keys

Main cost constants

Cost	MySQL 5.7	MySQL 8.0
Read a random disk page	1.0	1.0
Read a data page from memory buffer	1.0	0.25
Evaluate query condition	0.2	0.1
Compare keys/records	0.1	0.05

Range analysis

Compare cost of table scan to index scan

- Cost of table scan
 - Based on number of pages to read
- Cost of index scan alternatives
 - Index scan (covering)
 - Index range scan
 - Index merge
 - Index for grouping
 - Skip scan (new in 8.0)
 - Loose-index scan

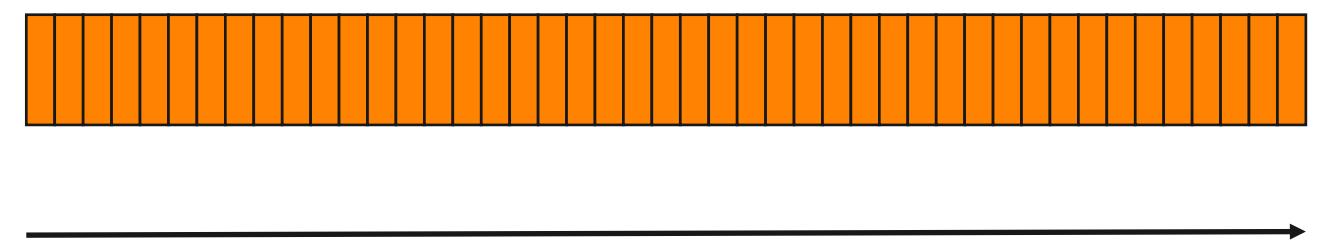
Range Analysis Example

Table scan vs Index range scan

```
SELECT SUM(o_totalprice) FROM orders  
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```

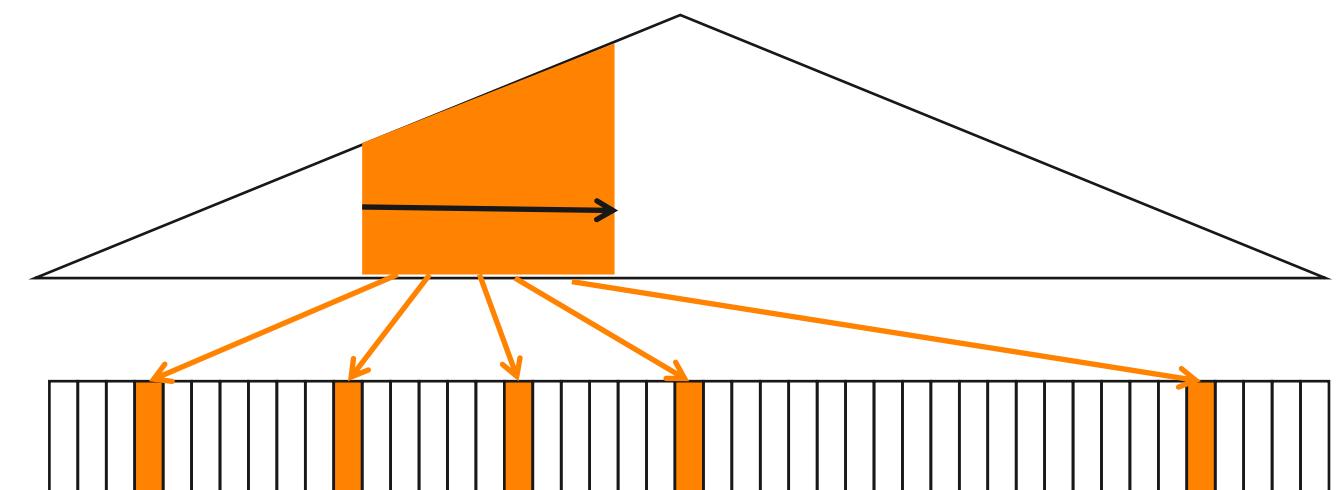
Table scan:

- IO-cost: #pages in table * IO_BLOCK_READ_COST
- CPU cost: #rows * ROW_EVALUATE_COST



Range scan (on secondary index):

- IO-cost: #rows_in_range * IO_BLOCK_READ_COST
- CPU cost: #rows_in_range * ROW_EVALUATE_COST



Range Analysis Example

Table scan vs index range scan cnt.

**EXPLAIN SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';**

id	select type	table	type	possible keys	key	key len	rows	filtered	Extra
1	SIMPLE	orders	ALL	i_o_orderdate	NULL	NULL	1480845	31.13	Using where

**EXPLAIN SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-06-30';**

Id	select type	table	type	possible keys	key	key len	rows	filtered	Extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	222102	100.00	Using index condition

Range Analysis Example

Optimizer trace

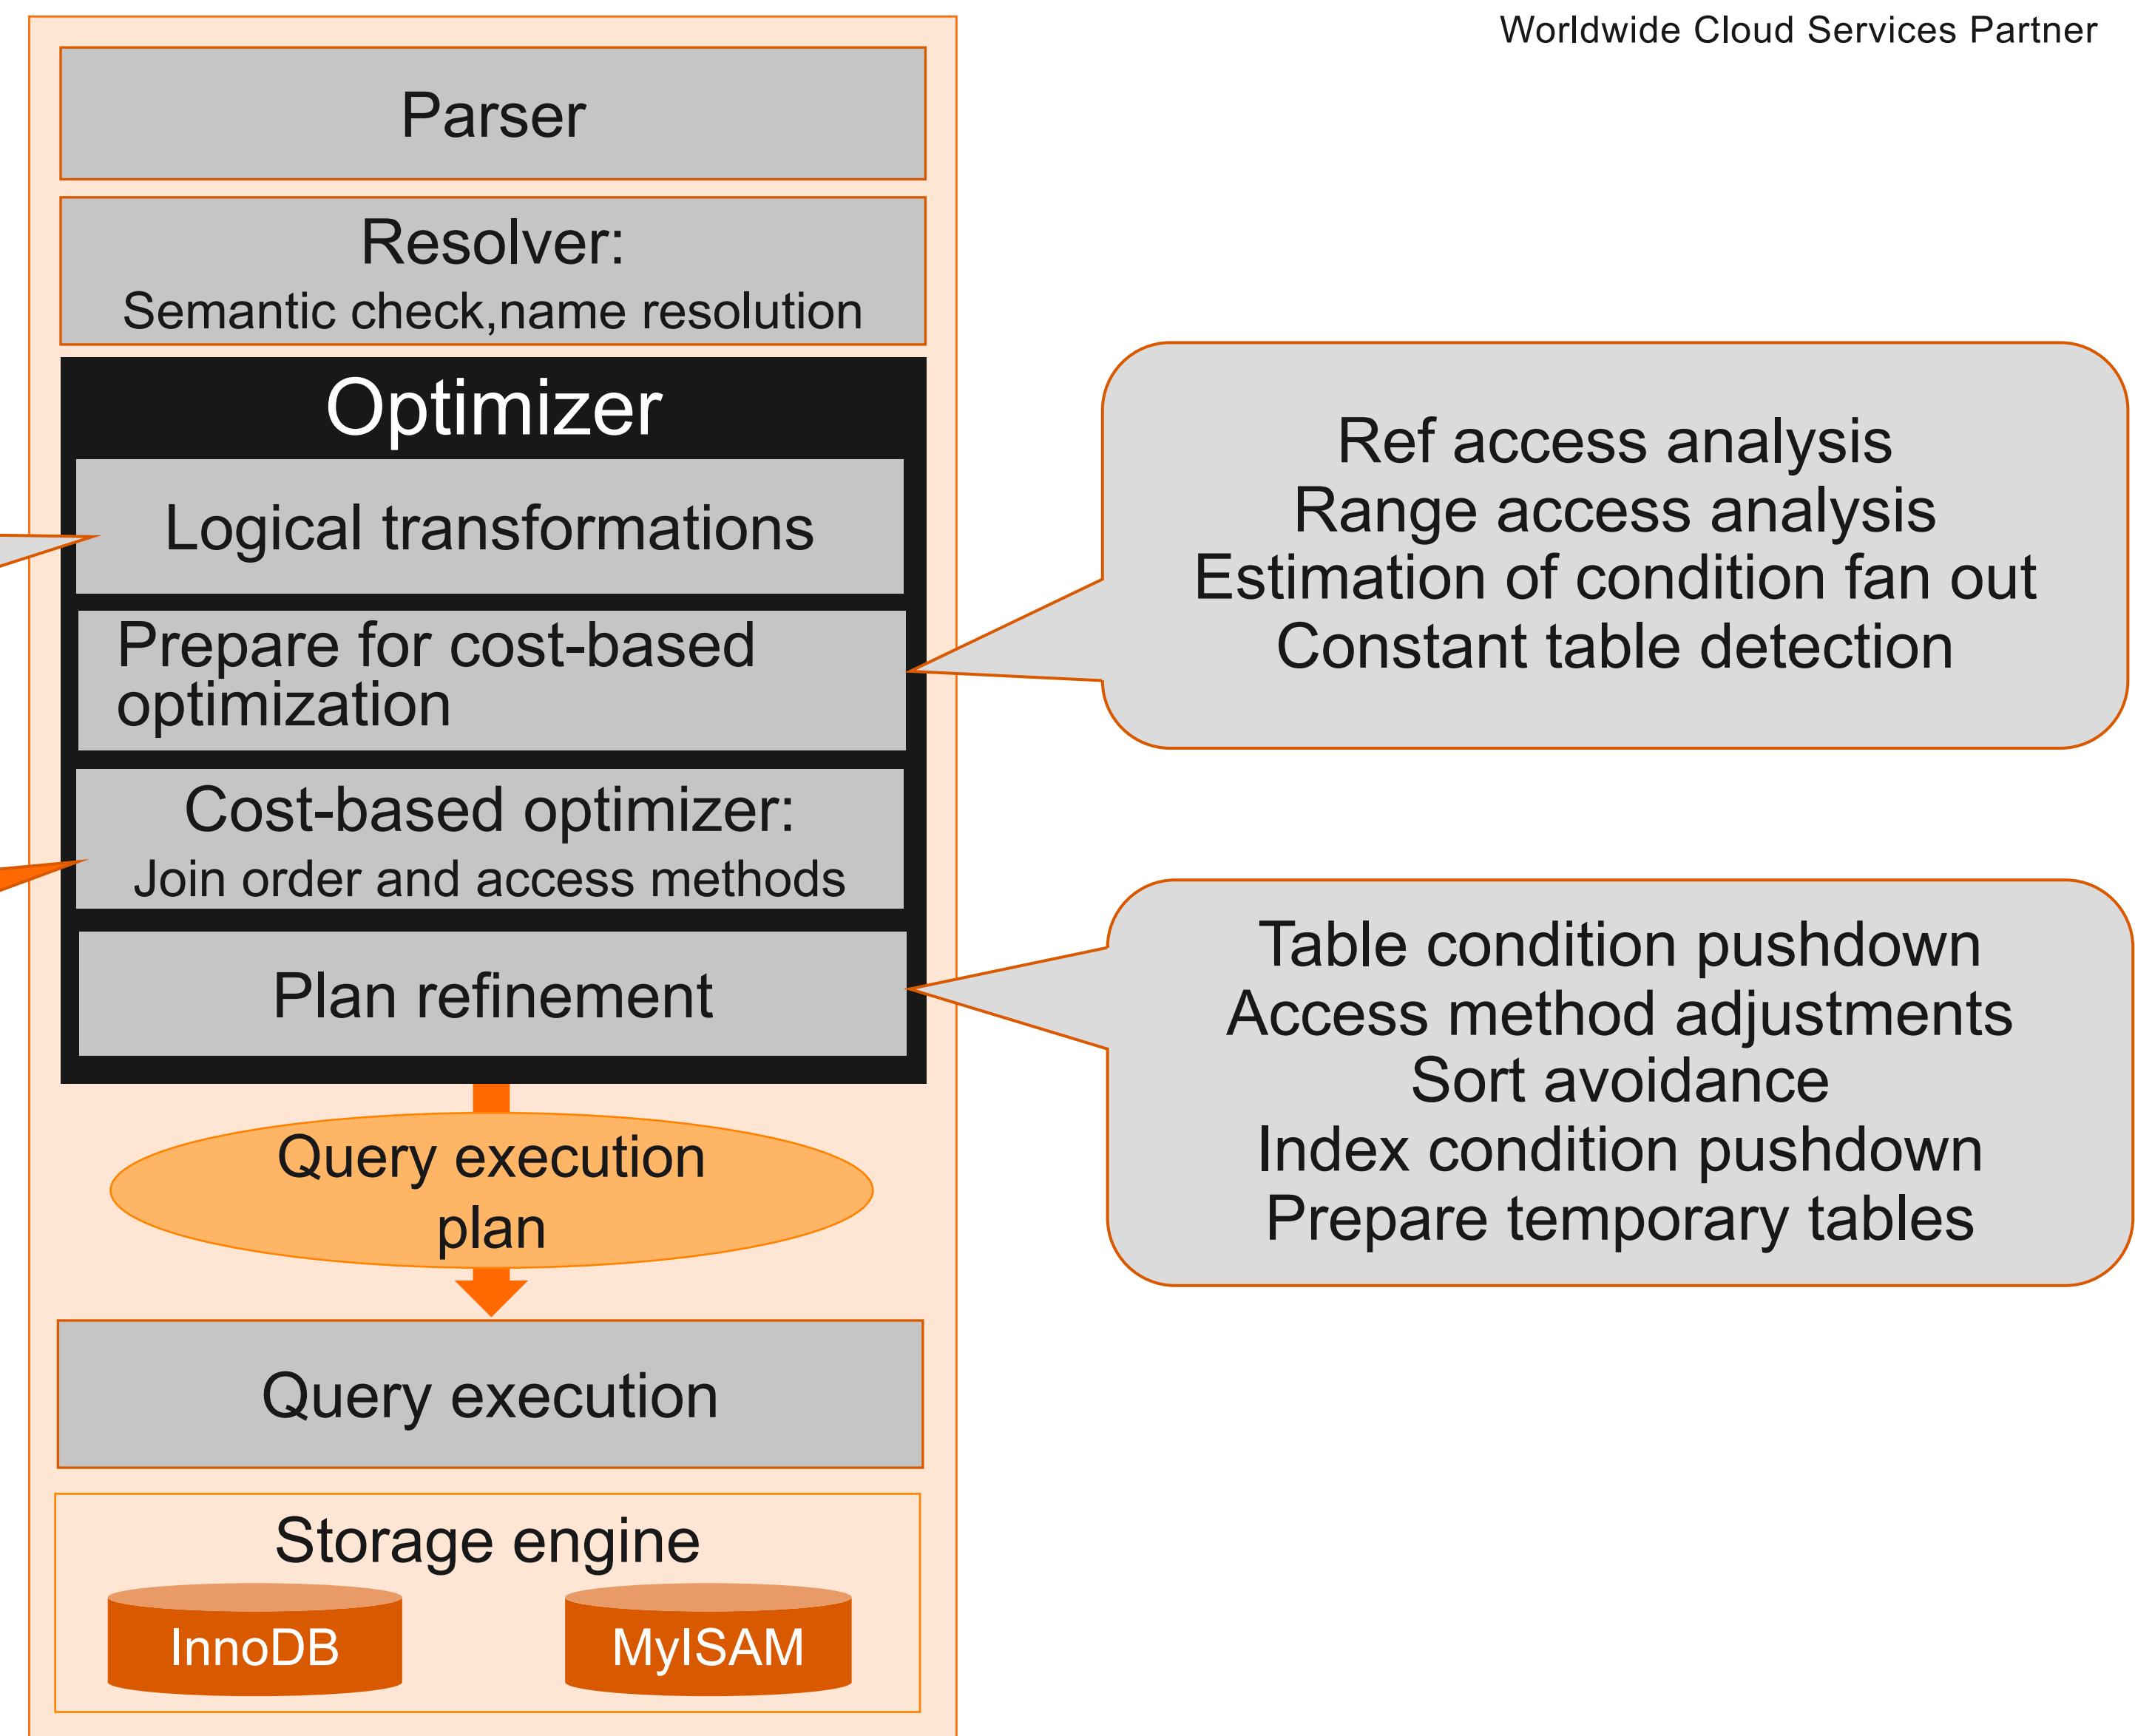
```
"rows_estimation": [
  {
    "table": "`orders`",
    "range_analysis": {
      "table_scan": {
        "rows": 1480845,
        "cost": 151138
      },
      "potential_range_indexes": [
        {
          "index": "PRIMARY",
          "usable": false,
          "cause": "not_applicable"
        },
        {
          "index": "i_o_orderdate",
          "usable": true,
          "key_parts": [
            "o_orderDATE",
            "o_orderkey"
          ]
        }
      ],
      "setup_range_conditions": [
        ],
      "group_index_range": {
        "chosen": false,
        "cause": "not_group_by_or_distinct"
      },
      "skip_scan_range": {
        "potential_skip_scan_indexes": [
          {
            "index": "i_o_orderdate",
            "usable": false,
            "cause": "query_references_nonkey_column"
          }
        ],
        "analyze_range_alternatives": [
          {
            "range_scan_alternatives": [
              {
                "index": "i_o_orderdate",
                "ranges": [
                  "0x21940f <= o_orderDATE <= 0x9f950f"
                ],
                "index_dives_for_eq_ranges": true,
                "rowid_ordered": false,
                "using_mrr": false,
                "index_only": false,
                "rows": 460938,
                "cost": 161329,
                "chosen": false,
                "cause": "cost"
              }
            ],
            "analyze_roworder_intersect": [
              {
                "usable": false,
                "cause": "too_few_roworder_scans"
              }
            }
          }
        ]
      }
    }
  }
]
```

Query Optimization

Main phases

Negation elimination
Equality and constant propagation
Evaluation of constant expressions
Substitution of generated columns

Join order
Access method selection

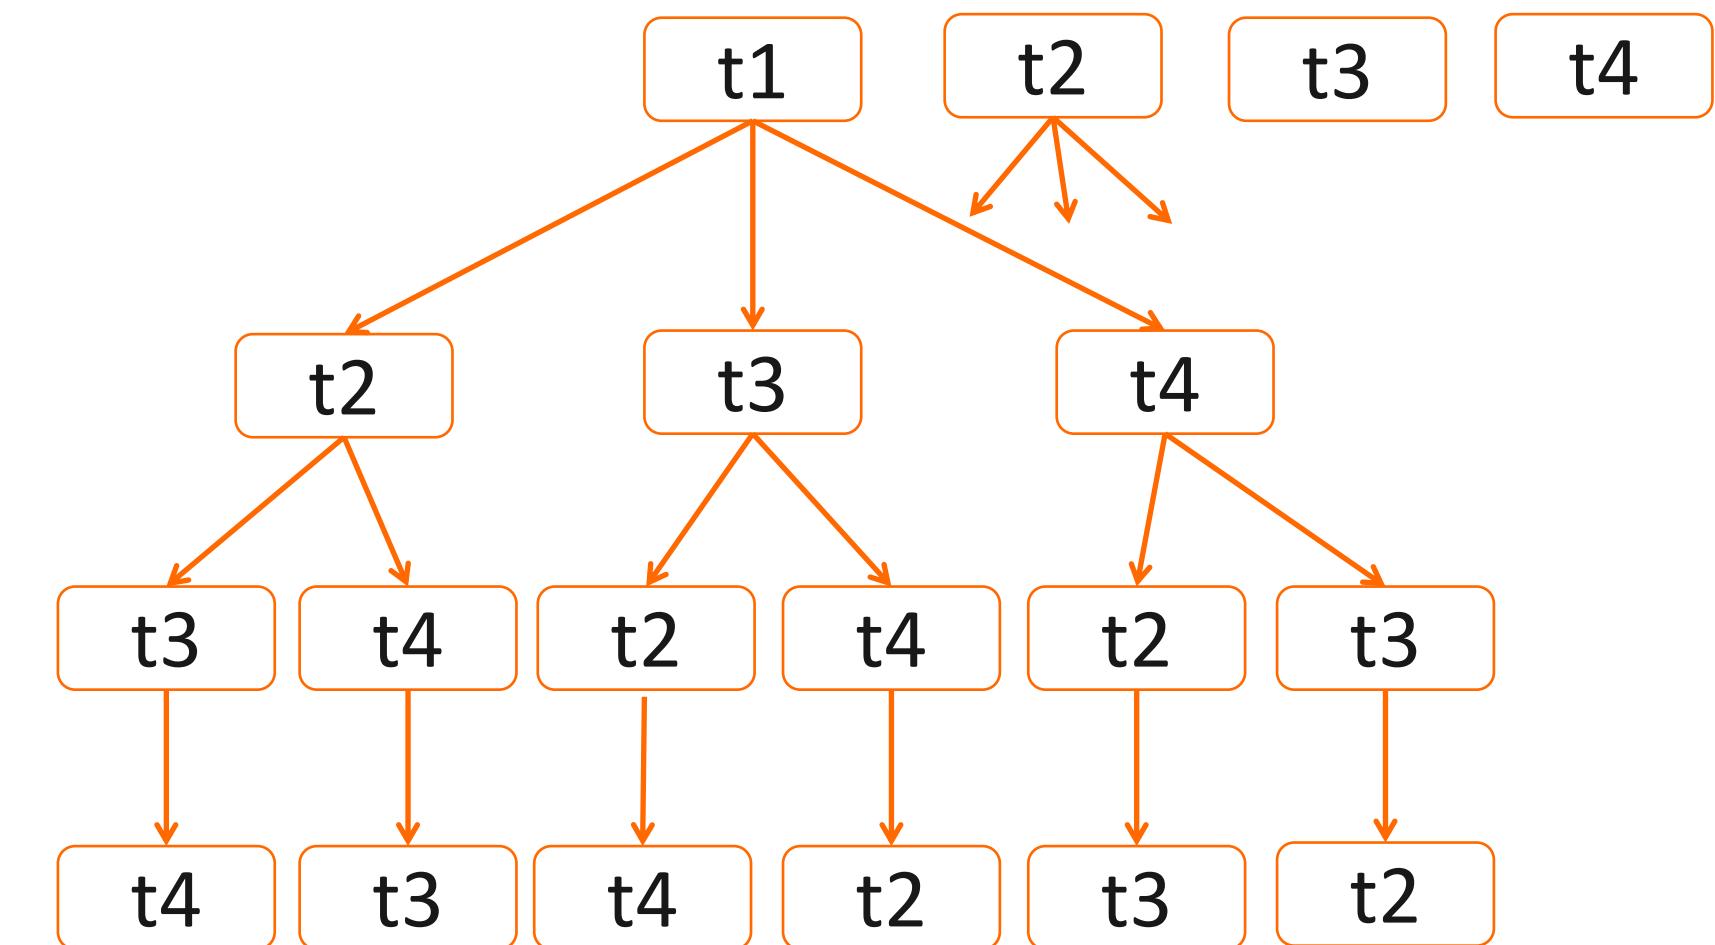


Join Optimizer

"Greedy search strategy"

- Goal: Given a JOIN of N tables, find the best JOIN ordering
- Only considers left-deep plans
- Strategy:
 - Start with all 1-table plans
 - Sorted based on size and *key dependency*
 - Expand each plan with remaining tables
 - Depth-first
 - If “cost of partial plan” > “cost of best plan”:
 - “Prune” plan
 - Heuristic pruning:
 - Prune less promising partial plans
 - May in rare cases miss most optimal plan (turn off with **set optimizer_prune_level = 0**)

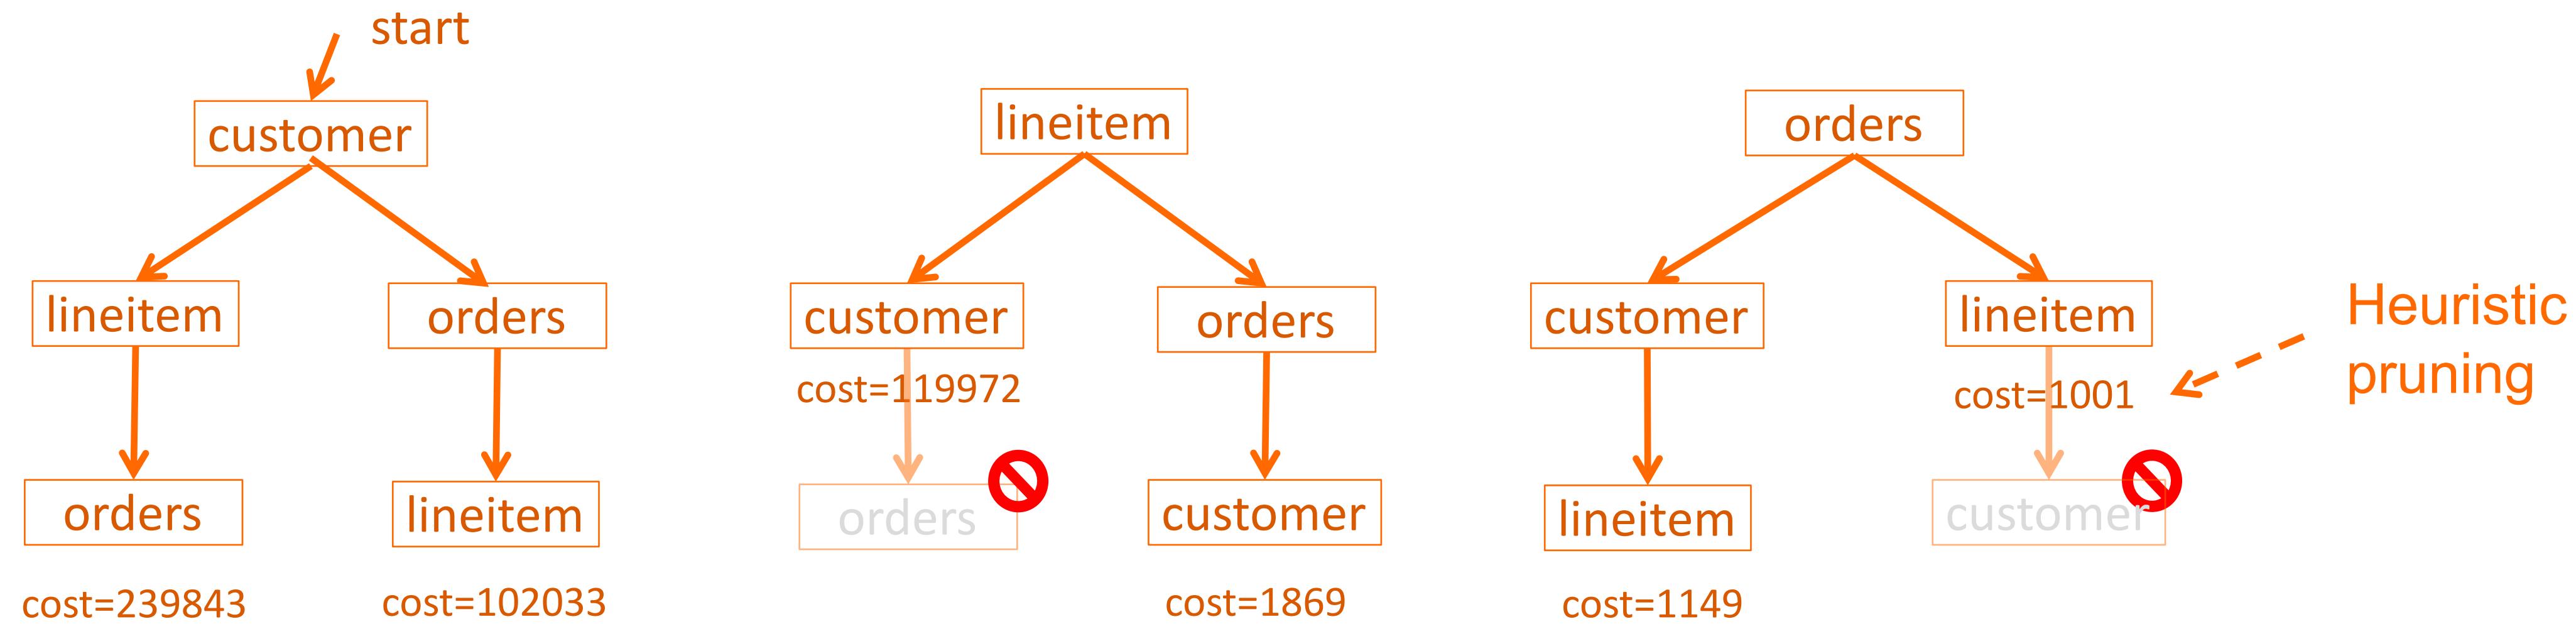
$N!$ possible
plans



Join Optimizer Illustrated

DBT3-Q3

```
SELECT l_orderkey, sum(l_extendedprice * (1 - l_discount)) AS revenue,  
       o_orderdate, o_shippriority  
FROM customer JOIN orders ON c_custkey = o_custkey  
               JOIN lineitem ON l_orderkey = o_orderkey  
WHERE c_mktsegment = 'FURNITURE'  
  AND o_orderdate < '1997-04-15' AND l_shipdate > '1997-04-15'  
...
```



Join Optimizer

Optimizer trace

```
▼ considered_execution_plans: [
  ▼ {
    plan_prefix: [ ],
    table: "`customer`",
    ▶ best_access_path: { ... },
    condition_filtering_pct: 100,
    rows_for_plan: 478.9,
    cost_for_plan: 244.2,
    ▶ rest_of_plan: [
      ▼ {
        ▼ plan_prefix: [
          "`customer`"
        ],
        table: "`lineitem`",
        ▶ best_access_path: { ... },
        condition_filtering_pct: 100,
        rows_for_plan: 1.19e6,
        cost_for_plan: 119735,
        ▶ rest_of_plan: [
          ▼ {
            ▼ plan_prefix: [
              "`customer`",
              "`lineitem`"
            ],
            table: "`orders`",
            ▶ best_access_path: { ... },
            condition_filtering_pct: 5,
            rows_for_plan: 59742,
            cost_for_plan: 239843,
            chosen: true
          }
        ]
      }
    ]
  }
  ▼ {
    plan_prefix: [
      "`customer`"
    ],
    table: "`orders`",
    ▶ best_access_path: { ... },
    condition_filtering_pct: 0.0471,
    rows_for_plan: 478.9,
    cost_for_plan: 101865,
    ▶ rest_of_plan: [
      ▼ {
        ▼ plan_prefix: [
          "`customer`",
          "`orders`"
        ],
        table: "`lineitem`",
        ▶ best_access_path: { ... },
        condition_filtering_pct: 100,
        rows_for_plan: 486.3,
        cost_for_plan: 102033,
        chosen: true
      }
    ]
  },
  ▼ {
    plan_prefix: [ ],
    table: "`lineitem`",
    ▶ best_access_path: { ... },
    condition_filtering_pct: 100,
    rows_for_plan: 486.3,
    cost_for_plan: 102033
  }
],
```

Choosing Access Path

Ref access vs. table/index scan

- Ref Access (index look-up)
 - Read all records with a given key value using an index
 - Examples:
SELECT * FROM t1 WHERE t1.key = 7;
SELECT * FROM t1, t2 WHERE t1.key = t2.key;
 - “eq_ref”:
 - Reading from a unique index, max one record returned
 - “ref”:
 - Reading from a non-unique index or a prefix of an unique index, possibly multiple records returned
 - The record estimate is based on cardinality number from index statistics
 - Indexes available for ref access depends on the join prefix
- Table/index scan
 - The “winner” of the earlier “range_analysis”
- Also considers:
 - Join buffering (BNL/BKA)
 - Filtering effects of conditions

Choosing the Access Path

customer

```
"plan_prefix": [
],
"table": "customer",
"best_access_path": {
  "considered_access_paths": [
    {
      "access_type": "ref",
      "index": "PRIMARY",
      "usable": false,
      "chosen": false
    },
    {
      "rows_to_scan": 2367,
      "filtering_effect": [
        {
          "condition": "(`customer`.`c_mktsegment` = 'FURNITURE')",
          "histogram_selectivity": 0.2023
        }
      ],
      "final_filtering_effect": 0.2023,
      "access_type": "scan",
      "resulting_rows": 478.9,
      "cost": 244.2,
      "chosen": true
    }
  ],
  "condition_filtering_pct": 100,
  "rows_for_plan": 478.9,
  "cost_for_plan": 244.2
}
```

Choosing the Access Path

customer → lineitem

```
"rest_of_plan": [
  {
    "plan_prefix": [
      "customer"
    ],
    "table": "`lineitem`",
    "best_access_path": {
      "considered_access_paths": [
        {
          "access_type": "ref",
          "index": "PRIMARY",
          "usable": false,
          "chosen": false
        },
        {
          "access_type": "ref",
          "index": "i_l_orderkey",
          "usable": false,
          "chosen": false
        }
      ]
    }
  }
]
```

Choosing the Access Path

customer → lineitem → orders

```
"rest_of_plan": [
  {
    "plan_prefix": [
      `customer`,
      `lineitem`
    ],
    "table": "`orders`",
    "best_access_path": {
      "considered_access_paths": [
        {
          "access_type": "eq_ref",
          "index": "PRIMARY",
          "rows": 1,
          "cost": 120108,
          "chosen": true,
          "cause": "clustered_pk_chosen_by_heuristics"
        },
      ]
    }
  }
]
```

```
{
  "rows_to_scan": 2502,
  "filtering_effect": [
  ],
  "final_filtering_effect": 0.8477,
  "access_type": "scan",
  "using_join_cache": true,
  "buffers_needed": 292,
  "resulting_rows": 2121,
  "cost": 2.53e8,
  "chosen": false
}
],
"condition_filtering_pct": 5,
"rows_for_plan": 59742,
"cost_for_plan": 239843,
"chosen": true
}
]
```

Condensed Trace For the Join Optimizer

- With many tables, the trace for join optimization may be thousands of lines
 - The trace for DBT3 Q8 (8 tables) is 16000 lines
 - Beware: Size for optimizer trace is limited by session variable `optimizer_trace_max_mem_size`
 - Default MySQL 5.7: 16 kB
 - Default MySQL 8.0: 1 MB
 - If `information_schema.optimizer_trace.missing_bytes_beyond_max_mem_size > 0`, increase `optimizer_trace_max_mem_size`
- `joinopttrace.js`
 - A script to present the trace in a condensed form
 - Available at <https://github.com/ogrovlen/opttrace>
 - Usage

```
node joinopttrace.js tracefile
```
- Please:
 - Test it out
 - Suggest/submit improvements
 - Report issues

joinopttrace

Example output (DBT3-Q3)

Table AccessType:IndexName Rows/Cost TotalRows/TotalCost

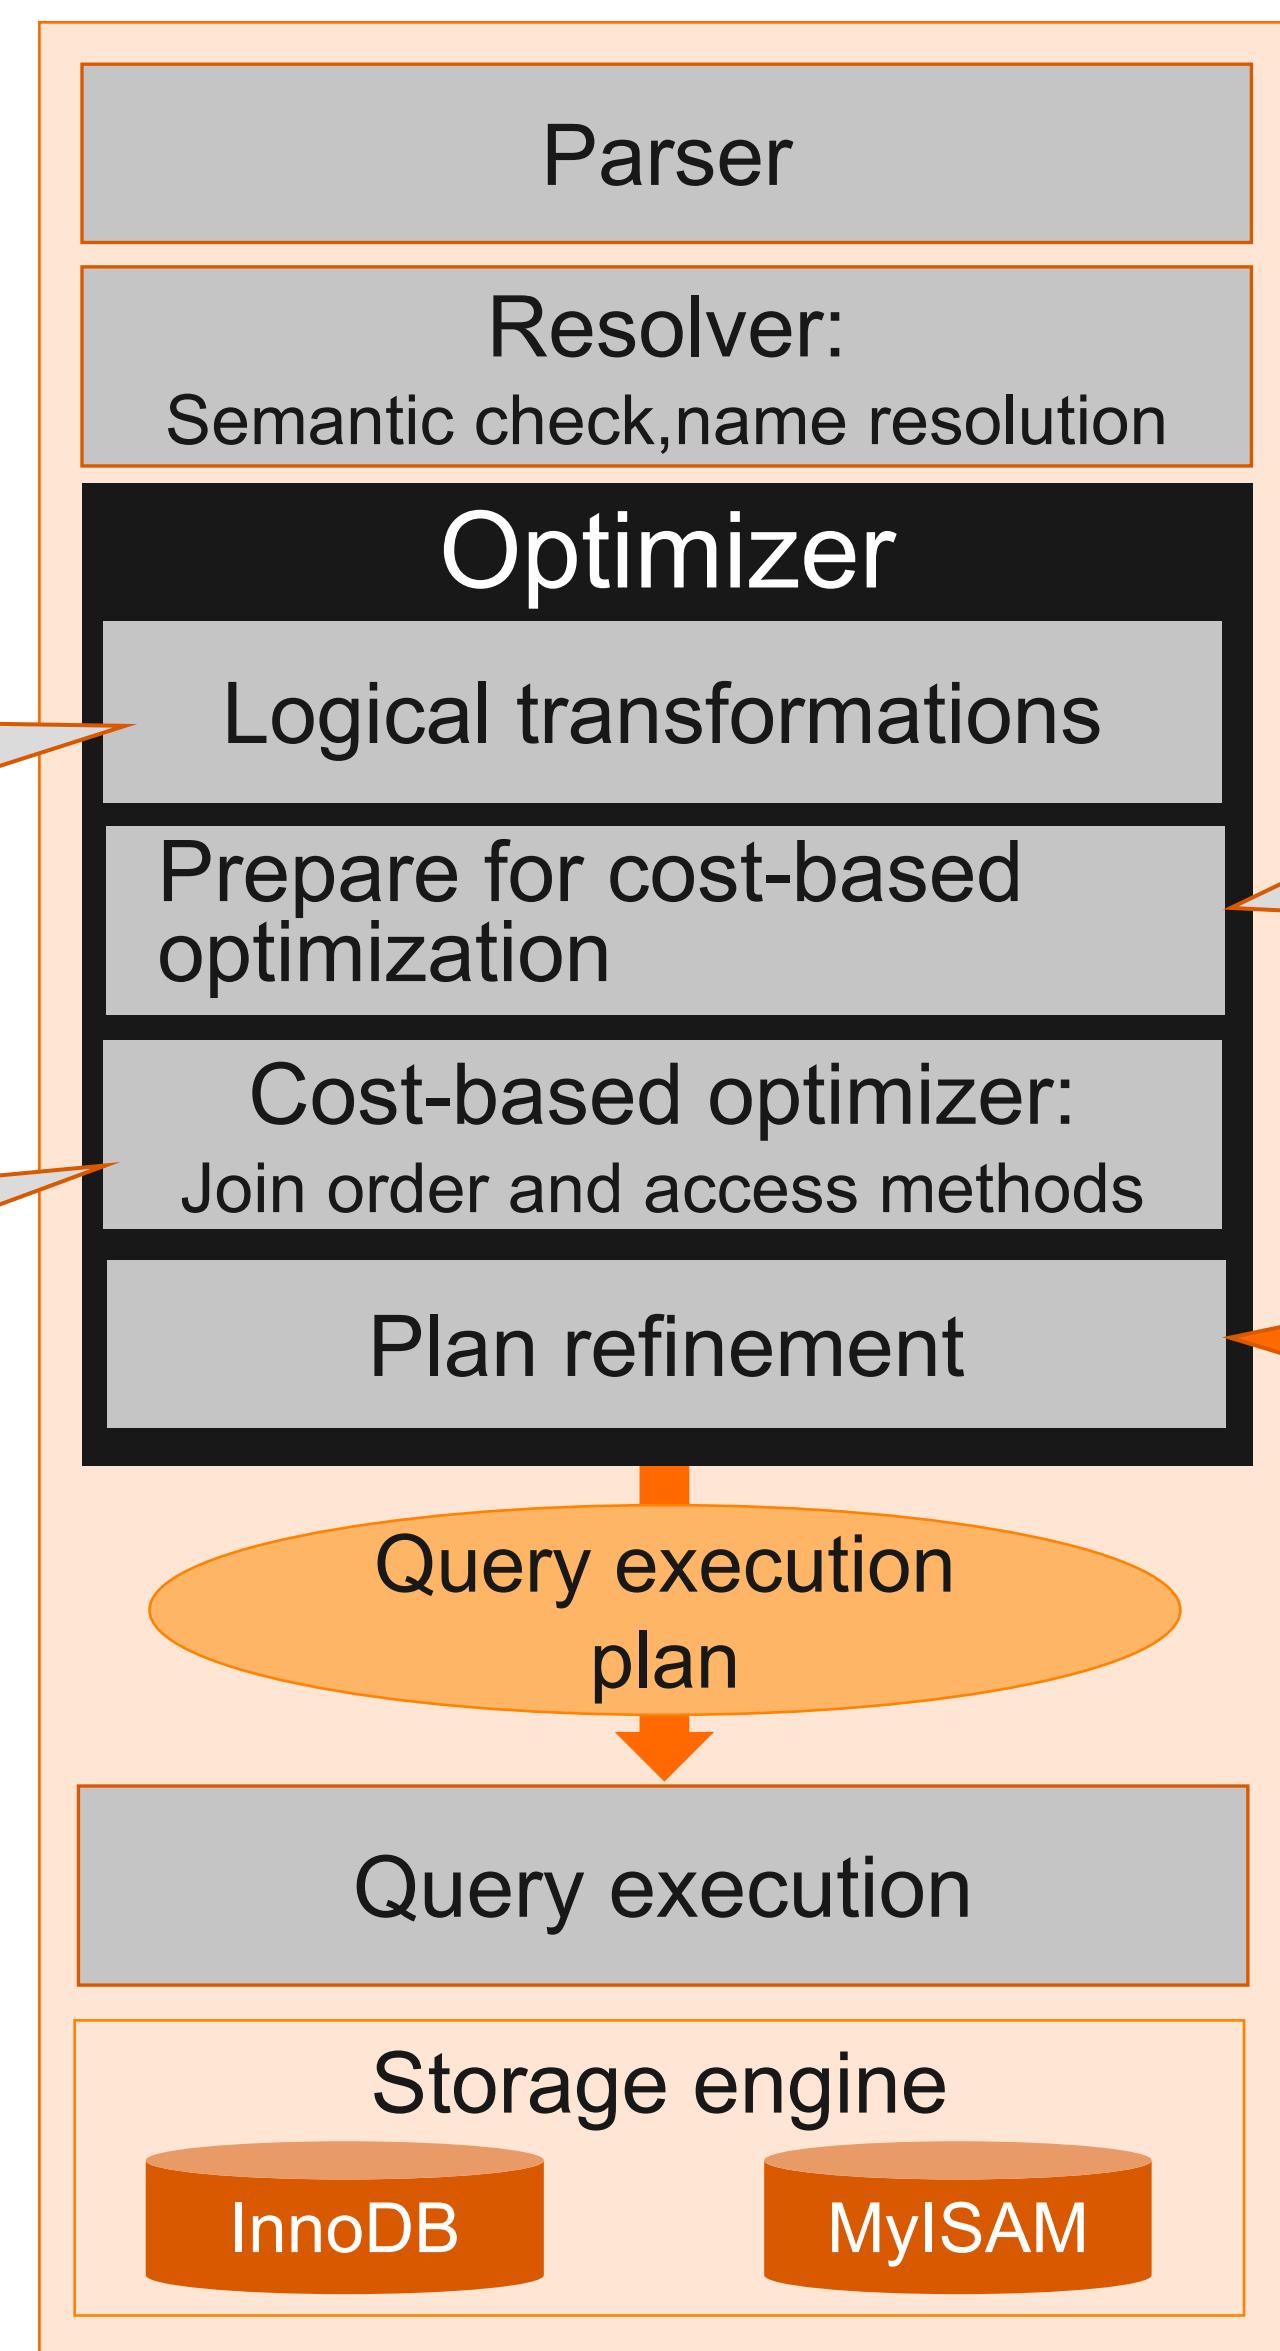
```
-----  
`customer` scan 2367/244.2 478.9/244.2  
`lineitem` scan 2495/119491 1190000/119735  
  `orders` eq_ref:PRIMARY 1/120108 59742/239843 *** NEW BEST PLAN ***  
`orders` scan 2502/101621 478.9/101865  
  `lineitem` ref:PRIMARY 1.0155/168.35 486.3/102033 *** NEW BEST PLAN ***  
`lineitem` scan 2495/255.5 2495/255.5  
`customer` scan 2367/119716 1190000/119972 PRUNED(cost)  
`orders` eq_ref:PRIMARY 1/873.25 2115.1/1128.8  
  `customer` eq_ref:PRIMARY 1/740.27 427.92/1869 *** NEW BEST PLAN ***  
`orders` scan 2502/255.7 2121/255.7  
  `customer` eq_ref:PRIMARY 1/742.35 429.12/998.05  
    `lineitem` ref:PRIMARY 1.0155/150.86 435.76/1148.9 *** NEW BEST PLAN ***  
  `lineitem` ref:PRIMARY 1.0155/745.63 2153.8/1001.3 PRUNED(heuristic)
```

Query Optimization

Main phases

Negation elimination
Equality and constant propagation
Evaluation of constant expressions
Substitution of generated columns

Access method selection
Join order



Ref access analysis
Range access analysis
Estimation of condition fan out
Constant table detection

Table condition pushdown
Access method adjustments
Sort avoidance
Index condition pushdown
Prepare temporary tables

Assigning Query Conditions to Tables

Evaluate conditions as early as possible in join order

```
"attaching_conditions_to_tables": {  
    "original_condition": "((`customer`.`c_custkey` = `orders`.`o_custkey`) and (`lineitem`.`l_orderkey` = `orders`.`o_orderkey`) and  
(`customer`.`c_mktsegment` = 'FURNITURE') and (`orders`.`o_orderDATE` < DATE'1997-04-15') and (`lineitem`.`l_shipDATE` > DATE'1997-  
04-15'))",  
    "attached_conditions_computation": [  
        {  
            "table": "`orders`",  
            "rechecking_index_usage": {  
                "recheck_reason": "low_limit",  
                "limit": 10,  
                "row_estimate": 2121 } }],  
    "attached_conditions_summary": [  
        {  
            "table": "`orders`",  
            "attached": "((`orders`.`o_orderDATE` < DATE'1997-04-15') and (`orders`.`o_custkey` is not null))",  
            {  
                "table": "`customer`",  
                "attached": "((`customer`.`c_custkey` = `orders`.`o_custkey`) and (`customer`.`c_mktsegment` = 'FURNITURE'))",  
                {  
                    "table": "`lineitem`",  
                    "attached": "((`lineitem`.`l_orderkey` = `orders`.`o_orderkey`) and (`lineitem`.`l_shipDATE` > DATE'1997-04-15'))" } ] }]
```

Assigning Query Conditions to Tables

Remove conditions satisfied by ref access

```
"finalizing_table_conditions": [  
    {  
        "table": "`orders`",  
        "original_table_condition": "((`orders`.`o_orderDATE` < DATE'1997-04-15') and (`orders`.`o_custkey` is not null))",  
        "final_table_condition": "((`orders`.`o_orderDATE` < DATE'1997-04-15') and (`orders`.`o_custkey` is not null))",  
        {  
            "table": "`customer`",  
            "original_table_condition": "((`customer`.`c_custkey` = `orders`.`o_custkey`) and (`customer`.`c_mktsegment` = 'FURNITURE'))",  
            "final_table_condition": "(`customer`.`c_mktsegment` = 'FURNITURE')",  
            {  
                "table": "`lineitem`",  
                "original_table_condition": "((`lineitem`.`l_orderkey` = `orders`.`o_orderkey`) and (`lineitem`.`l_shipDATE` > DATE'1997-04-15'))",  
                "final_table_condition": "(`lineitem`.`l_shipDATE` > DATE'1997-04-15')"  
            }  
        }  
    }  
]
```

ORDER BY Optimization

Avoid sorting, if possible

- Change to a different index that provides result in sorted order
- Read in descending order
- Example:

```
EXPLAIN SELECT * FROM orders WHERE o_totalprice > 400000
ORDER BY o_orderdate DESC LIMIT 10;
```

id	select type	table	type	possible keys	key	key len	rows	filtered	Extra
1	SIMPLE	orders	index	NULL	i_o_orderdate	4	10	0.57	Using where; Backward index scan

- Join queries:
 - Join order is already fixed.
 - Sorting can only be avoided if ordering is provided by an index from the first table in the join order

Assigning Query Conditions to Tables

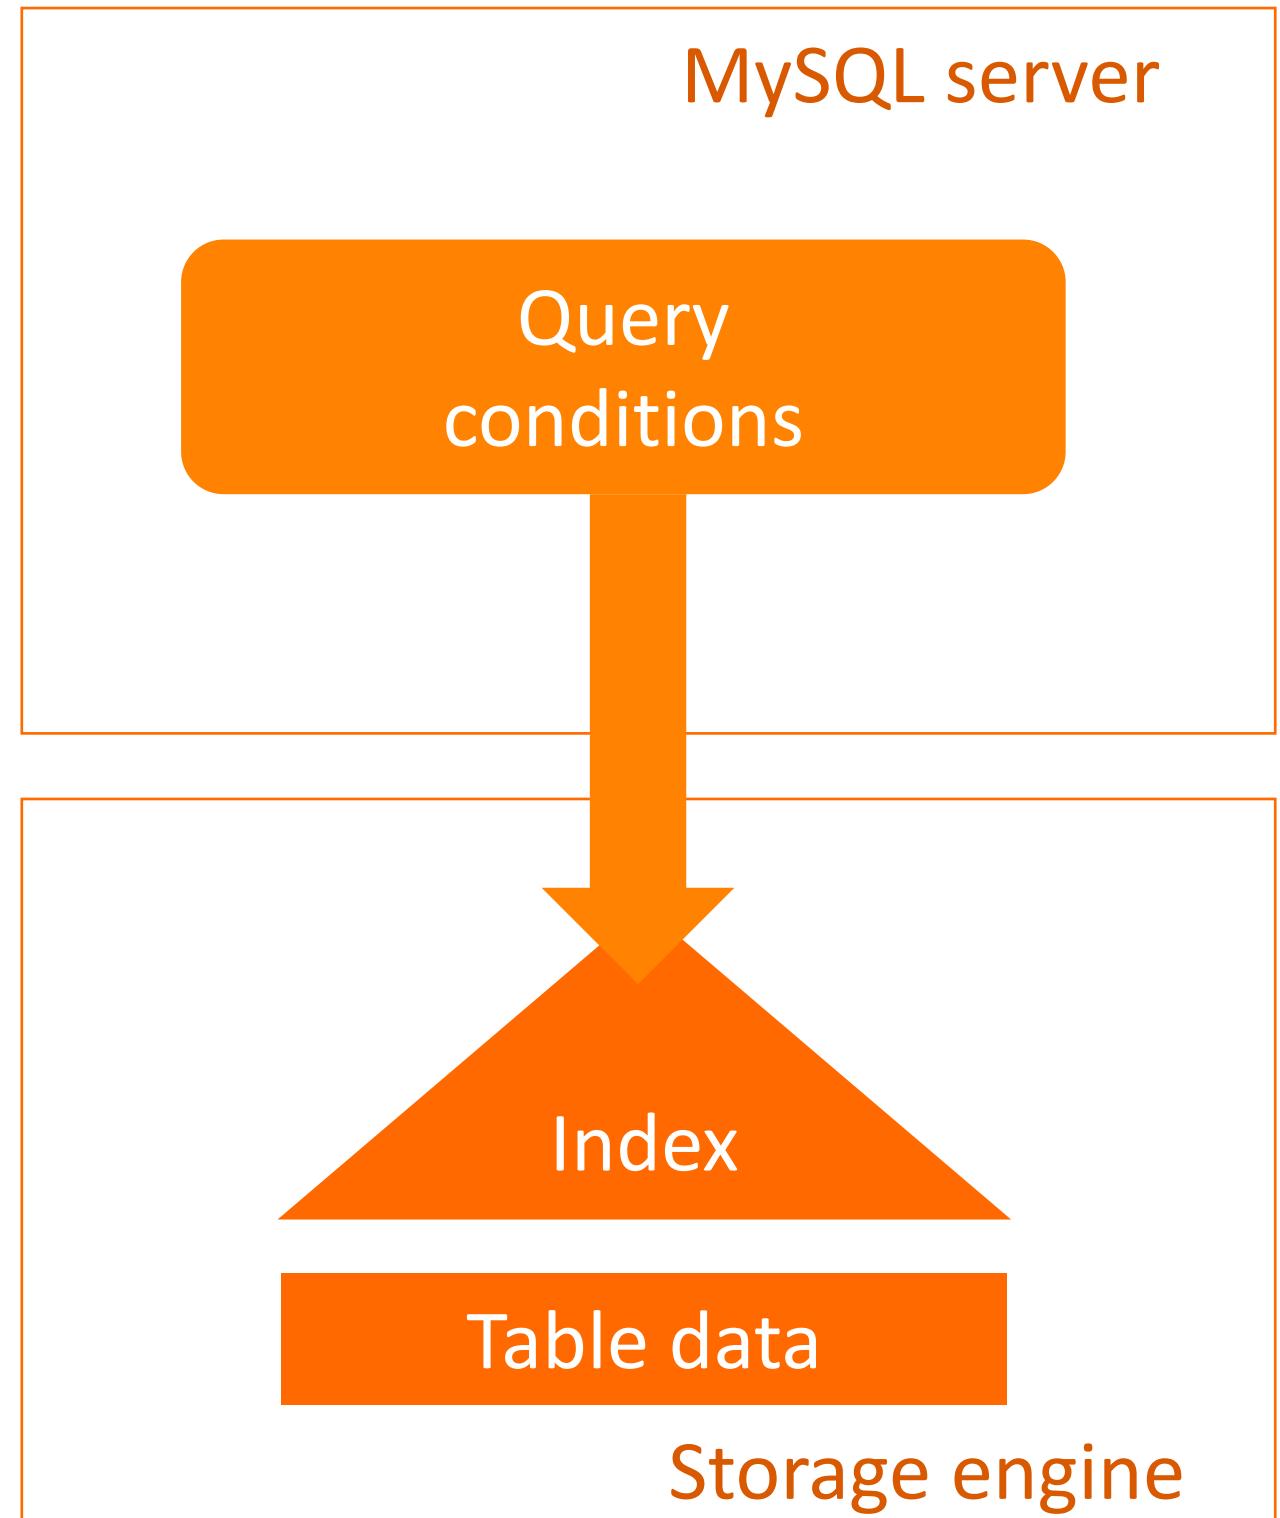
Evaluate conditions as early as possible in join order

```
{  
  "optimizing_distinct_group_by_order_by": {  
    "simplifying_order_by": {  
      "original_clause": "`orders`.`o_orderDATE` desc",  
      "items": [  
        {  
          "item": "`orders`.`o_orderDATE`"  
        }  
      ],  
      "resulting_clause_is_simple": true,  
      "resulting_clause": "`orders`.`o_orderDATE` desc"  
    }  
  },  
  {  
    "reconsidering_access_paths_for_index_ordering": {  
      "clause": "ORDER BY",  
      "steps": [  
      ],  
      "index_order_summary": {  
        "table": "`orders`",  
        "index_provides_order": true,  
        "order_direction": "desc",  
        "index": "i_o_orderdate",  
        "plan_changed": true,  
        "access_type": "index"  
      }  
    }  
  },  
},  
{  
  "reconsidering_access_paths_for_index_ordering": {  
    "clause": "ORDER BY",  
    "steps": [  
    ],  
    "index_order_summary": {  
      "table": "`orders`",  
      "index_provides_order": true,  
      "order_direction": "desc",  
      "index": "i_o_orderdate",  
      "plan_changed": true,  
      "access_type": "index"  
    }  
  }  
},  
}
```

- Unfortunately, no information on cost calculations!

Index Condition Pushdown

- Pushes conditions that can be evaluated on the index down to storage engine
 - Works only on indexed columns
- Goal: evaluate conditions without having to access the actual record
 - Reduces number of disk/block accesses
 - Reduces CPU usage



Index Condition Pushdown

How it works

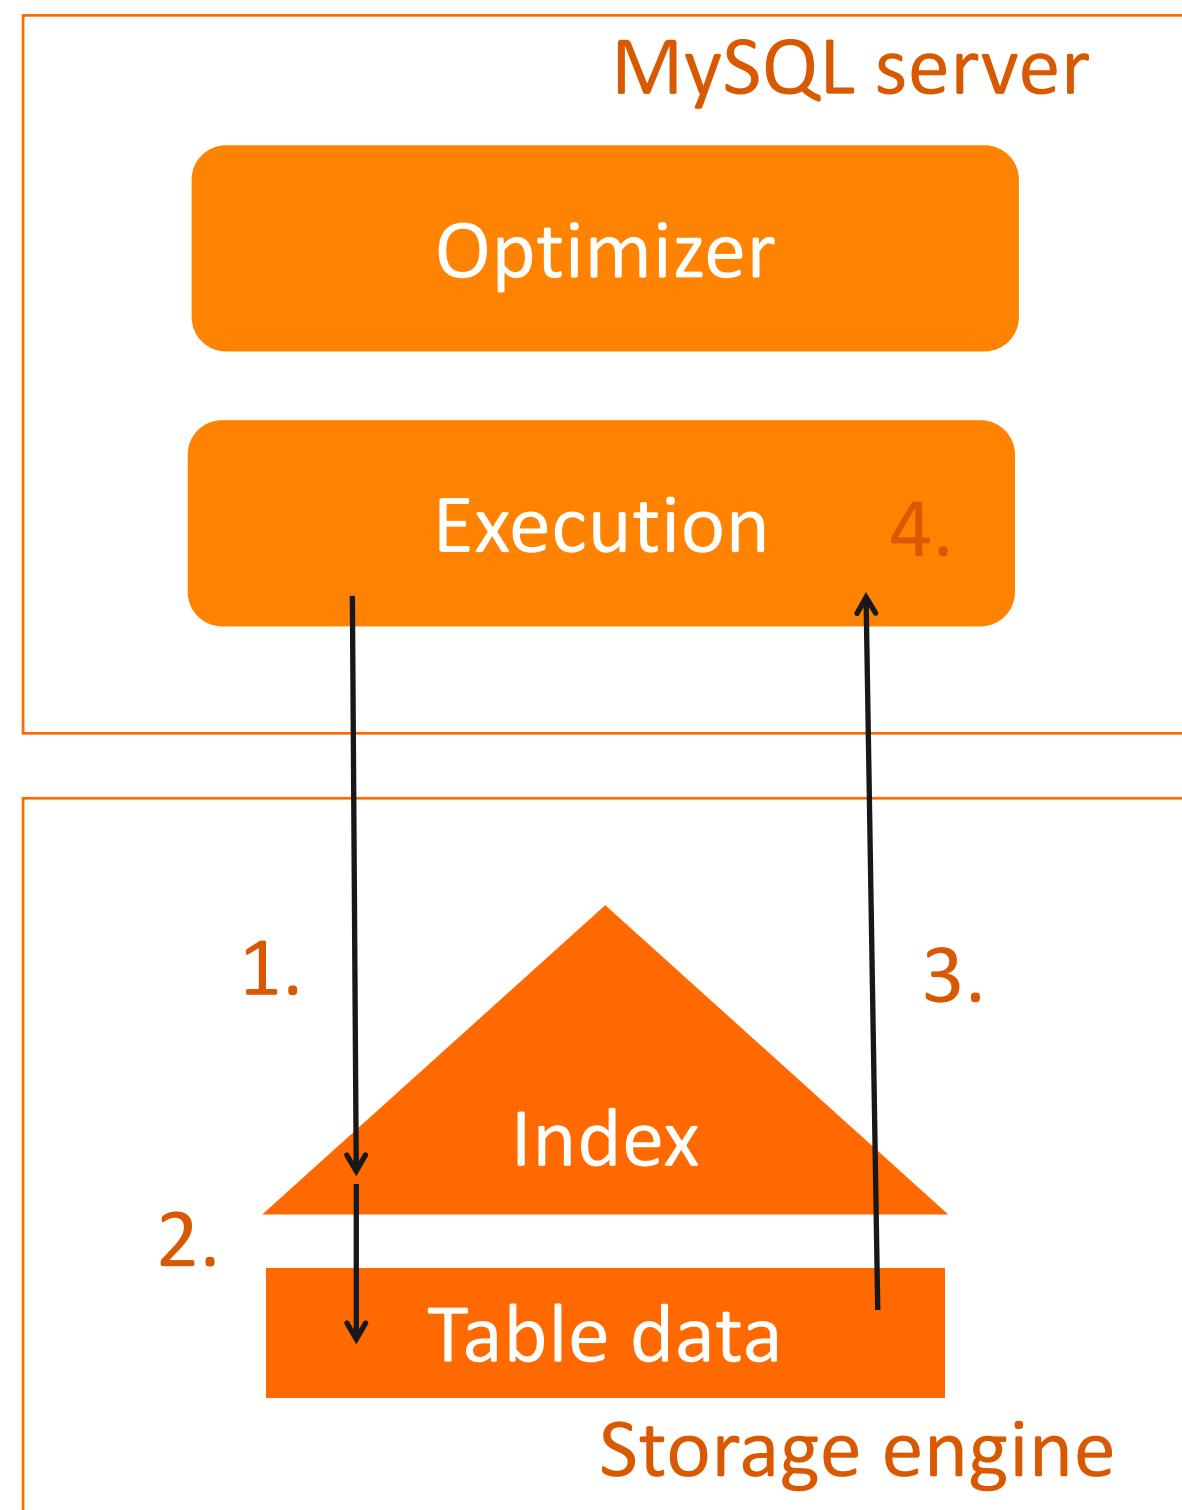
Without ICP:

Storage Engine:

1. Reads index
2. Reads record
3. Returns record

Server:

4. Evaluates condition



With ICP:

Storage Engine:

1. Reads index and **evaluates pushed index condition**
2. Reads record
3. Returns record

Server:

4. Evaluates **rest of condition**

Index Condition Pushdown

Optimizer trace

```
{  
  "refine_plan": [  
    {  
      "table": "part"  
    },  
    {  
      "table": "lineitem` FORCE INDEX ('i_l_partkey_plus')",  
      "pushed_index_condition": "(`lineitem`.`l_quantity` >= 1.00) and (`lineitem`.`l_quantity` <= <cache>((1 + 10))) and  
      (`lineitem`.`l_shipmode` in ('AIR','AIR REG'))",  
      "table_condition_attached": "(`lineitem`.`l_shipinstruct` = 'DELIVER IN PERSON')"  
    }  
  ],  
},  
Index: lineitem(l_partkey,l_quantity,l_shipmode)
```

Prepare Temporary Tables

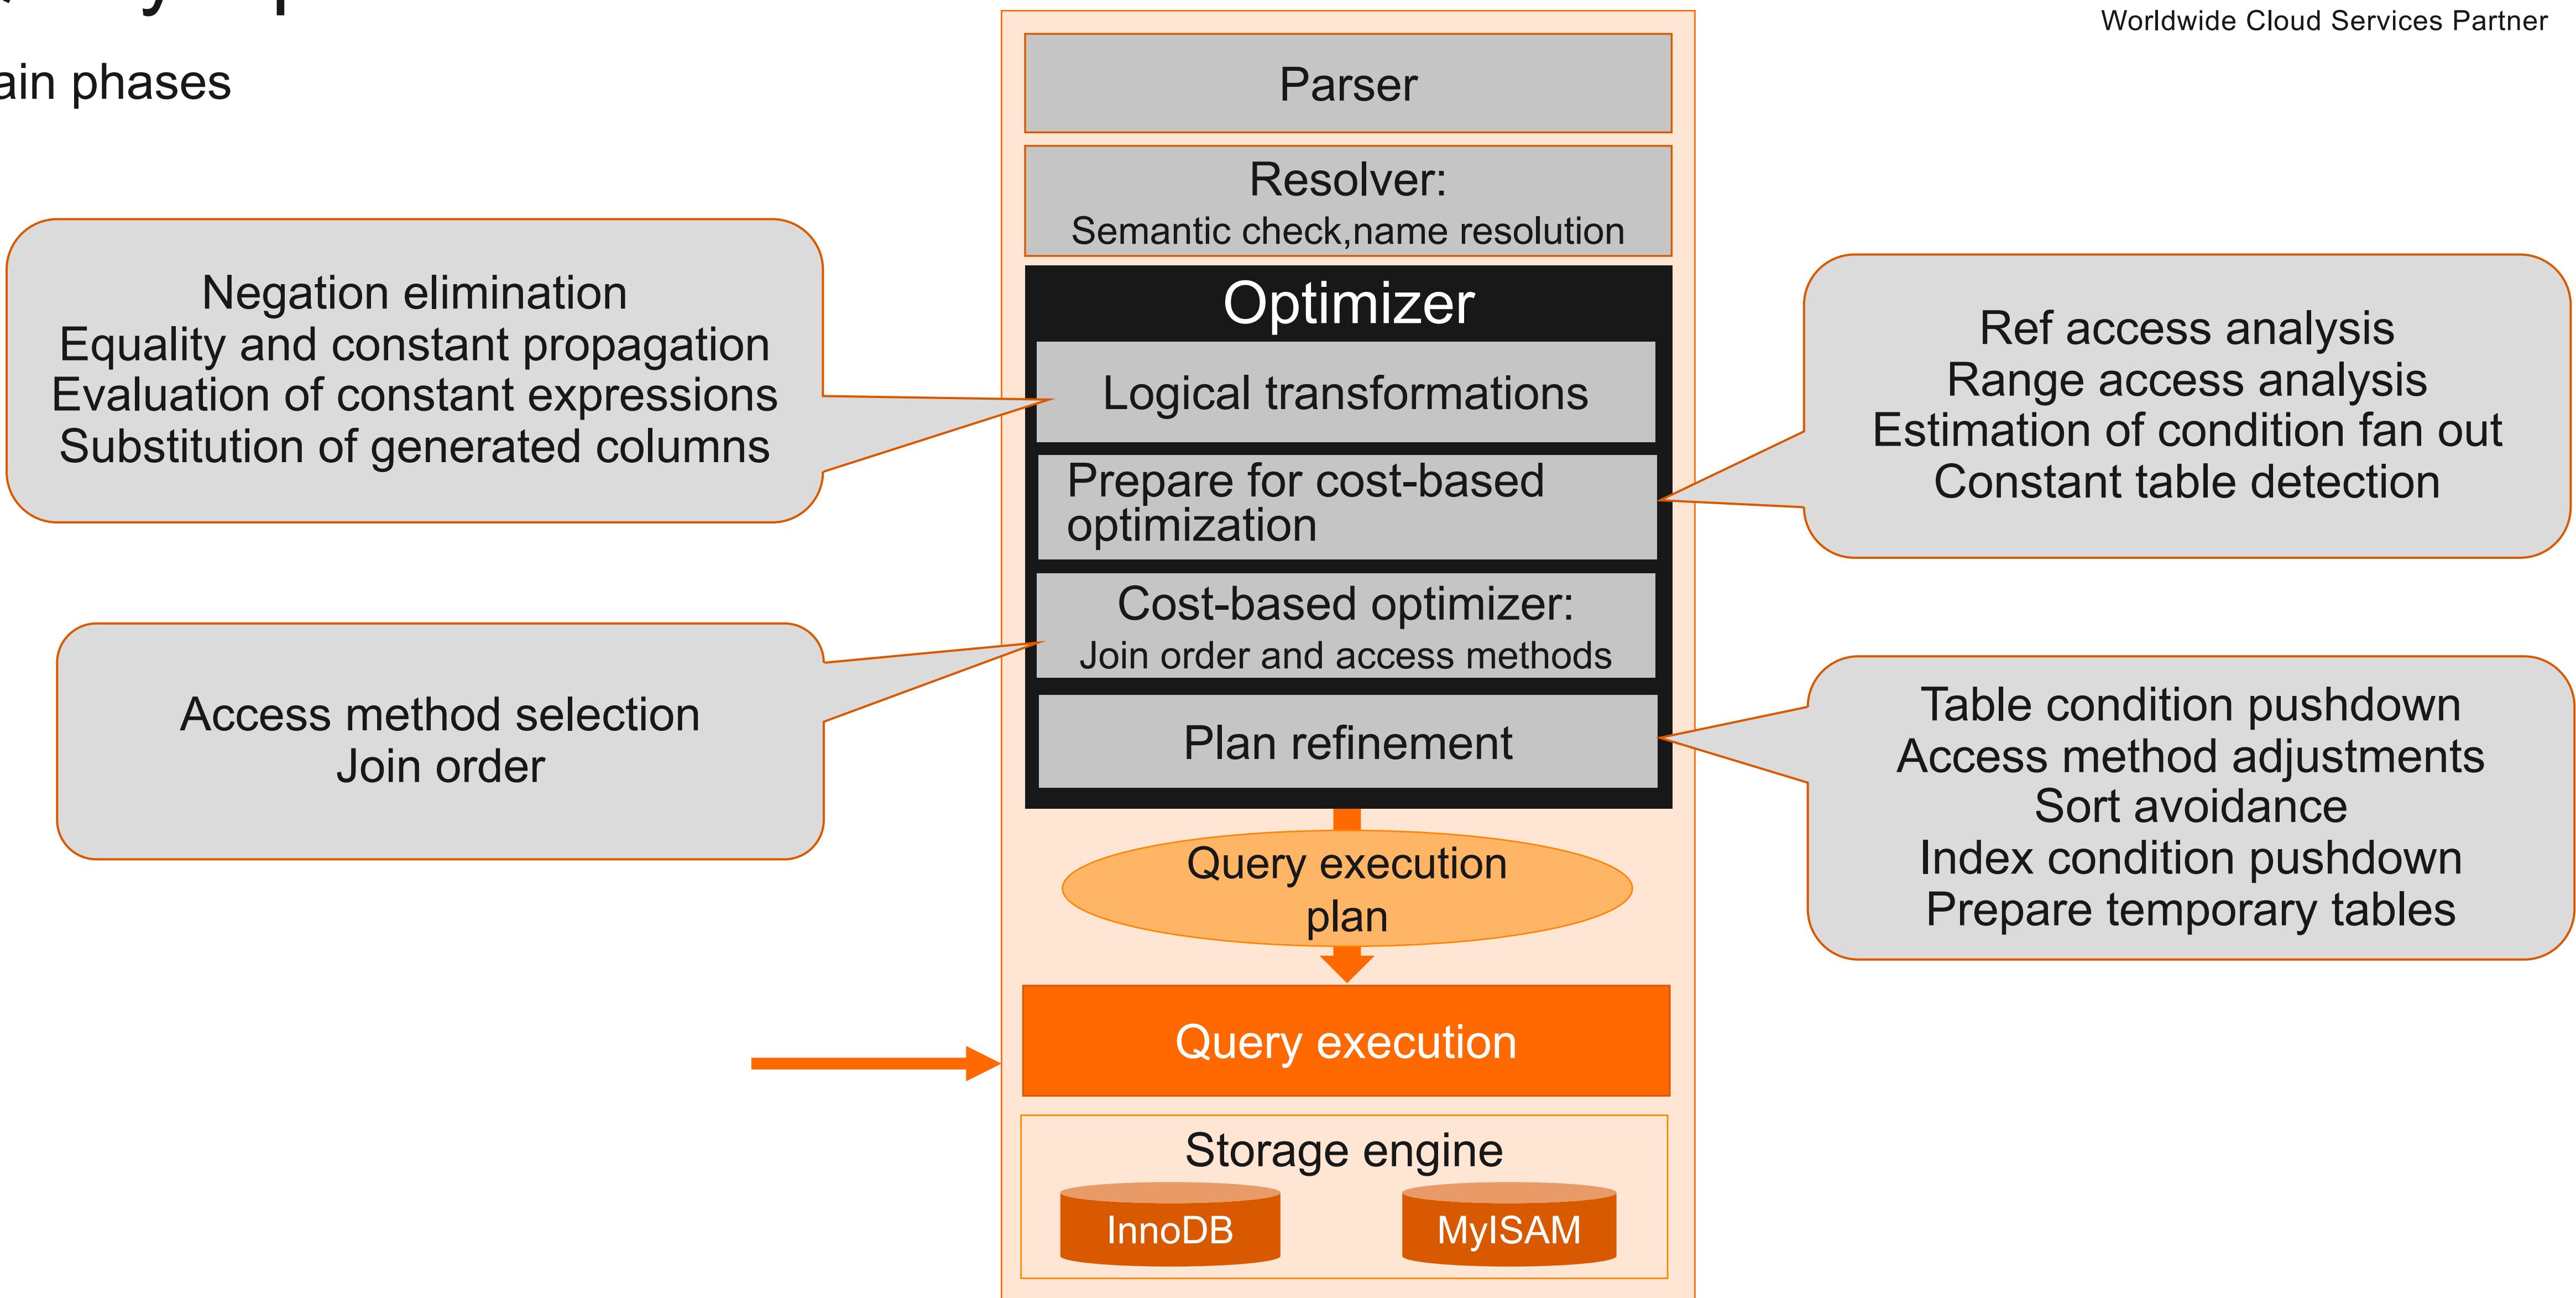
DBT3-Q3

```
SELECT l_orderkey, sum(l_extendedprice * (1 - l_discount)) AS revenue,
       o_orderdate, o_shippriority
  FROM customer JOIN orders ON c_custkey = o_custkey
                  JOIN lineitem ON l_orderkey = o_orderkey
 WHERE c_mktsegment = 'FURNITURE'
   AND o_orderdate < '1997-04-15' AND l_shipdate > '1997-04-15'
 GROUP by l_orderkey, o_orderdate, o_shippriority
 ORDER by revenue desc, o_orderdate
 LIMIT 10;
```

```
{
  "considering_tmp_tables": [
    {
      "adding_tmp_table_in_plan_at_position": 3,
      "write_method": "continuously_update_group_row"
    },
    {
      "adding_sort_to_table_in_plan_at_position": 3
    }
  ]
}
```

Query Optimization

Main phases



Query Execution Trace

- Information on temporary table creation
- Sort parameters and summary
- Subquery execution
- Dynamic range optimization
- May create large volumes of trace
 - Not available after EXPLAIN
 - Volume may be reduced by setting session variable:

```
SET optimizer_trace_features = "dynamic_range=off, repeated_subselect=off";
```

Query Execution Trace

Temporary table information

```
{  
  "join_execution": {  
    "select#": 1,  
    "steps": [  
      {  
        "temp_table_aggregate": {  
          "select#": 1,  
          "steps": [  
            {  
              "creating_tmp_table": {  
                "tmp_table_info": {  
                  "in_plan_at_position": 3,  
                  "columns": 4,  
                  "row_length": 34,  
                  "key_length": 13,  
                  "unique_constraint": false,  
                  "makes_grouped_rows": true,  
                  "cannot_insert_duplicates": false,  
                  "location": "TempTable"  
                } } ] } }
```

Query Execution Trace

Sort information

```
{  
  "sorting_table_in_plan_at_position": 3,  
  "filesort_information": [  
    {  
      "direction": "desc",  
      "field": "revenue"  
    },  
    {  
      "direction": "asc",  
      "field": "o_orderDATE"  
    }  
  ],  
  "filesort_priority_queue_optimization": {  
    "limit": 10,  
    "chosen": true  
  },  
  "filesort_execution": [  
  ],  
  "filesort_summary": {  
    "memory_available": 262144,  
    "key_size": 33,  
    "row_size": 33,  
    "max_rows_per_buffer": 11,  
    "num_rows_estimate": 452,  
    "num_rows_found": 442,  
    "num_initial_chunks_spilled_to_disk": 0,  
    "peak_memory_used": 451,  
    "sort_algorithm": "std::sort",  
    "unpacked_addon_fields": "using_priority_queue",  
    "sort_mode": "<fixed_sort_key, rowid>"  
  }  
}
```

A Last Tip

Use MySQL JSON functionality to process the optimizer trace

- Example: Present the filesort summary from the last execution:

```
SELECT JSON_PRETTY(trace->'$.steps[*].join_execution.steps[*].filesort_summary')
FROM information_schema.optimizer_trace;
```

- Queries against `information_schema.optimizer_trace` will not generate a new trace, so you can execute multiple queries on the same trace.

More information

- The optimizer trace manual: <https://dev.mysql.com/doc/internals/en/optimizer-tracing.html>
 - Enable/disable trace of selected features of the optimizer
 - Tune trace purging (by default an optimizer trace overwrites the previous trace)
 - Make it more human readable (but no longer valid JSON)
- My blog: <https://oysteing.blogspot.com>
- For information about MySQL related work at Alibaba:
Attend the **AliSQL & POLARDB** track tomorrow

Thank You!

Please, remember to rate this session in the
Percona Live APP!

